

GANs for Image Processing & Computer Vision

Taeoh Kim
MVP Lab Seminar - Refined
Episode-4, 2017. 08. 23

Contents

1. GAN

2. GAN Papers

References

GAN based Papers

In CVPR 2017 (Main)

In ICML 2017

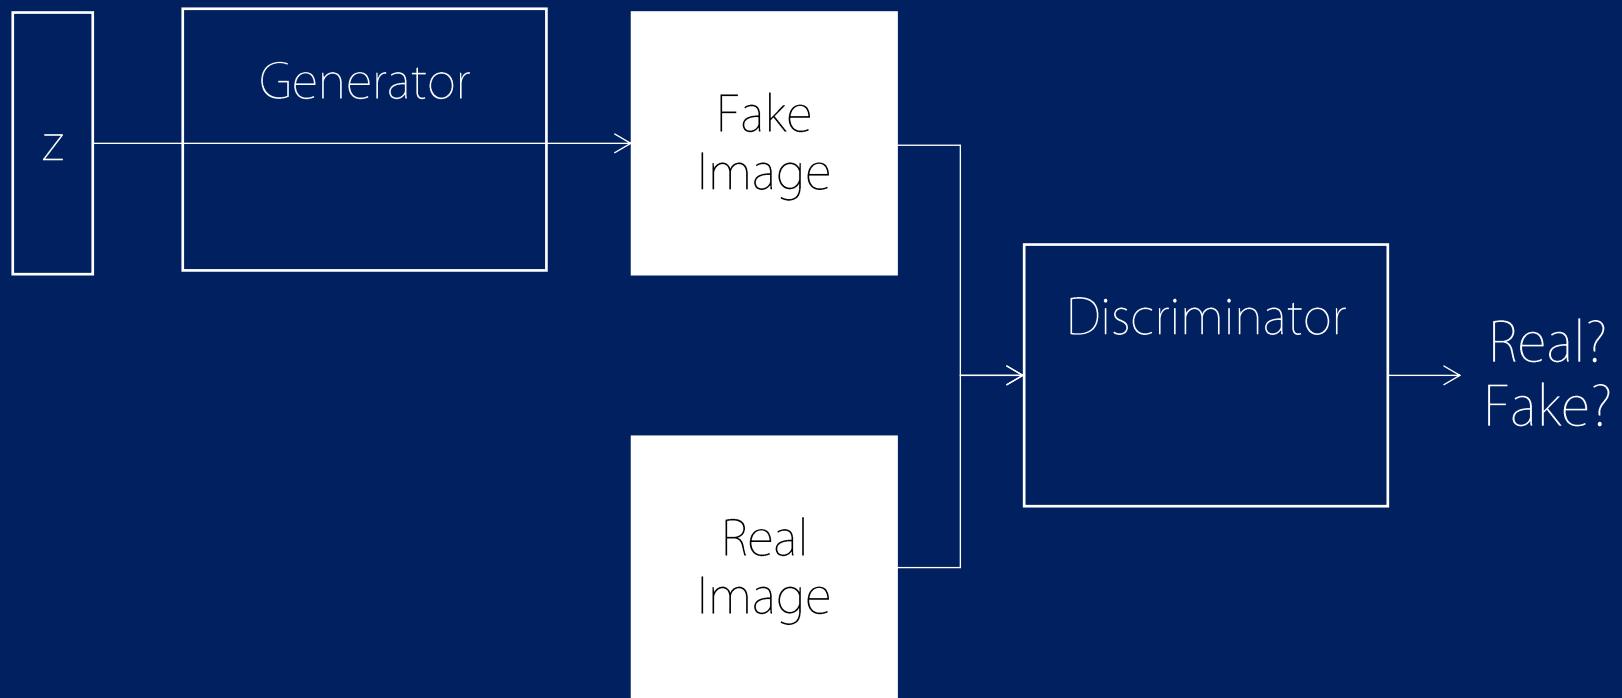
In ICCV 2017

Part 1. GAN

Generative

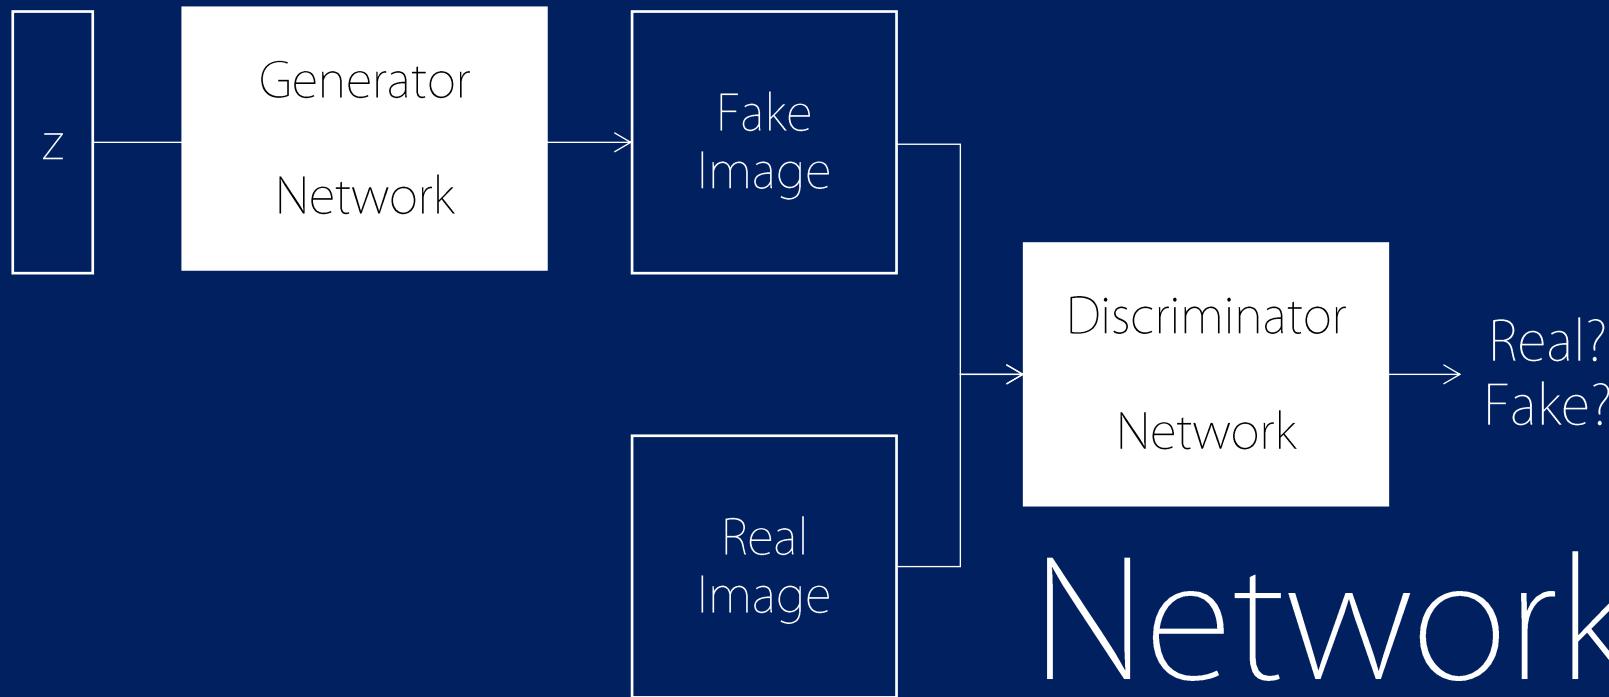


Generative



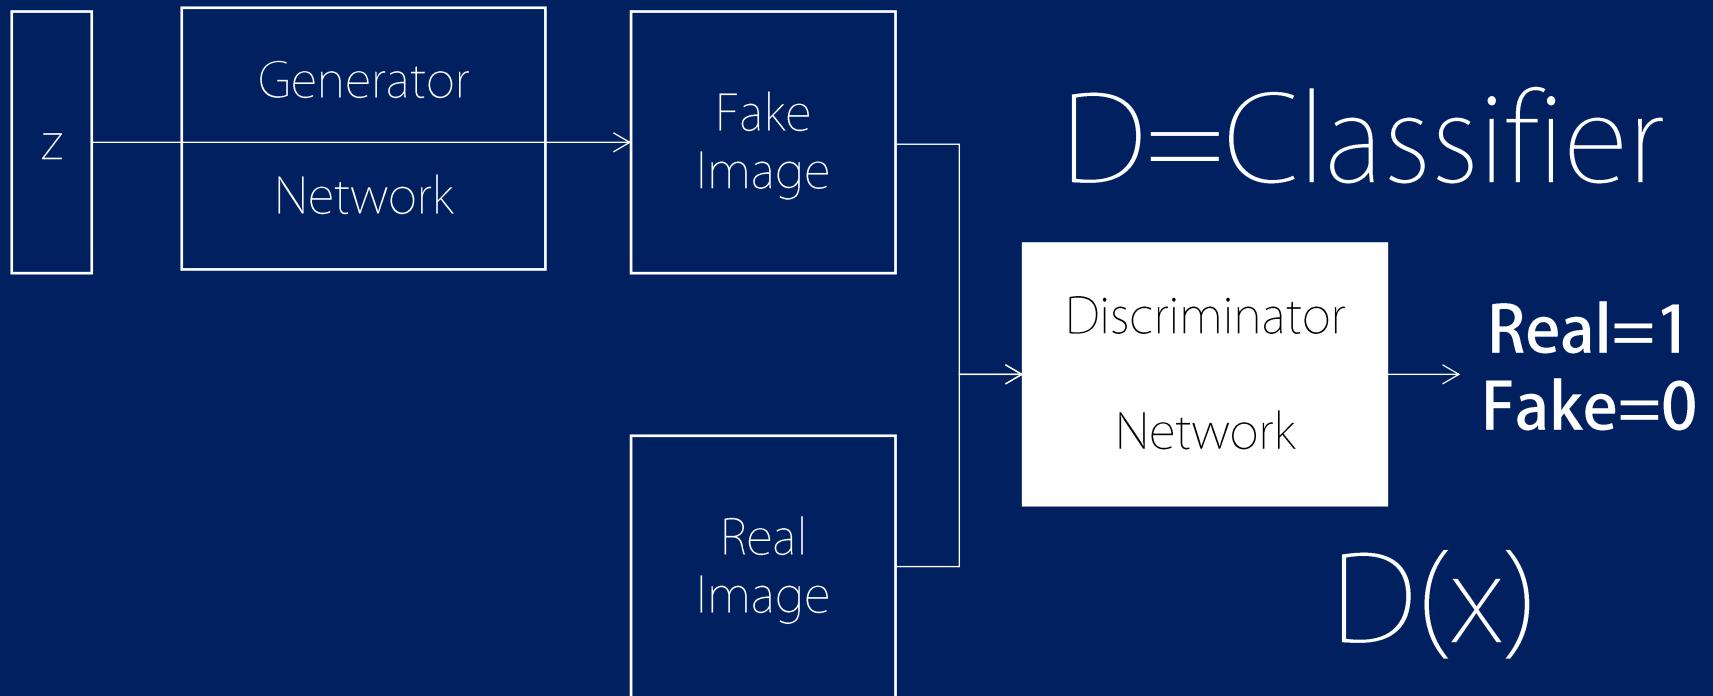
Adversarial

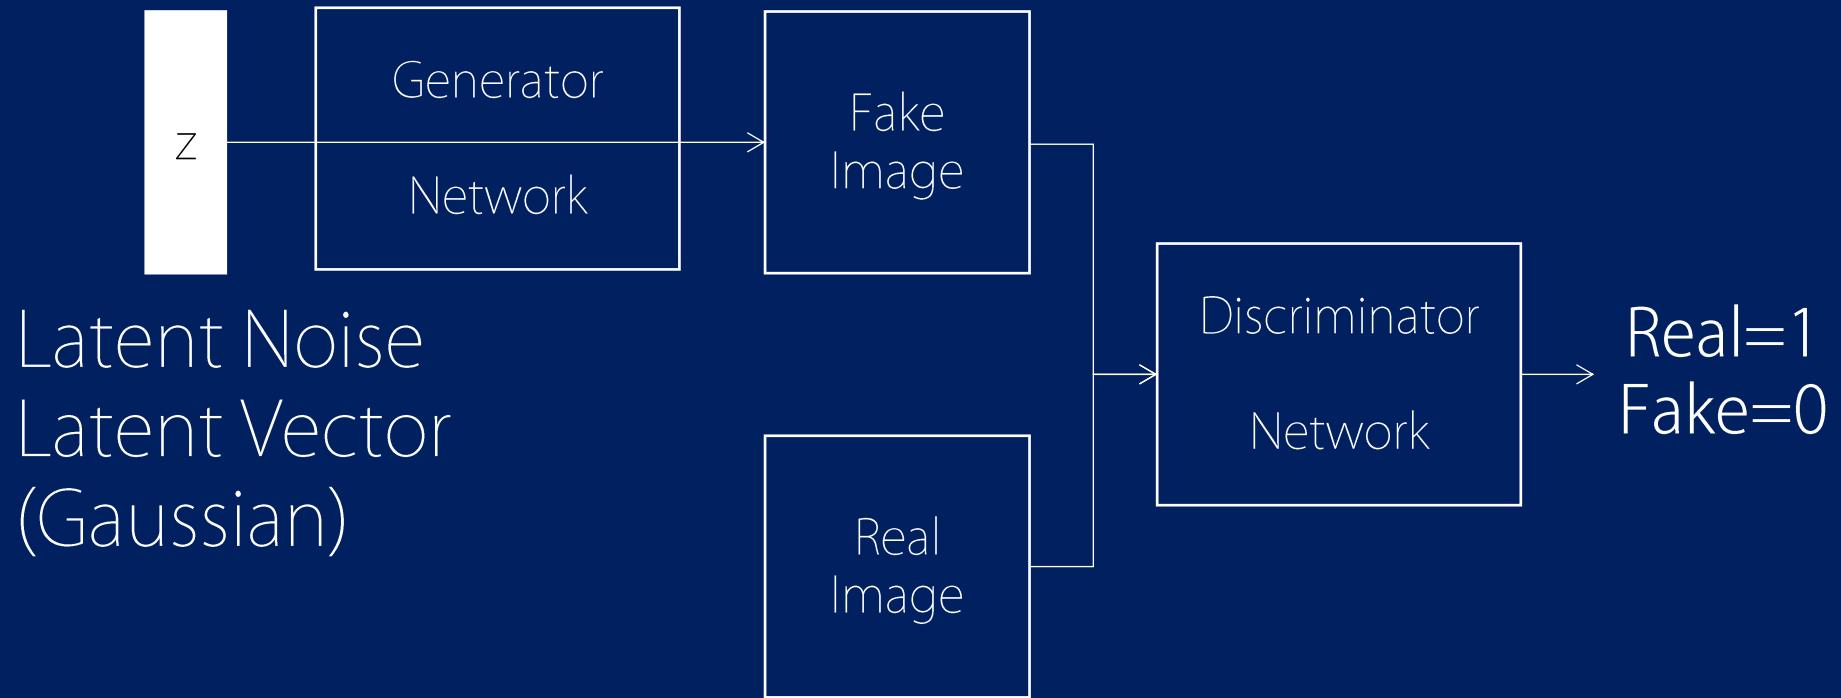
Generative



Adversarial

Networks

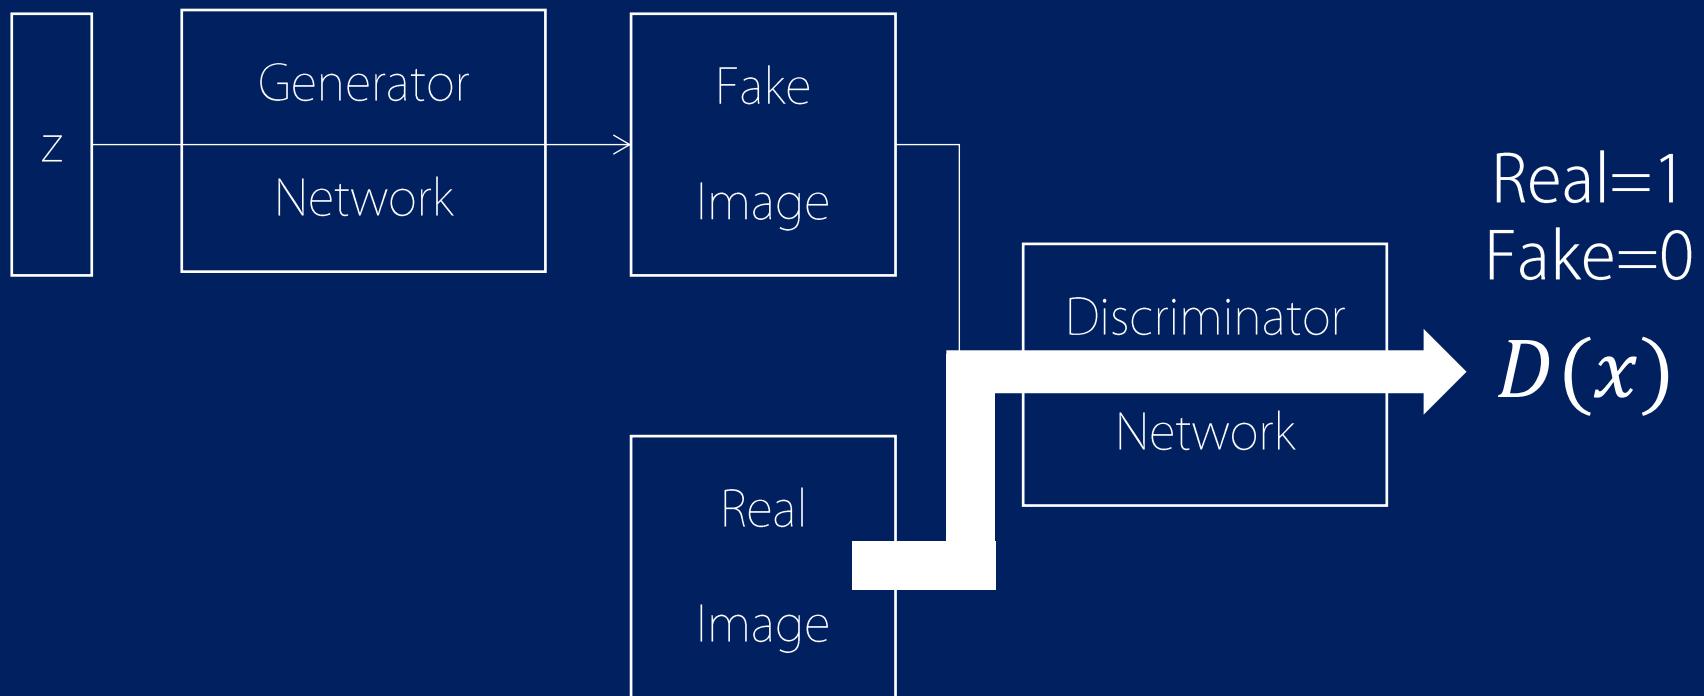




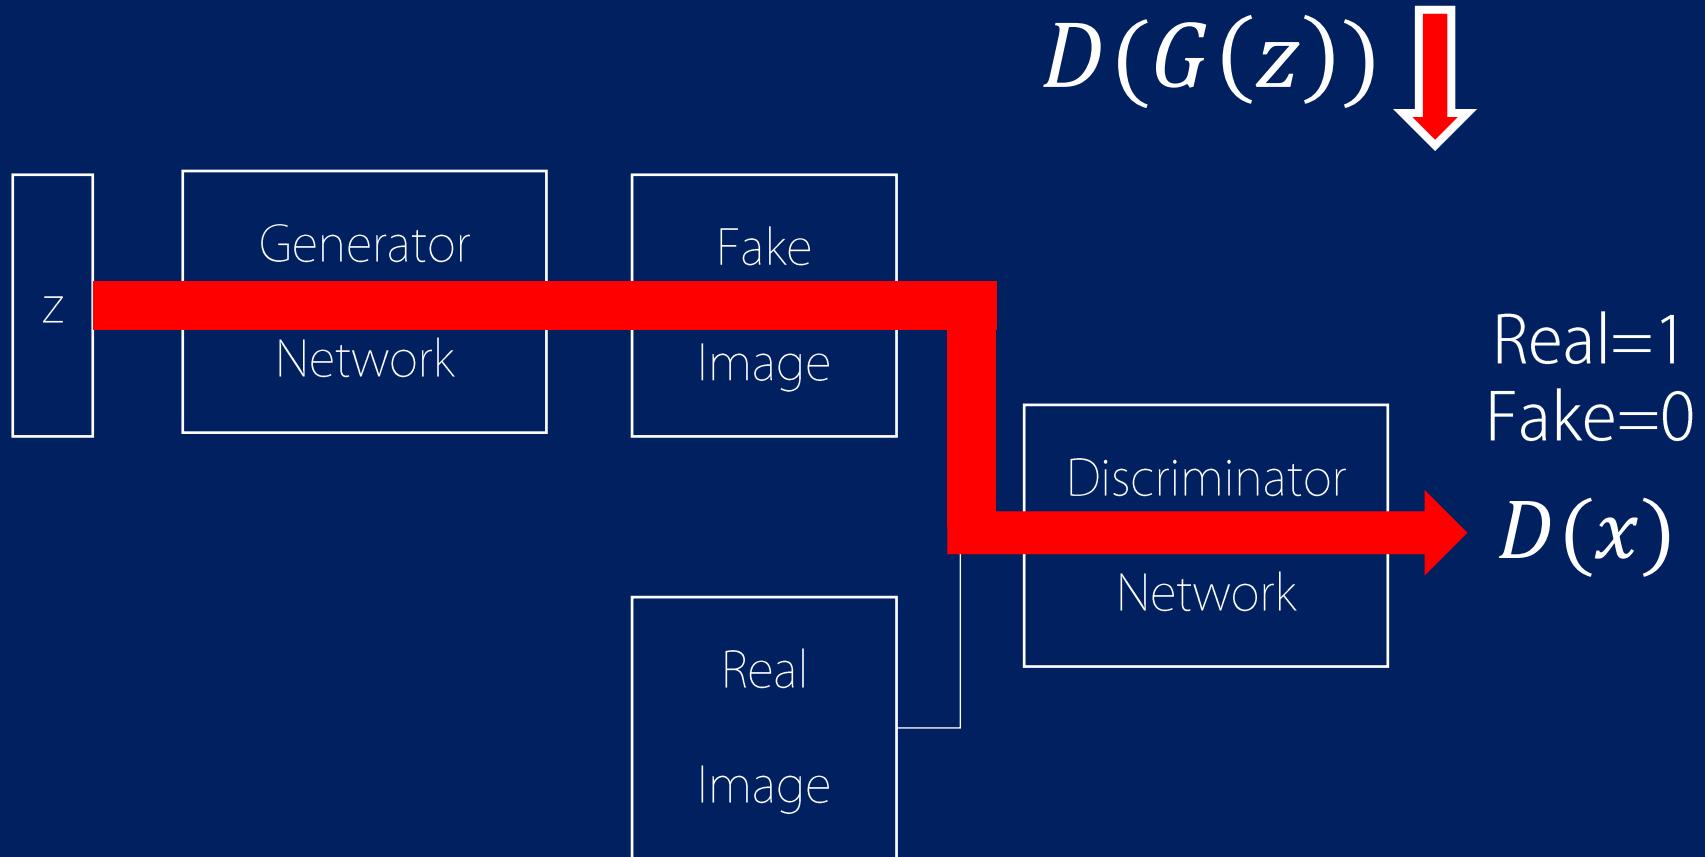
Training GAN

Training Discriminator

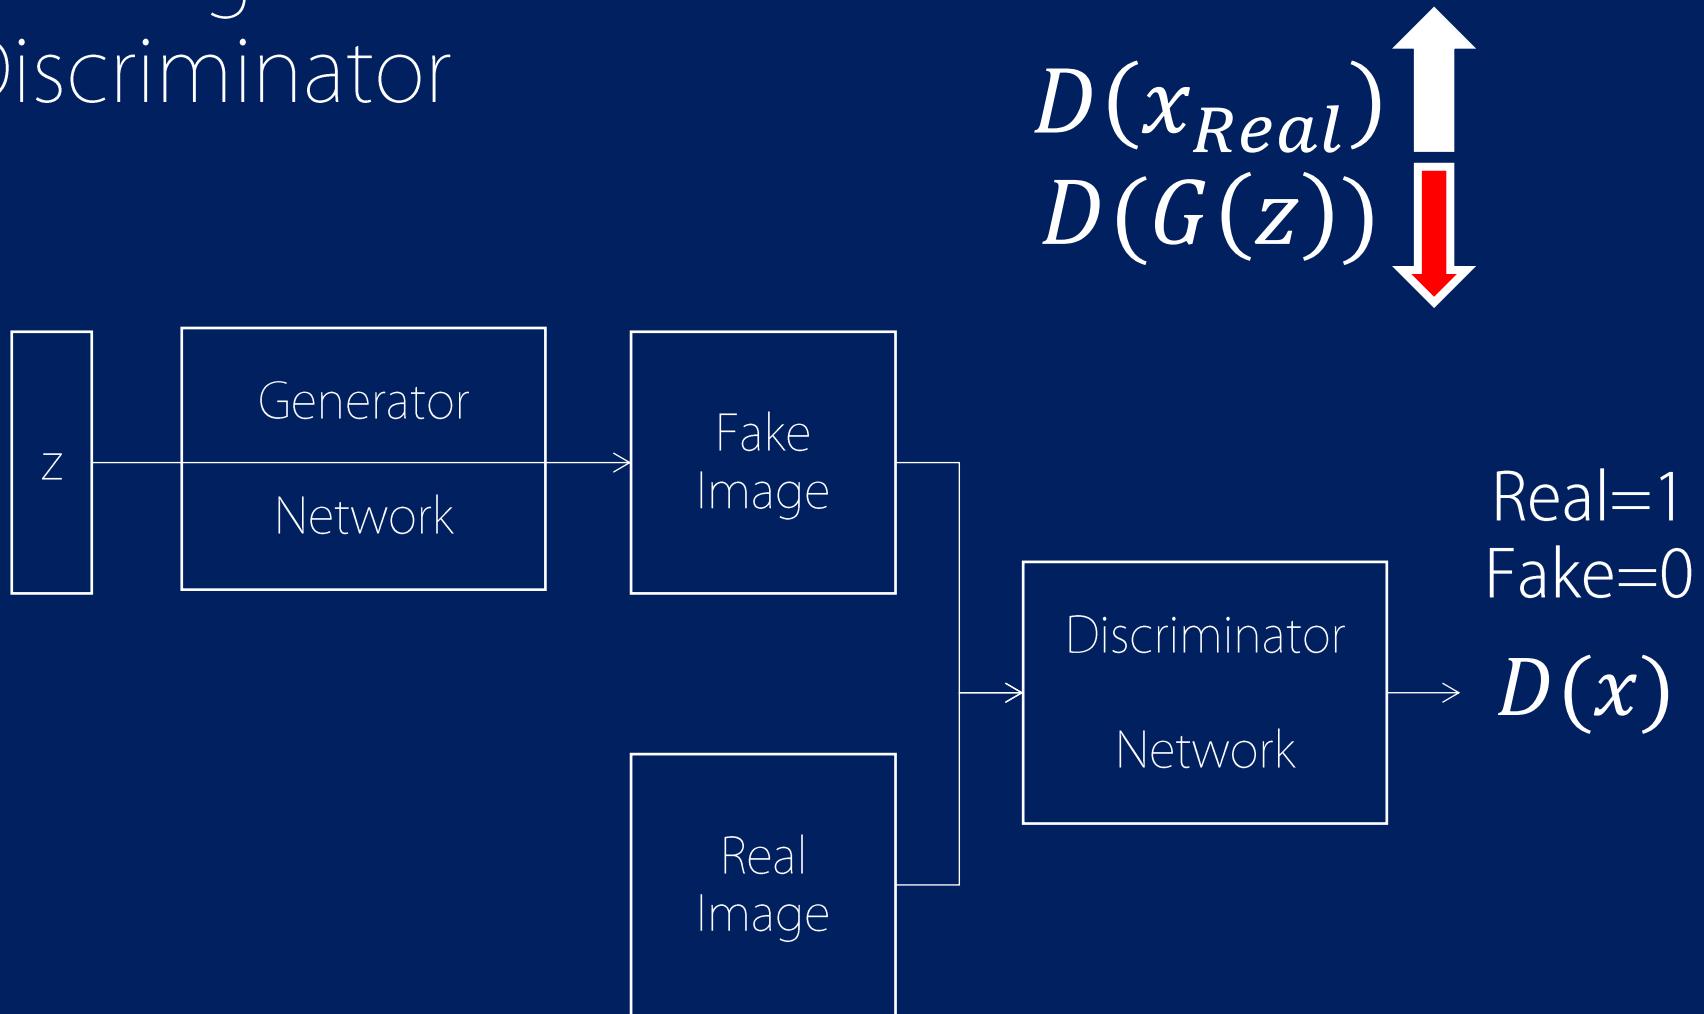
$$D(x_{Real}) \uparrow$$



Training Discriminator



Training Discriminator



$$D(x_{Real})$$

$$D(G(z))$$

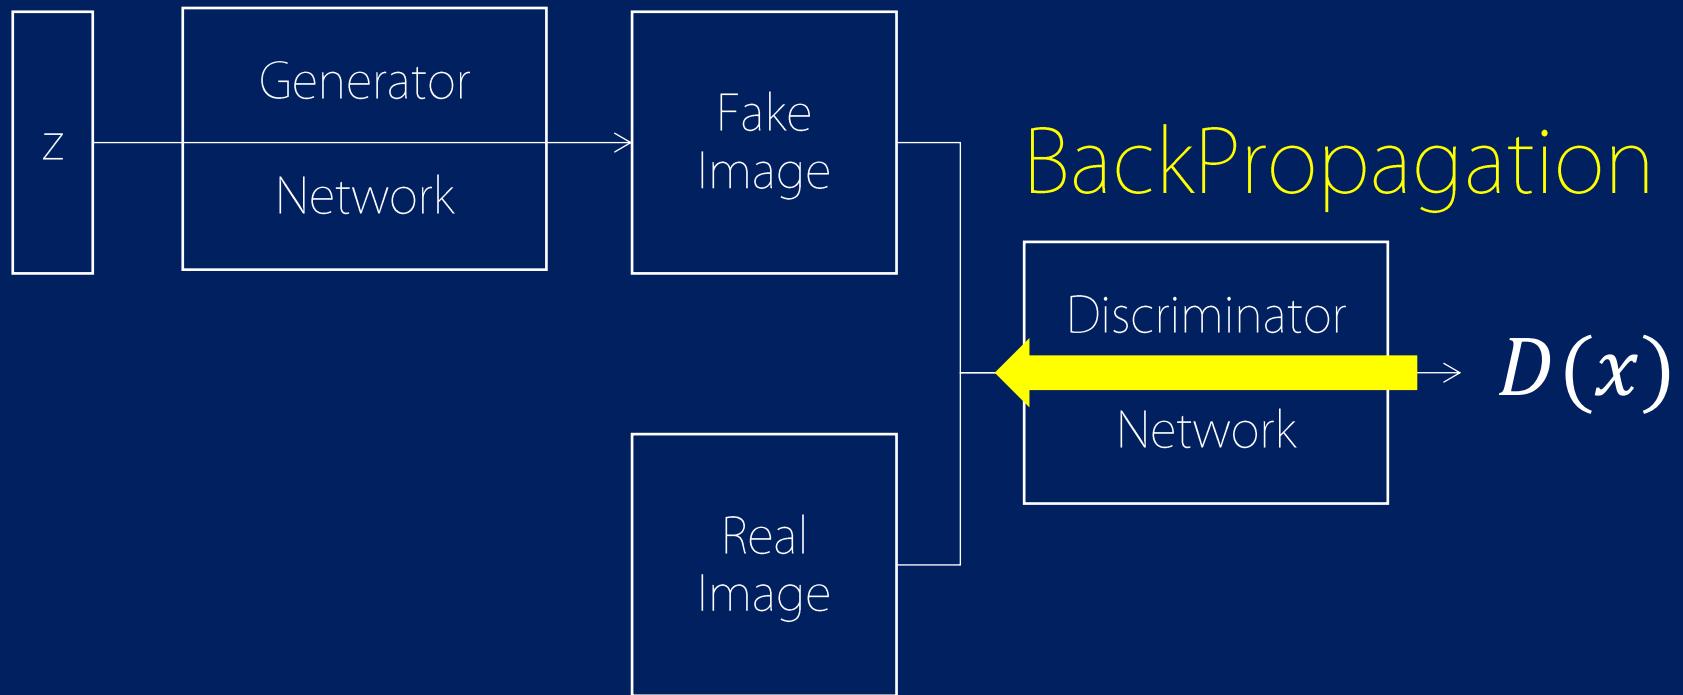


$$\text{Real}=1$$

$$\text{Fake}=0$$

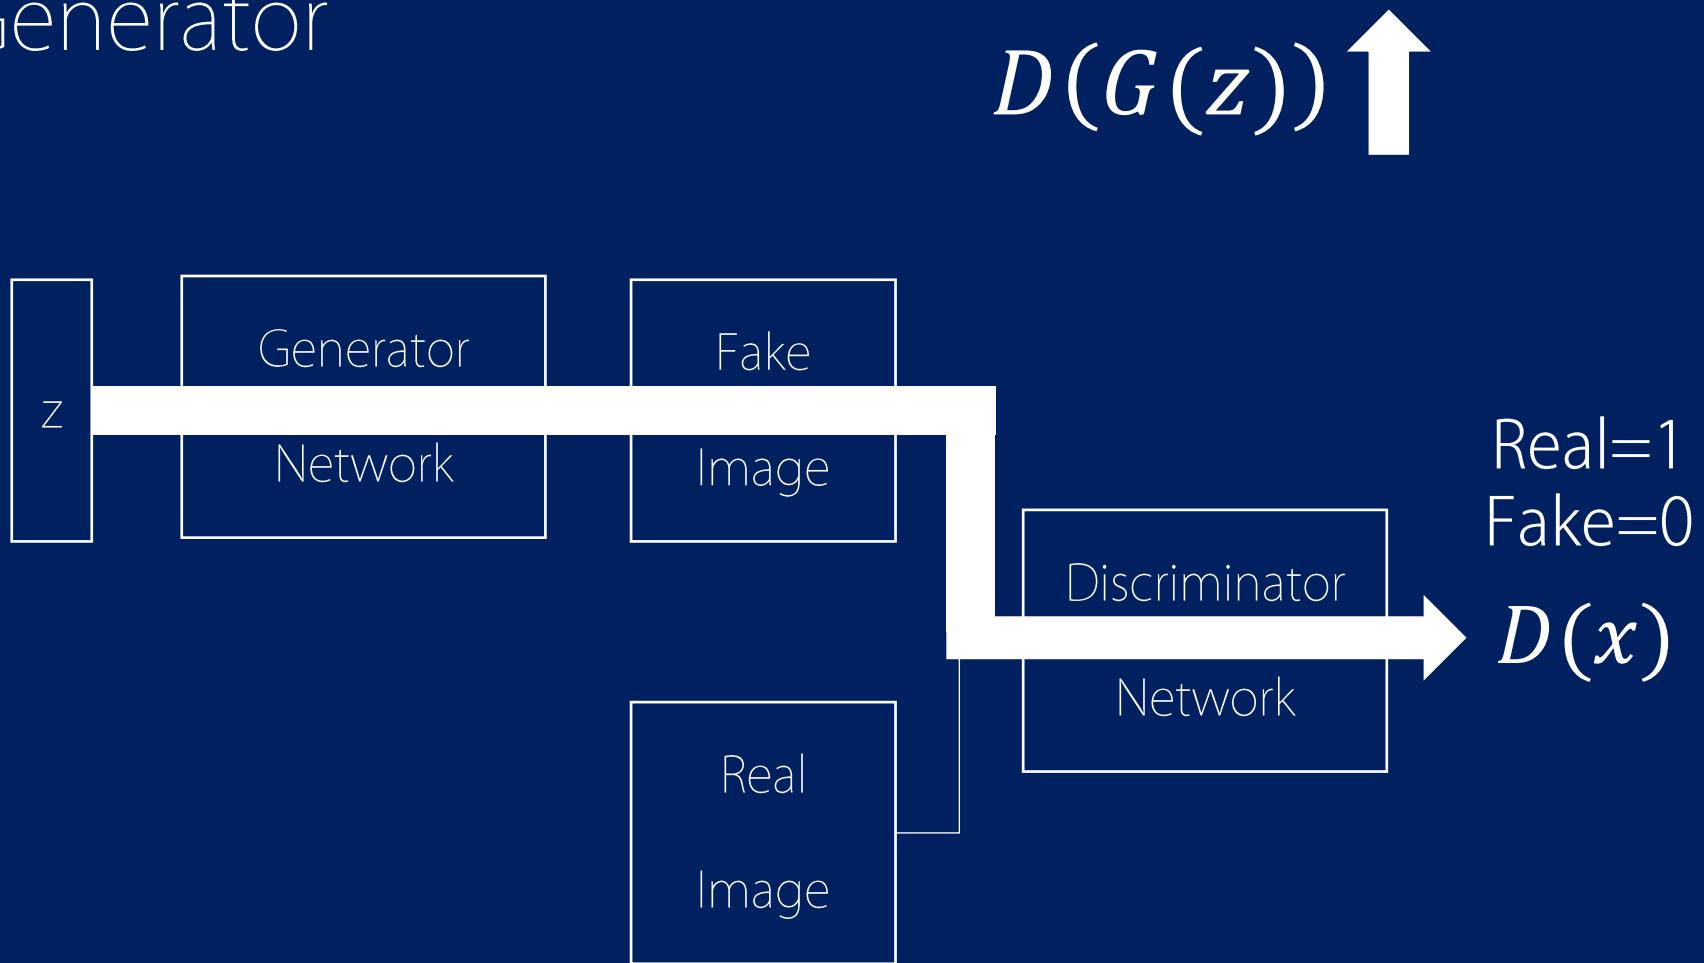
$$D(x)$$

Training Discriminator



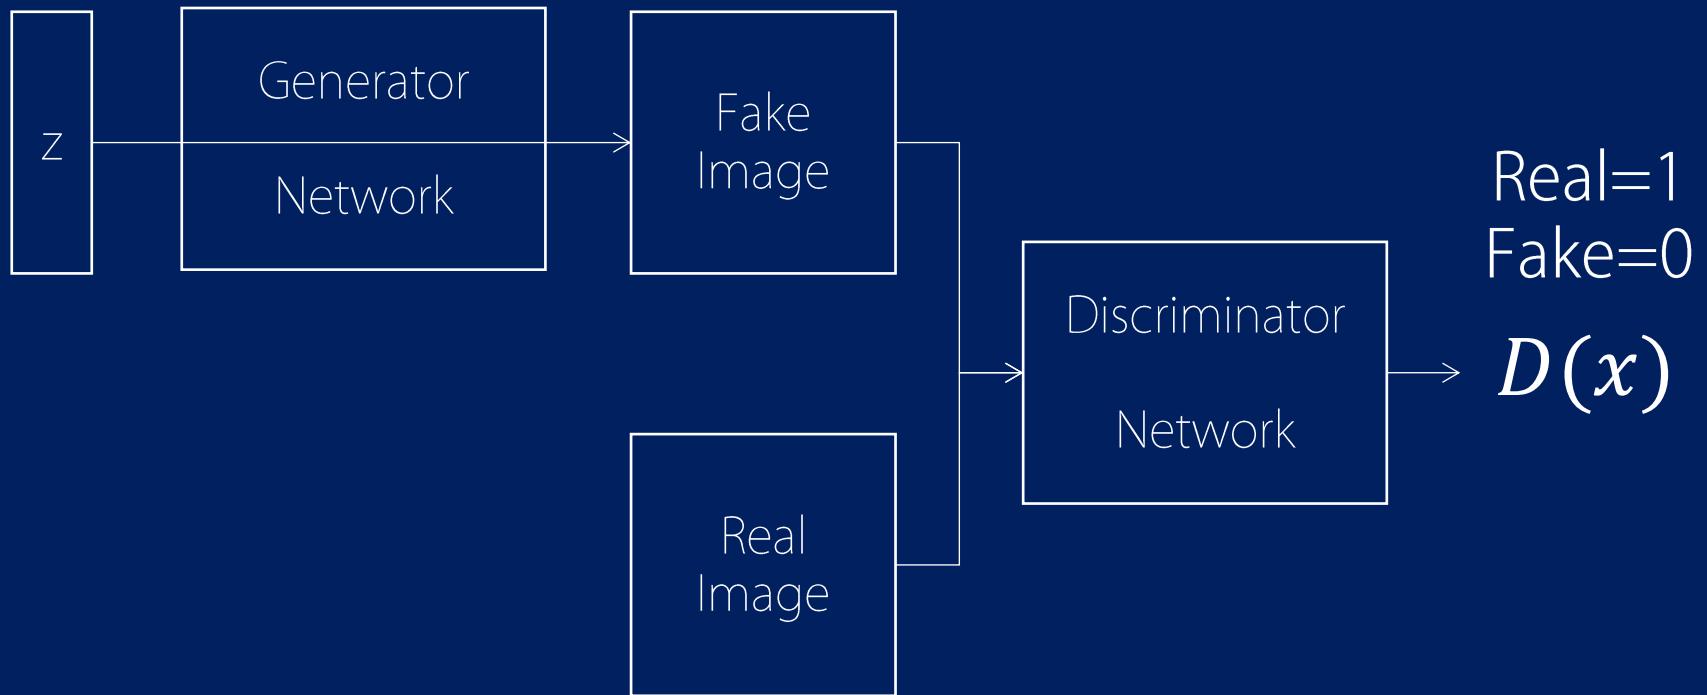
$$D_{Loss} = Error(D(x), 1) + Error(D(G(z)), 0)$$

Training Generator



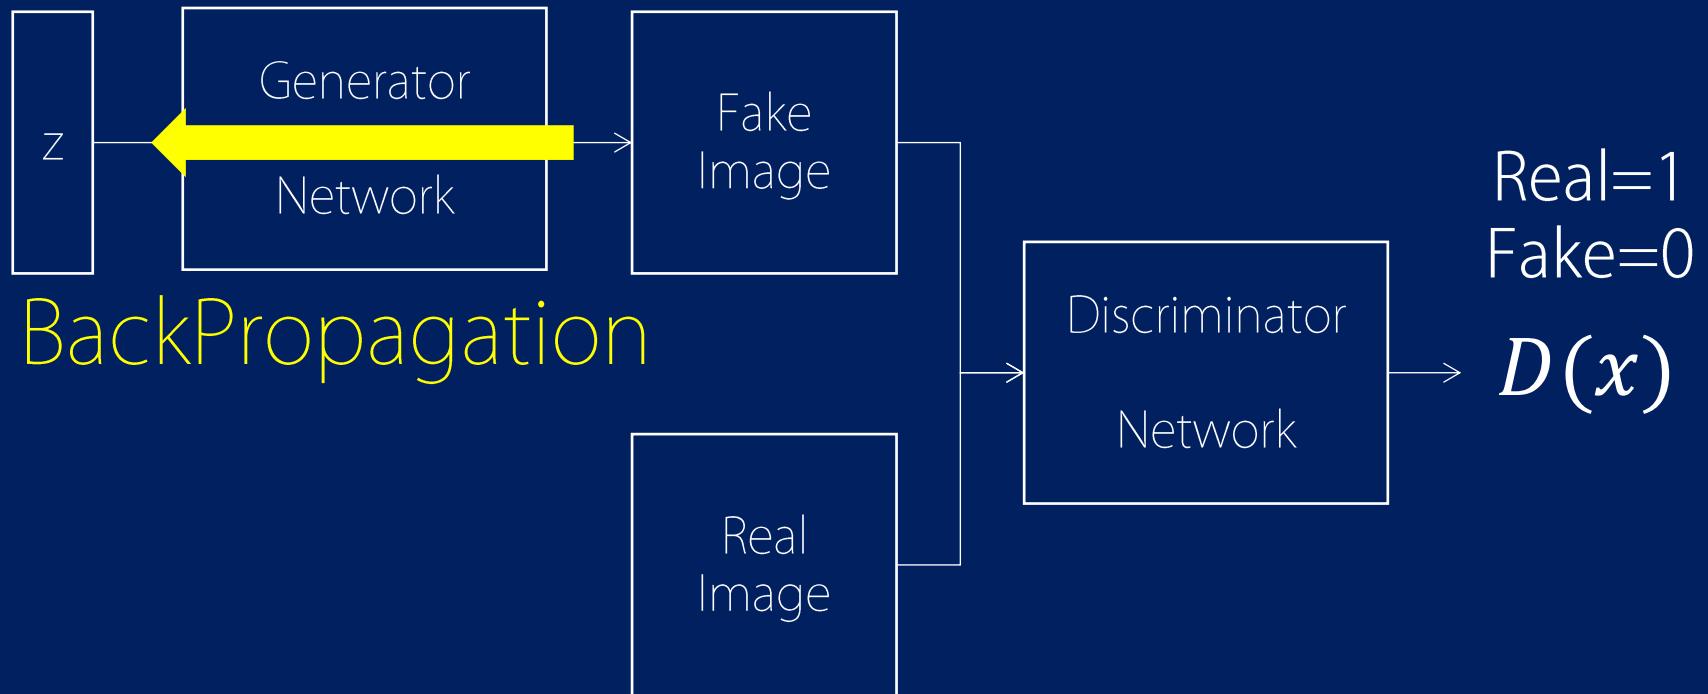
Training Generator

$D(G(z)) \uparrow$



$$G_{Loss} = Error(D(G(z)), 1)$$

Training Generator



$$G_{Loss} = Error(D(G(z)), 1)$$

Define Loss Function

Linear Regression Loss

: Mean Square Error (MSE)

Logistic Regression / Classification Loss

: Cross Entropy (CE)

$$-y_{true} \cdot \log(y_{pred}) - (1 - y_{true}) \cdot \log(1 - y_{pred})$$

Define Loss Function

$$-y_{true} \cdot \log(y_{pred}) - (1 - y_{true}) \cdot \log(1 - y_{pred})$$

$$\begin{aligned} D_{Loss} &= Error(D(x), 1) + Error(D(G(z)), 0) \\ &= -\log(D(x)) - \log(1 - D(G(z))) \end{aligned}$$

$$\begin{aligned} G_{Loss} &= Error(D(G(z)), 1) \\ &= -\log(D(G(z))) \end{aligned}$$

Define Loss Function

$$D_{Loss} = -\log(D(x)) - \log(1 - D(G(z)))$$

$$G_{Loss} = -\log(D(G(z)))$$

Value Function (Mathematical Form)

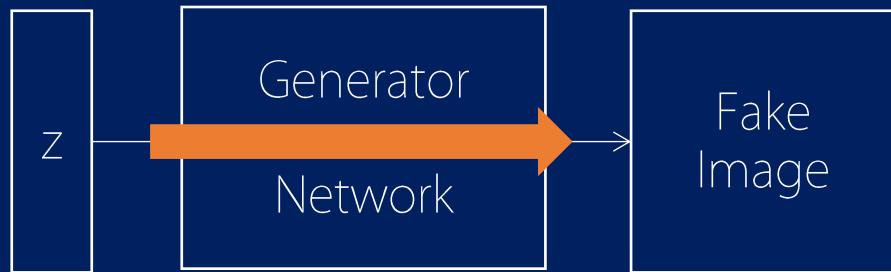
$$\min_G \max_D V(D, G) = E_{x \sim p_{data}} \log(D(x)) + E_{z \sim p_z} \log(1 - D(G(z)))$$

Min-Max Game
Nash Equilibrium

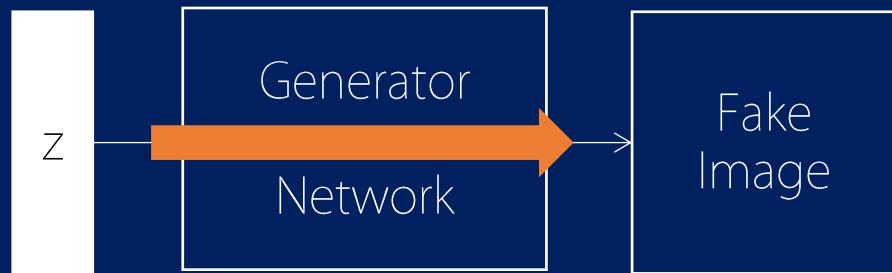
Testing
GAN

Testing GAN

= Testing Generator = Generate From Noise



Testing During Training



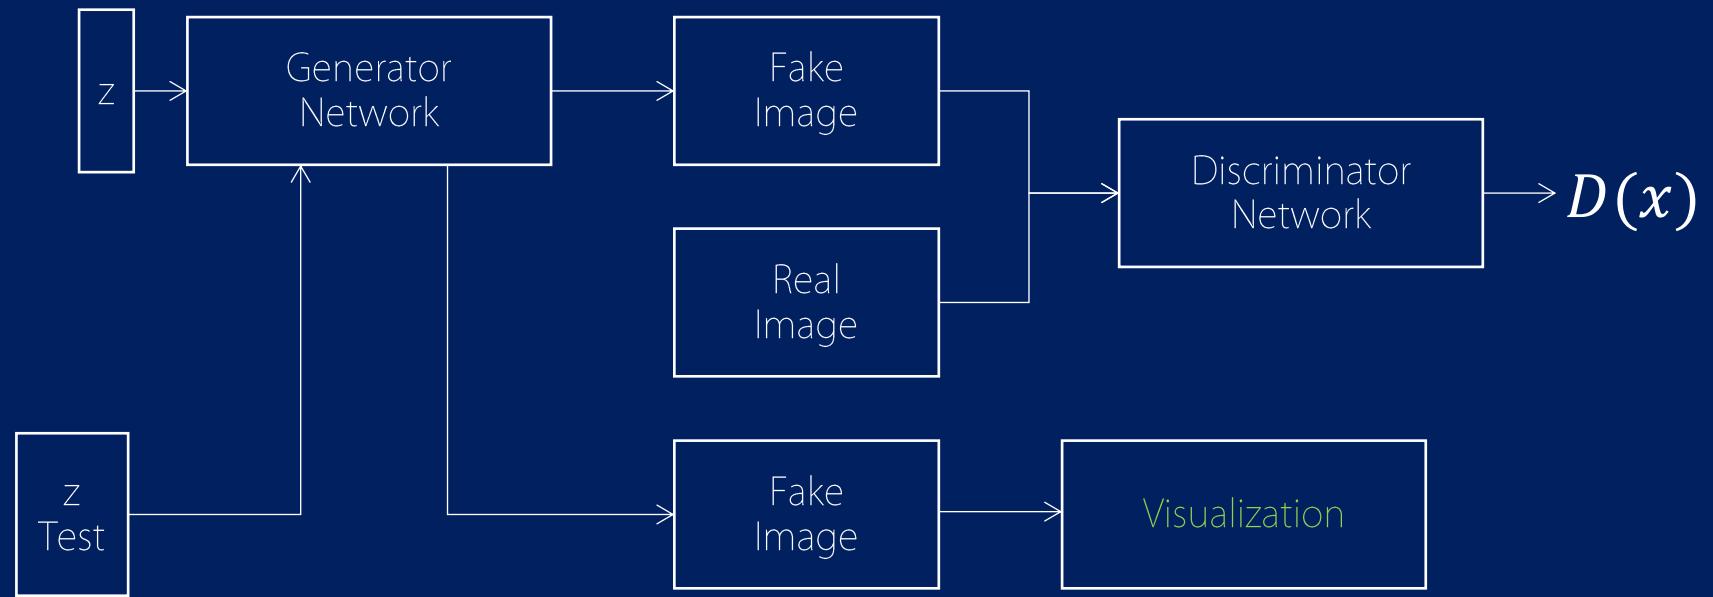
Fix Noise

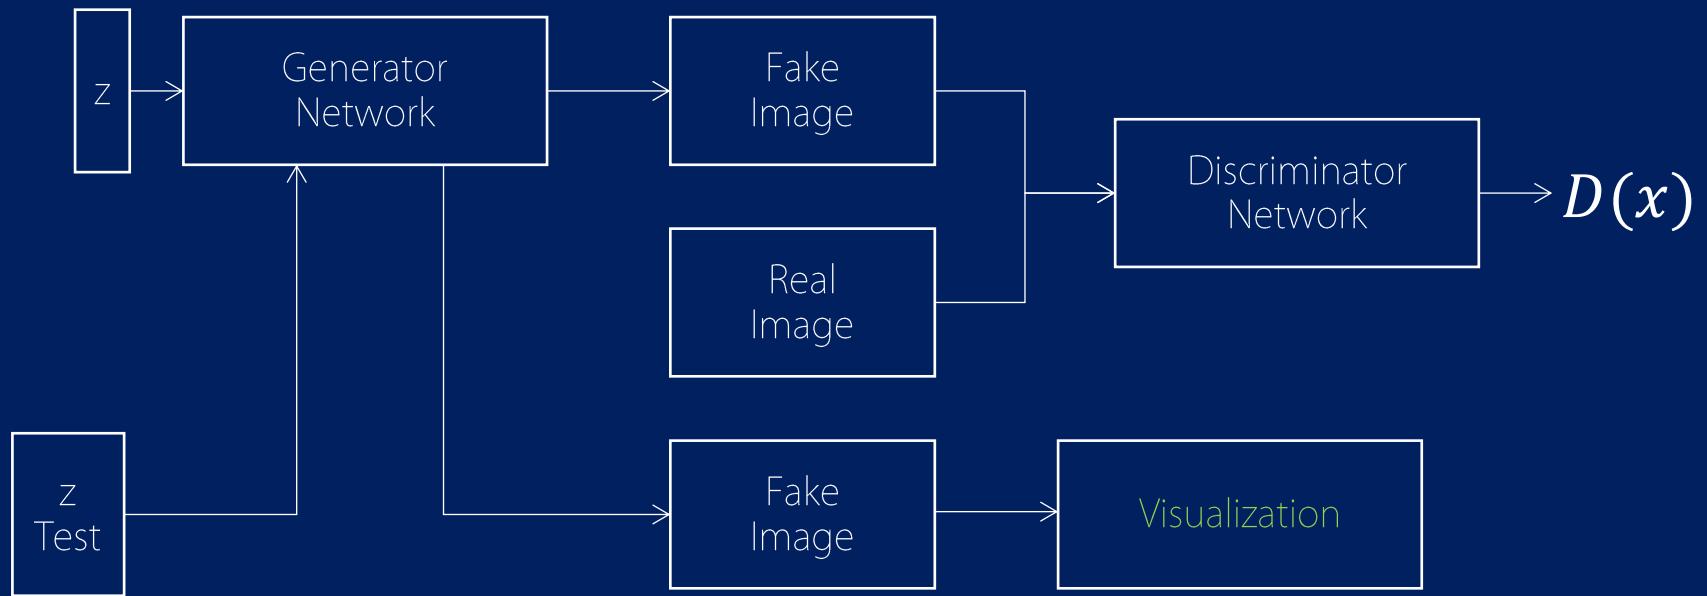
GAN

Implementation

Tensorflow

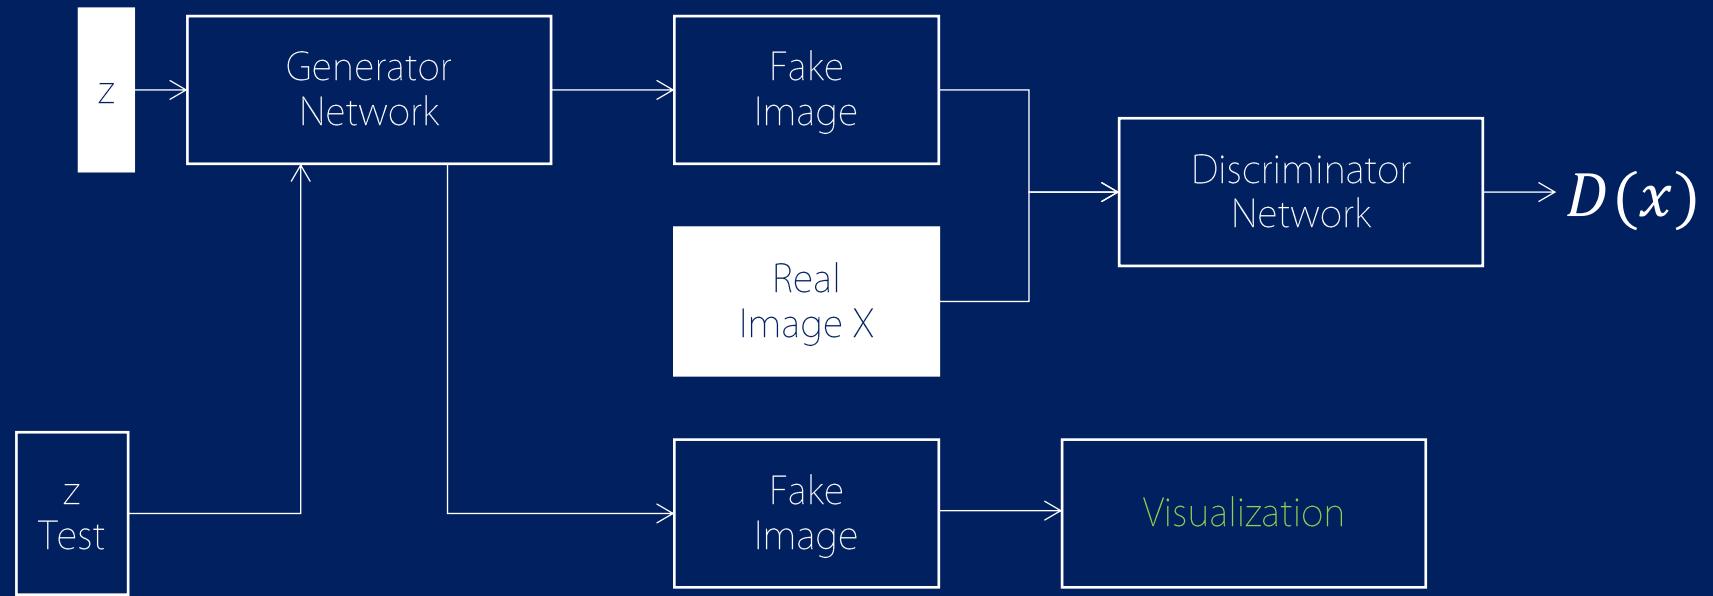
50 Lines



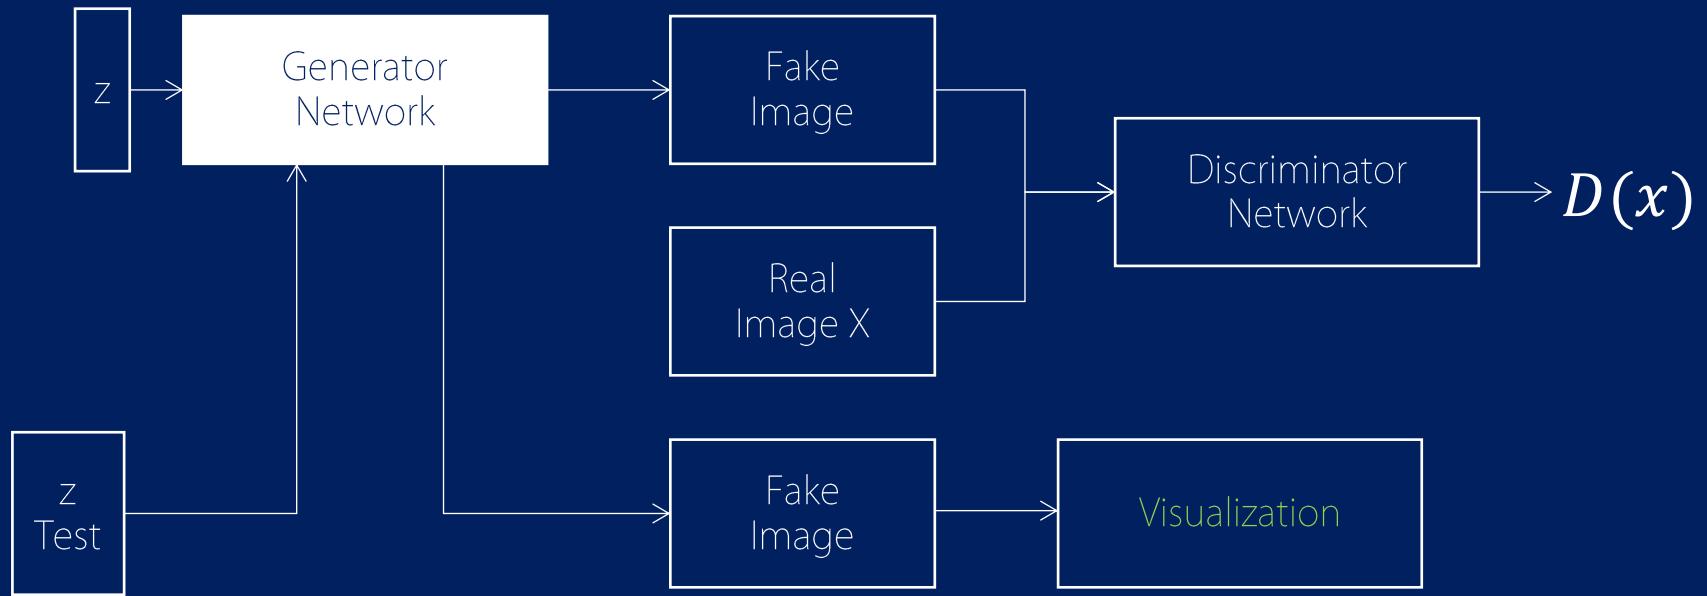


```
import tensorflow as tf
import matplotlib.pyplot as plt
import numpy as np
from tensorflow.examples.tutorials.mnist import input_data

mnist = input_data.read_data_sets("./mnist/data/", one_hot=True)
```



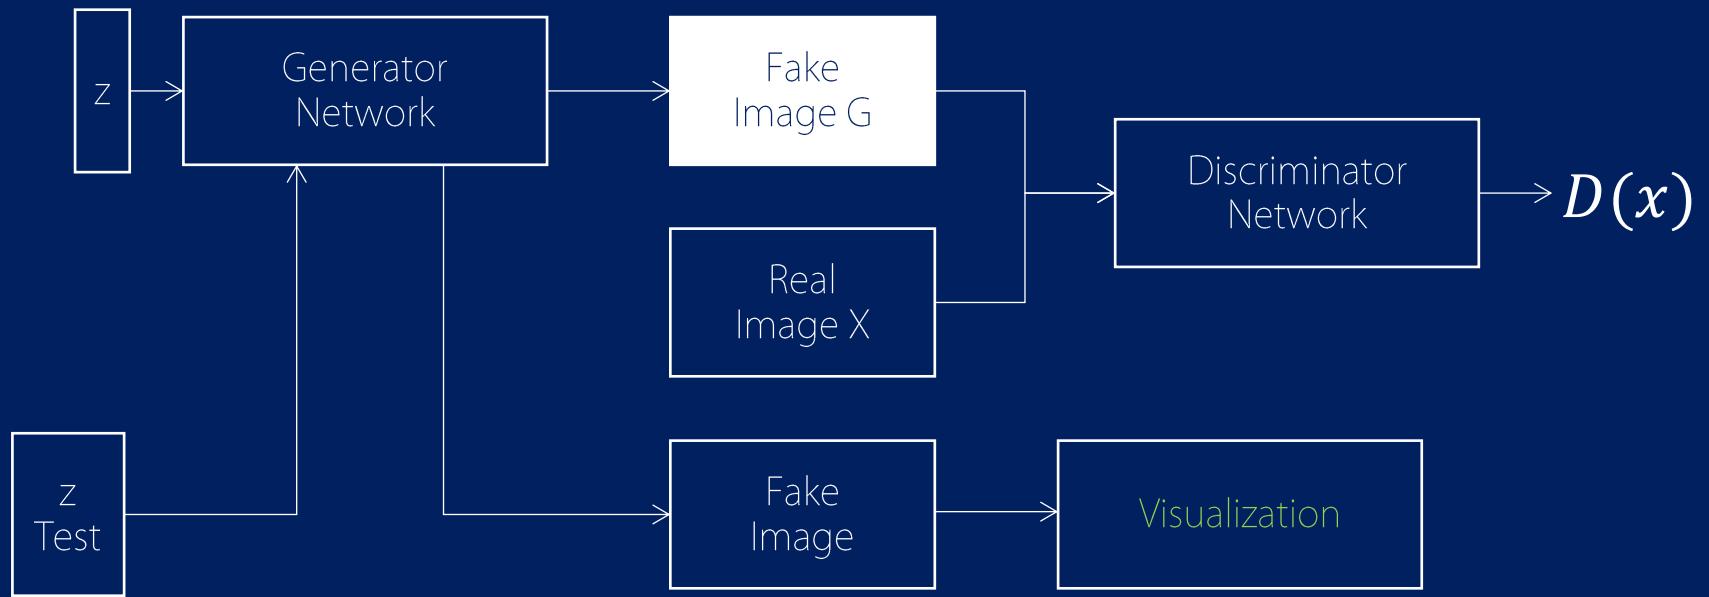
```
X = tf.placeholder(tf.float32, [None, 28 * 28])
Z = tf.placeholder(tf.float32, [None, 128])
```



```

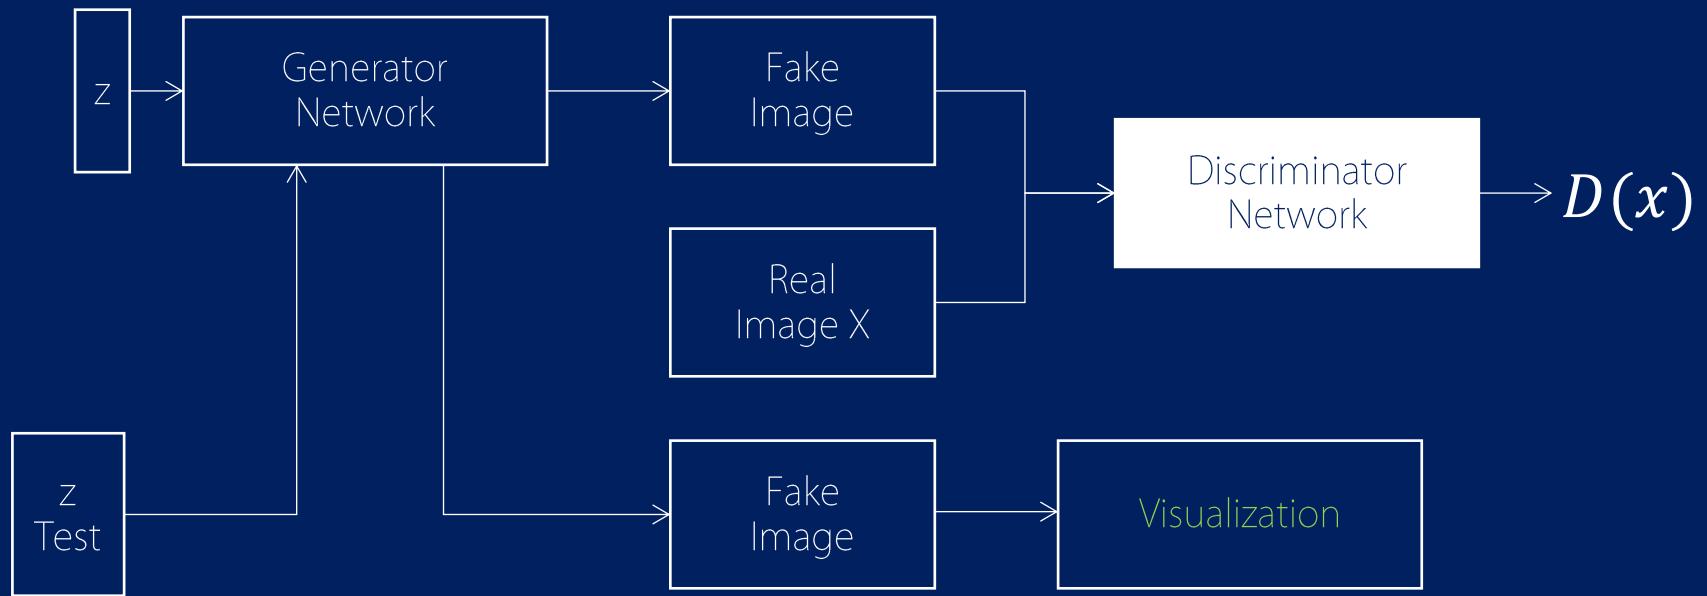
G_W1 = tf.Variable(tf.random_normal([128, 256], stddev=0.01))
G_W2 = tf.Variable(tf.random_normal([256, 28 * 28], stddev=0.01))
G_b1 = tf.Variable(tf.zeros([256]))
G_b2 = tf.Variable(tf.zeros([28 * 28]))

def generator(noise_z): # 128 -> 256 -> 28*28
    hidden = tf.nn.relu(tf.matmul(noise_z, G_W1) + G_b1)
    output = tf.nn.sigmoid(tf.matmul(hidden, G_W2) + G_b2)
    return output
  
```



```
G = generator(Z)
```

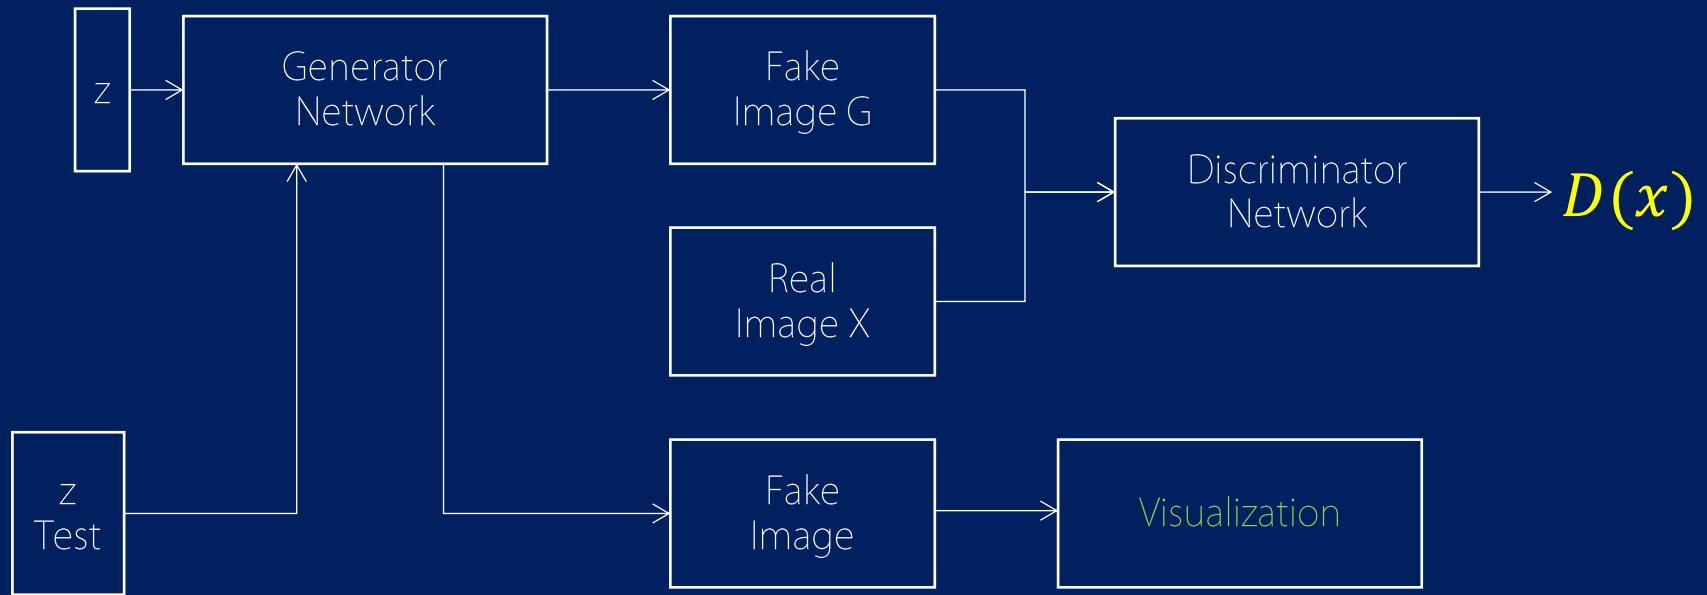
```
# Z->G
```



```

D_W1 = tf.Variable(tf.random_normal([28 * 28, 256], stddev=0.01))
D_W2 = tf.Variable(tf.random_normal([256, 1], stddev=0.01))
D_b1 = tf.Variable(tf.zeros([256]))
D_b2 = tf.Variable(tf.zeros([1]))

def discriminator(inputs): # 28*28 -> 256 -> 1
    hidden = tf.nn.relu(tf.matmul(inputs, D_W1) + D_b1)
    output = tf.nn.sigmoid(tf.matmul(hidden, D_W2) + D_b2)
    return output
  
```



```

loss_D = -tf.reduce_mean(
    tf.log(discriminator(X)) + tf.log(1 - discriminator(G)))
  
```

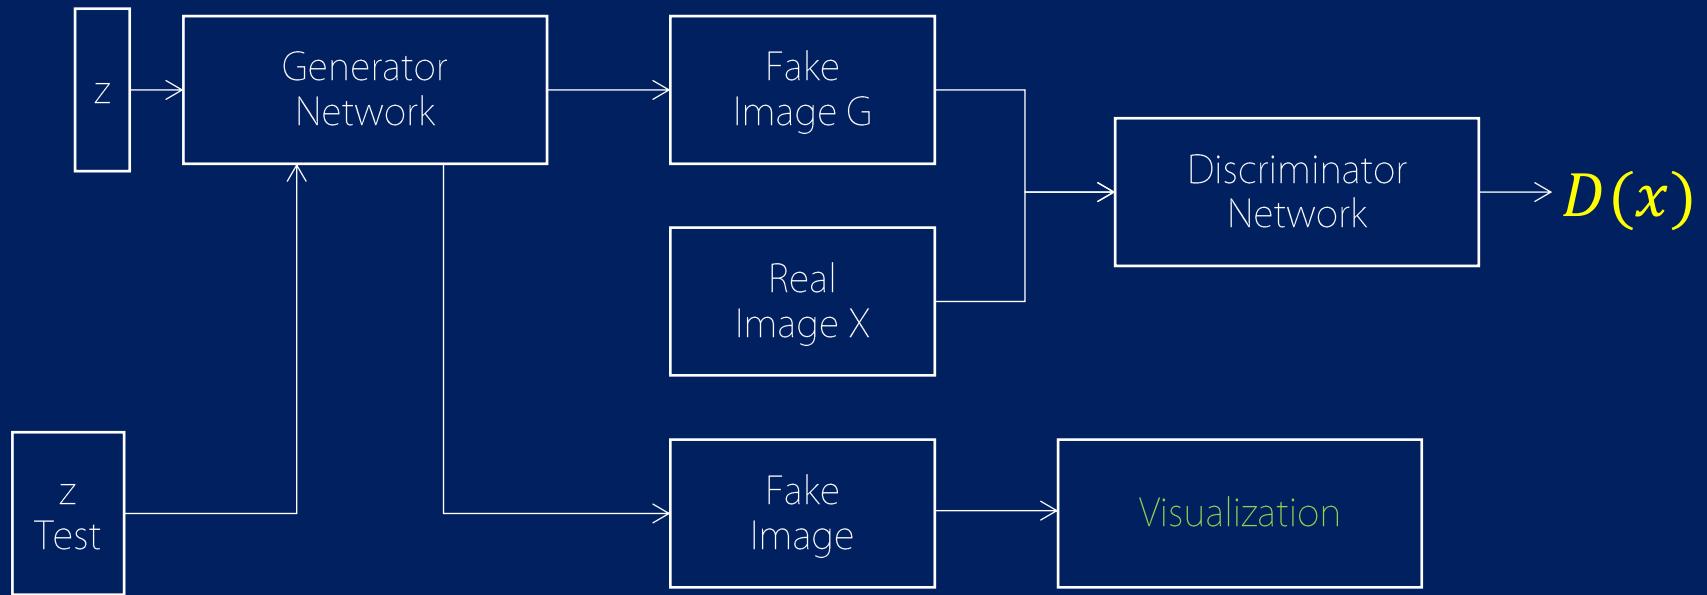
```

train_D = tf.train.AdamOptimizer(0.0002).minimize(
    loss_D, var_list=[D_W1, D_b1, D_W2, D_b2])
  
```

```

# Z->G->loss_D->train_D
# X->loss_D->train_D
  
```

$$\begin{aligned}
 D_{Loss} &= -\log(D(x)) - \log(1 - D(G(z))) \\
 G_{Loss} &= -\log(D(G(z)))
 \end{aligned}$$



```

loss_G = -tf.reduce_mean(tf.log(discriminator(G)))

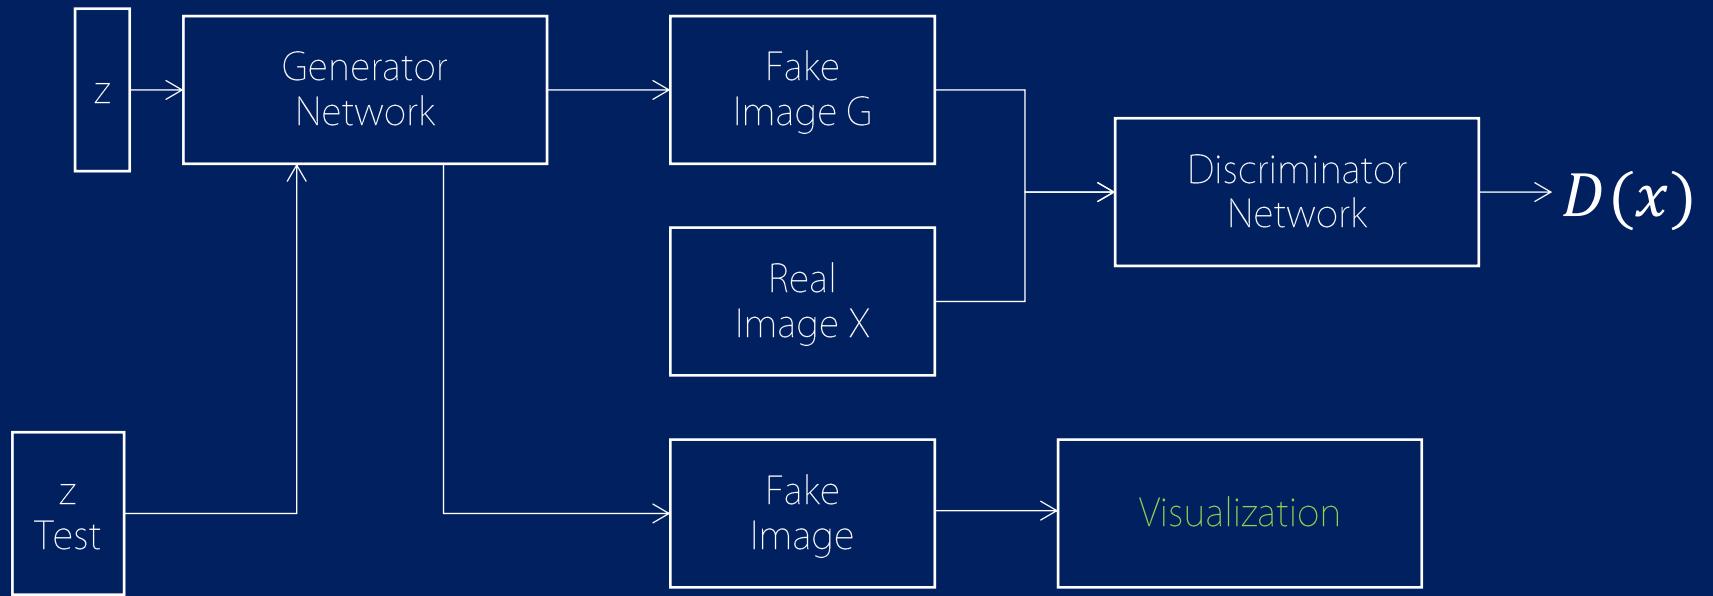
train_G = tf.train.AdamOptimizer(0.0002).minimize(
    loss_G, var_list=[G_W1, G_b1, G_W2, G_b2])

# Z->G->loss_G->train_G

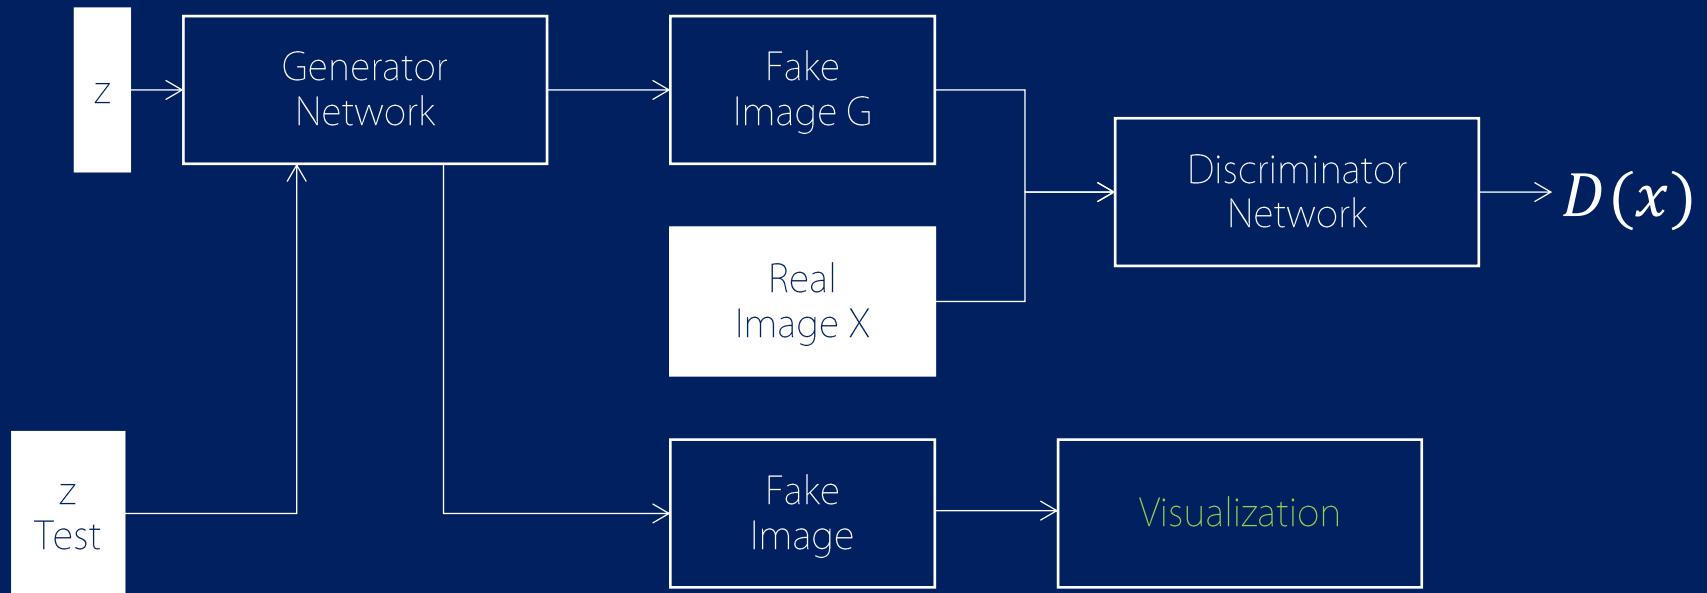
```

$$D_{Loss} = -\log(D(x)) - \log(1 - D(G(z)))$$

$$G_{Loss} = -\log(D(G(z)))$$



```
sess = tf.Session()  
  
sess.run(tf.global_variables_initializer())
```



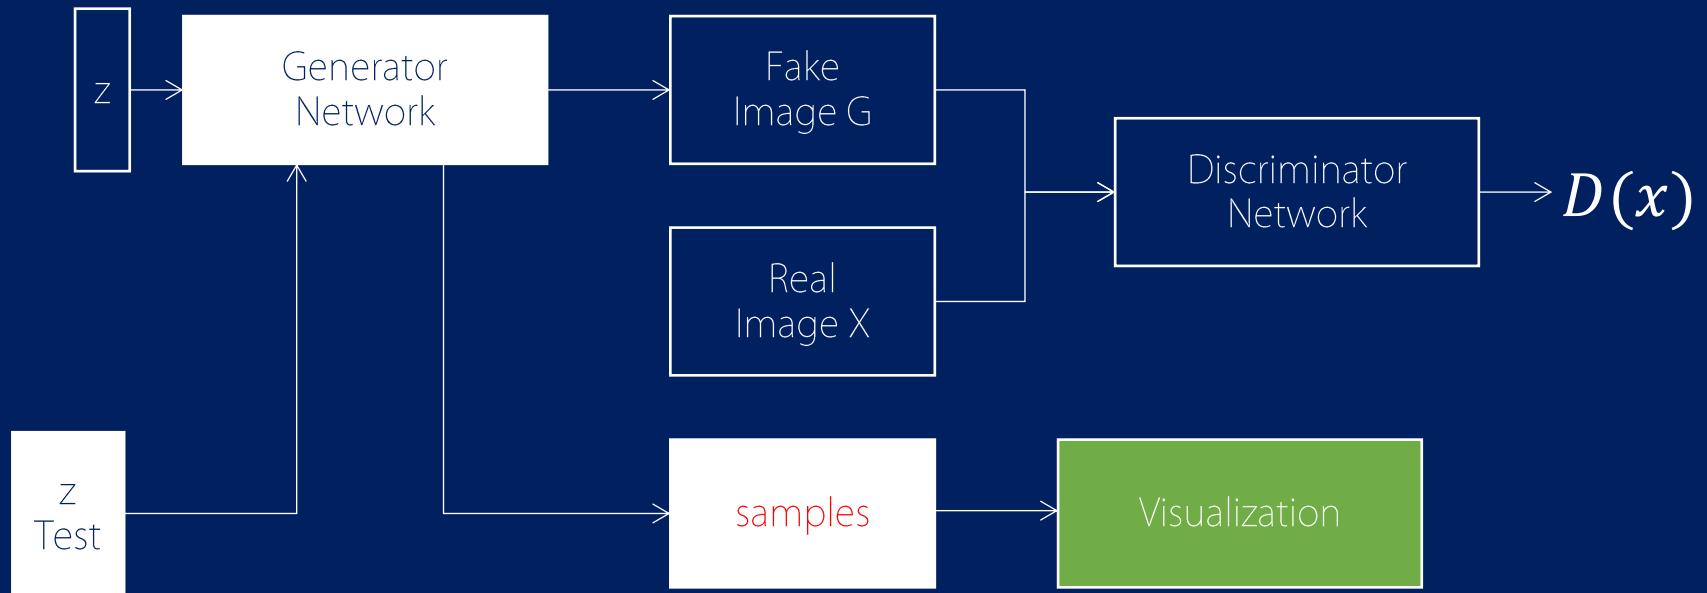
```

noise_test = np.random.normal(size=(10, 128))

for epoch in range(200):
    for i in range(int(mnist.train.num_examples / 100)):
        batch_xs, _ = mnist.train.next_batch(100)

        noise = np.random.normal(size=100, 128)

        sess.run(train_D, feed_dict={X: batch_xs, Z: noise})
        sess.run(train_G, feed_dict={Z: noise})
  
```



```

for epoch in range(200):
    ...
    if epoch == 0 or (epoch + 1) % 10 == 0:
        samples = sess.run(G, feed_dict={Z: noise_test})

        fig, ax = plt.subplots(1, 10, figsize=(10, 1))
        for i in range(10):
            ax[i].set_axis_off()
            ax[i].imshow(np.reshape(samples[i], (28, 28)))
        plt.savefig('samples_ex/{}.png'.format(str(epoch).zfill(3)), bbox_inches='tight')
        plt.close(fig)
    
```

Full Code

<https://github.com/HyeongminLEE/GANin50lines>

```
import tensorflow as tf
import matplotlib.pyplot as plt
import numpy as np
from tensorflow.examples.tutorials.mnist import input_data

mnist = input_data.read_data_sets("./mnist/data/", one_hot=True)
X = tf.placeholder(tf.float32, [None, 28 * 28]) # MNIST = 28*28
Z = tf.placeholder(tf.float32, [None, 128]) # Noise Dimension = 128

# ***** G-Network (Hidden Node # = 256)
G_W1 = tf.Variable(tf.random_normal([128, 256], stddev=0.01))
G_W2 = tf.Variable(tf.random_normal([256, 28 * 28], stddev=0.01))
G_b1 = tf.Variable(tf.zeros([256]))
G_b2 = tf.Variable(tf.zeros([28 * 28]))

def generator(noise_z): # 128 -> 256 -> 28*28
    hidden = tf.nn.relu(tf.matmul(noise_z, G_W1) + G_b1)
    output = tf.nn.sigmoid(tf.matmul(hidden, G_W2) + G_b2)
    return output

# ***** D-Network (Hidden Node # = 256)
D_W1 = tf.Variable(tf.random_normal([28 * 28, 256], stddev=0.01))
D_W2 = tf.Variable(tf.random_normal([256, 1], stddev=0.01))
D_b1 = tf.Variable(tf.zeros([256]))
D_b2 = tf.Variable(tf.zeros([1]))

def discriminator(inputs): # 28*28 -> 256 -> 1
    hidden = tf.nn.relu(tf.matmul(inputs, D_W1) + D_b1)
    output = tf.nn.sigmoid(tf.matmul(hidden, D_W2) + D_b2)
    return output

# ***** Generation, Loss, Optimization and Session Init.
G = generator(Z)
loss_D = -tf.reduce_mean(tf.log(discriminator(X)) + tf.log(1 - discriminator(G)))
loss_G = -tf.reduce_mean(tf.log(discriminator(G)))
train_D = tf.train.AdamOptimizer(learning_rate=0.0002).minimize(loss_D, var_list=[D_W1, D_b1, D_W2, D_b2])
train_G = tf.train.AdamOptimizer(learning_rate=0.0002).minimize(loss_G, var_list=[G_W1, G_b1, G_W2, G_b2])

sess = tf.Session()
sess.run(tf.global_variables_initializer())

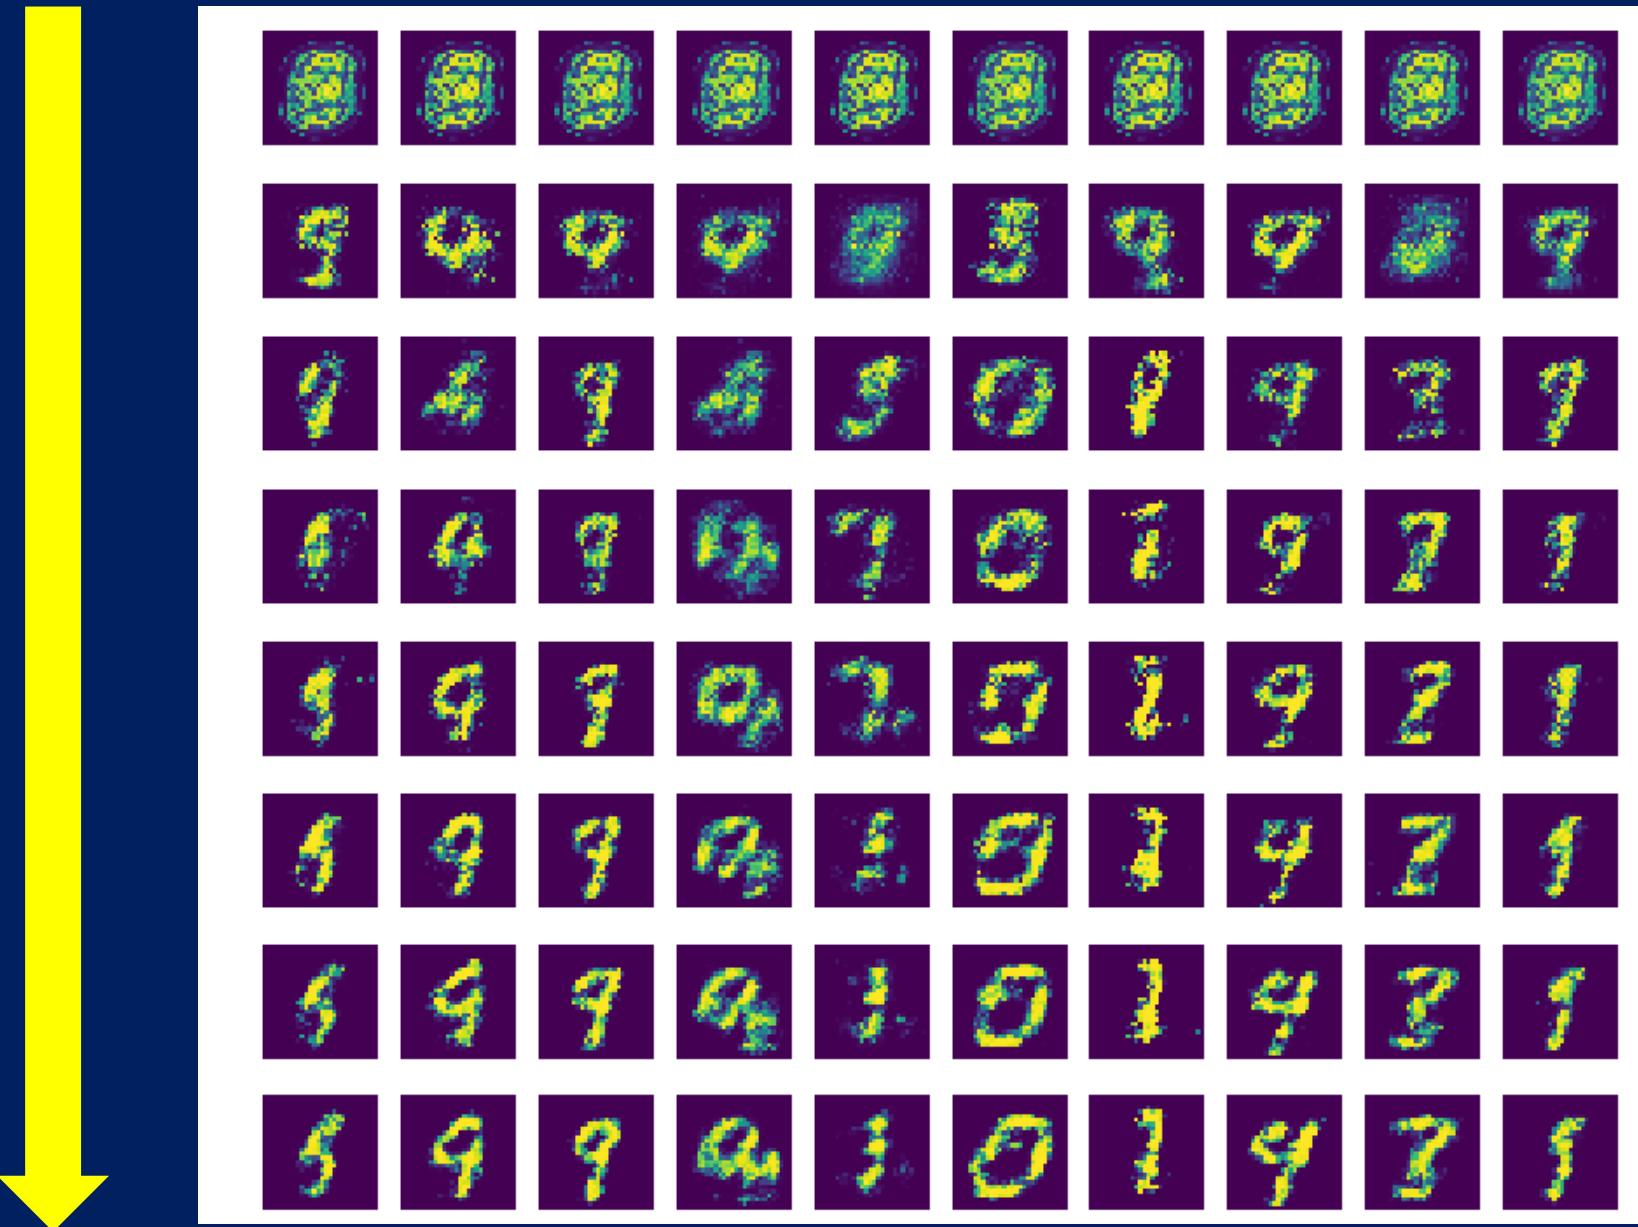
# ***** Training and Testing
noise_test = np.random.normal(size=(10, 128)) # 10 = Test Sample Size
for epoch in range(200): # 200 = Num. of Epoch
    for i in range(int(mnist.train.num_examples / 100)): # 100 = Batch Size
        batch_xs, _ = mnist.train.next_batch(100)
        noise = np.random.normal(size=(100, 128))

        sess.run(train_D, feed_dict={X: batch_xs, Z: noise})
        sess.run(train_G, feed_dict={Z: noise})

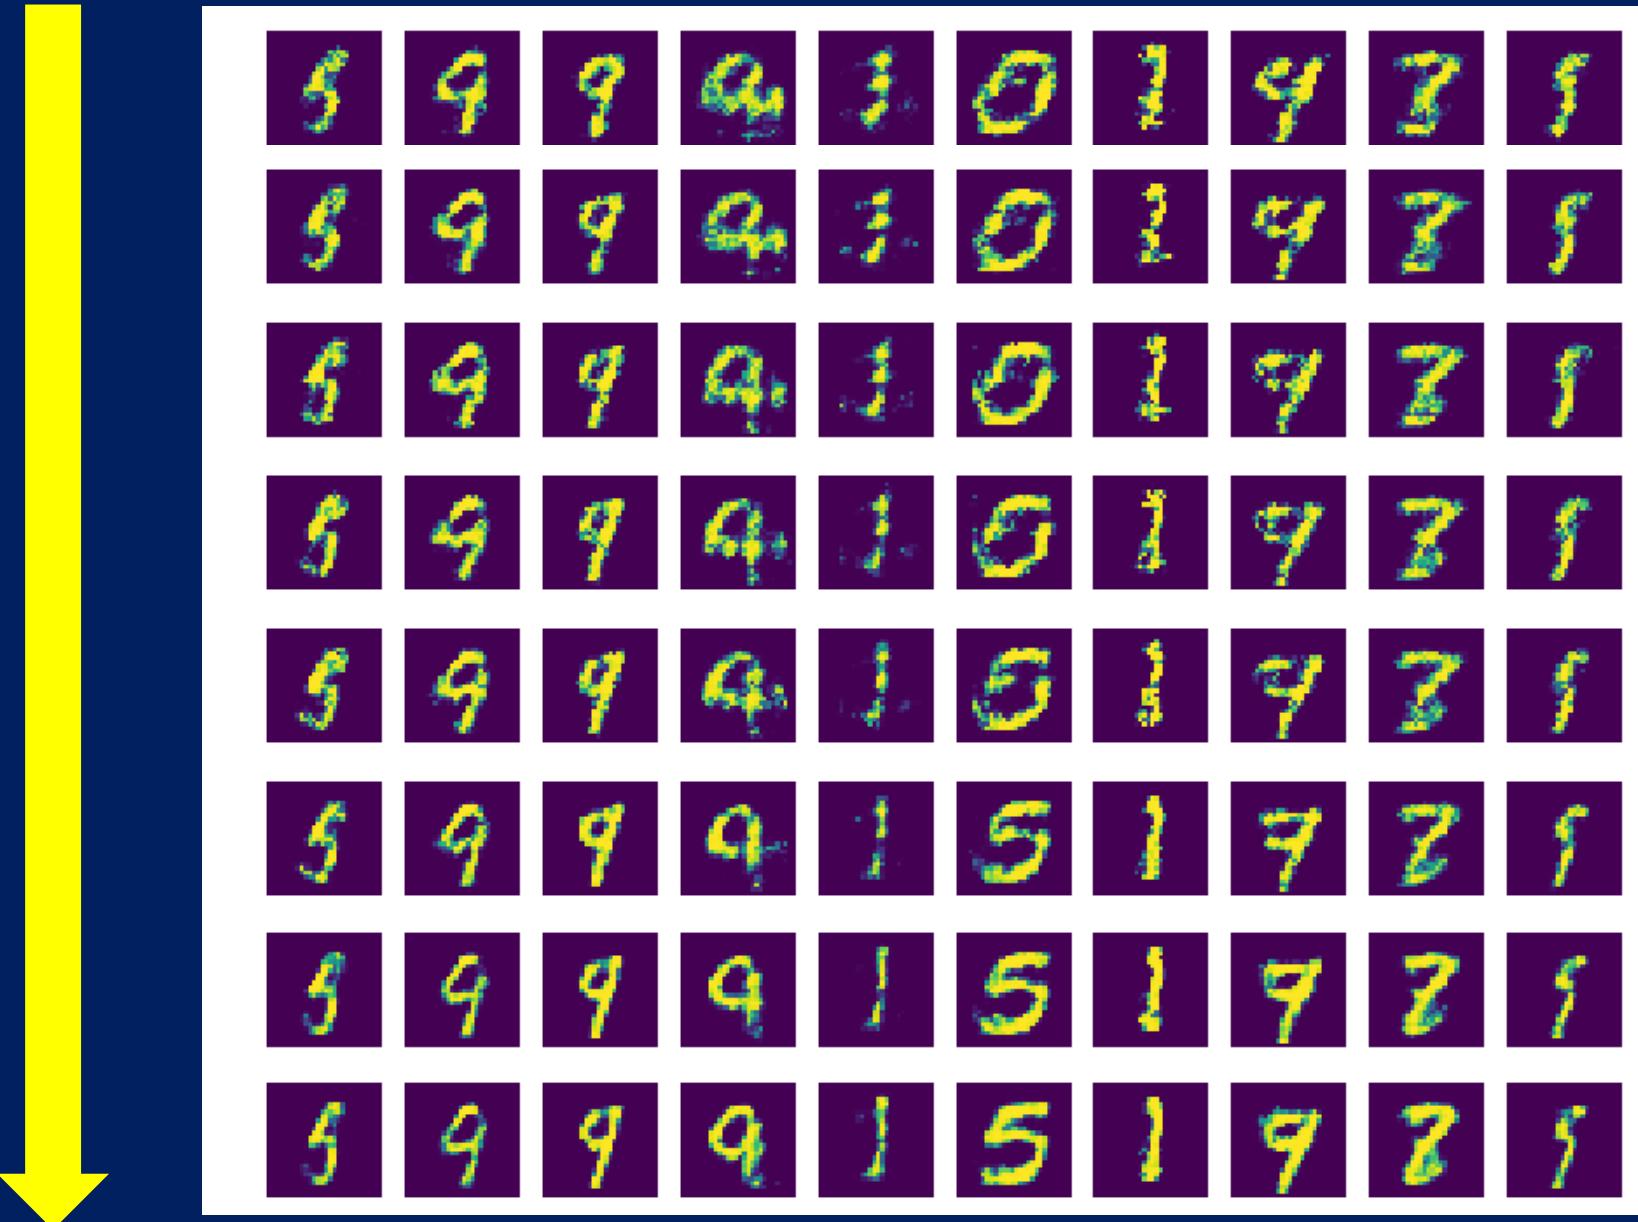
    if epoch == 0 or (epoch + 1) % 10 == 0: # 10 = Saving Period
        samples = sess.run(G, feed_dict={Z: noise_test})

        fig, ax = plt.subplots(1, 10, figsize=(10, 1))
        for i in range(10):
            ax[i].set_axis_off()
            ax[i].imshow(np.reshape(samples[i], (28, 28)))
        plt.savefig('samples_ex/{}.png'.format(str(epoch).zfill(3)),
                    bbox_inches='tight')
        plt.close(fig)
```

Results



Results



GAN Does Not
Easily Converge!

And Quality is Not Good!

→ DCGAN
Solve this Empirically

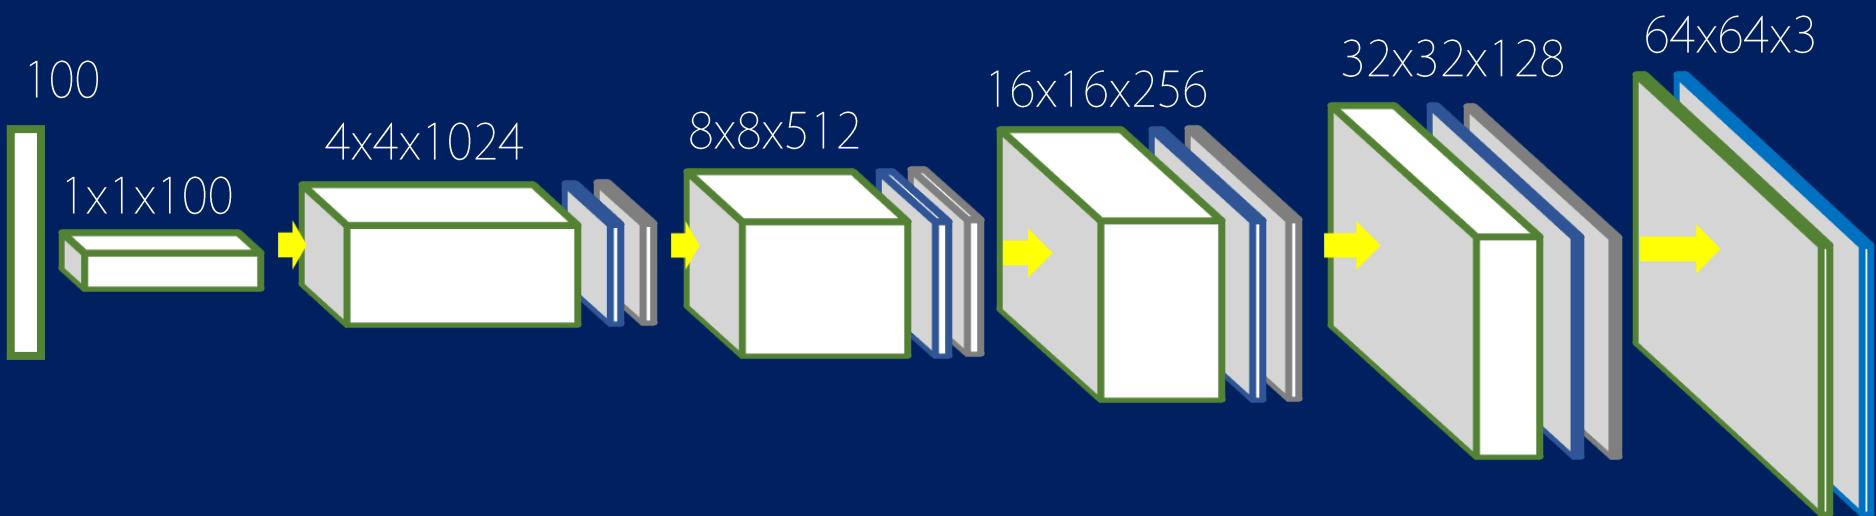
DCGAN

Deep Convolutional GAN

DCGAN's Solution

- Go Deep + CNN
- Batch Normalization
- Fully Convolutional
- No Pooling
- Leaky-ReLU

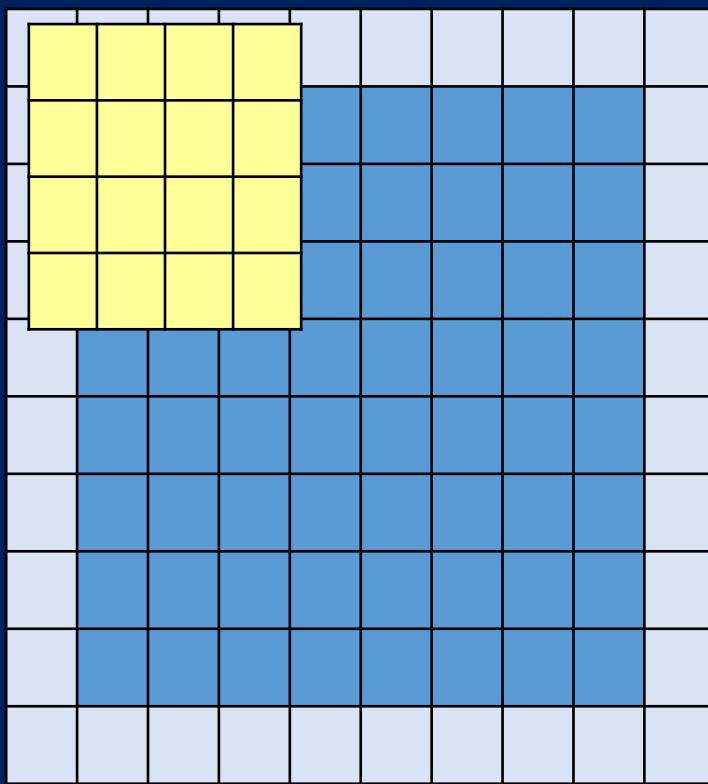
DCGAN Generator

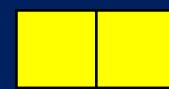
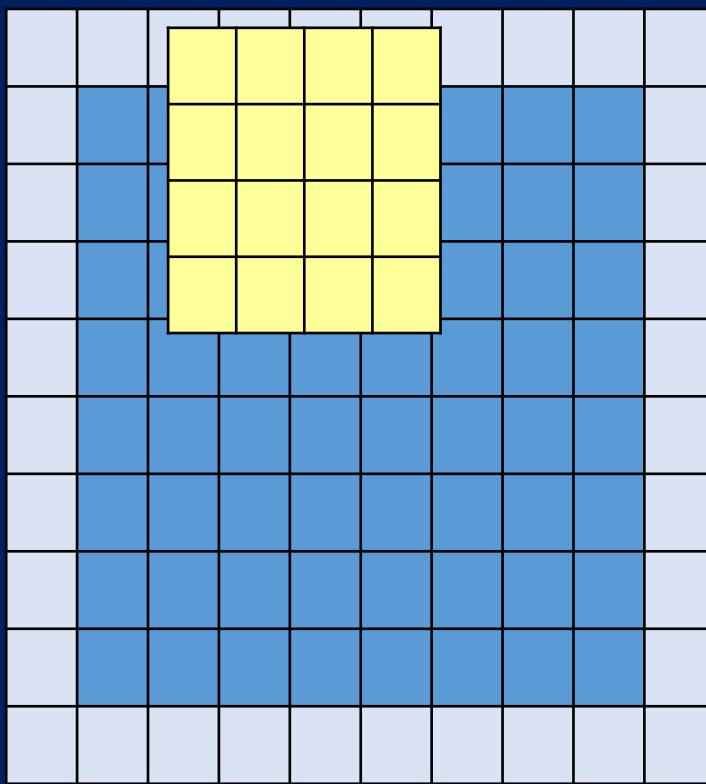


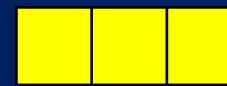
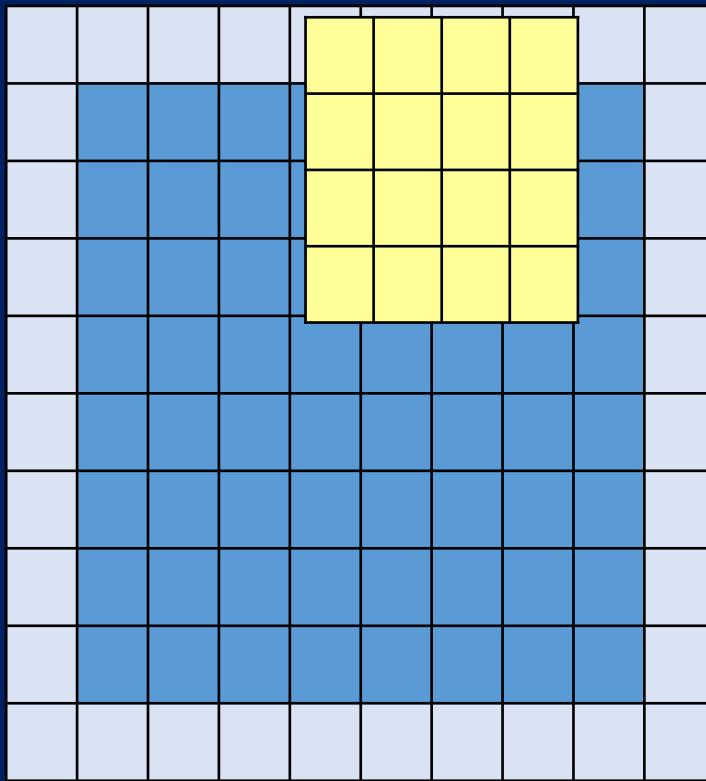
Filter Size = 4x4
Stride = 2
Padding = 1

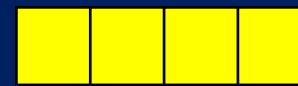
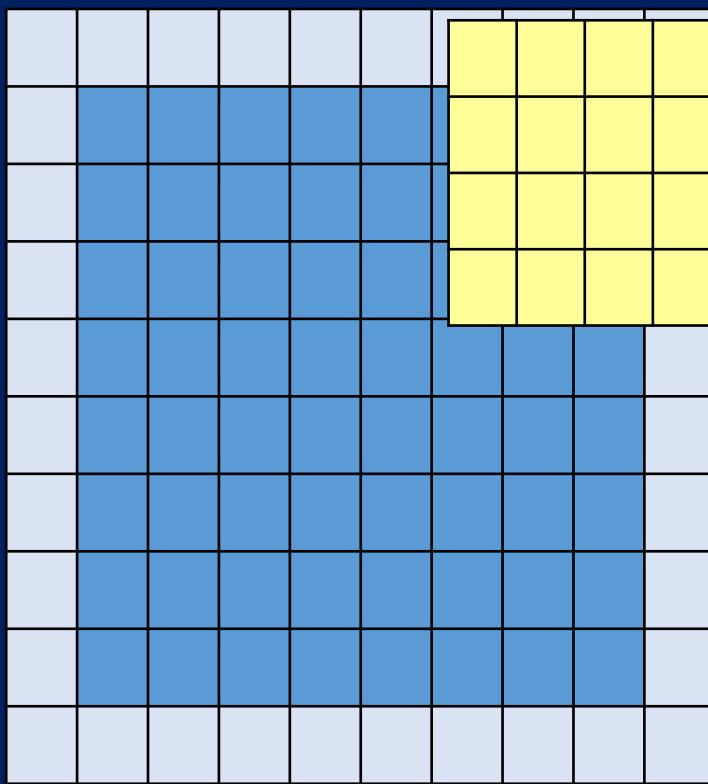


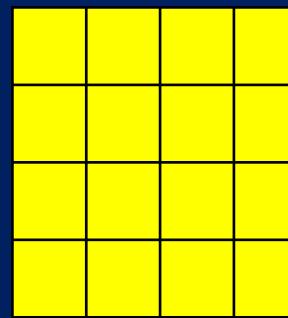
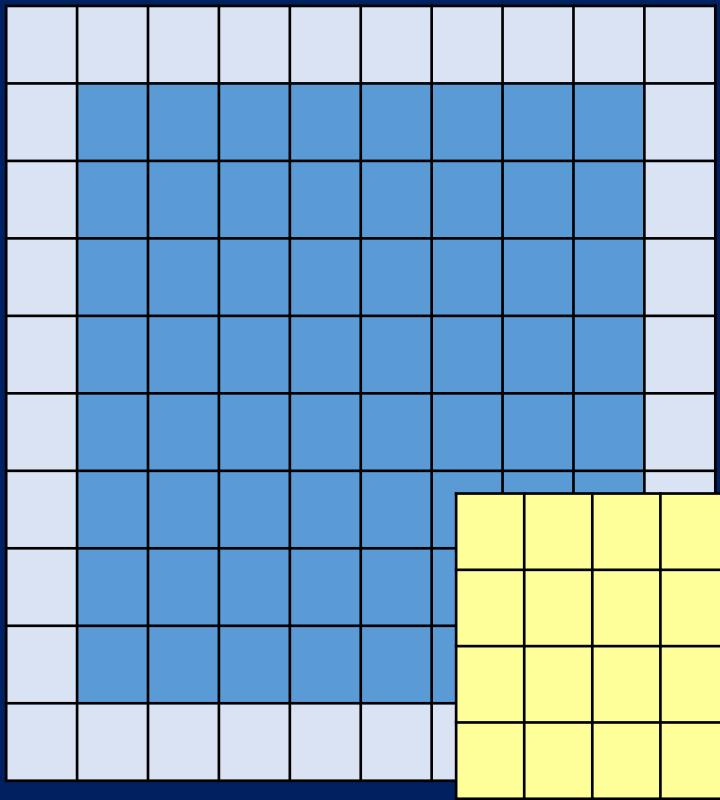
Image Should be Normalized [-1 , 1]



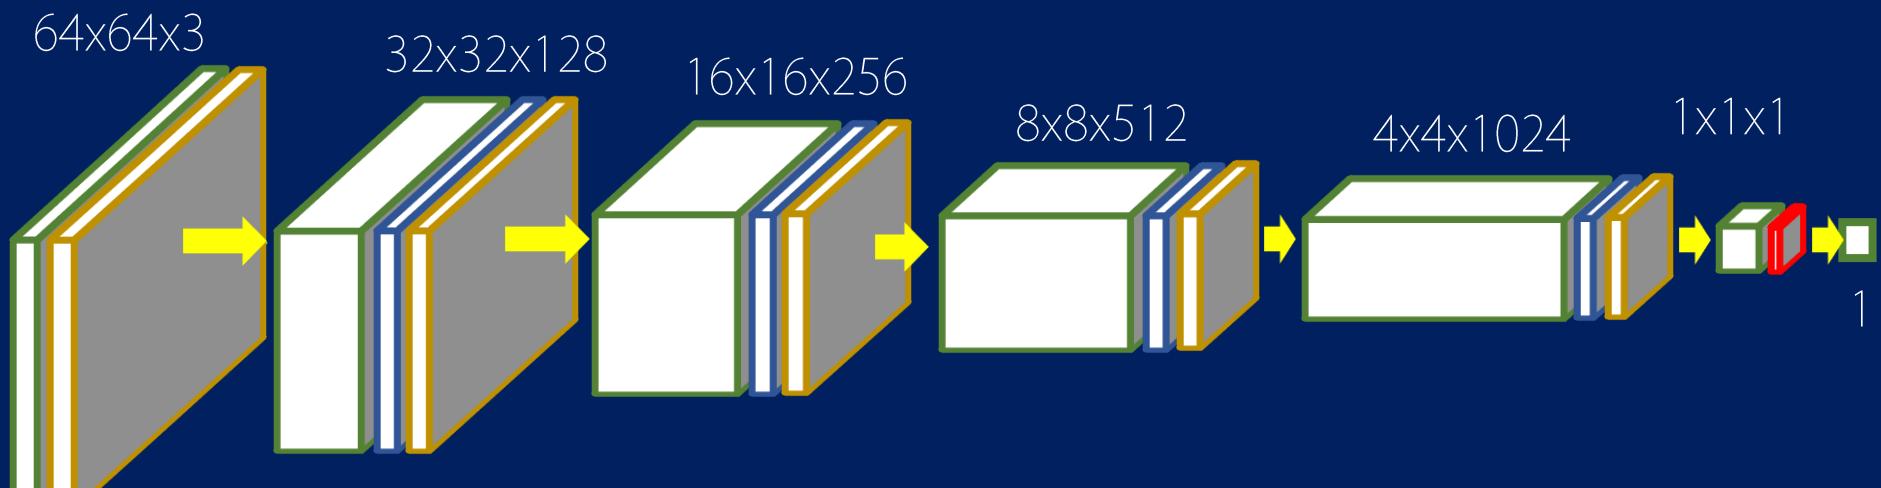








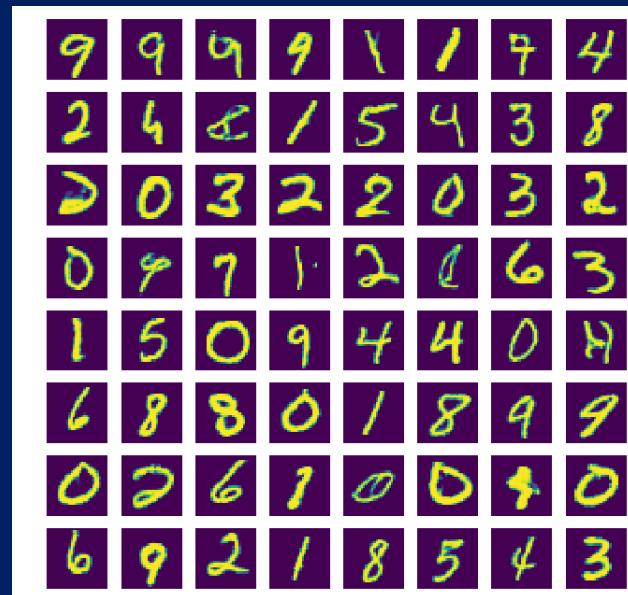
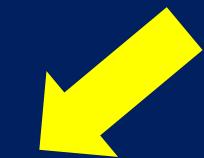
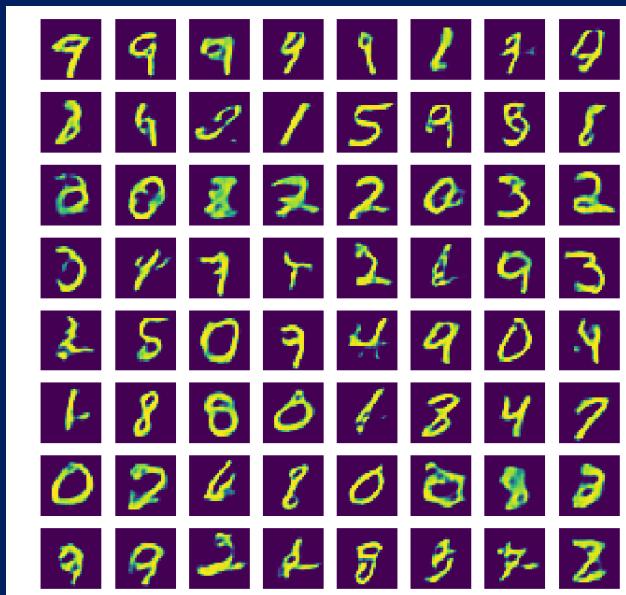
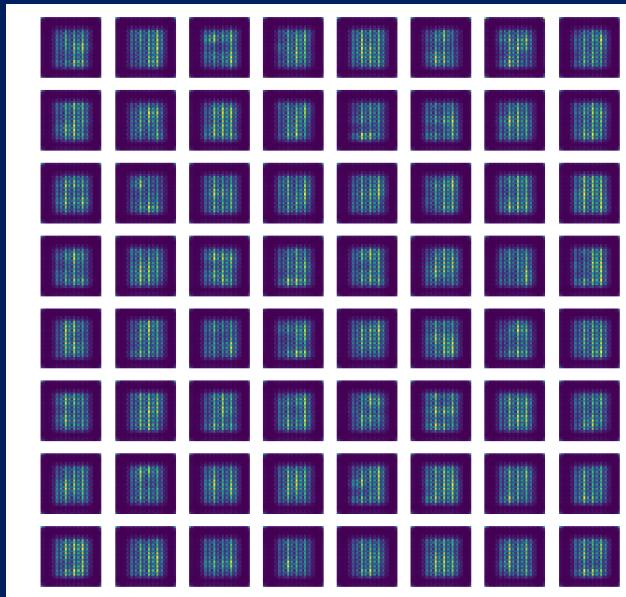
DCGAN Discriminator



Filter Size = 4x4
Stride = 2
Padding = 1



DCGAN Results



DCGAN Code

https://github.com/taeoh-kim/DCGAN_Pytorch

GAN History

Prerequisites



Random Process
Probability Theory
Machine Learning

Generative Models



Bayes, Naïve Bayes
Hidden Markov Model
Autoencoder, VAE

GAN
2014



DCGAN
2015

GAN Theory

f-GAN, WGAN, BEGAN

Latent Space

InfoGAN

Applications

Pix2Pix, CycleGAN
DiscoGAN, ...



InfoGAN

Exploit Latent Noise

Exploit Latent Vector z

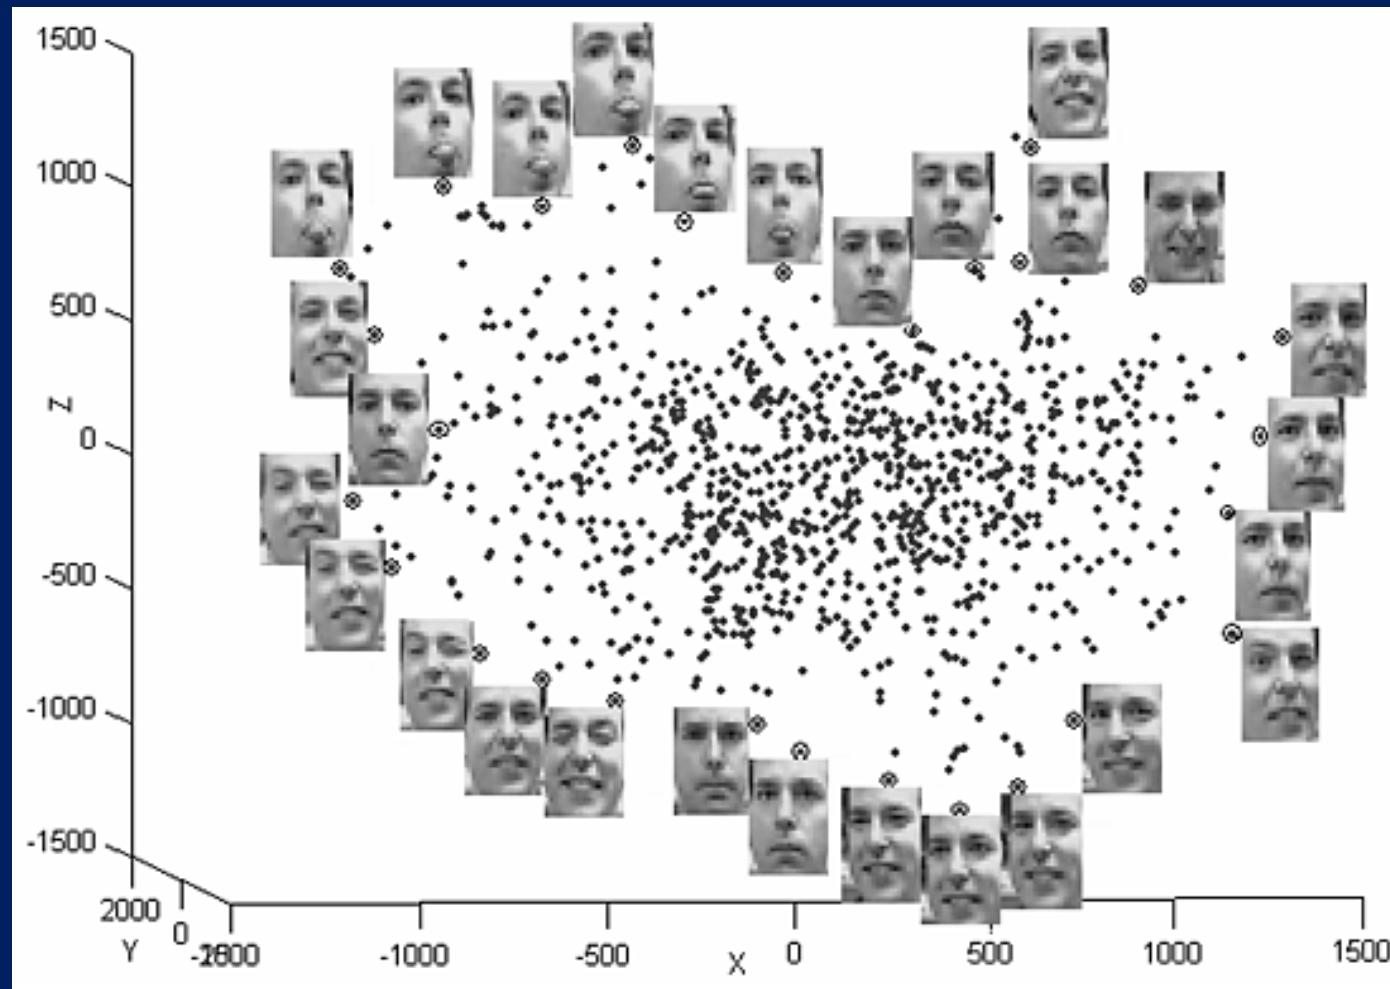
How to Exploit z ?

and I want Generate “Specific” Image!
ex) Smile Woman Face with Black Skin



- Conditional-GAN
with Label
- InfoGAN / VAE
with Meaningful z

Data Manifold Hypothesis

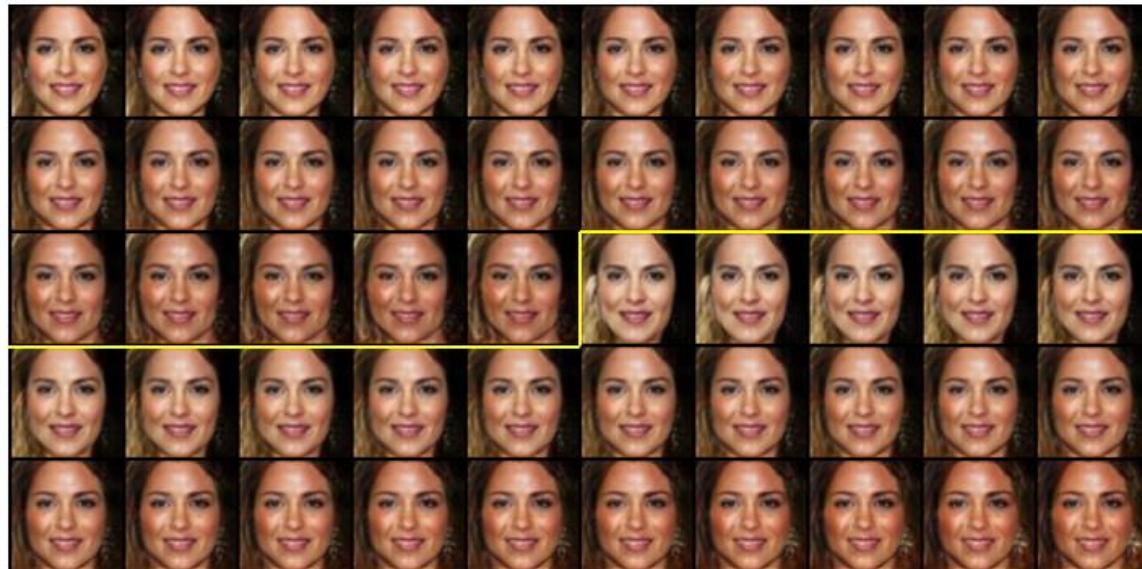


InfoGAN https://github.com/taeoh-kim/GAN_pytorch

1 Categorical Feature, 2 Continuous Feature, 256 Noise



Neutral



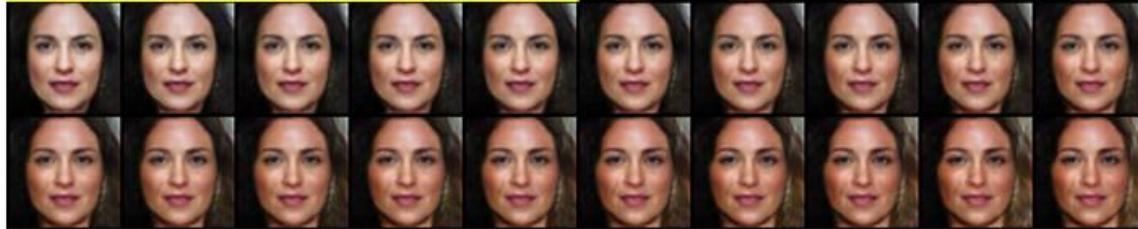
Smile

InfoGAN https://github.com/taeoh-kim/GAN_pytorch

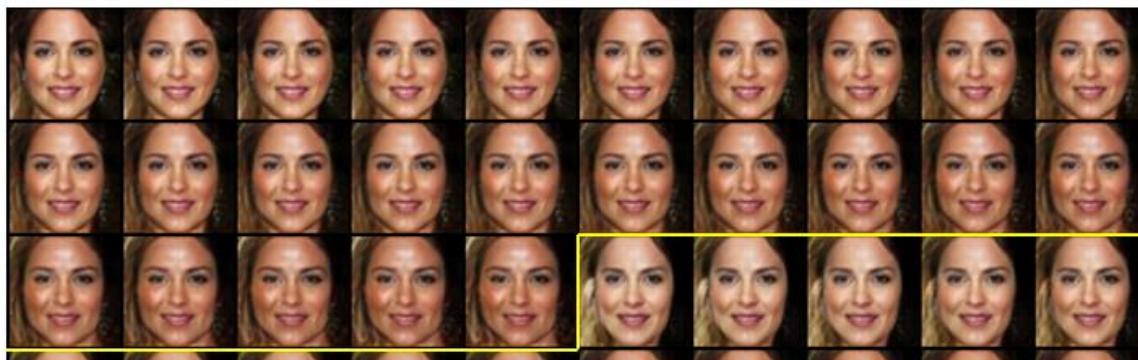
1 Categorical Feature, 2 Continuous Feature, 256 Noise



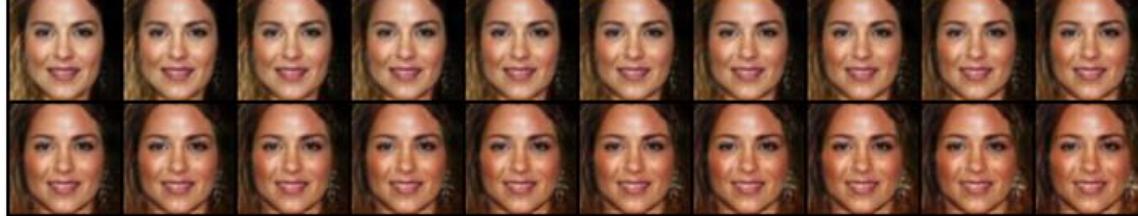
Aging



Skin/Hair Color



Aging



Skin/Hair Color

GAN Issue

- Convergence
- Realistic Visual
- Performance Measure
- Application

Part 2. GANs for CV

Overview

- Image Translation
 - Generation, Segmentation, Style Transfer
 - Edge Detection, Colorization, Multi-model Imaging
- Super Resolution
- Feature Upscaling: Object Detection
- 3D Reconstruction
- Semi-supervised Learning
 - Pose Estimation, Face Recognition, Pedestrian Detection
- Face Aging Modeling

Image Translation Domain Transfer

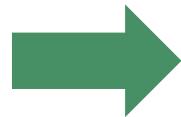
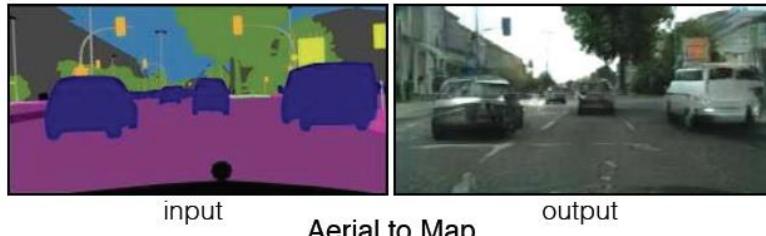


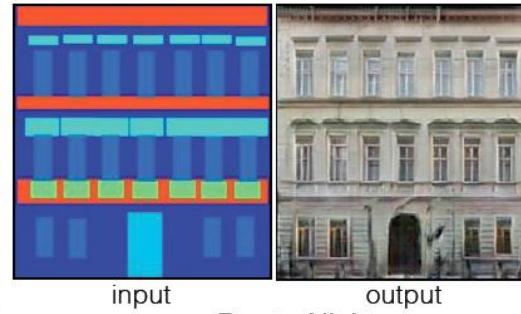
Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
a.k.a Pix2Pix, CVPR 2017

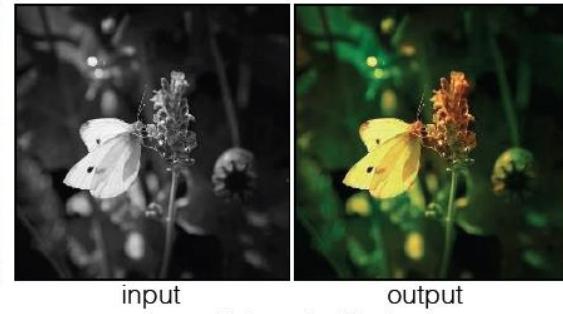
Labels to Street Scene



Labels to Facade



BW to Color



Aerial to Map



Day to Night



Edges to Photo



Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

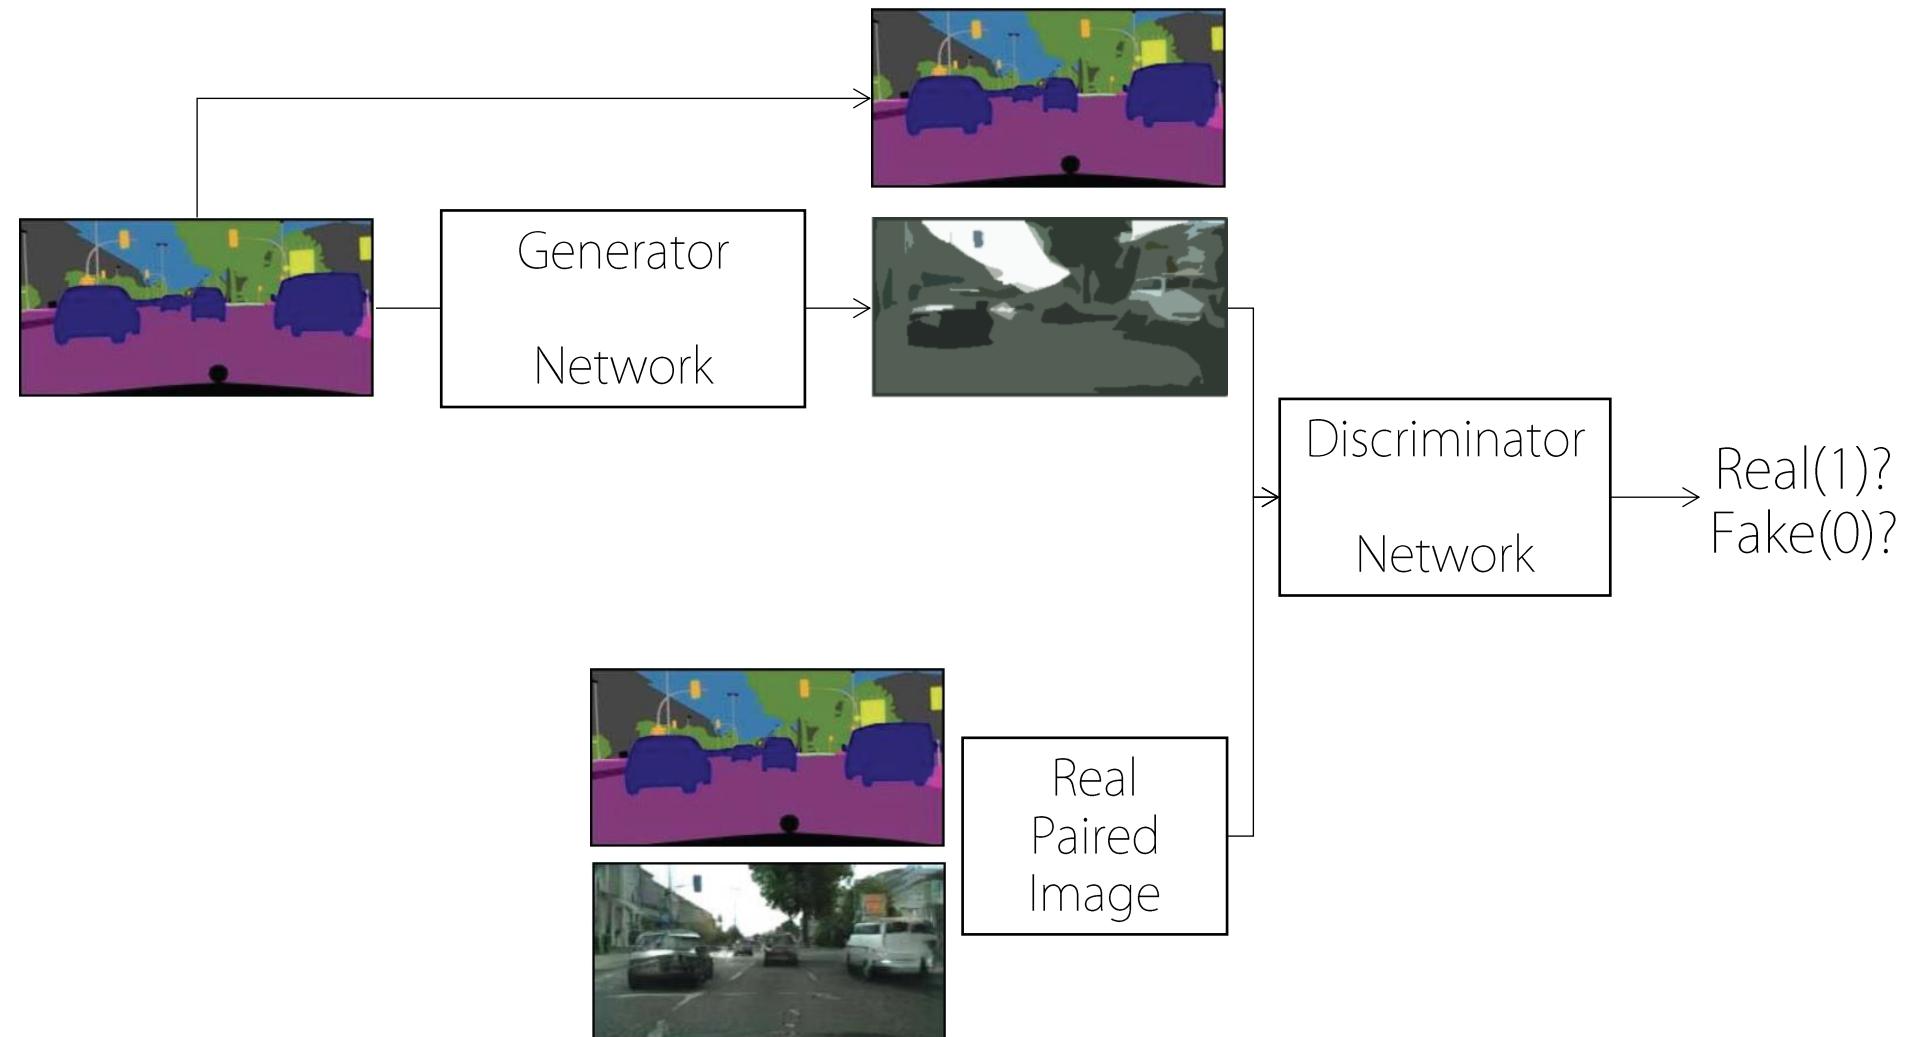


Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

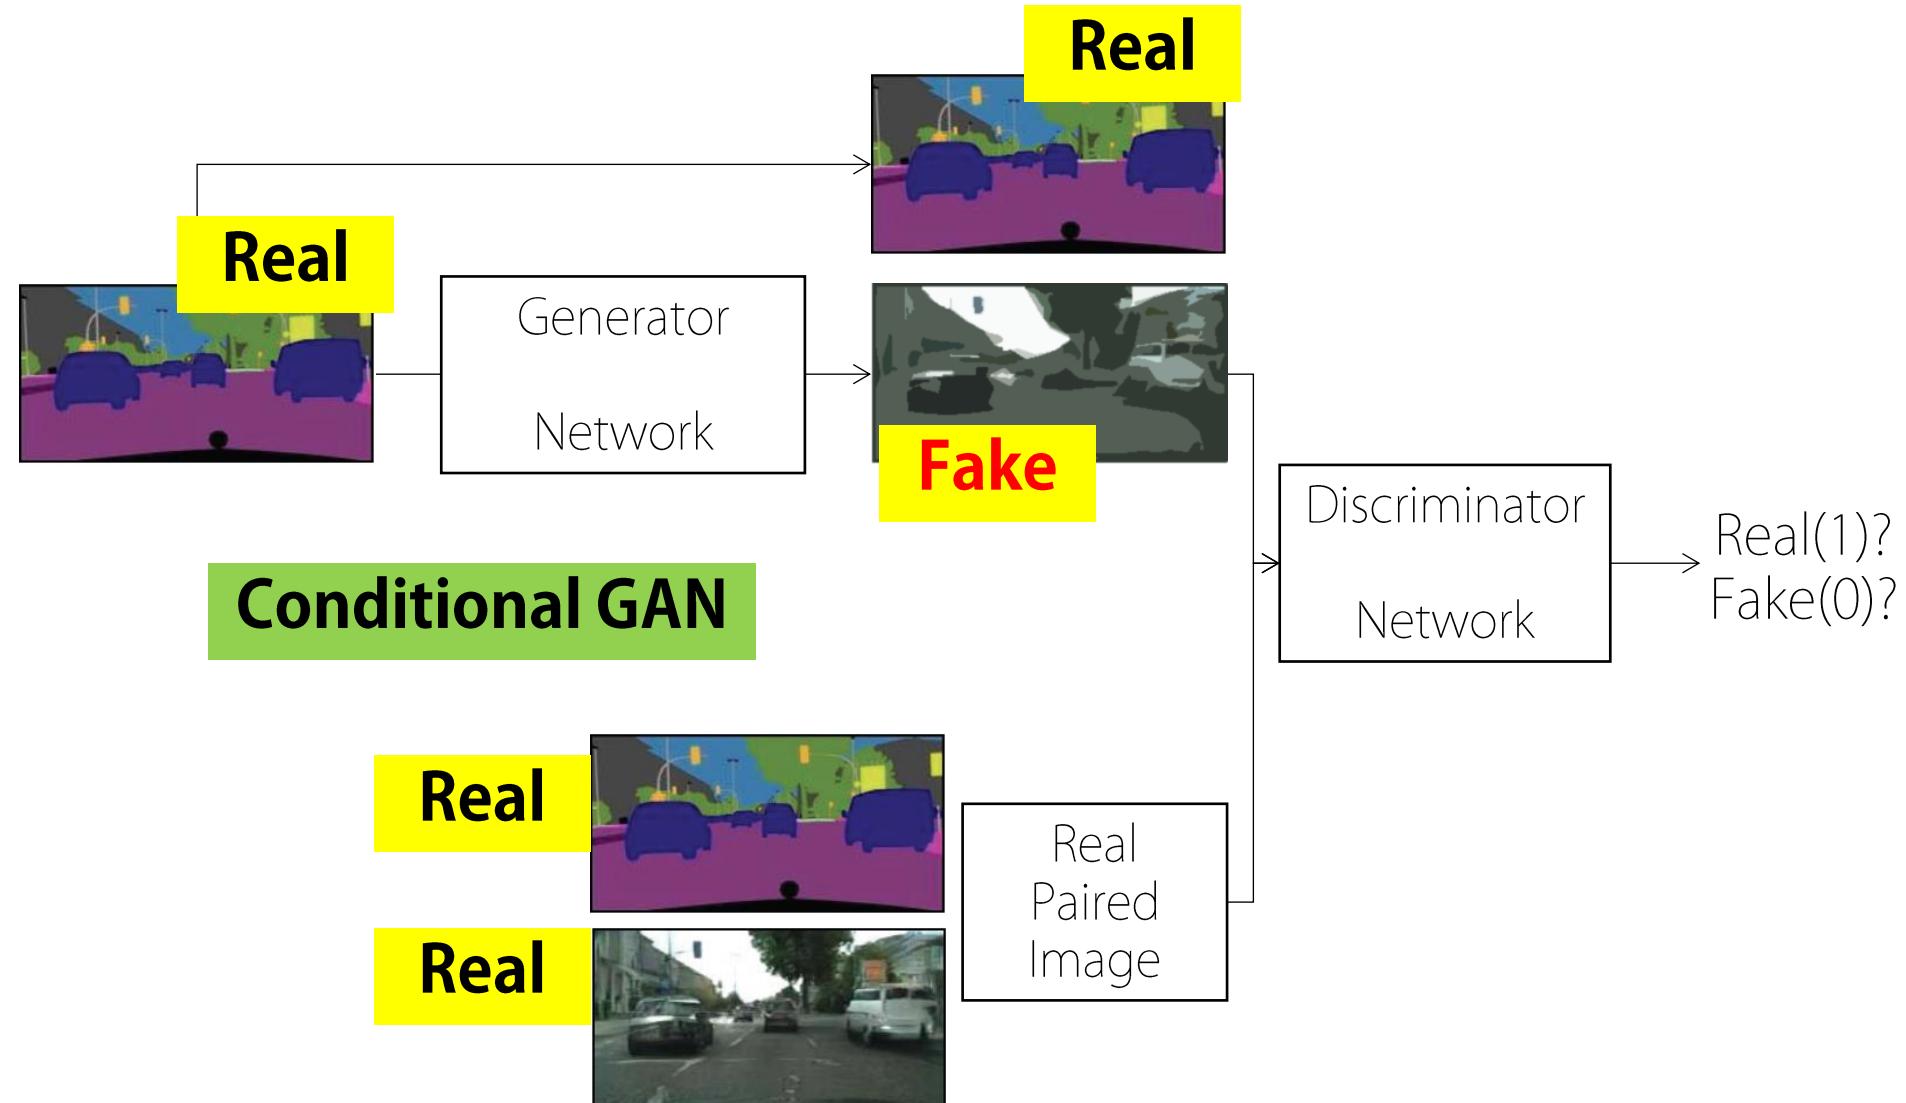


Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

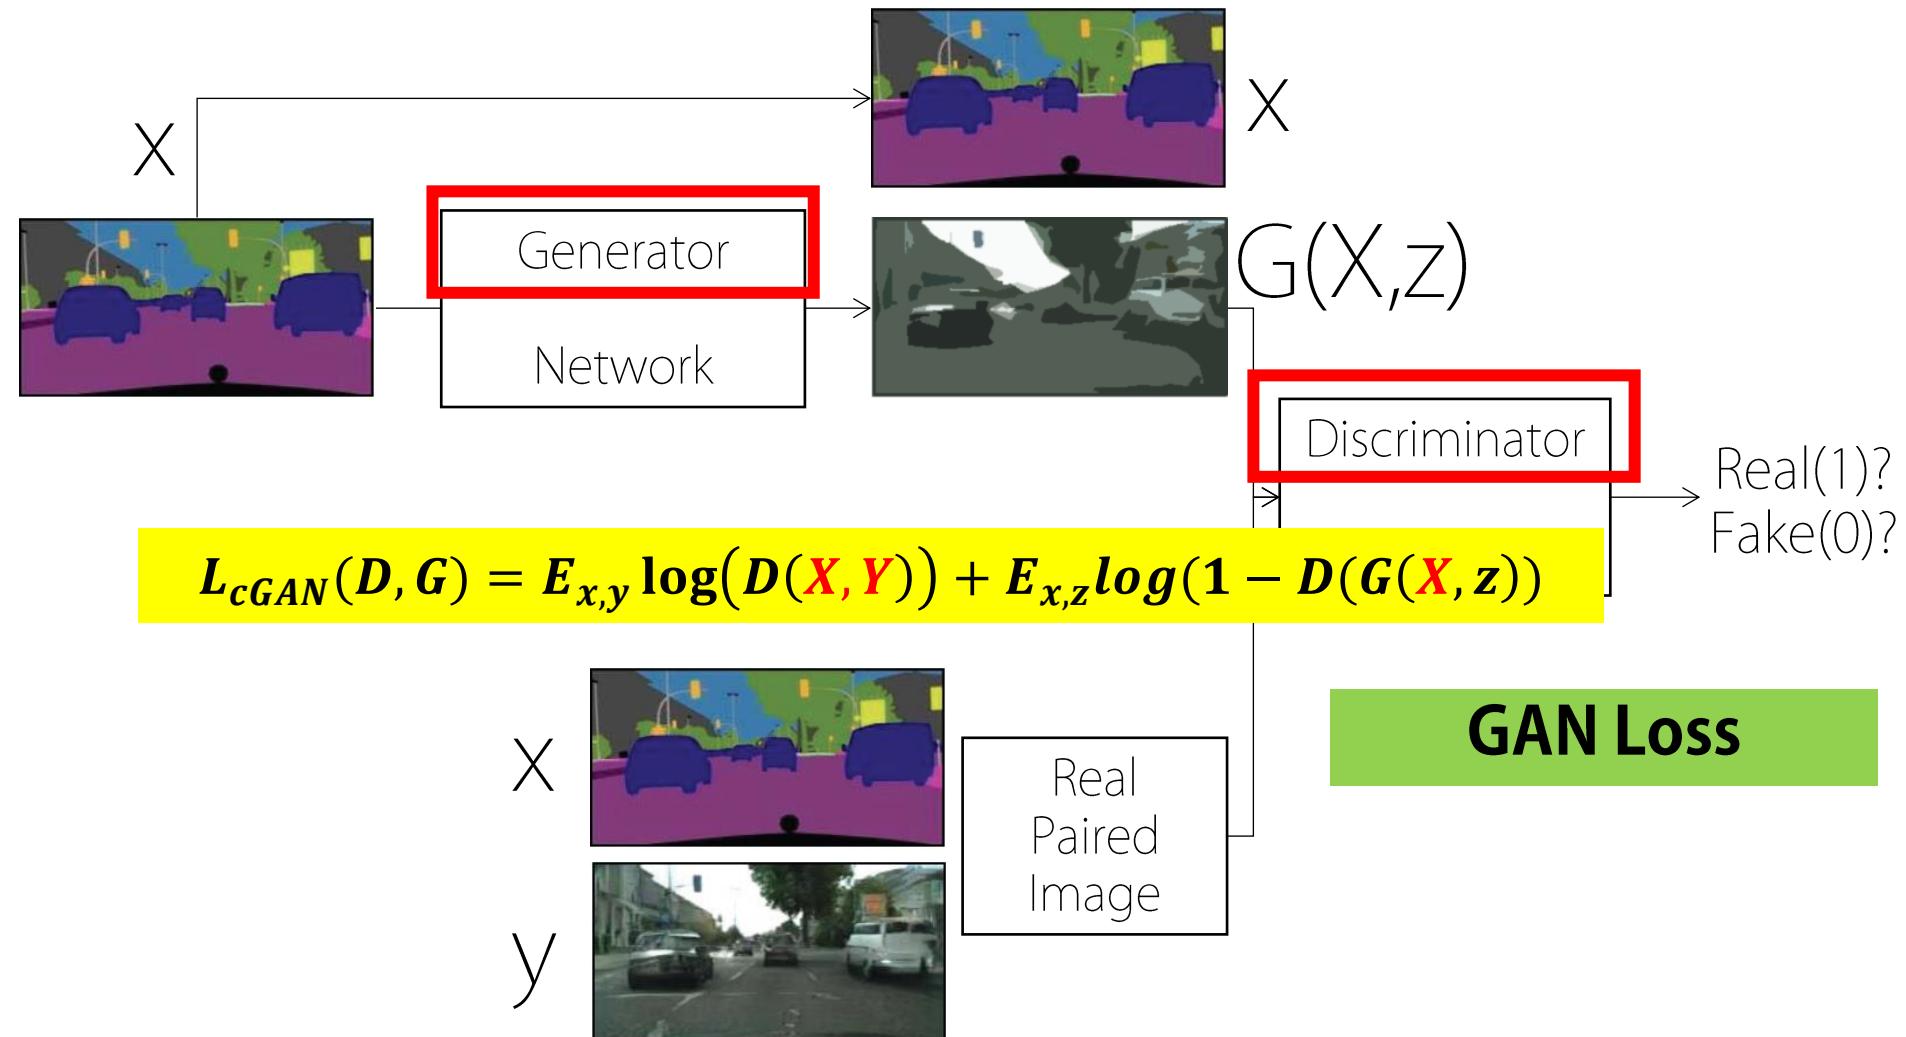


Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

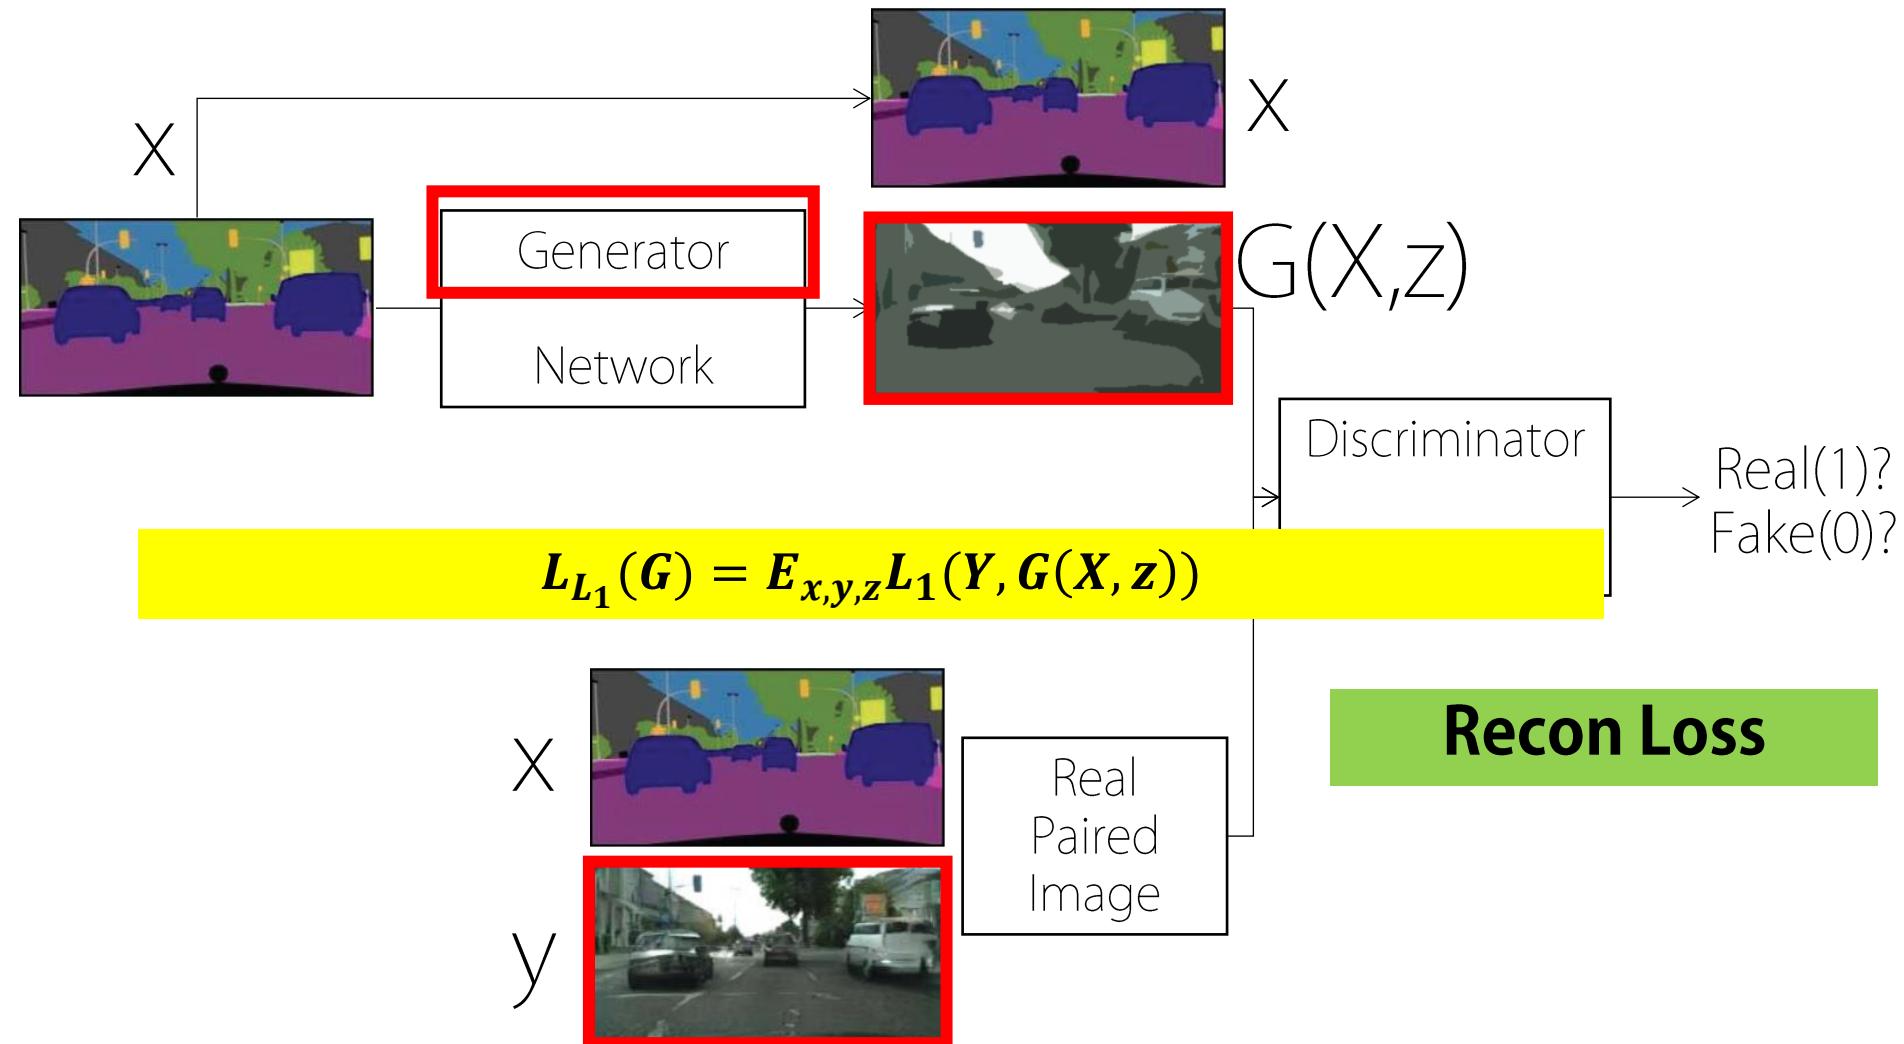


Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

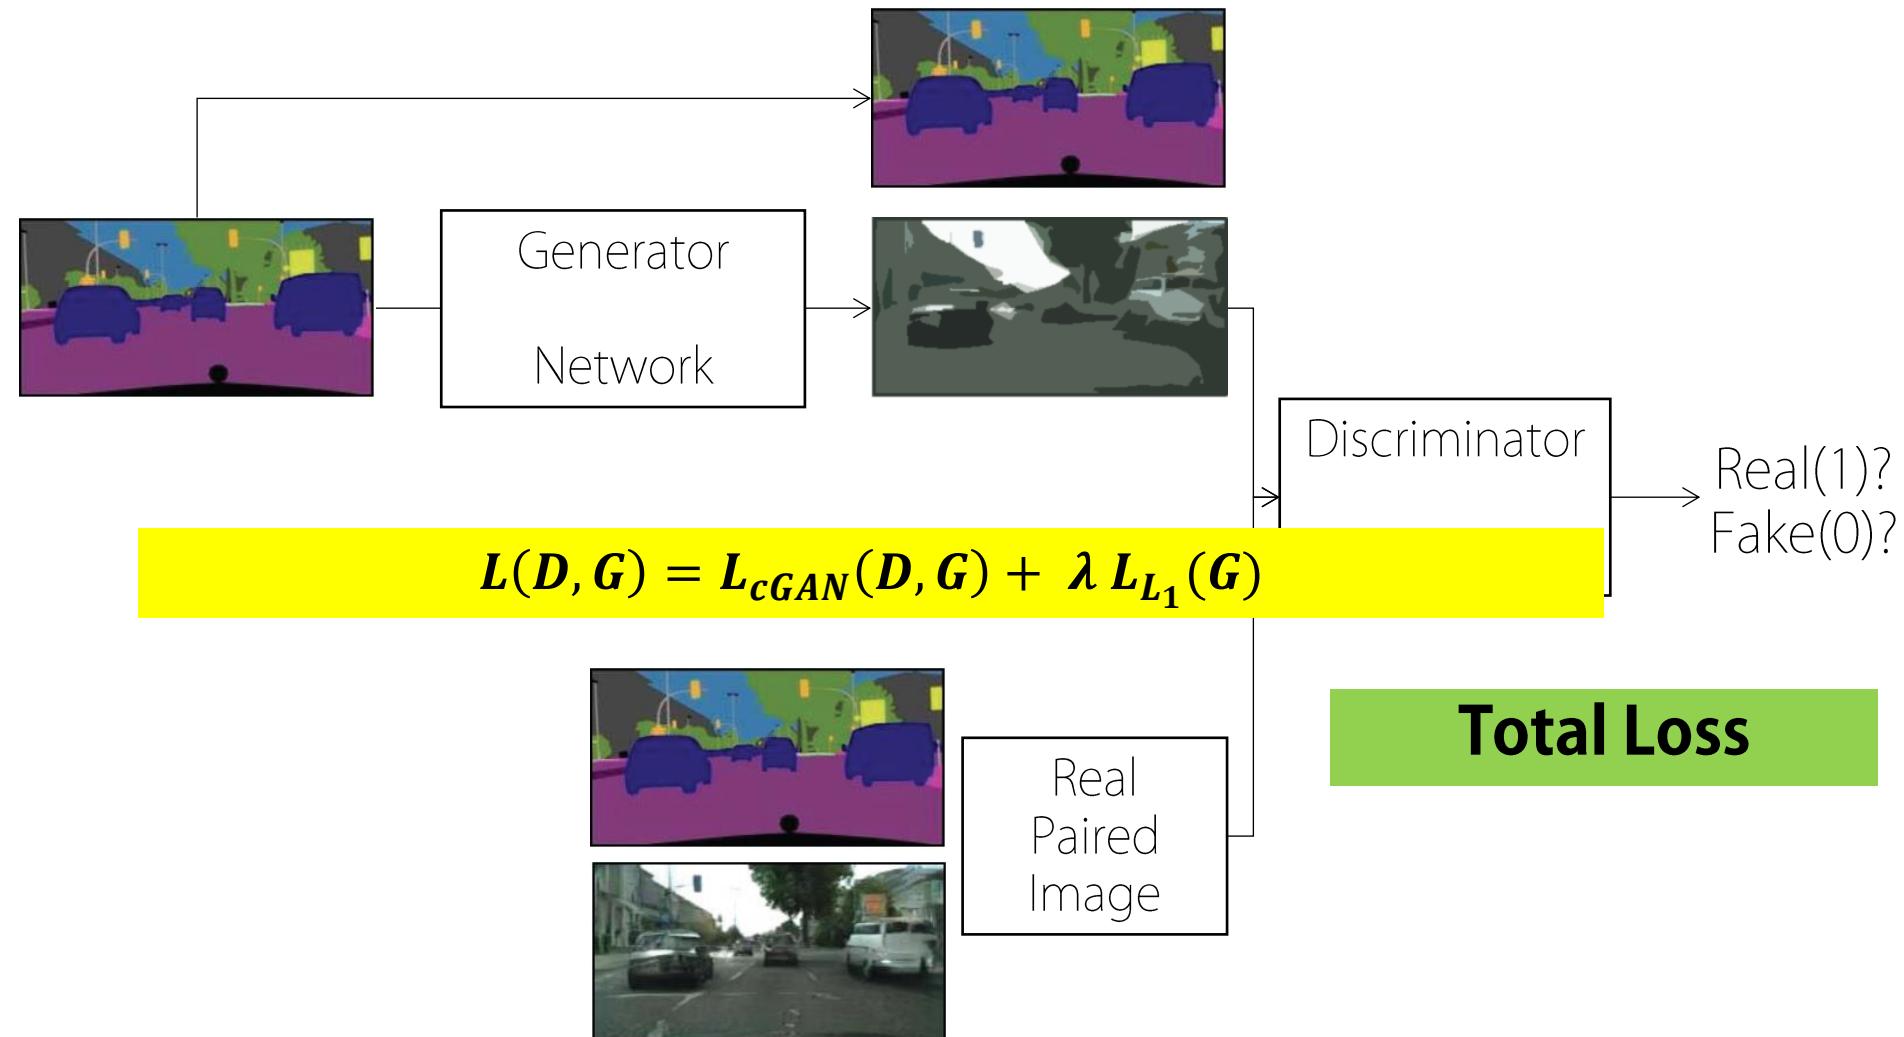


Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

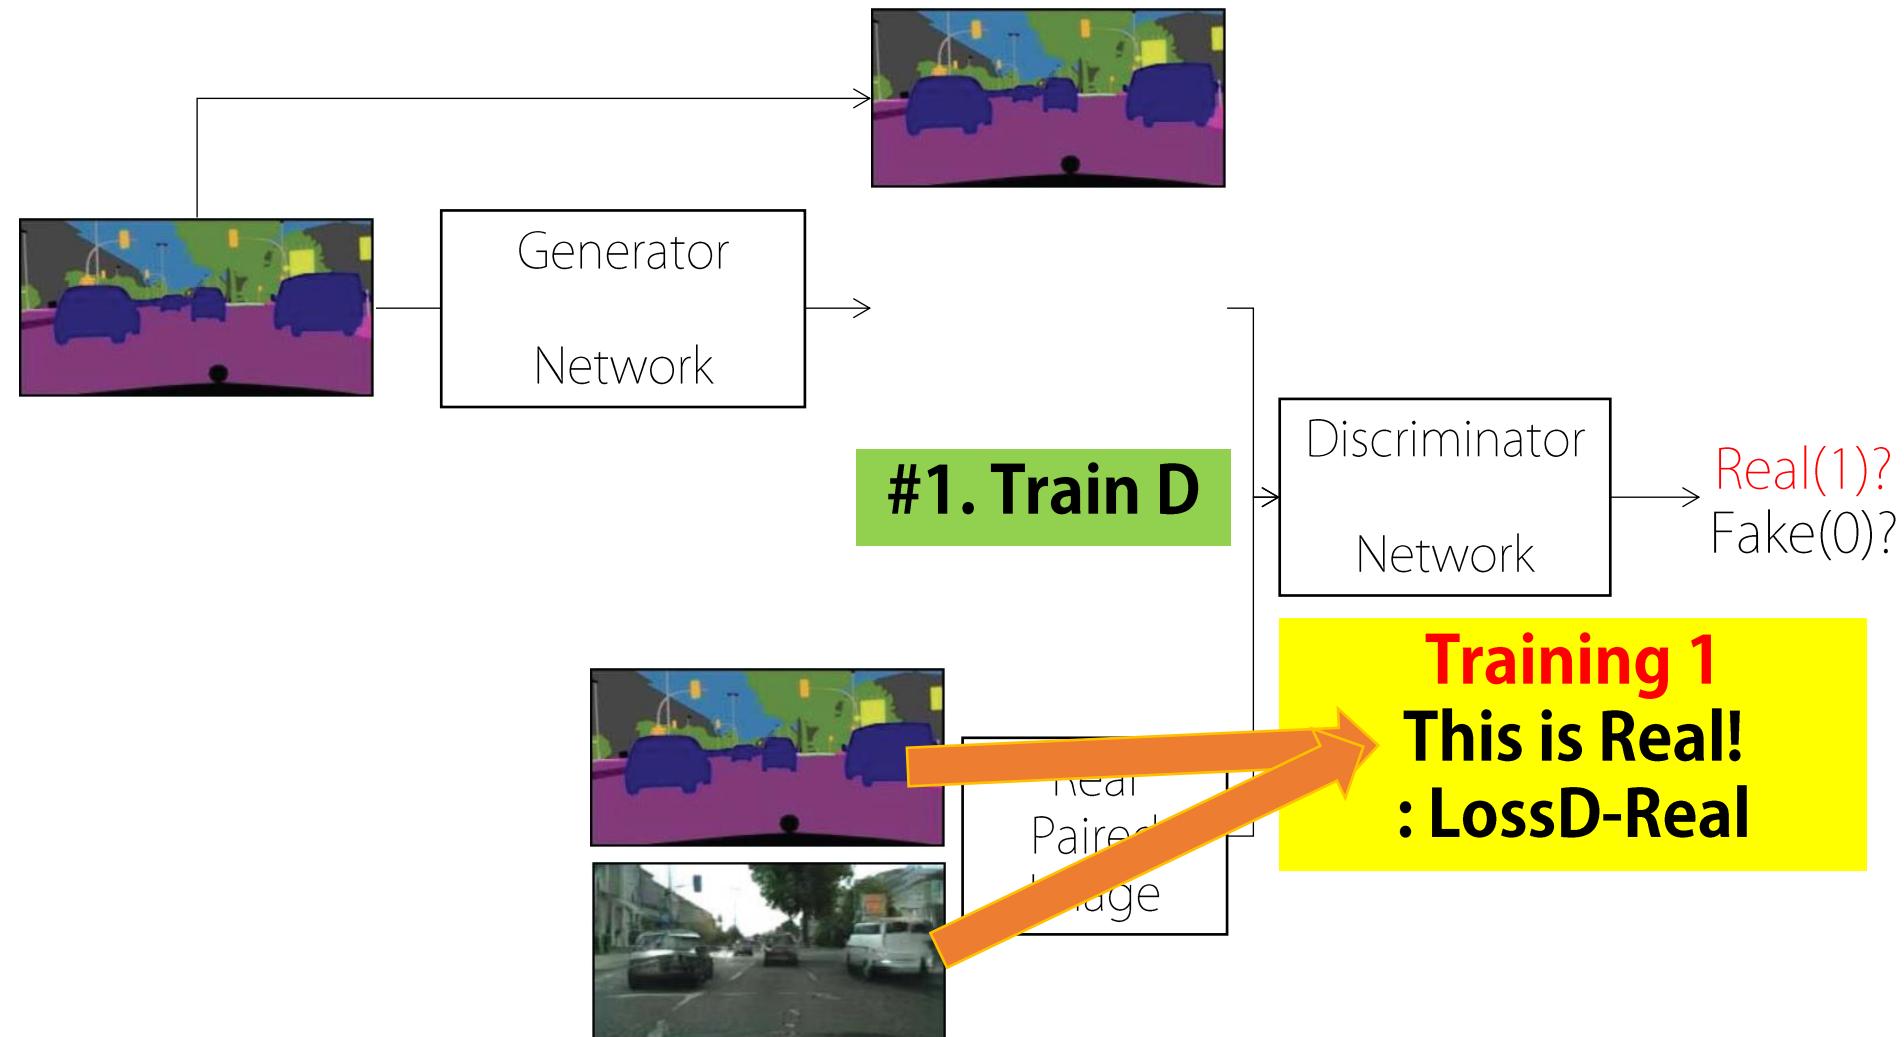


Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

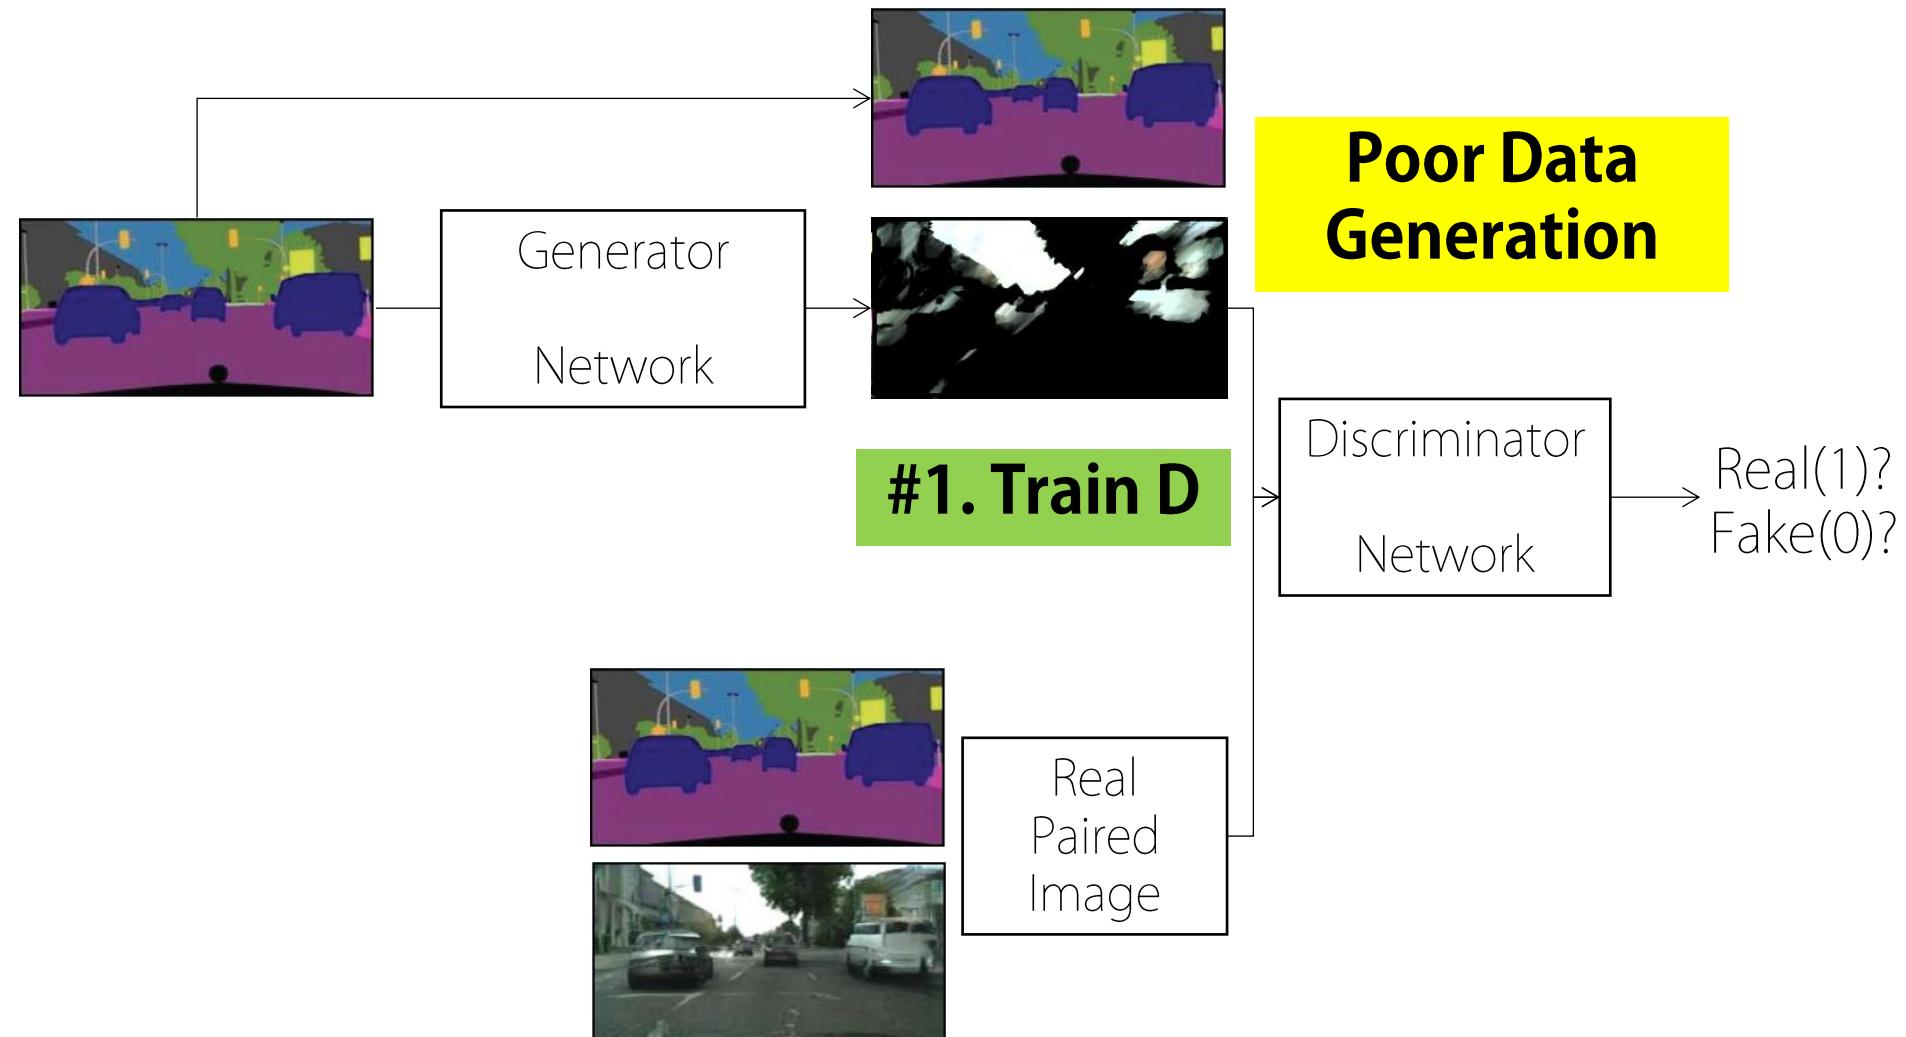


Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

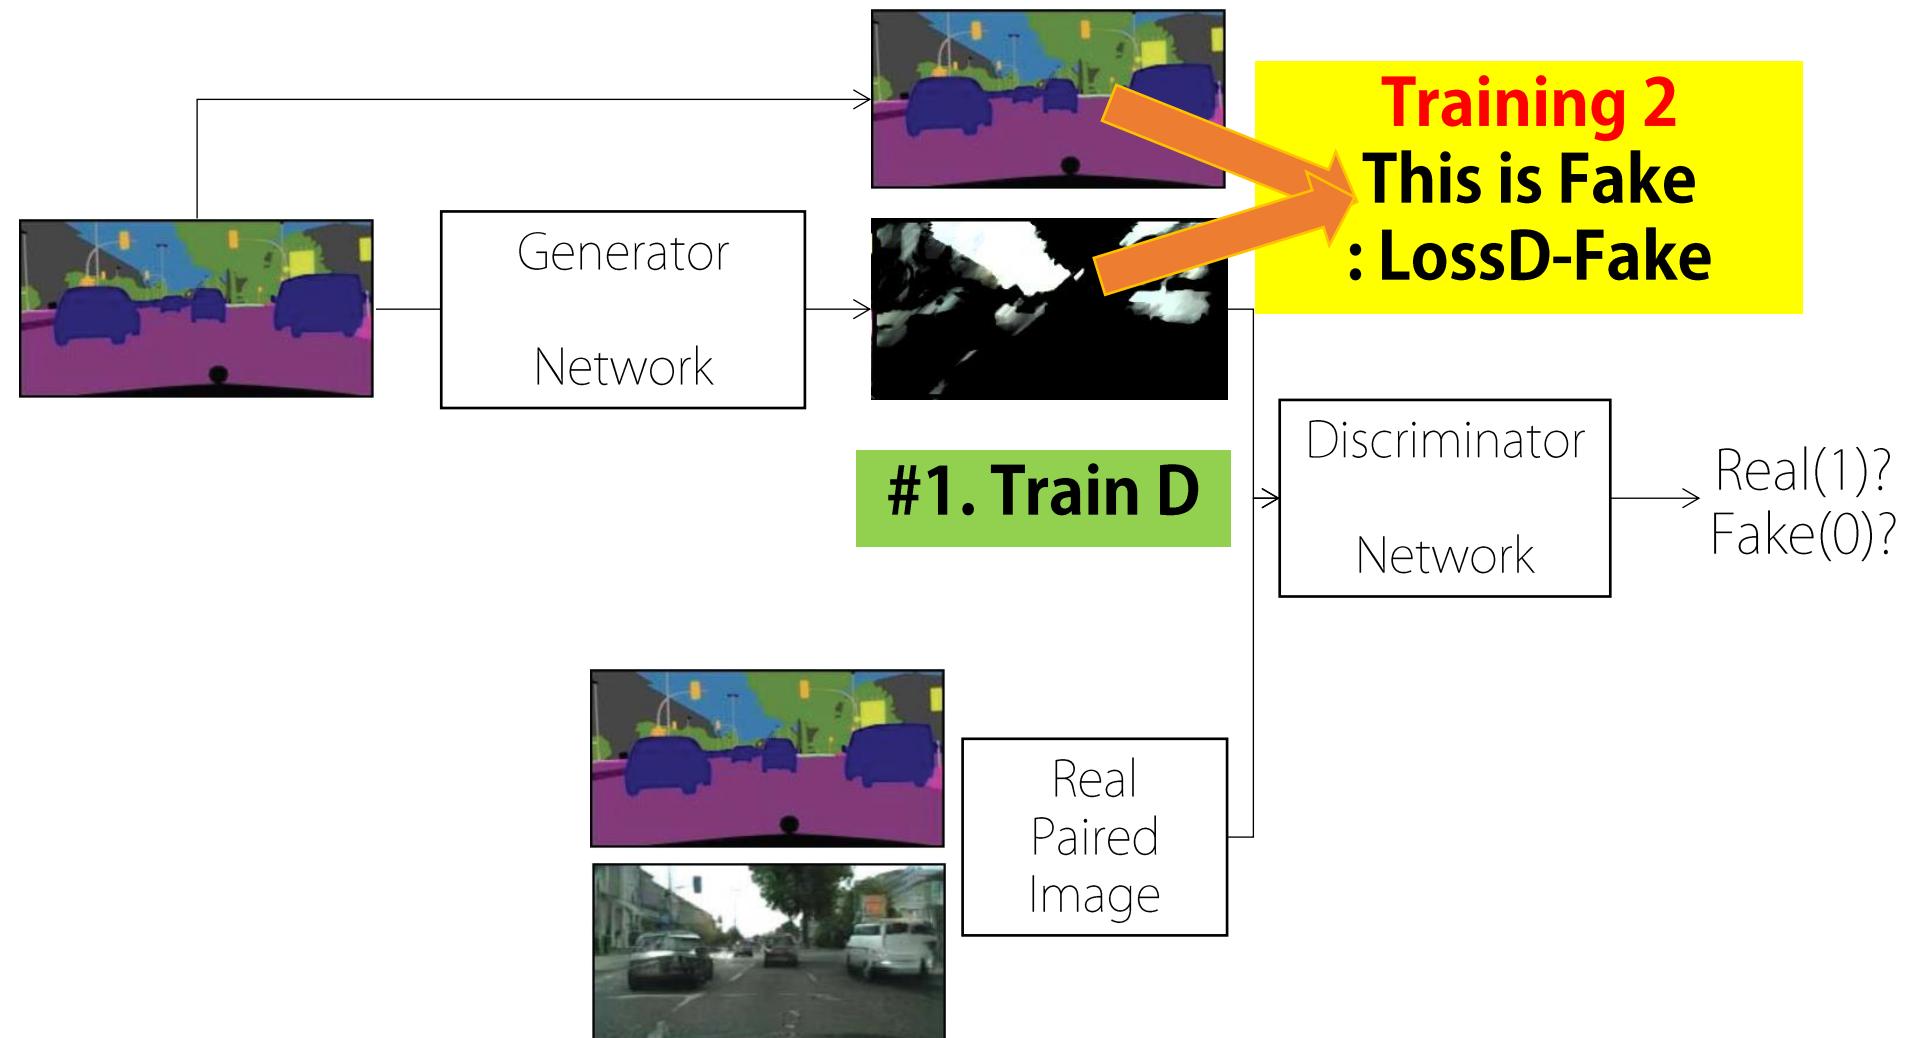


Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

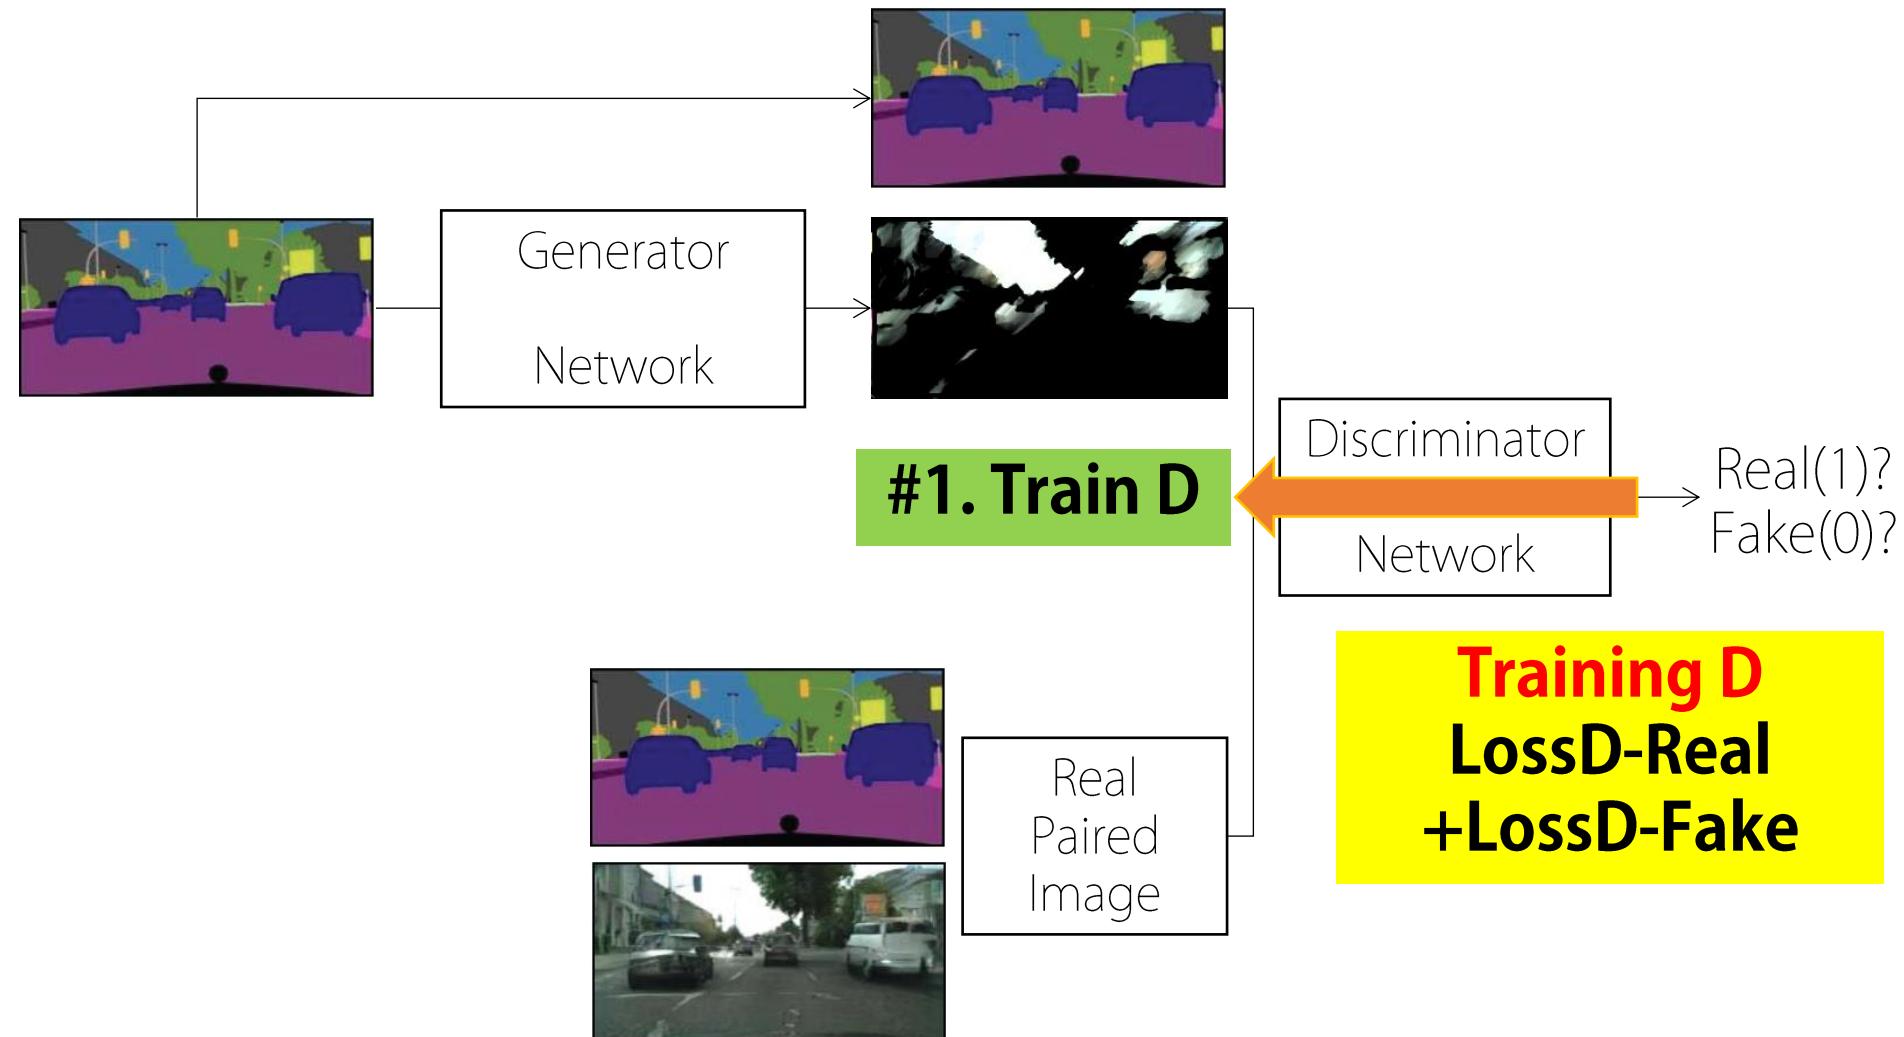


Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

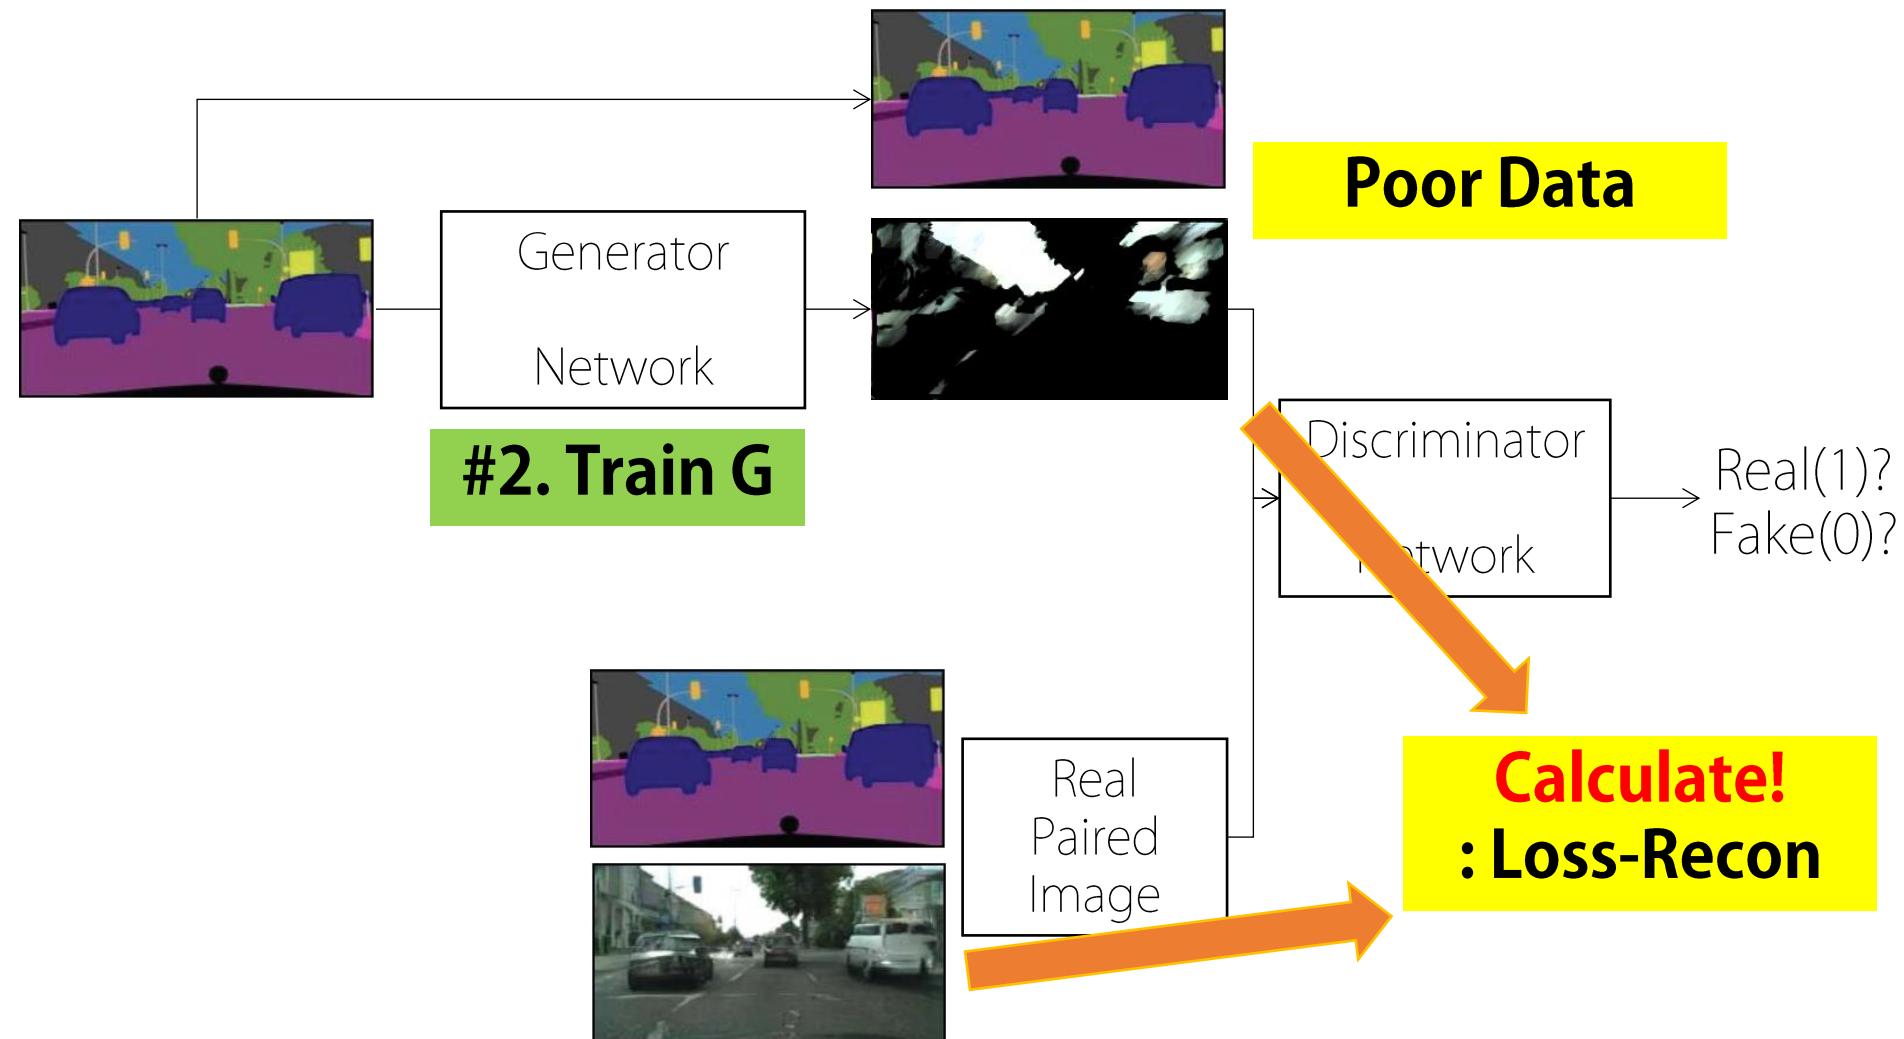


Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

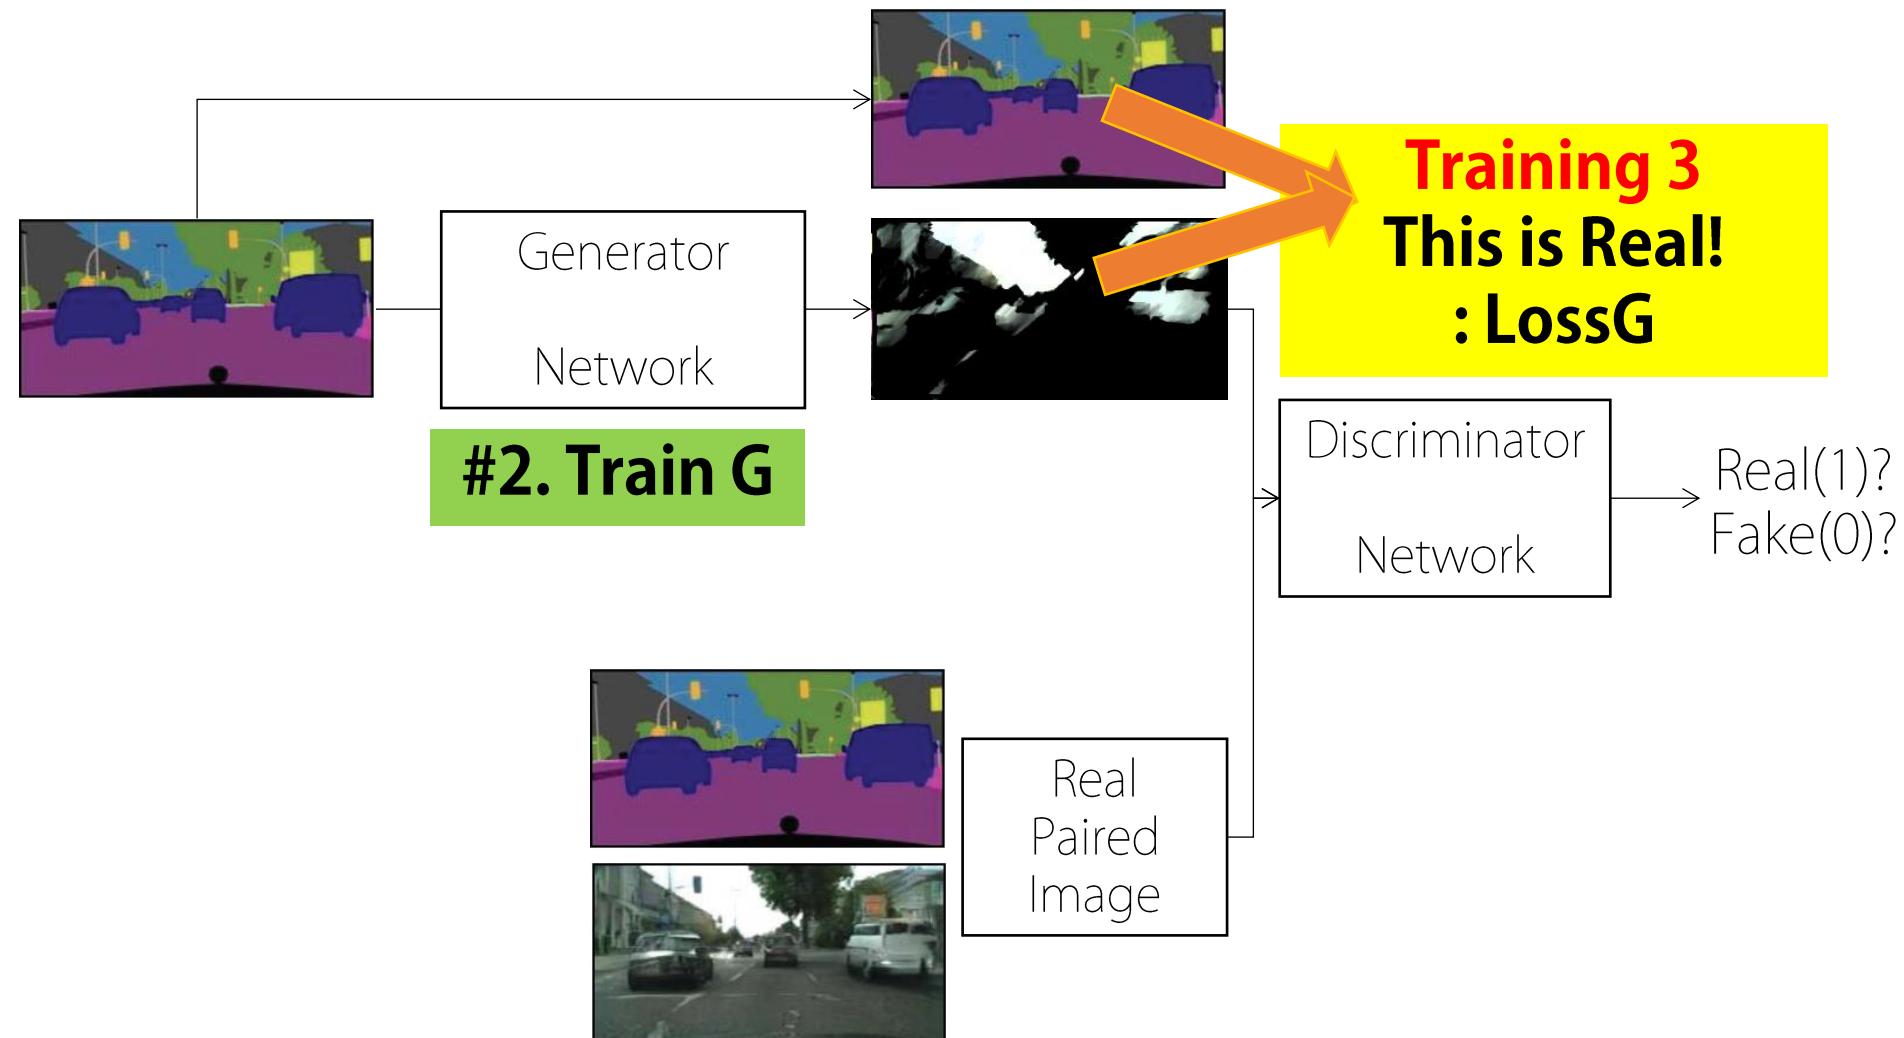


Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

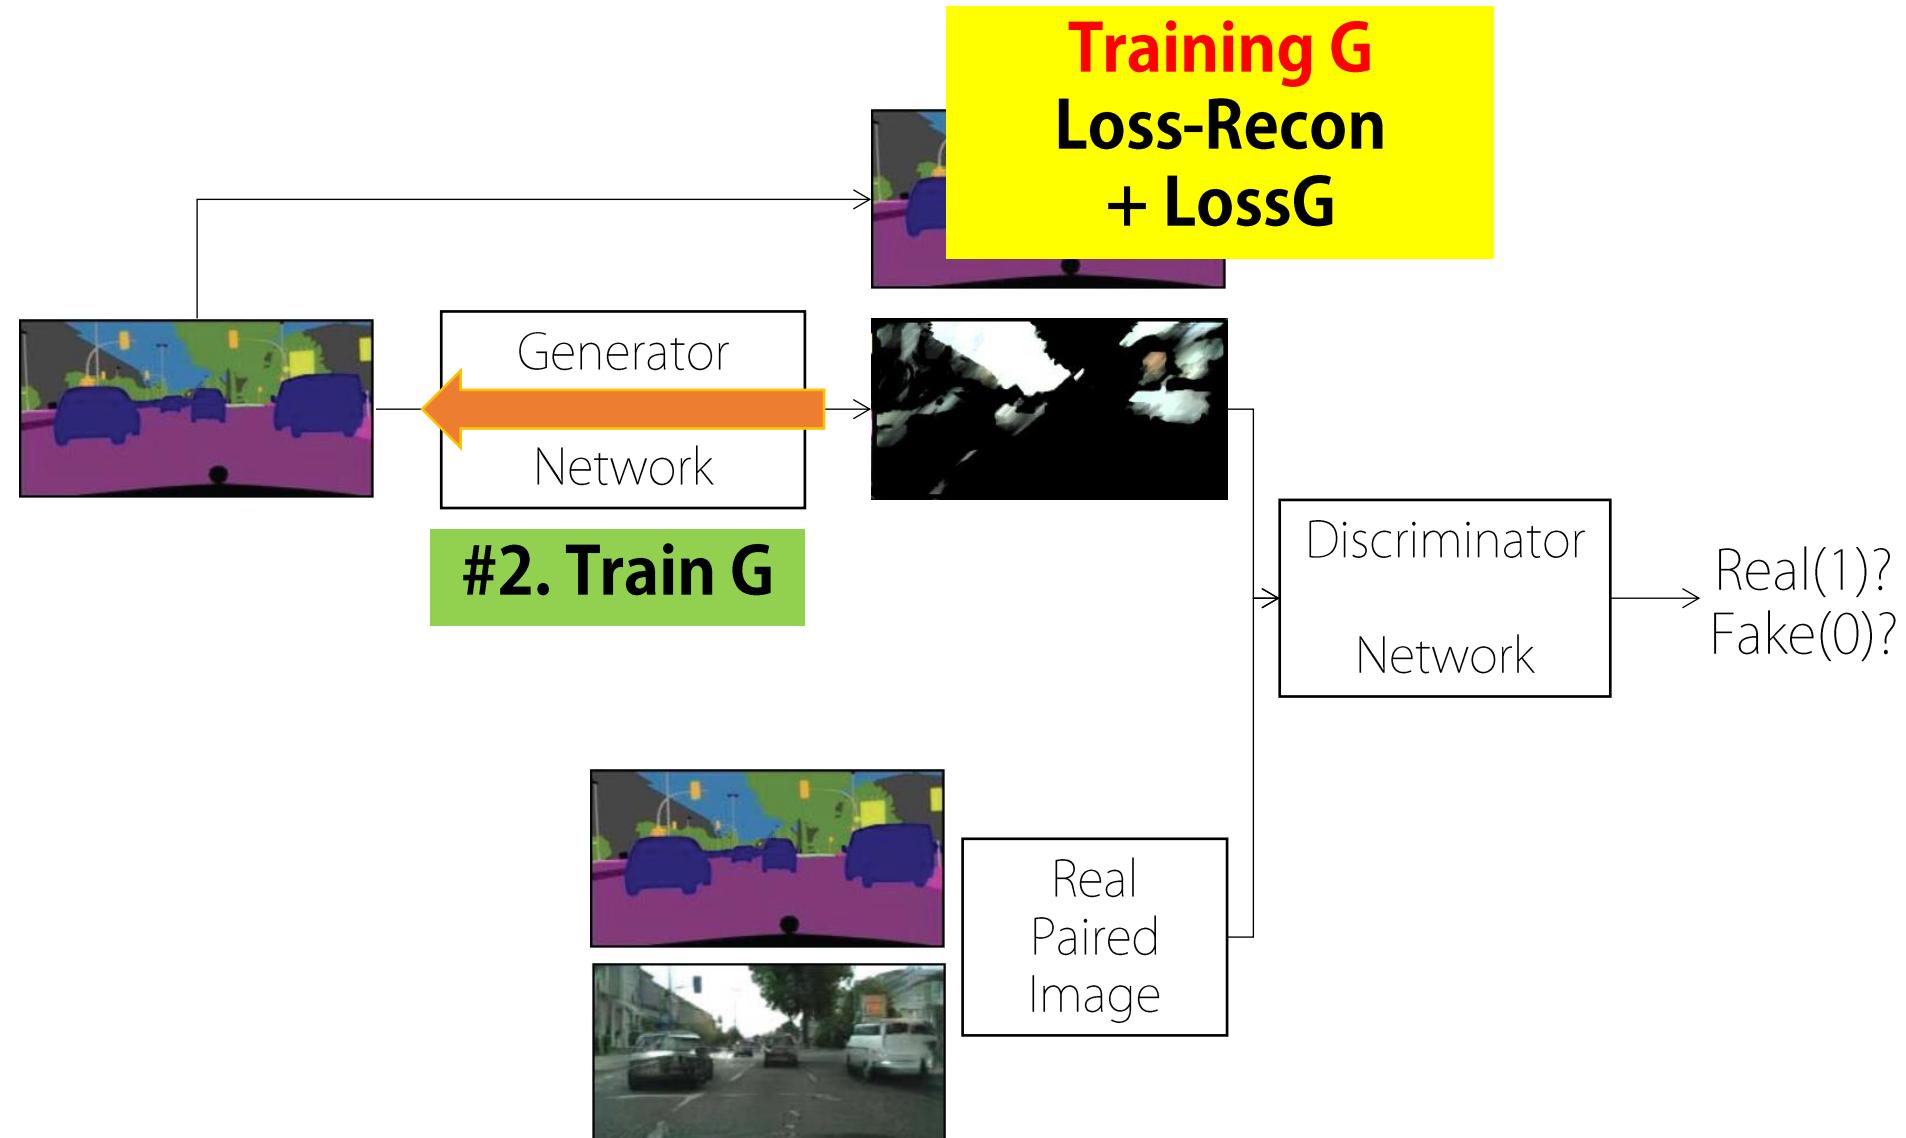


Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

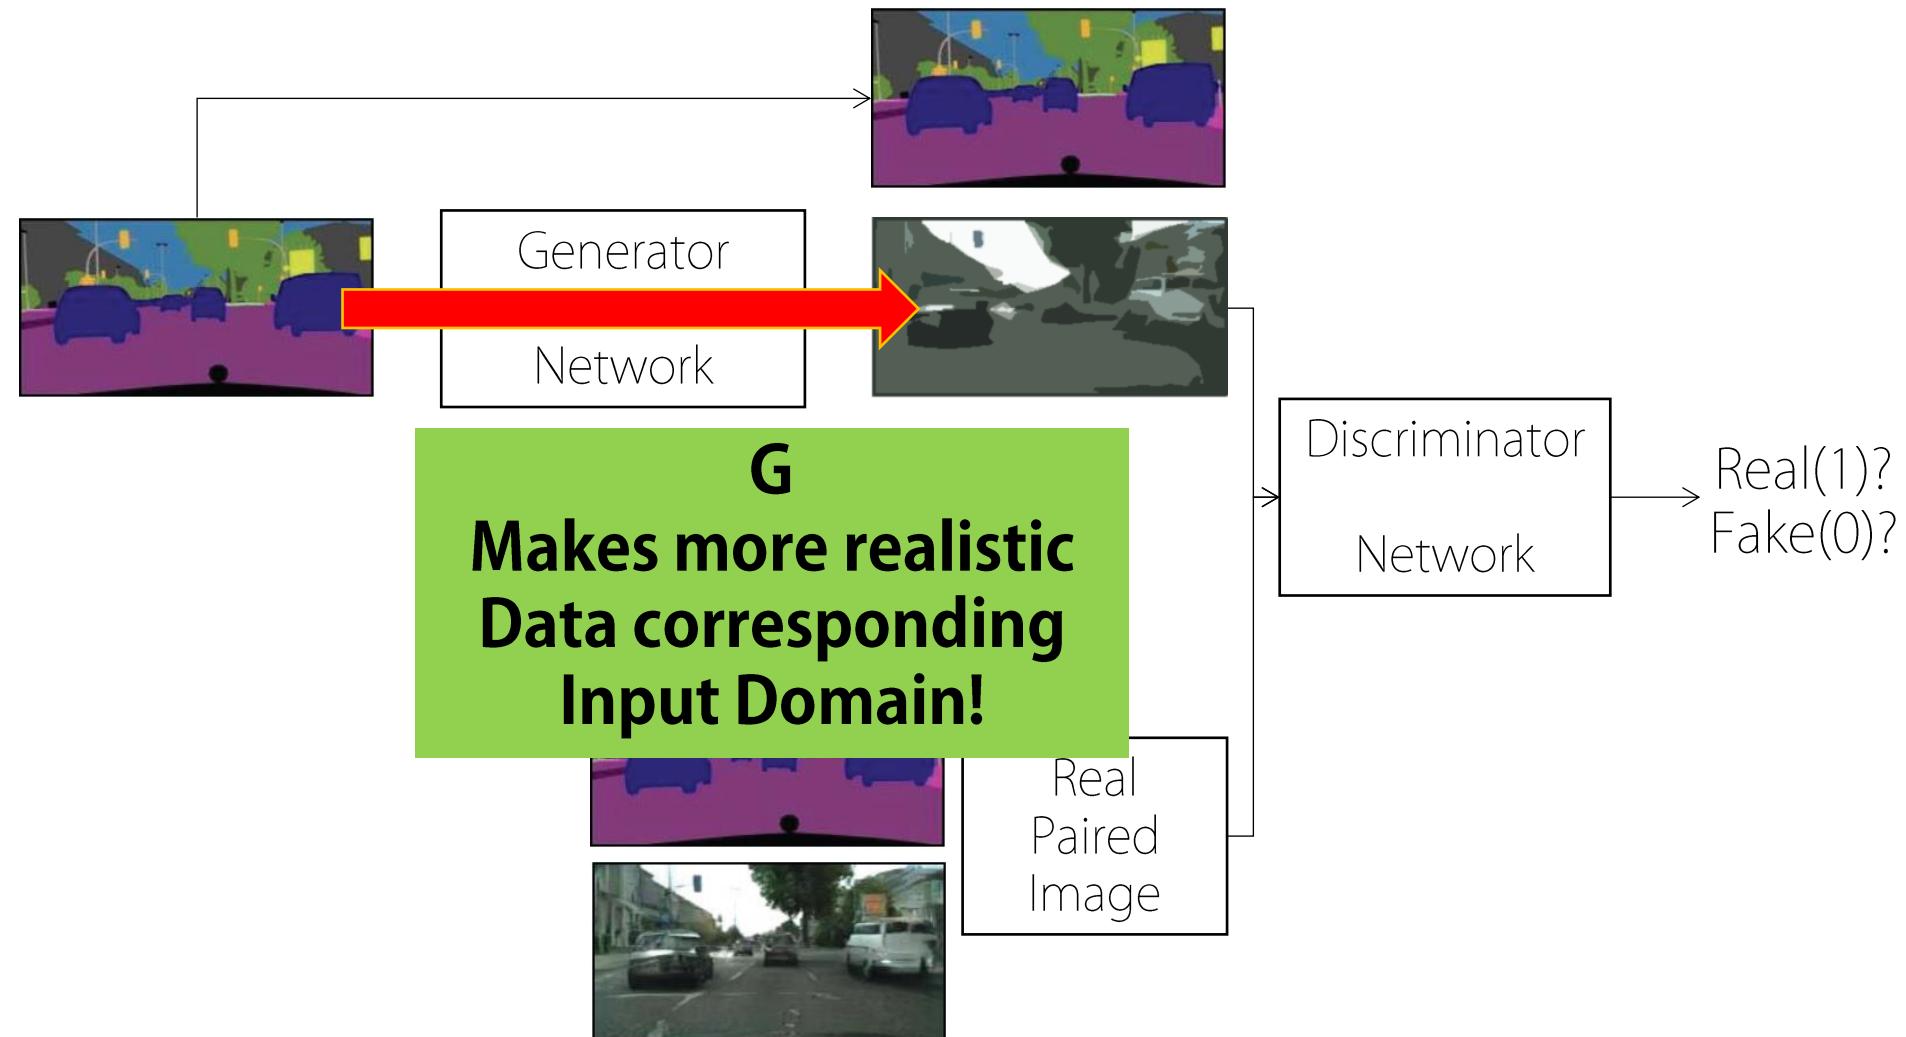
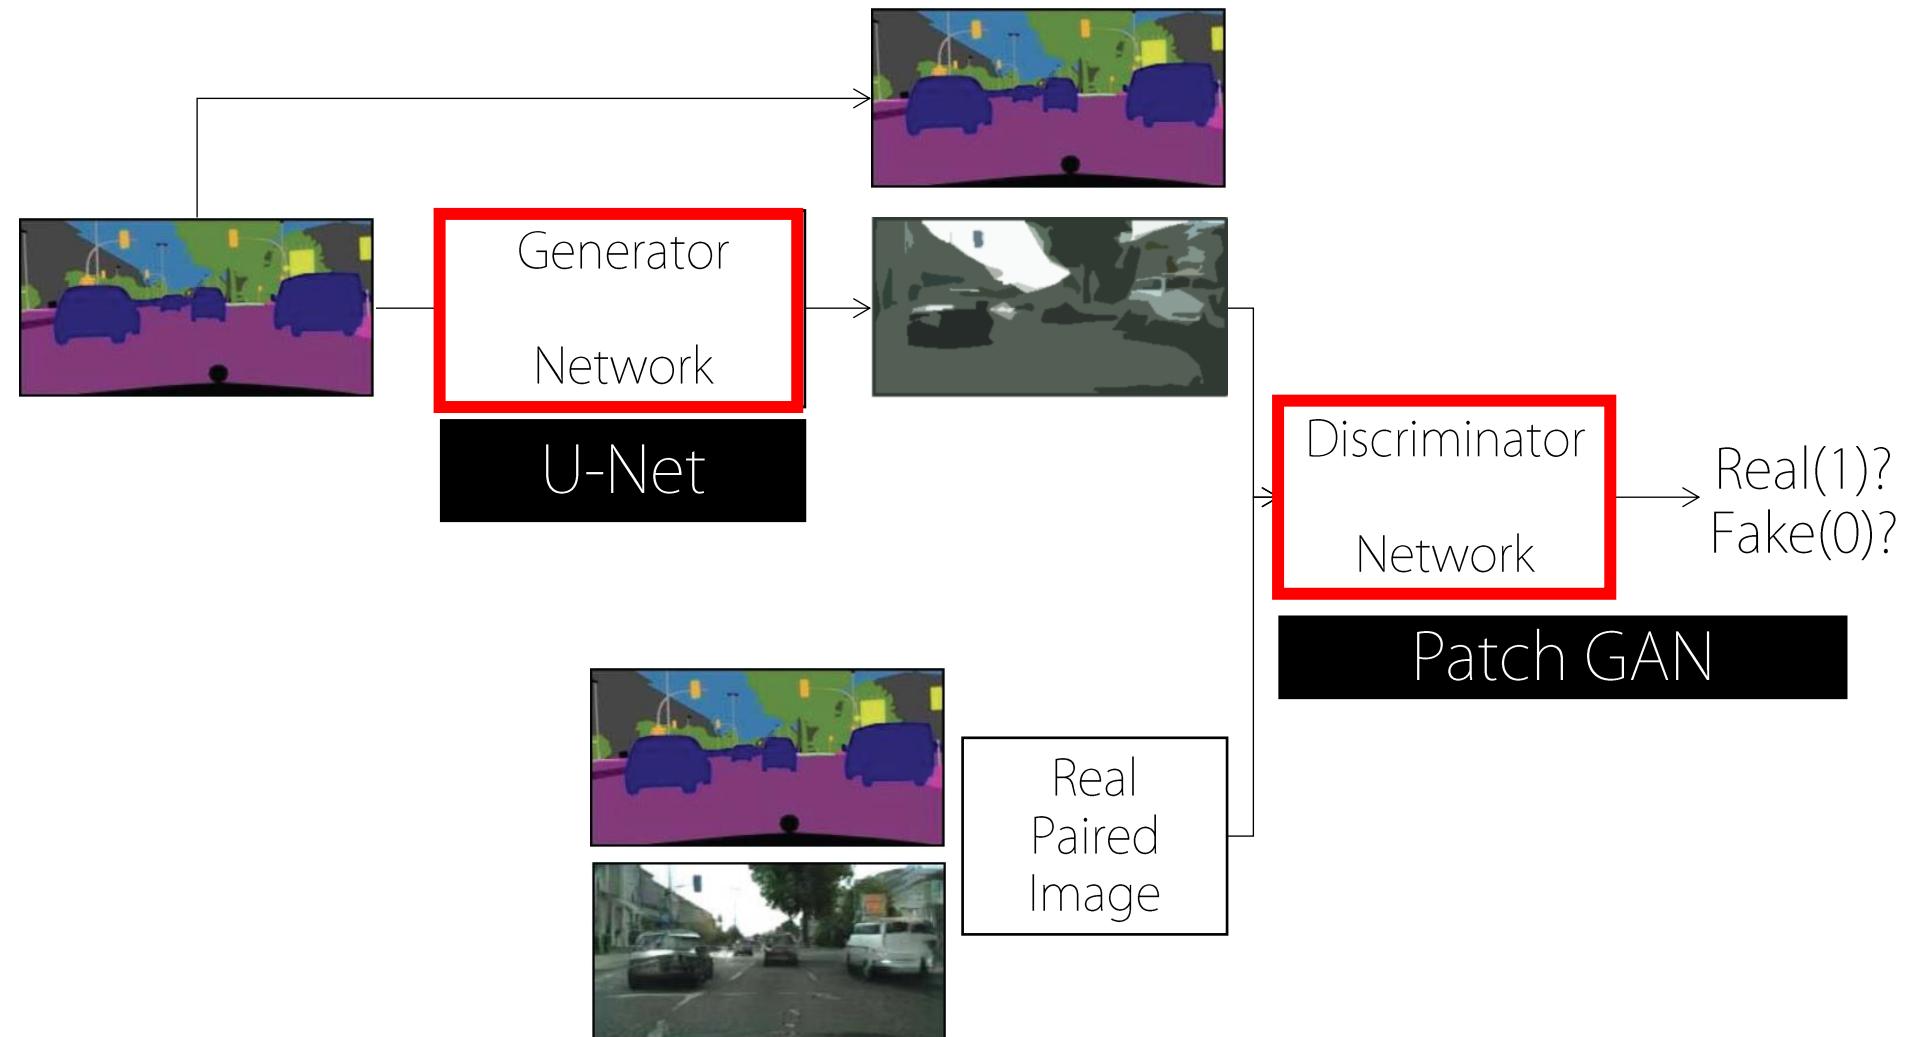


Image to Image Translation

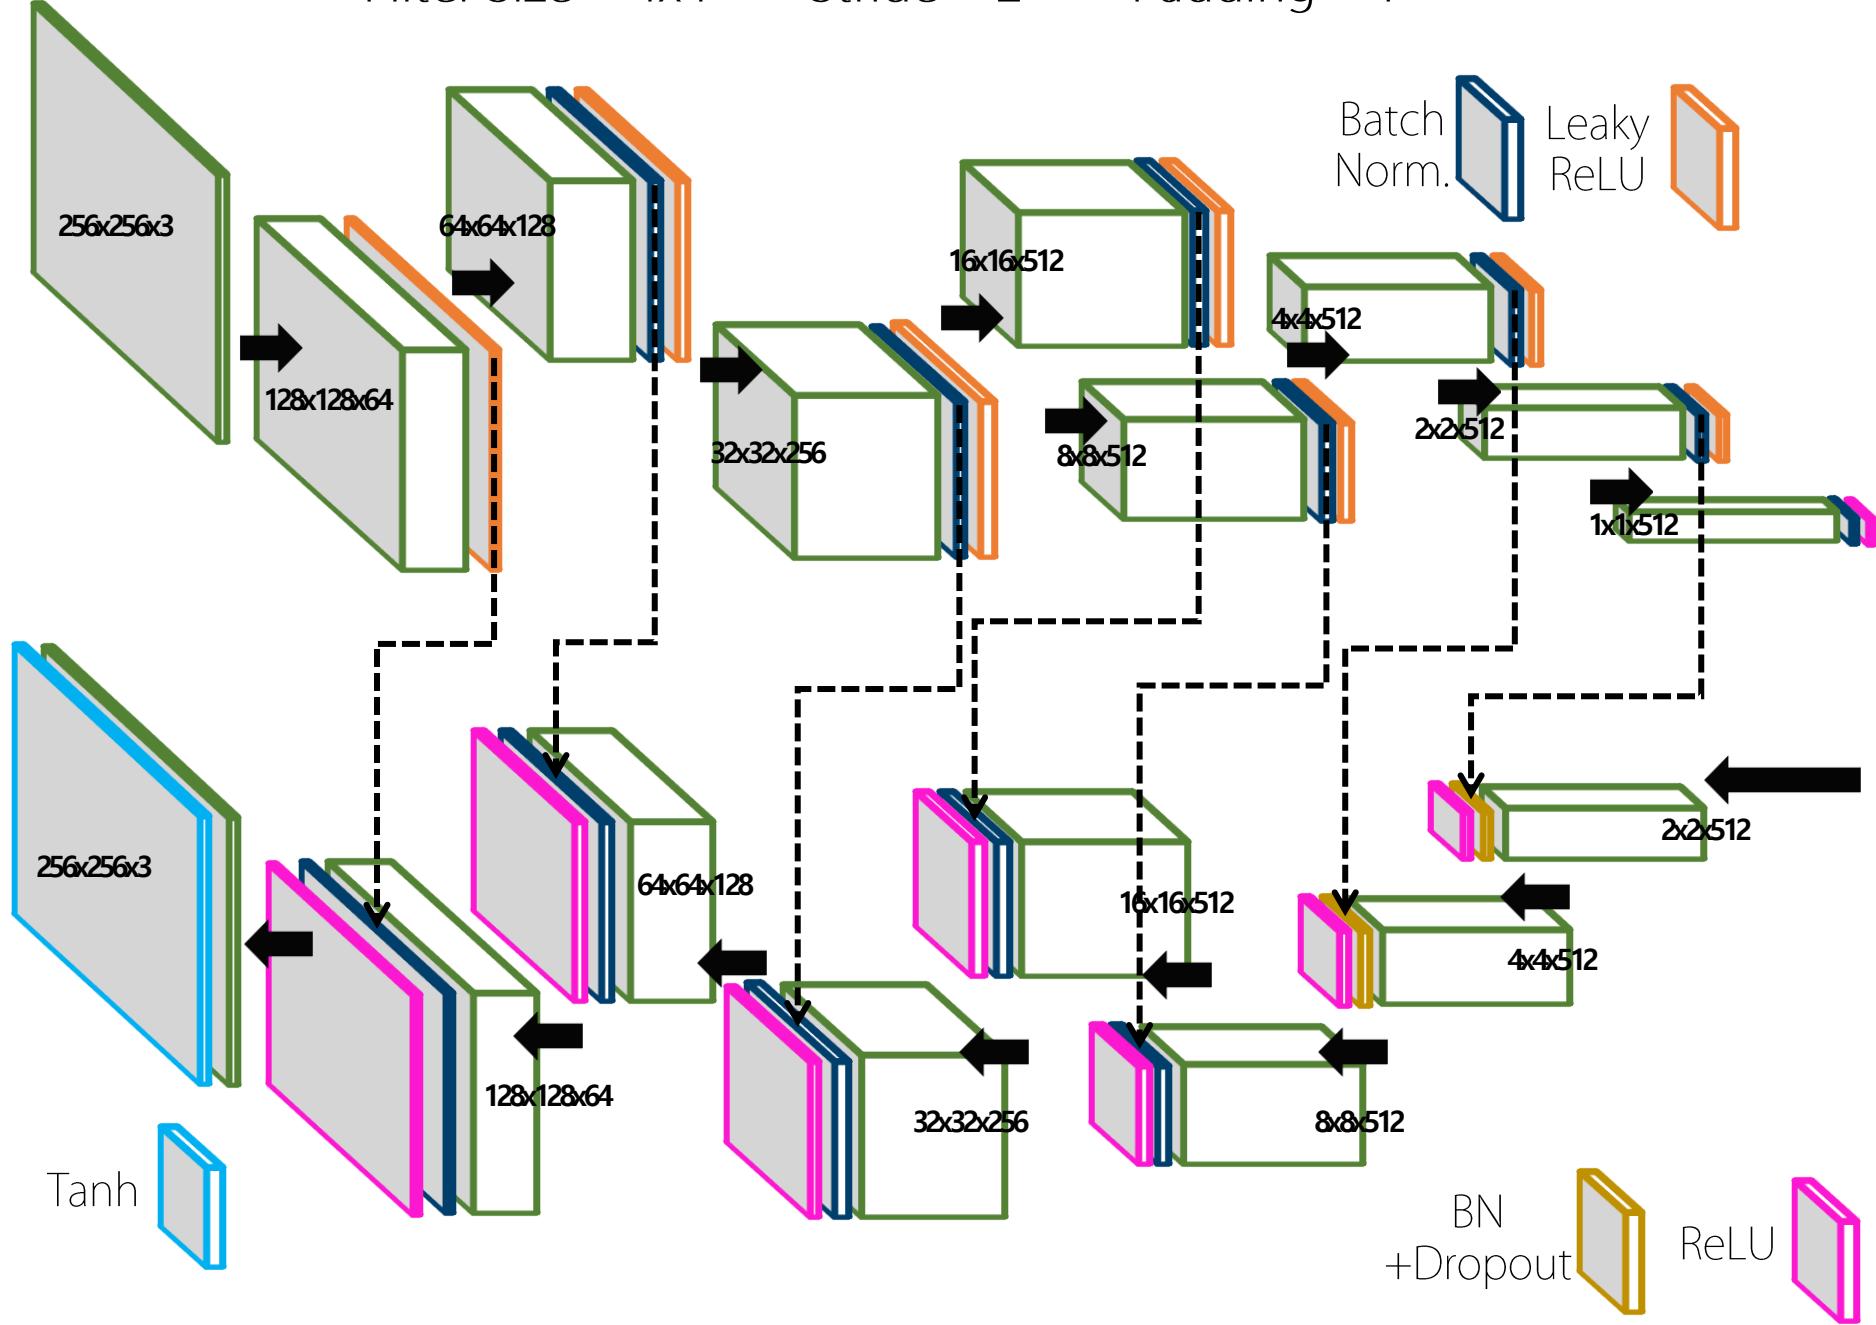
Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

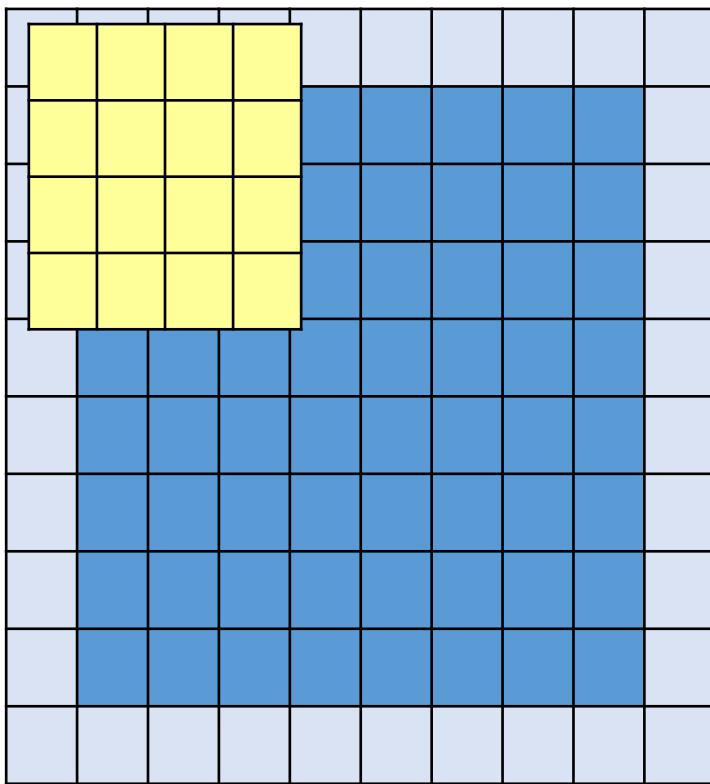


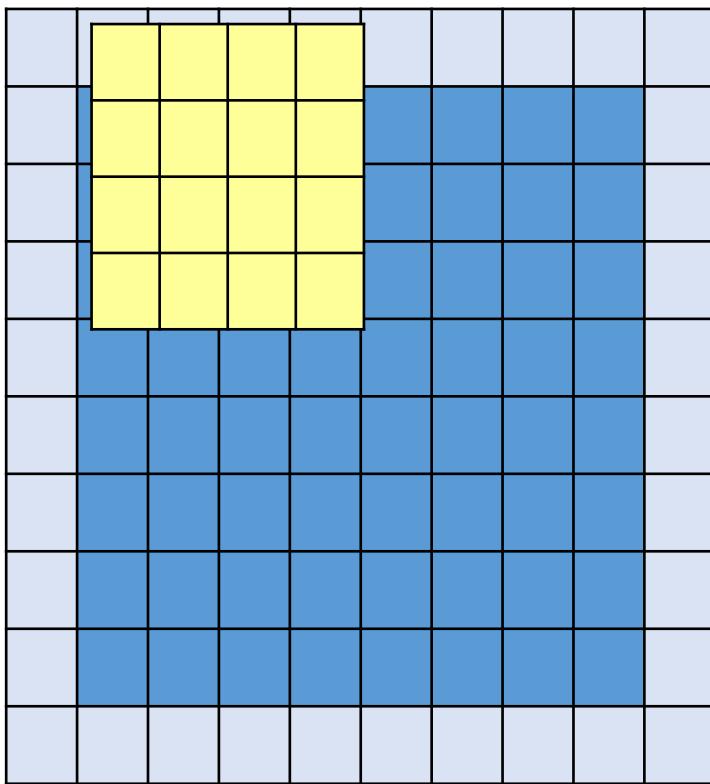
Filter Size = 4x4

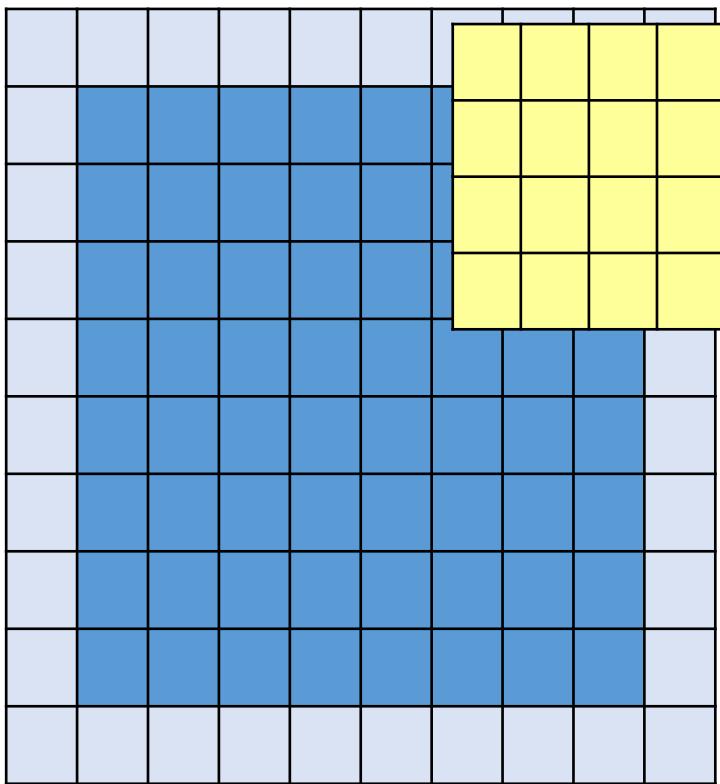
Stride = 2

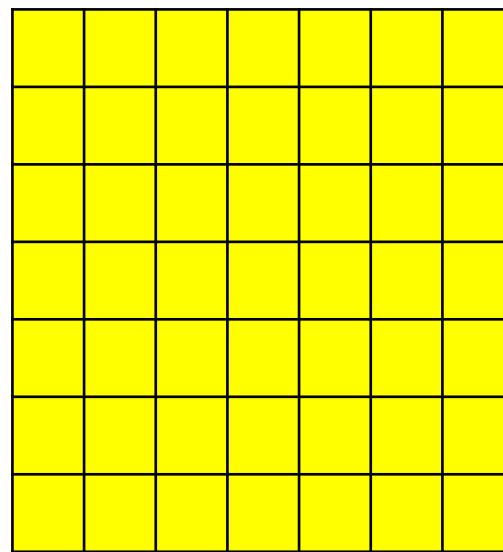
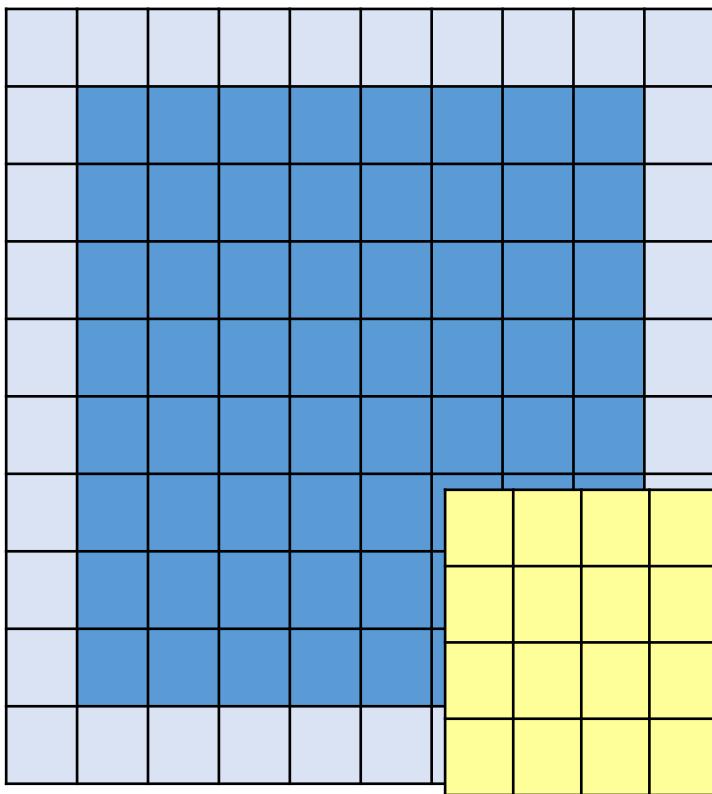
Padding = 1

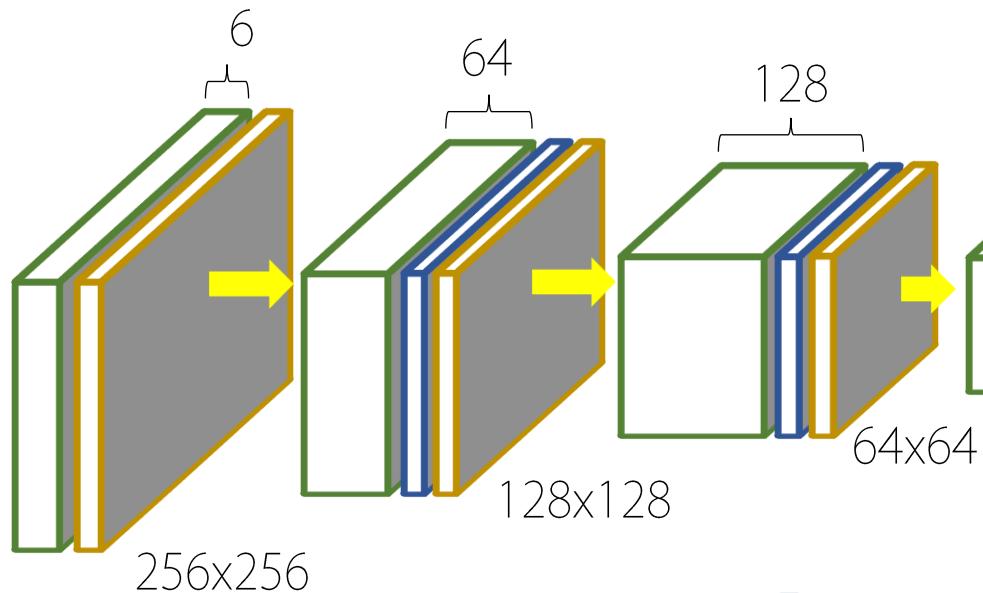












Conv Filter Size = 4x4
 Conv Stride = 2
 Conv Padding = 1

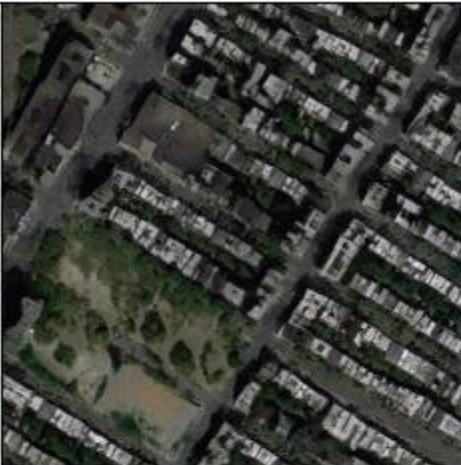
Batch Norm. Leaky ReLU Sigmoid

Conv Filter Size = 4x4
 Conv Stride = 1
 Conv Padding = 1

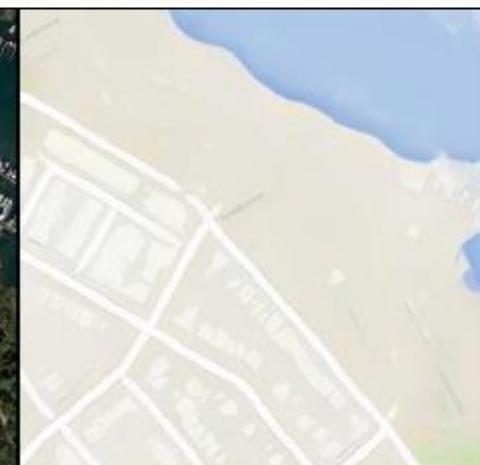
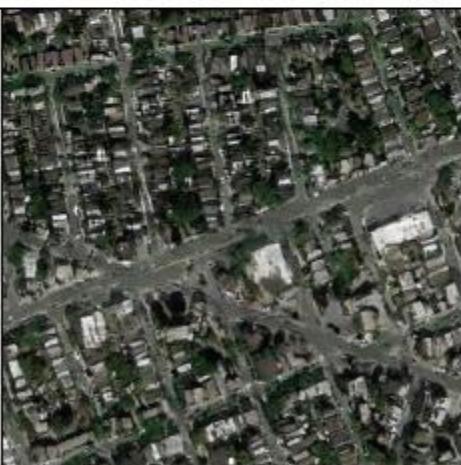
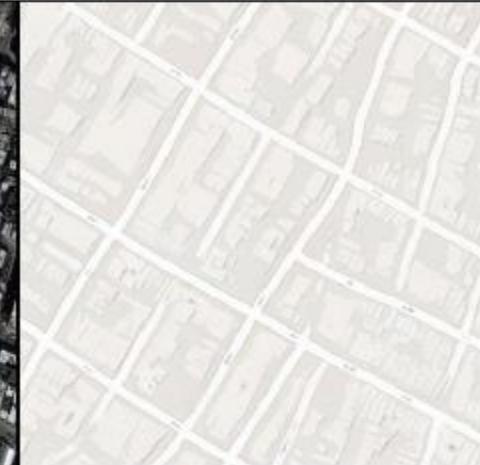
Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

Aerial photo to map



Map to aerial photo



input

output

input

output

Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017



Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017



Image to Image Translation

Image to Image Translation with Conditional Adversarial Networks,
CVPR 2017

Network Needs Paired Datasets

Image to Image Translation 2

Unpaired Image to Image Translation using Cycle-consistent Adversarial Networks,
a.k.a CycleGAN, ICCV 2017

Monet \curvearrowright Photos



Monet \rightarrow photo

Zebras \curvearrowright Horses



zebra \rightarrow horse

Summer \curvearrowright Winter



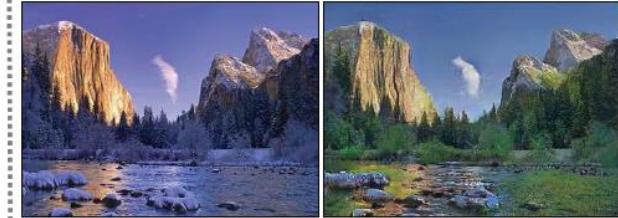
summer \rightarrow winter



photo \rightarrow Monet



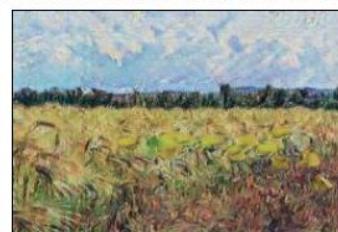
horse \rightarrow zebra



winter \rightarrow summer



Monet



Van Gogh



Cezanne



Ukiyo-e

Image to Image Translation 2

Unpaired Image to Image Translation using Cycle-consistent Adversarial Networks, ICCV 2017

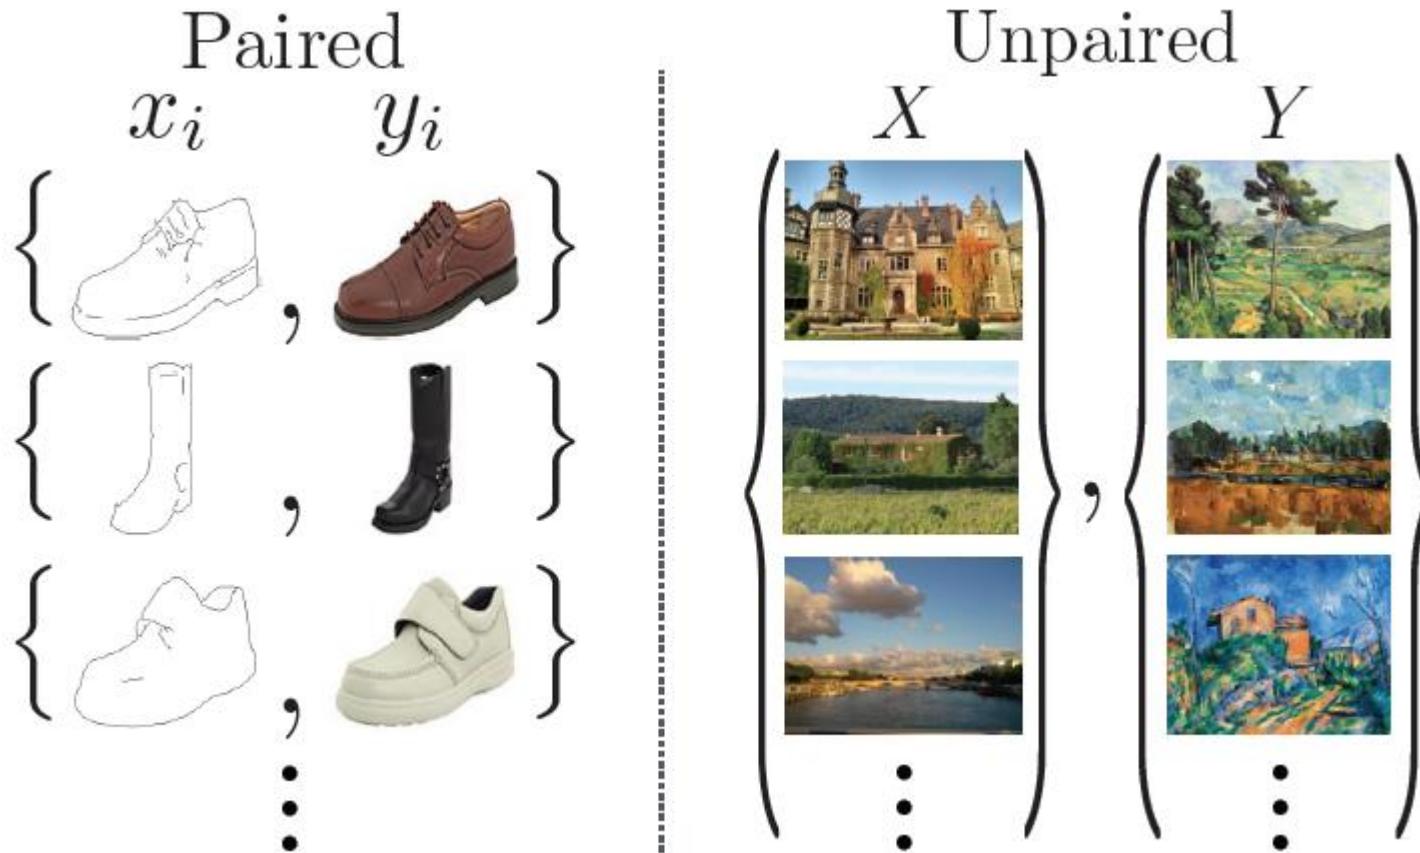
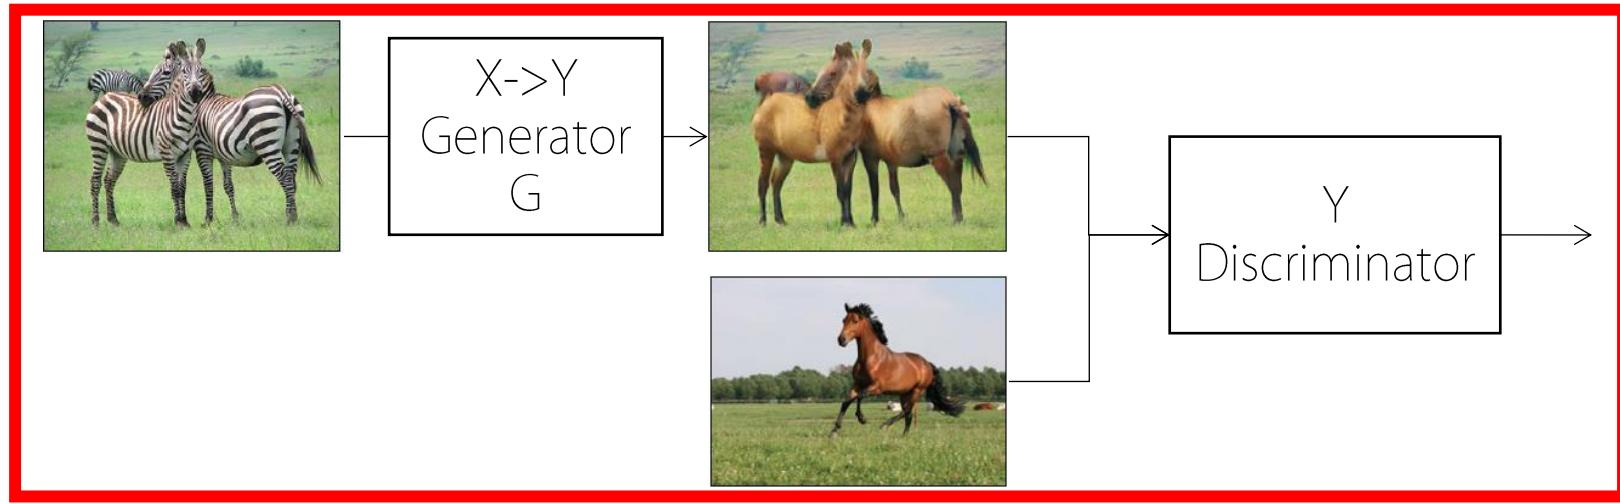


Image to Image Translation 2

Unpaired Image to Image Translation using Cycle-consistent Adversarial Networks, ICCV 2017



$Y \rightarrow X$



Adversarial Loss $X \rightarrow Y$

$$L_{GAN_Y} = E_y(\log D_Y(y)) + E_x(\log(1 - D_Y(G(x))))$$

$$L_{LSGAN_Y} = E_y((D_Y(y) - 1)^2) + E_x((D_Y(G(x)))^2)$$



Image to Image Translation 2

Unpaired Image to Image Translation using Cycle-consistent Adversarial Networks, ICCV 2017

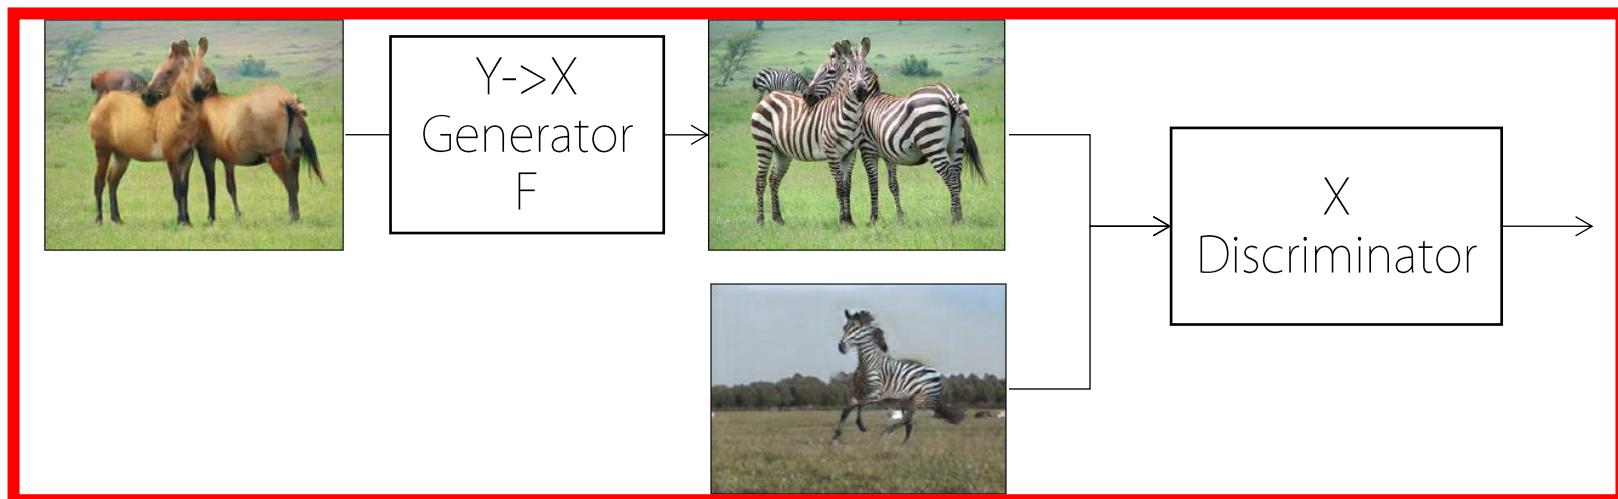
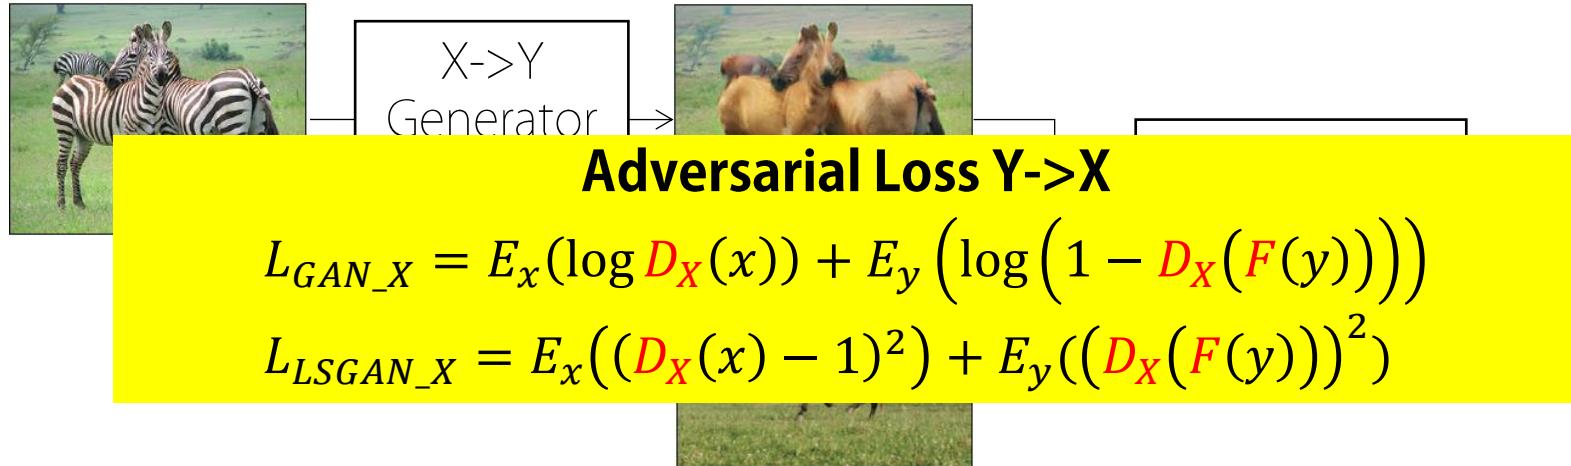


Image to Image Translation 2

Unpaired Image to Image Translation using Cycle-consistent Adversarial Networks, ICCV 2017

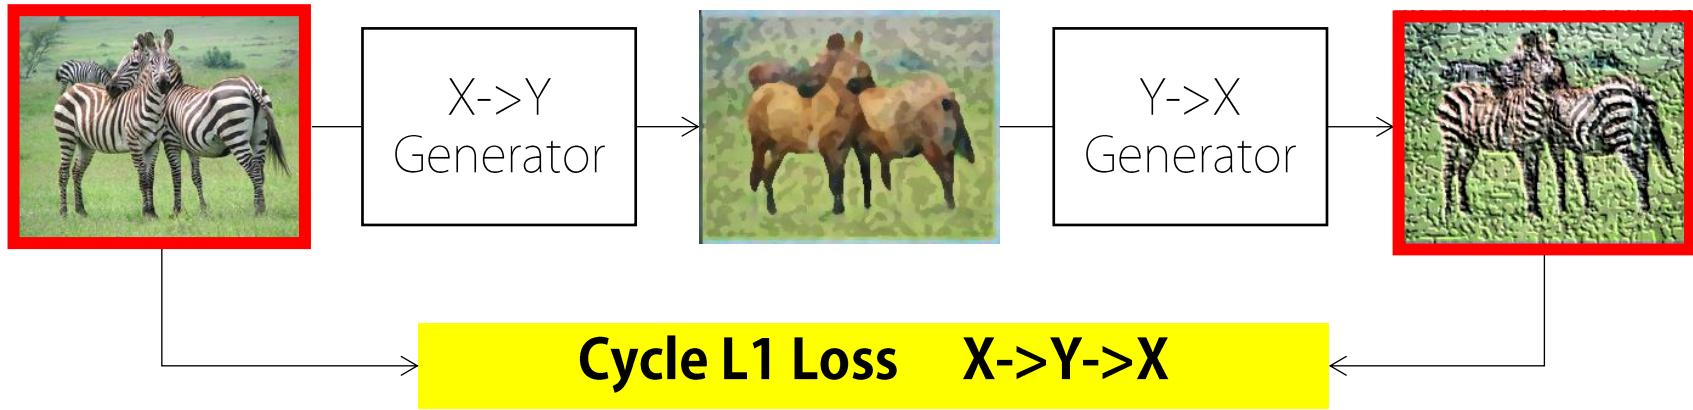


Image to Image Translation 2

Unpaired Image to Image Translation using Cycle-consistent Adversarial Networks, ICCV 2017

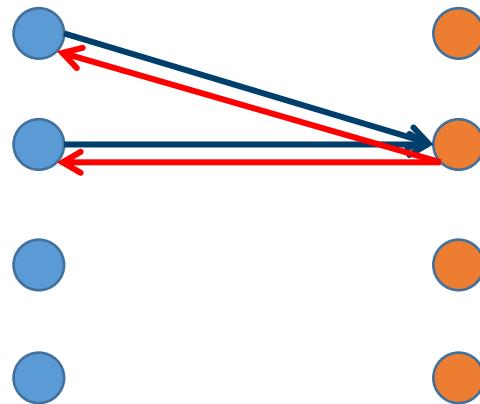
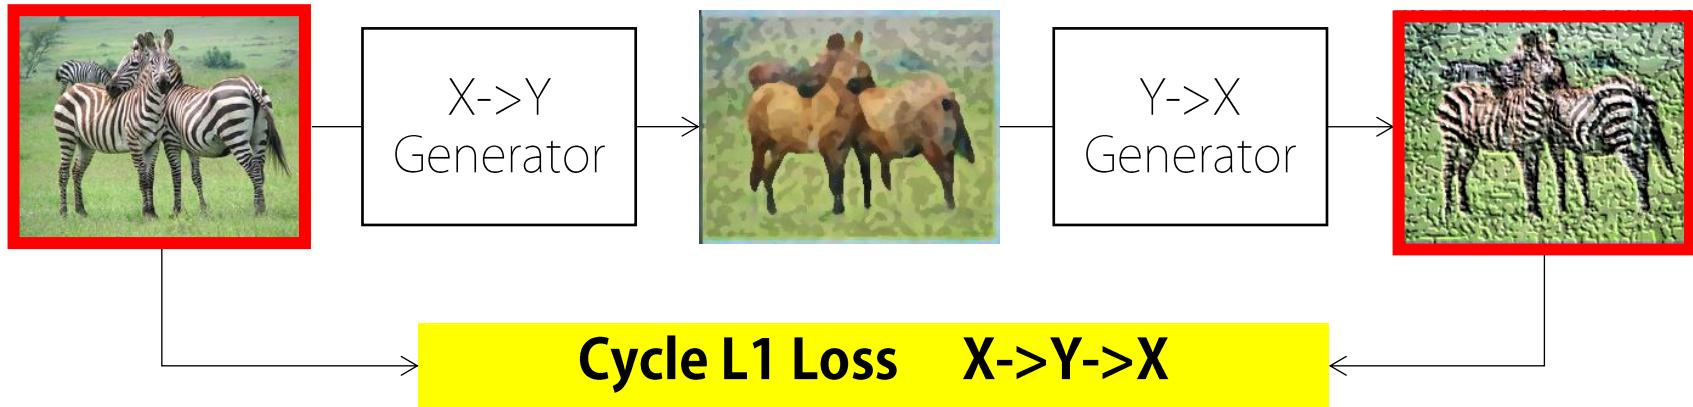


Image to Image Translation 2

Unpaired Image to Image Translation using Cycle-consistent Adversarial Networks, ICCV 2017

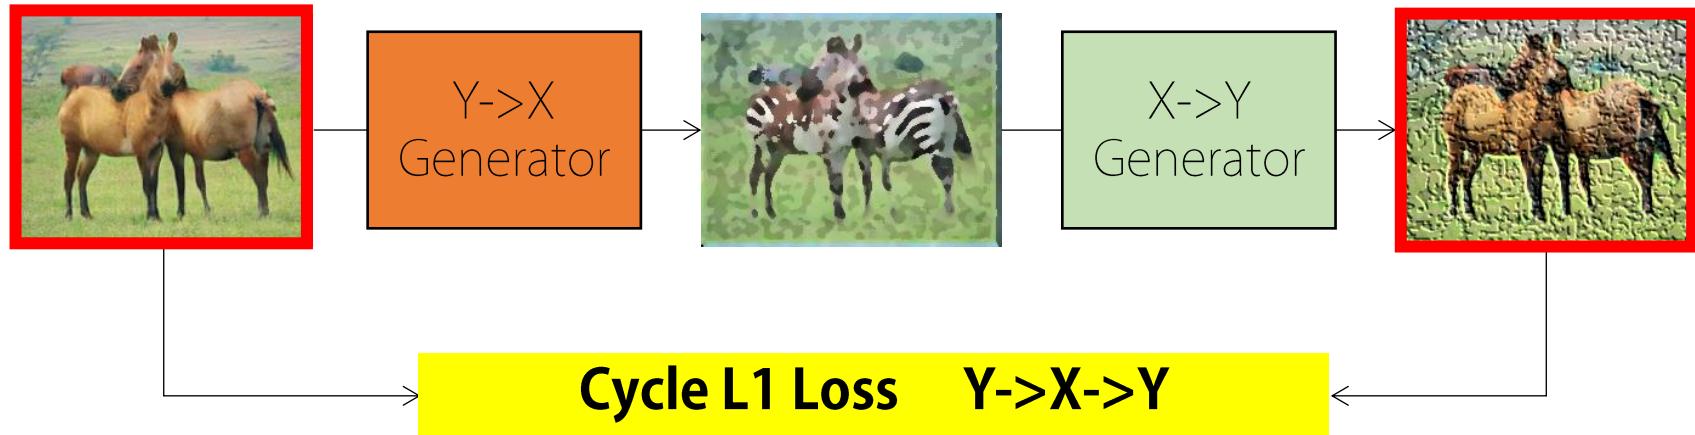
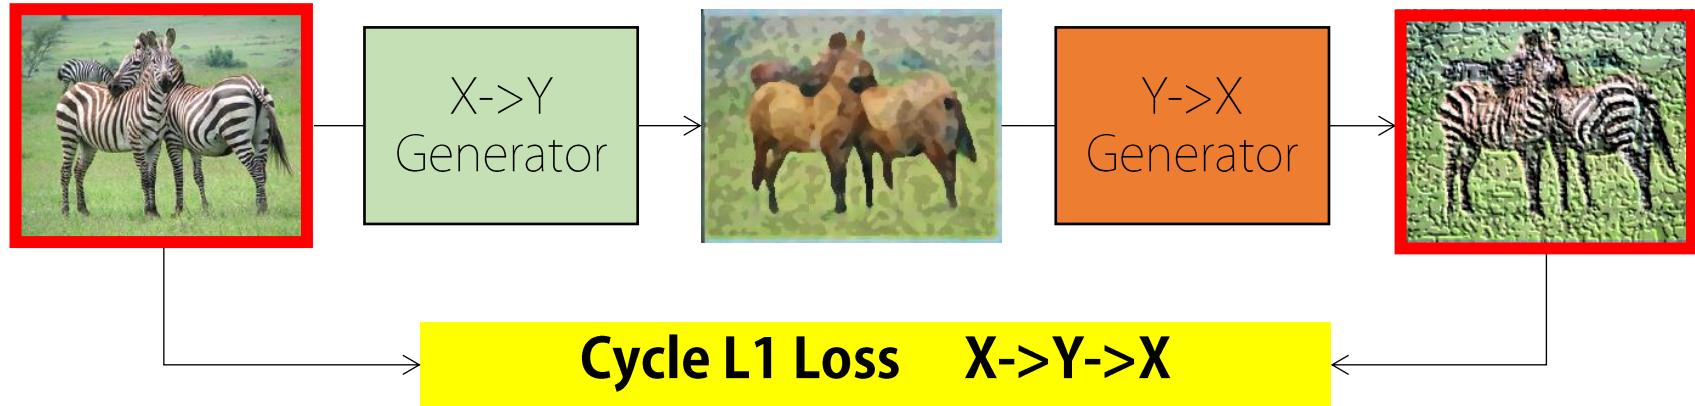
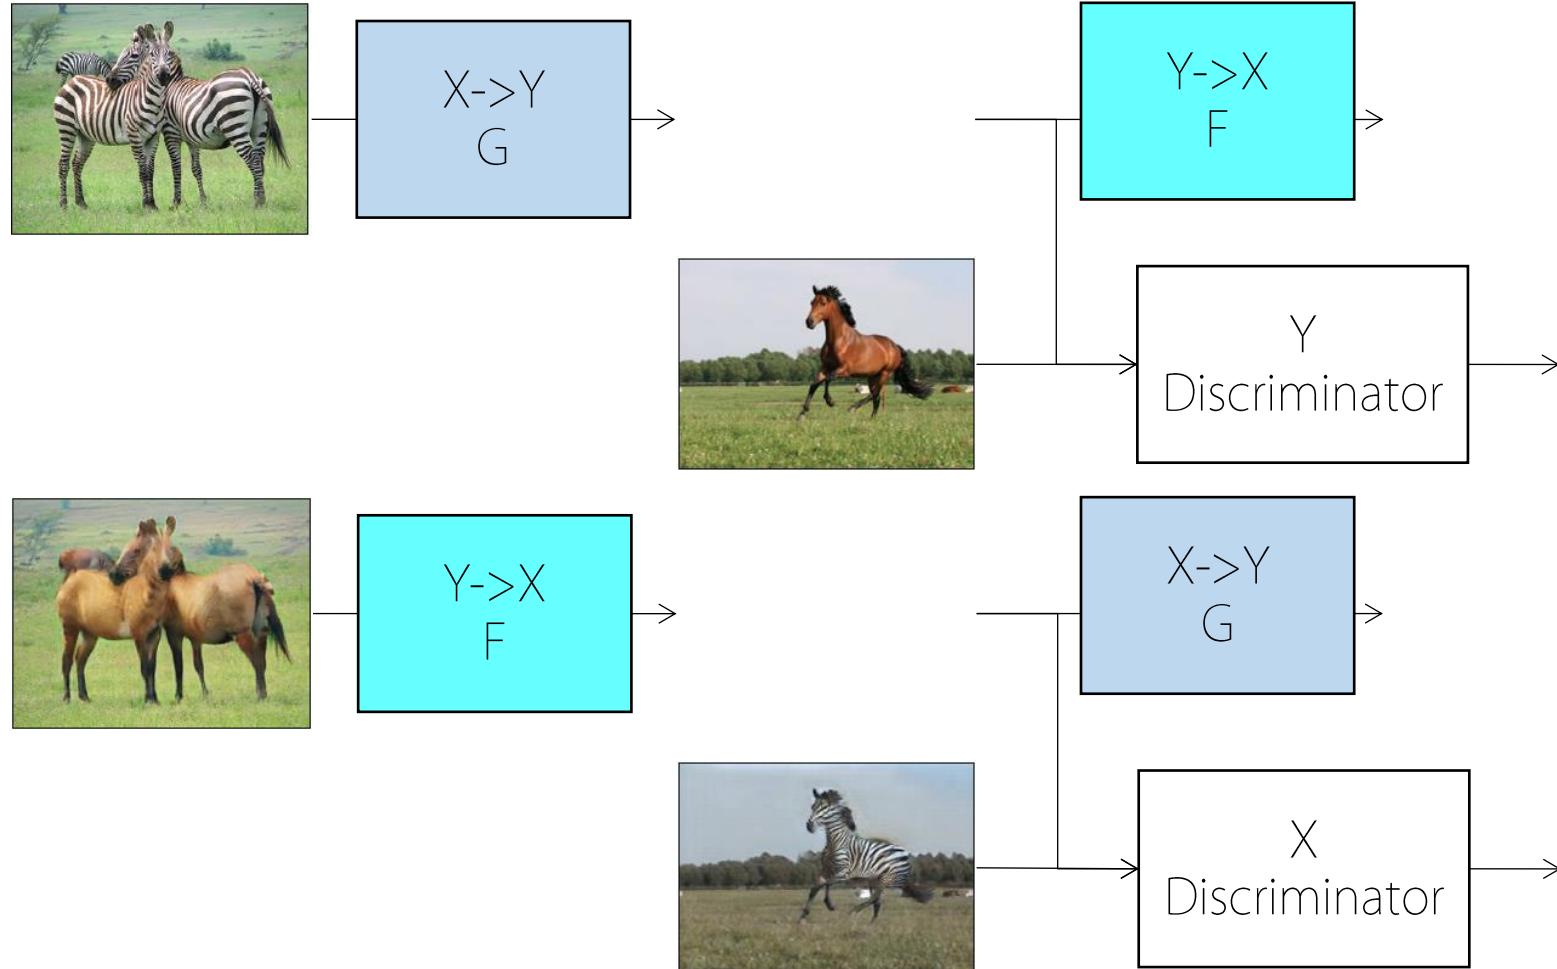
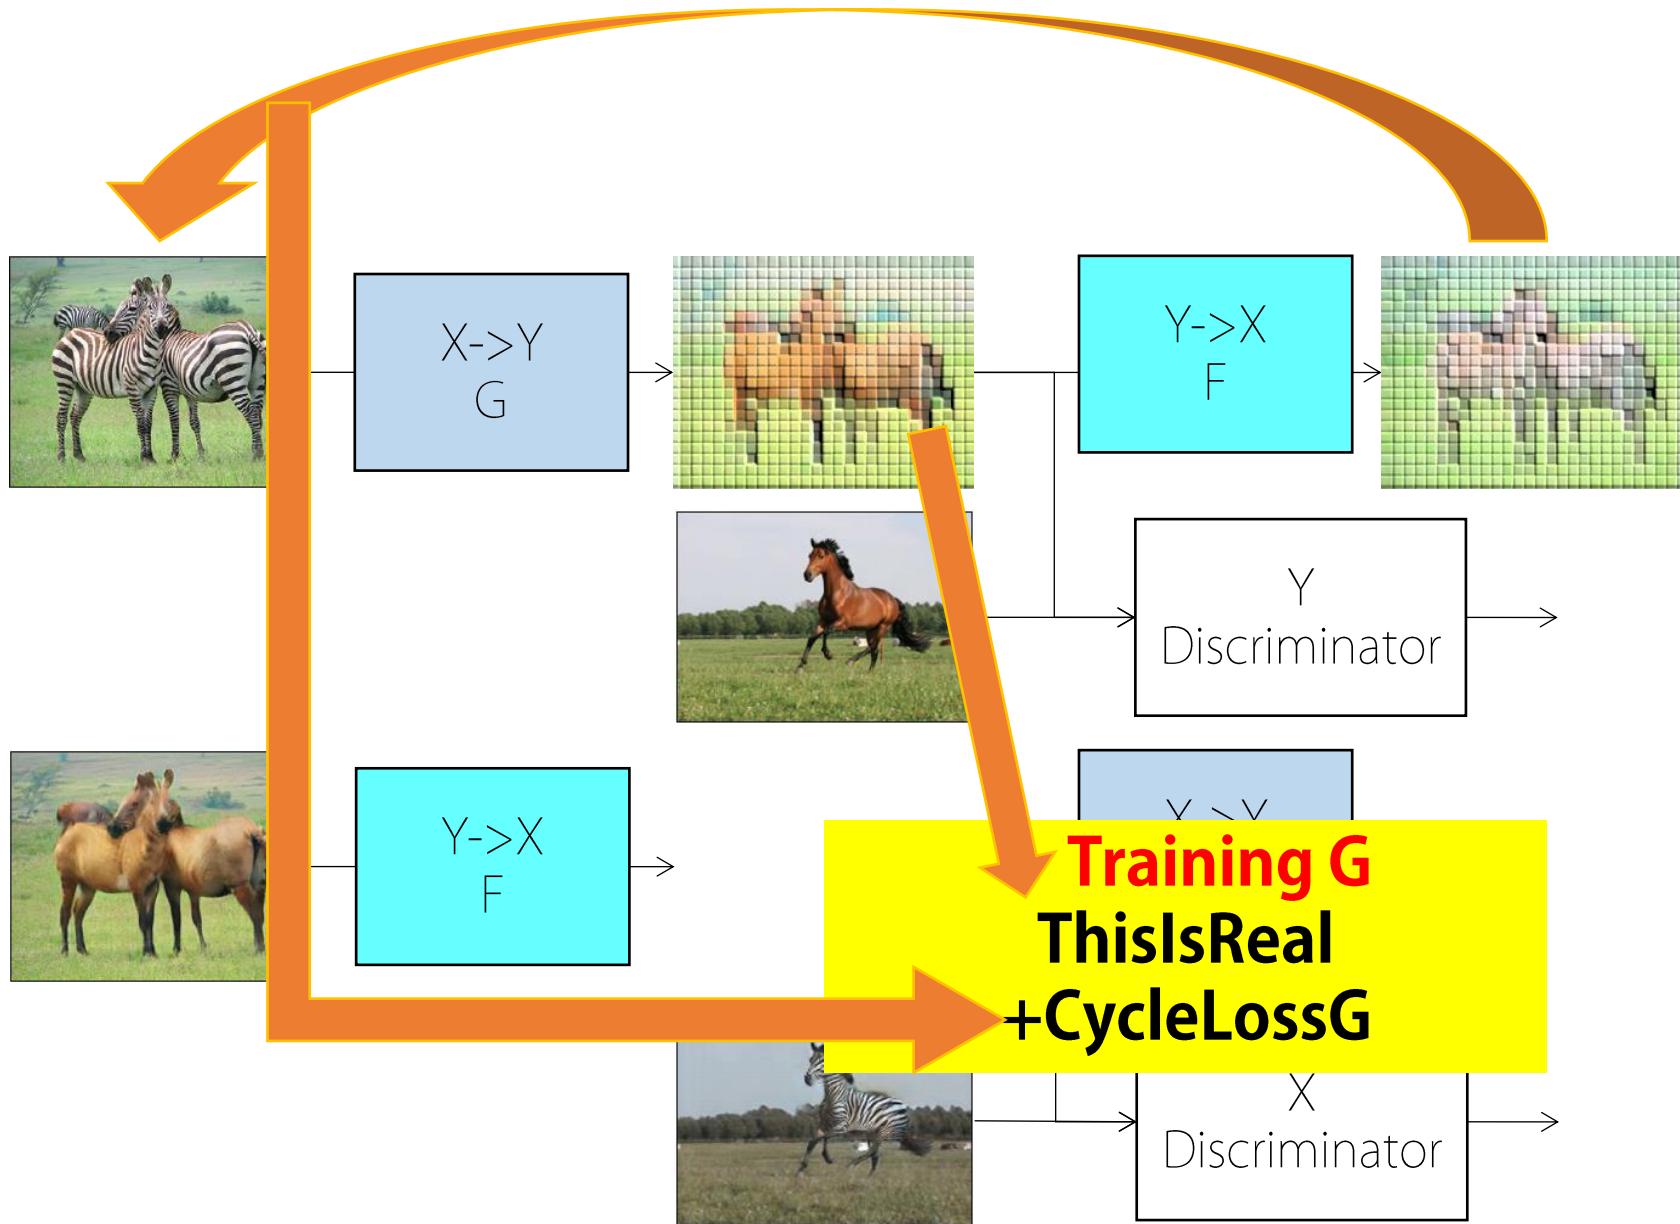


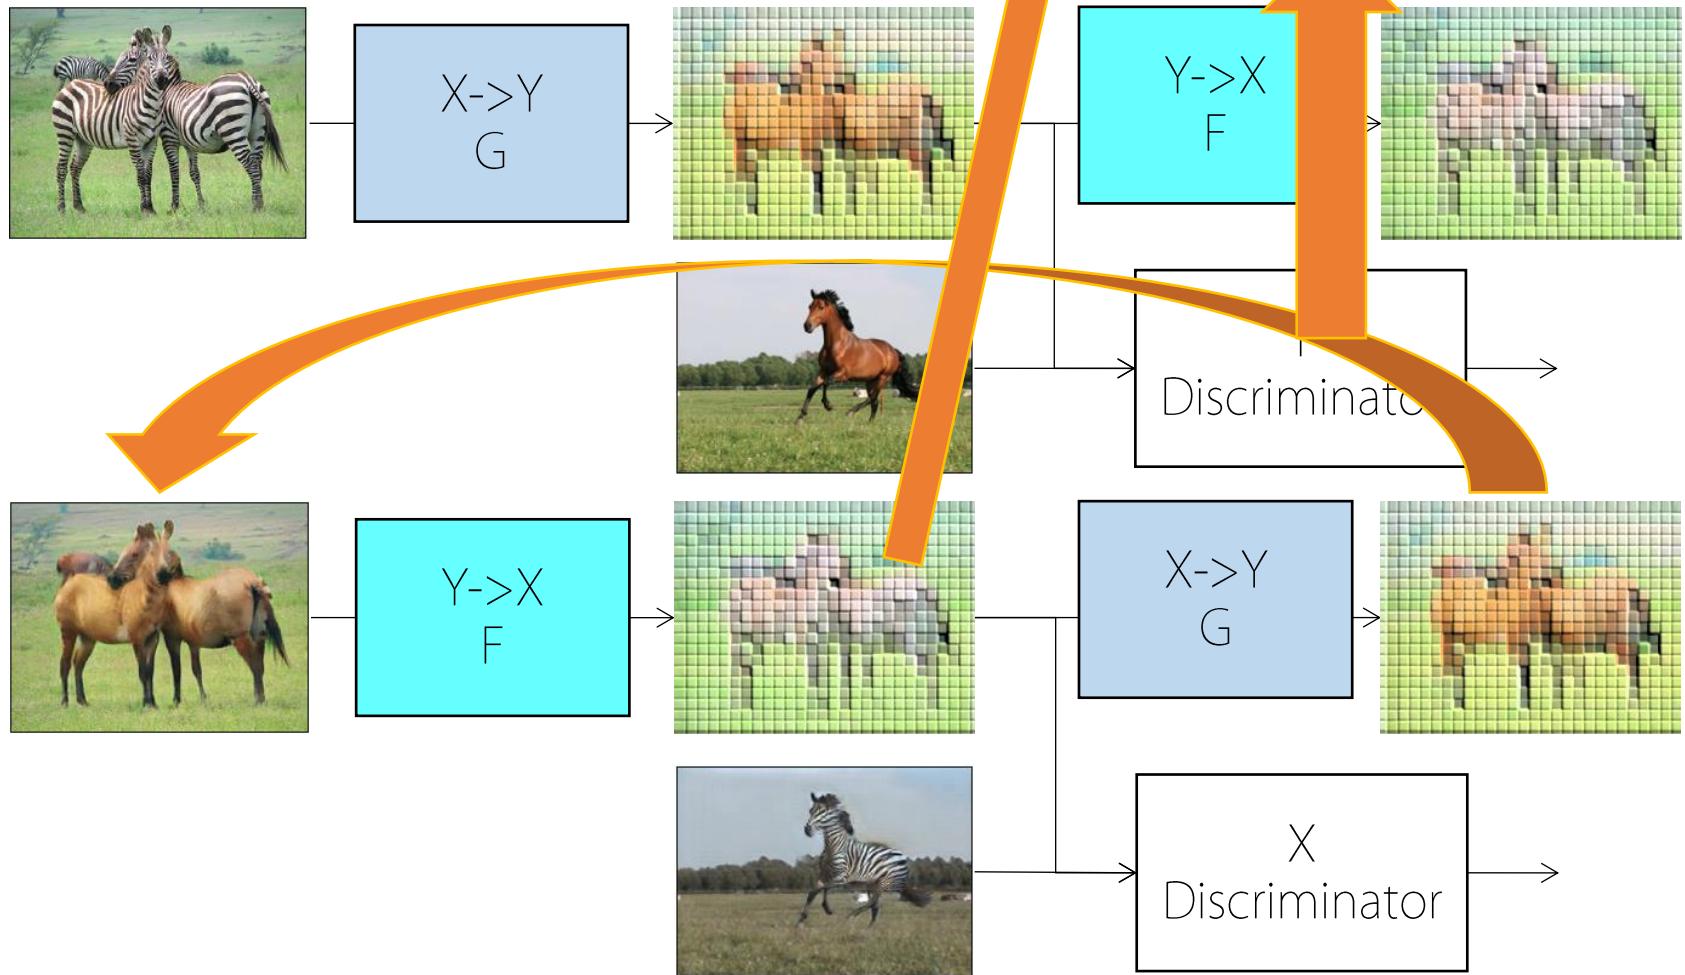
Image to Image Translation 2

Unpaired Image to Image Translation using Cycle-consistent Adversarial Networks, ICCV 2017



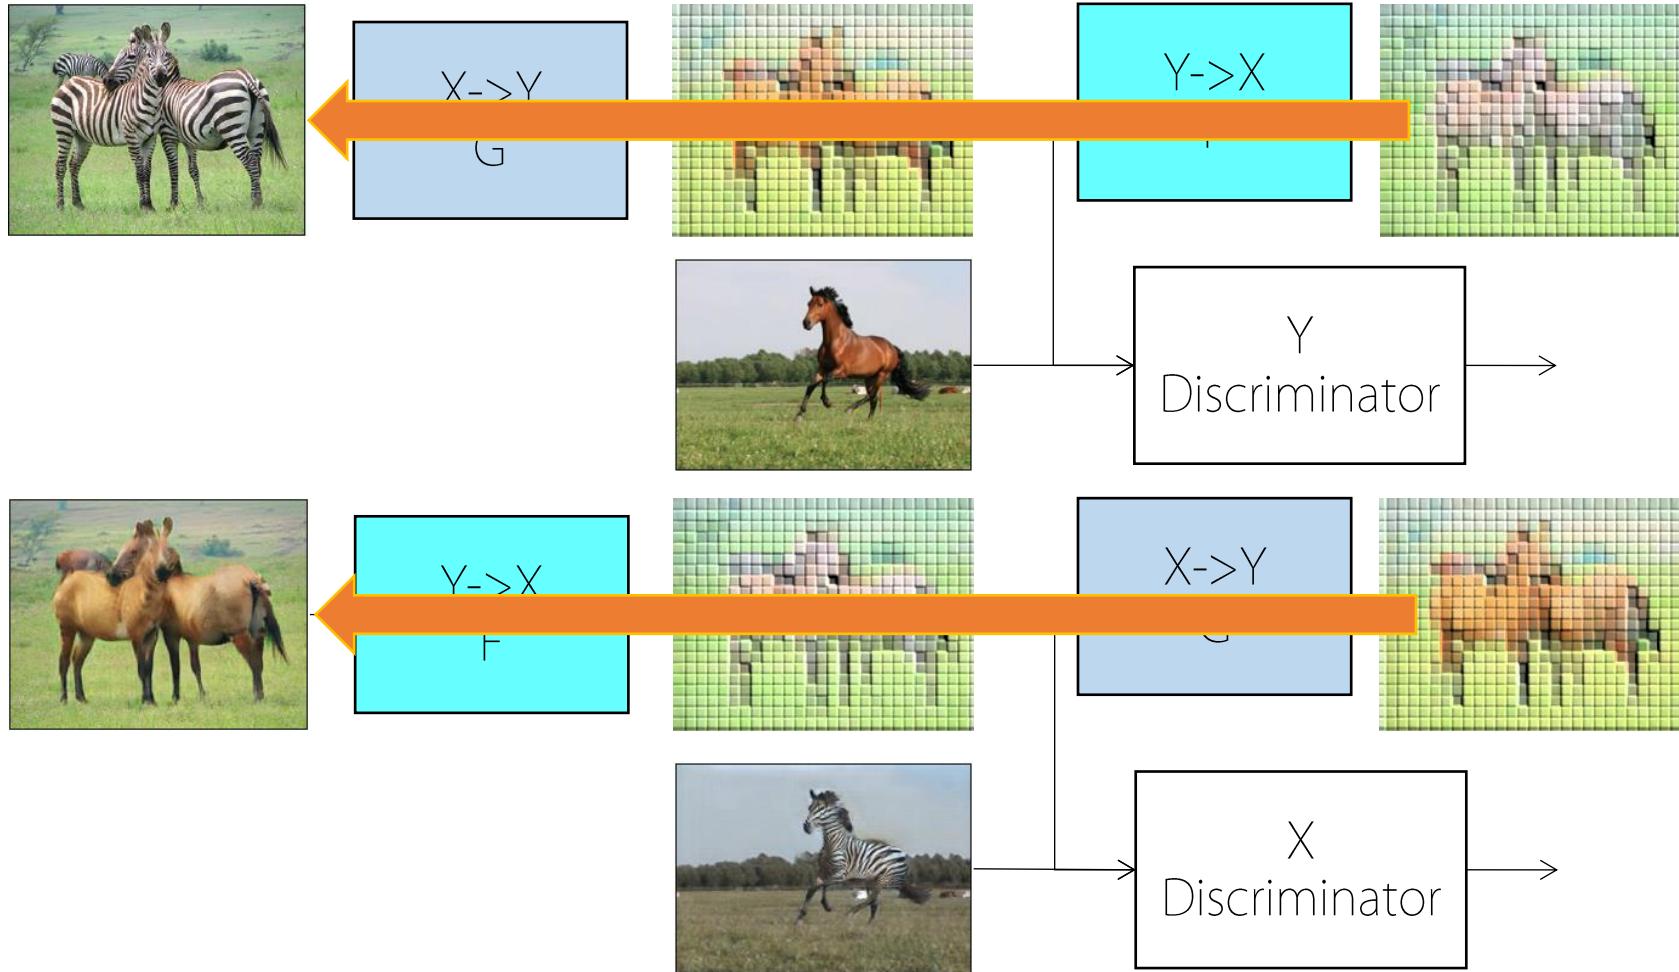


**Training F
ThisIsReal
+CycleLossF**



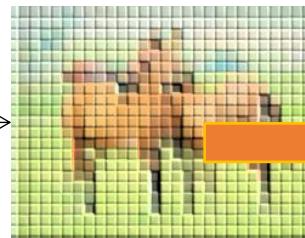
Train G+F

LossG+LossF





$X \rightarrow Y$
G

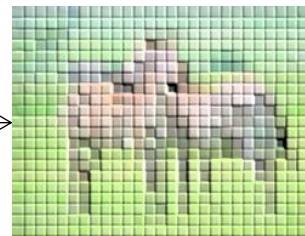


Train D_Y
This is Real+
This is Fake

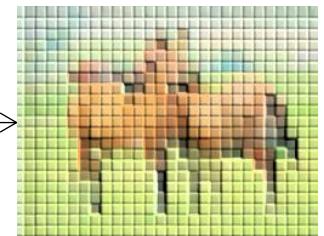
Y
Discriminator



$Y \rightarrow X$
F



$X \rightarrow Y$
G



X
Discriminator



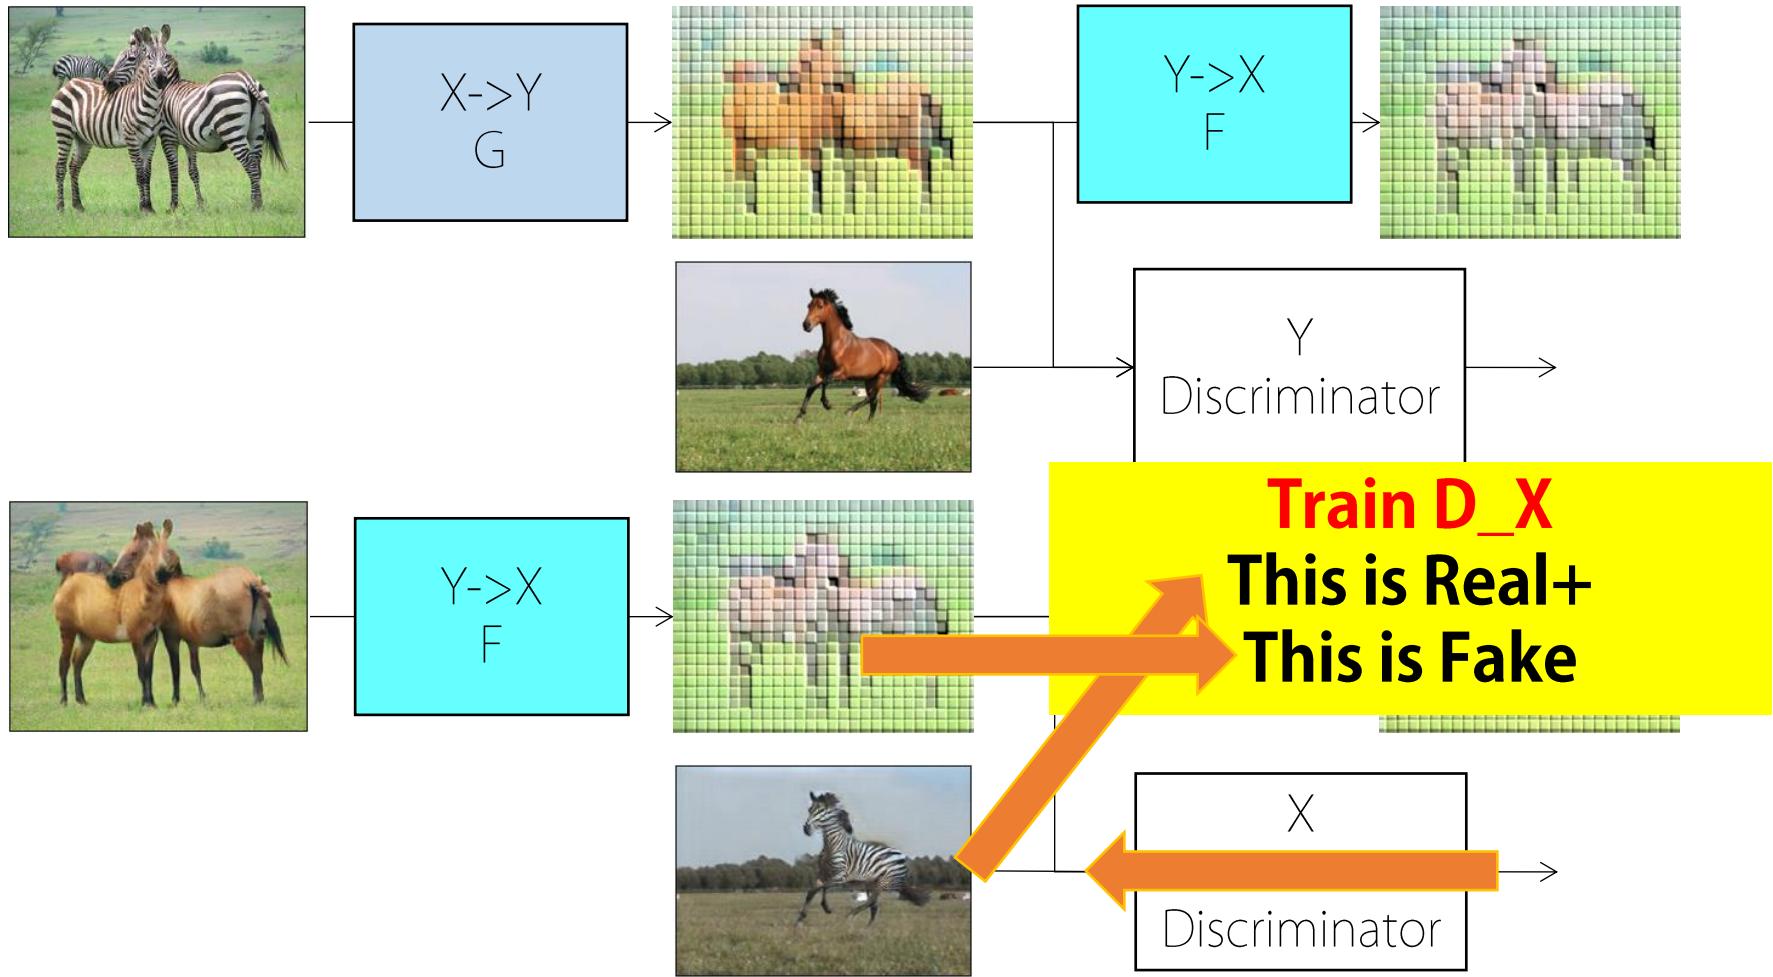


Image to Image Translation 2

Unpaired Image to Image Translation using Cycle-consistent Adversarial Networks, ICCV 2017

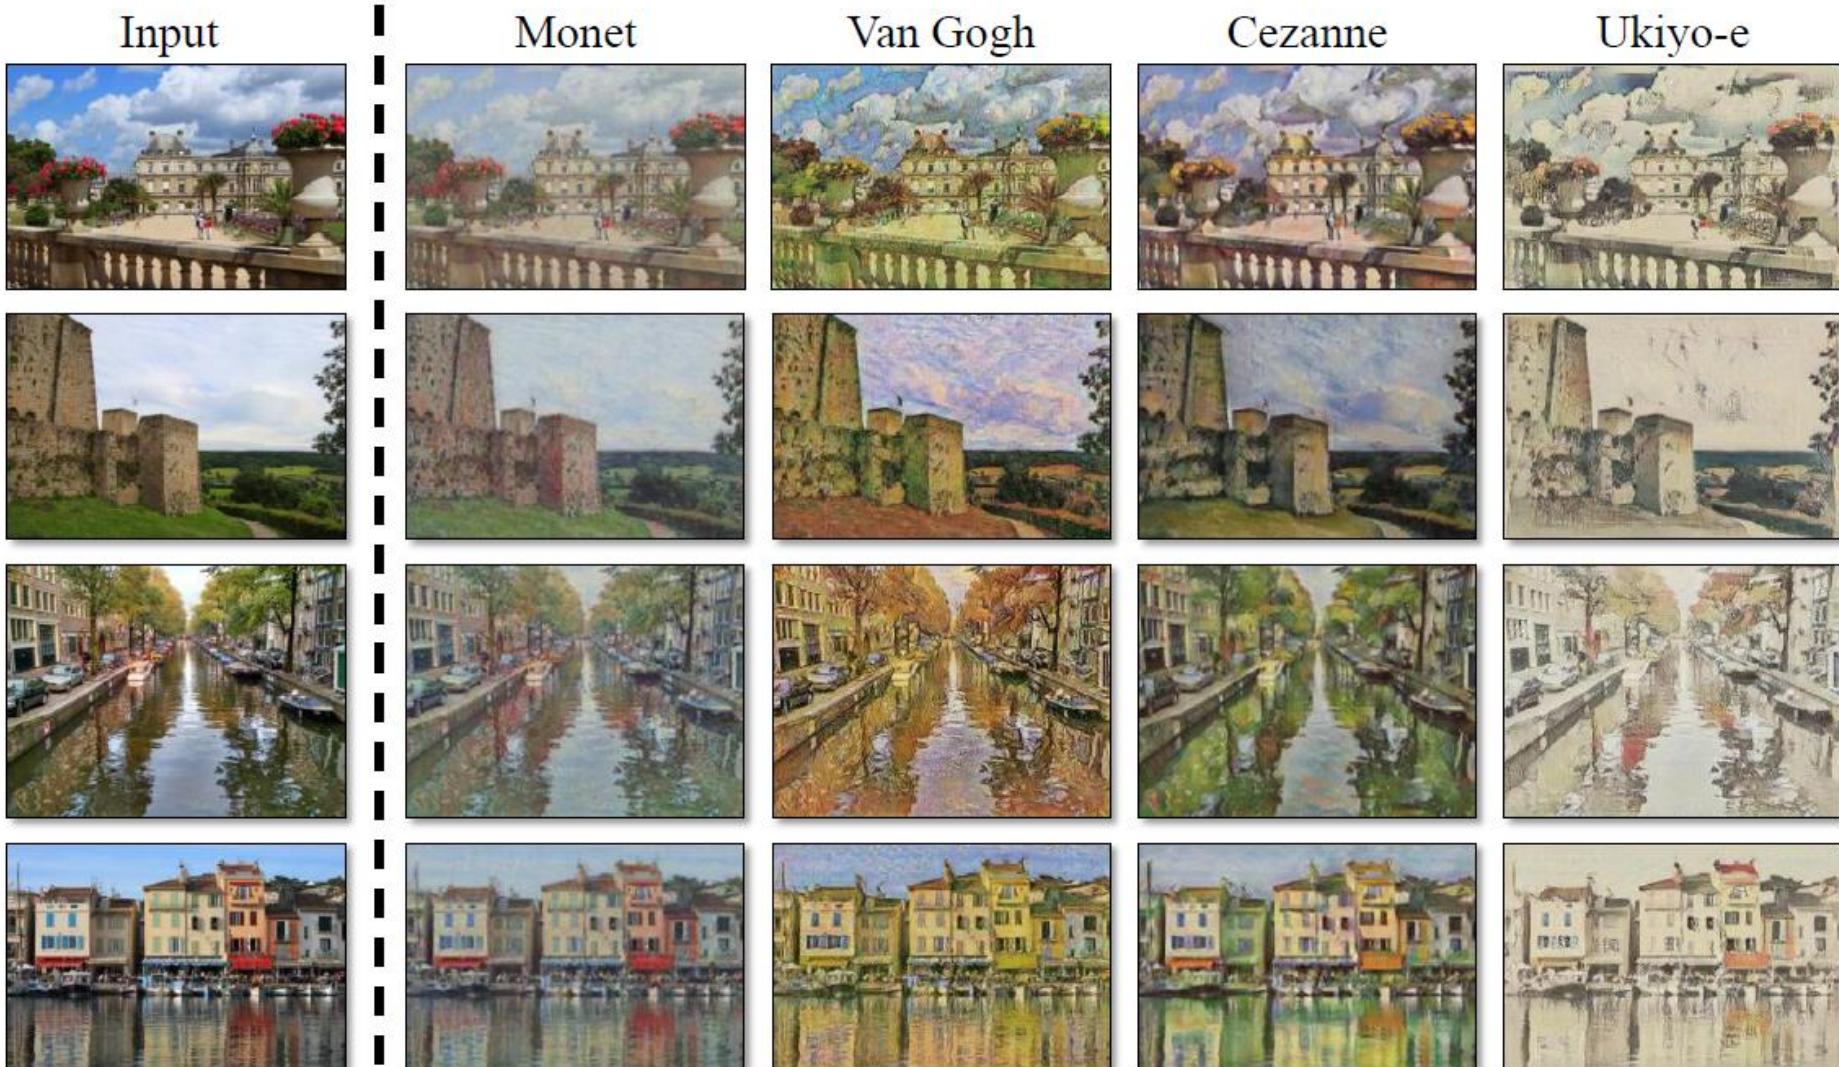


Image to Image Translation 2

Unpaired Image to Image Translation using Cycle-consistent Adversarial Networks, ICCV 2017

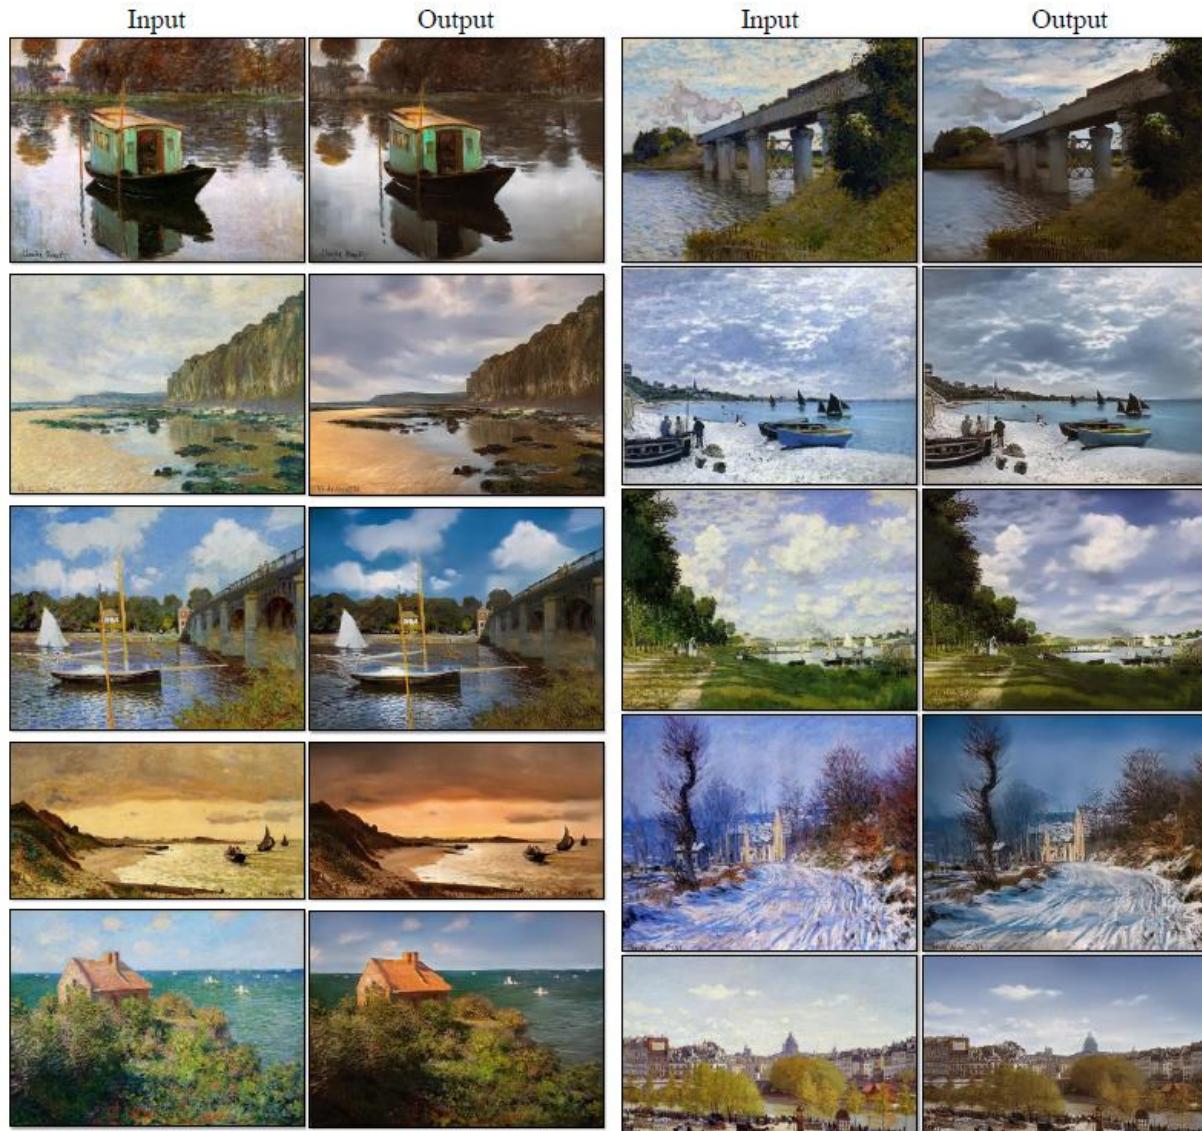
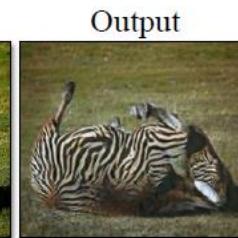
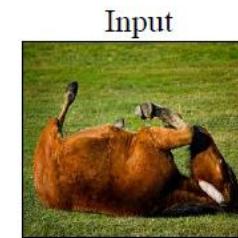
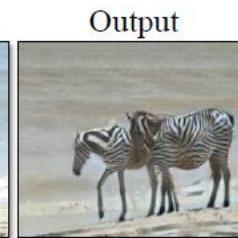
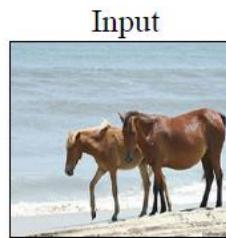
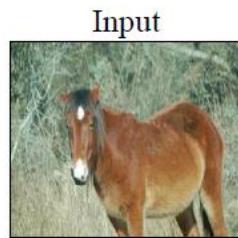


Image to Image Translation 2

Unpaired Image to Image Translation using Cycle-consistent Adversarial Networks, ICCV 2017



horse → zebra



zebra → horse



winter Yosemite → summer Yosemite



summer Yosemite → winter Yosemite

Image to Image Translation 2

Unpaired Image to Image Translation using Cycle-consistent Adversarial Networks, ICCV 2017



apple → orange



orange → apple

Input



Output



Input



Output



Input



Output



Image to Image Translation 3

Learning to Discover Cross-domain Relations with Generative Adversarial Networks,
a.k.a DiscoGAN, ICML 2017



(b) Handbag images (input) & **Generated** shoe images (output)



(c) Shoe images (input) & **Generated** handbag images (output)

Image to Image Translation 3

Learning to Discover Cross-domain Relations with Generative Adversarial Networks,
a.k.a DiscoGAN, ICML 2017



Image to Image Translation 3

Learning to Discover Cross-domain Relations with Generative Adversarial Networks,
a.k.a DiscoGAN, ICML 2017

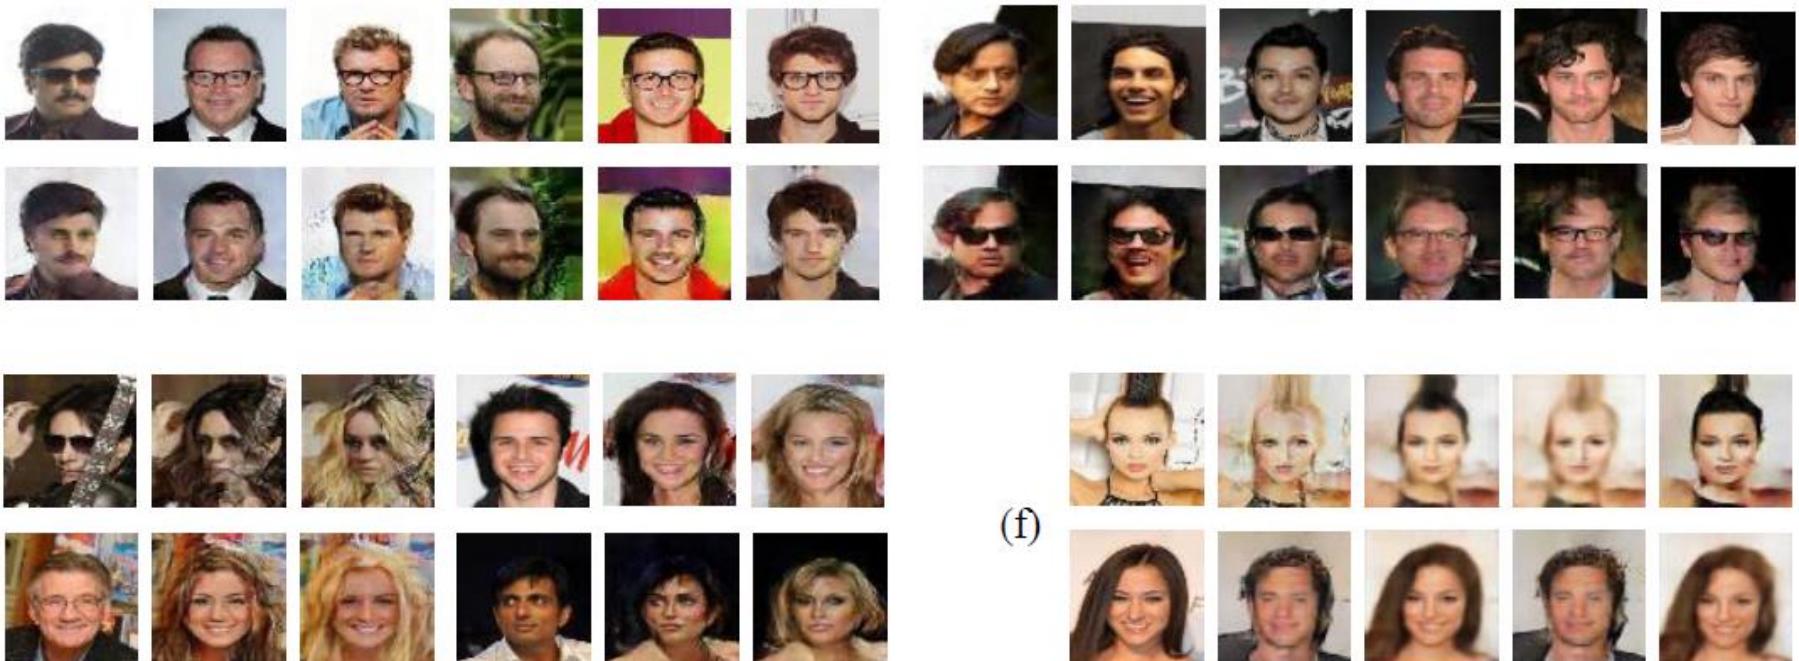


Image to Image Translation 3

Learning to Discover Cross-domain Relations with Generative Adversarial Networks,
a.k.a DiscoGAN, ICML 2017

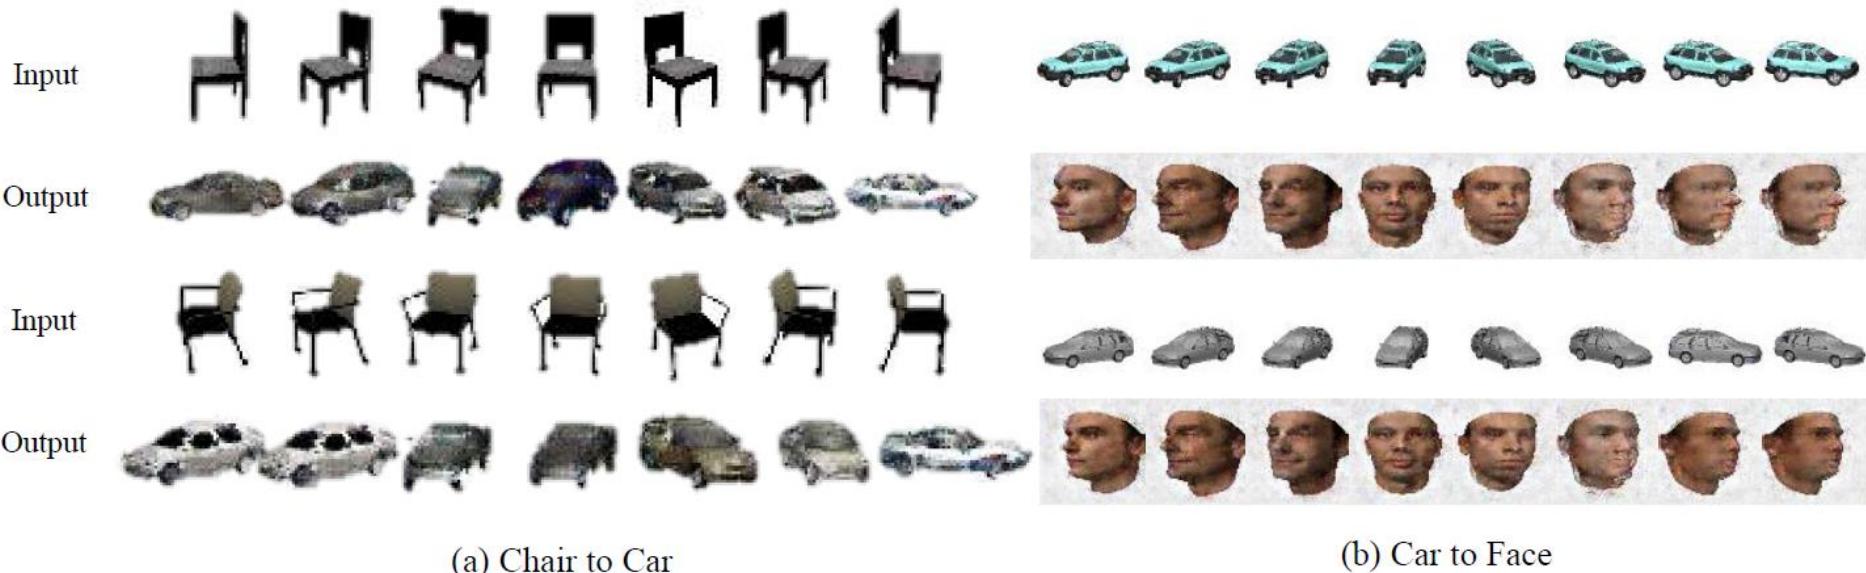


Image to Image Translation 3

Learning to Discover Cross-domain Relations with Generative Adversarial Networks, ICML 2017

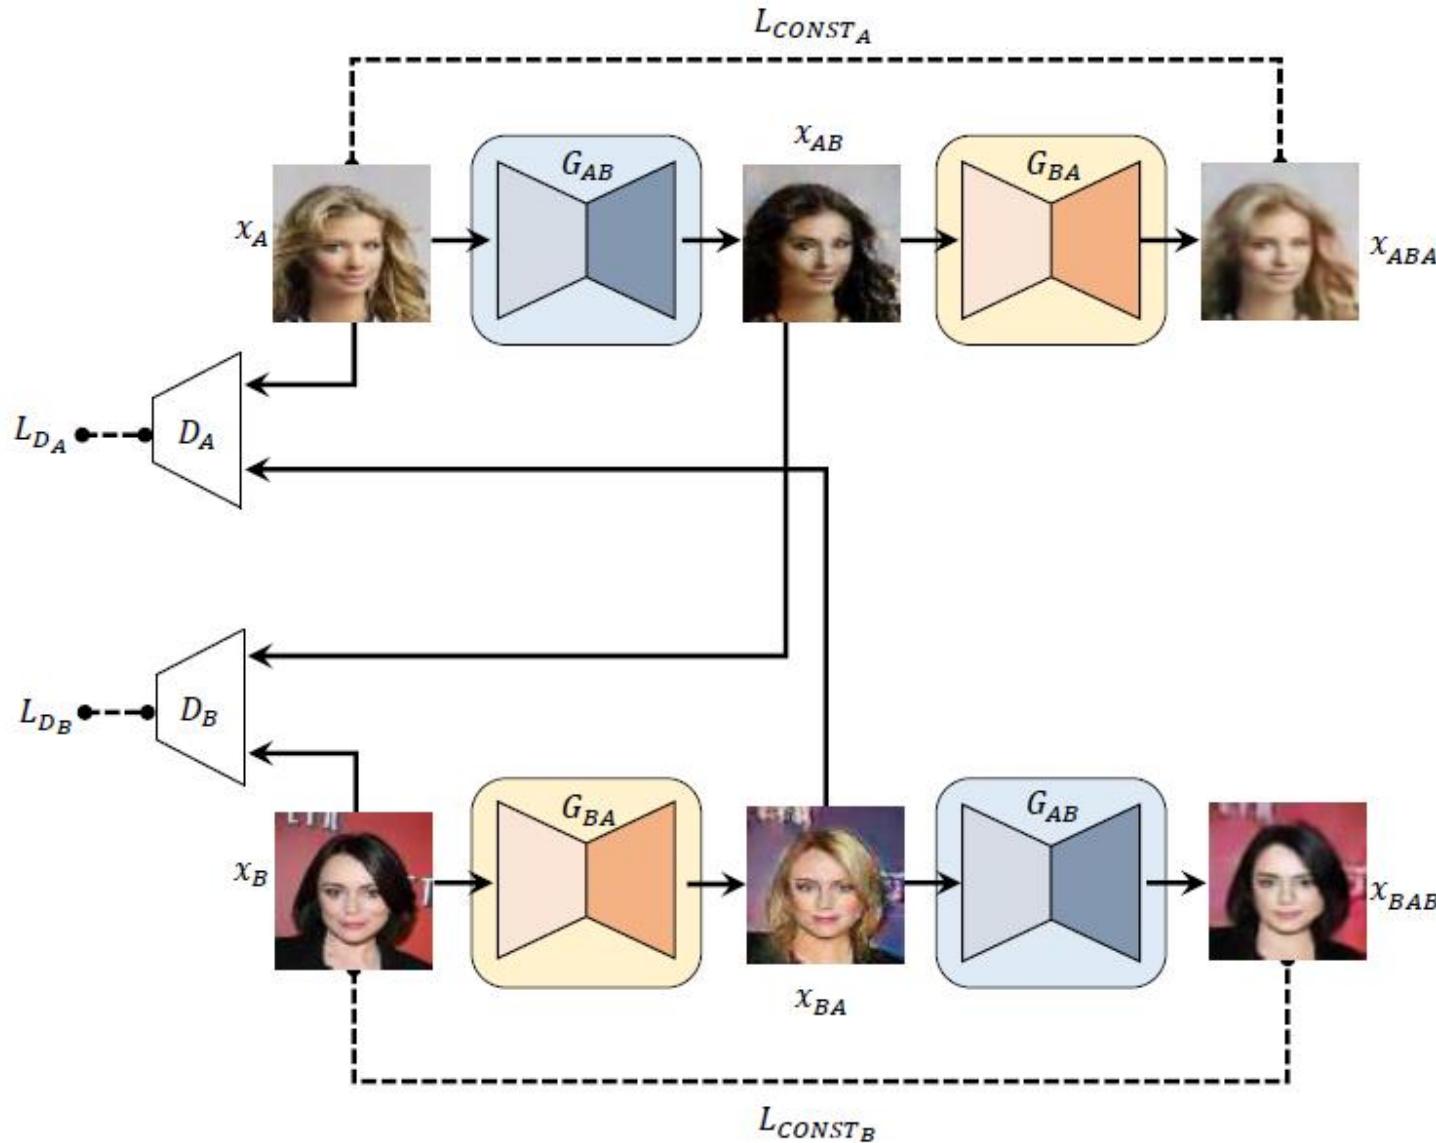


Image to Image Translation 3

Learning to Discover Cross-domain Relations with Generative Adversarial Networks, ICML 2017

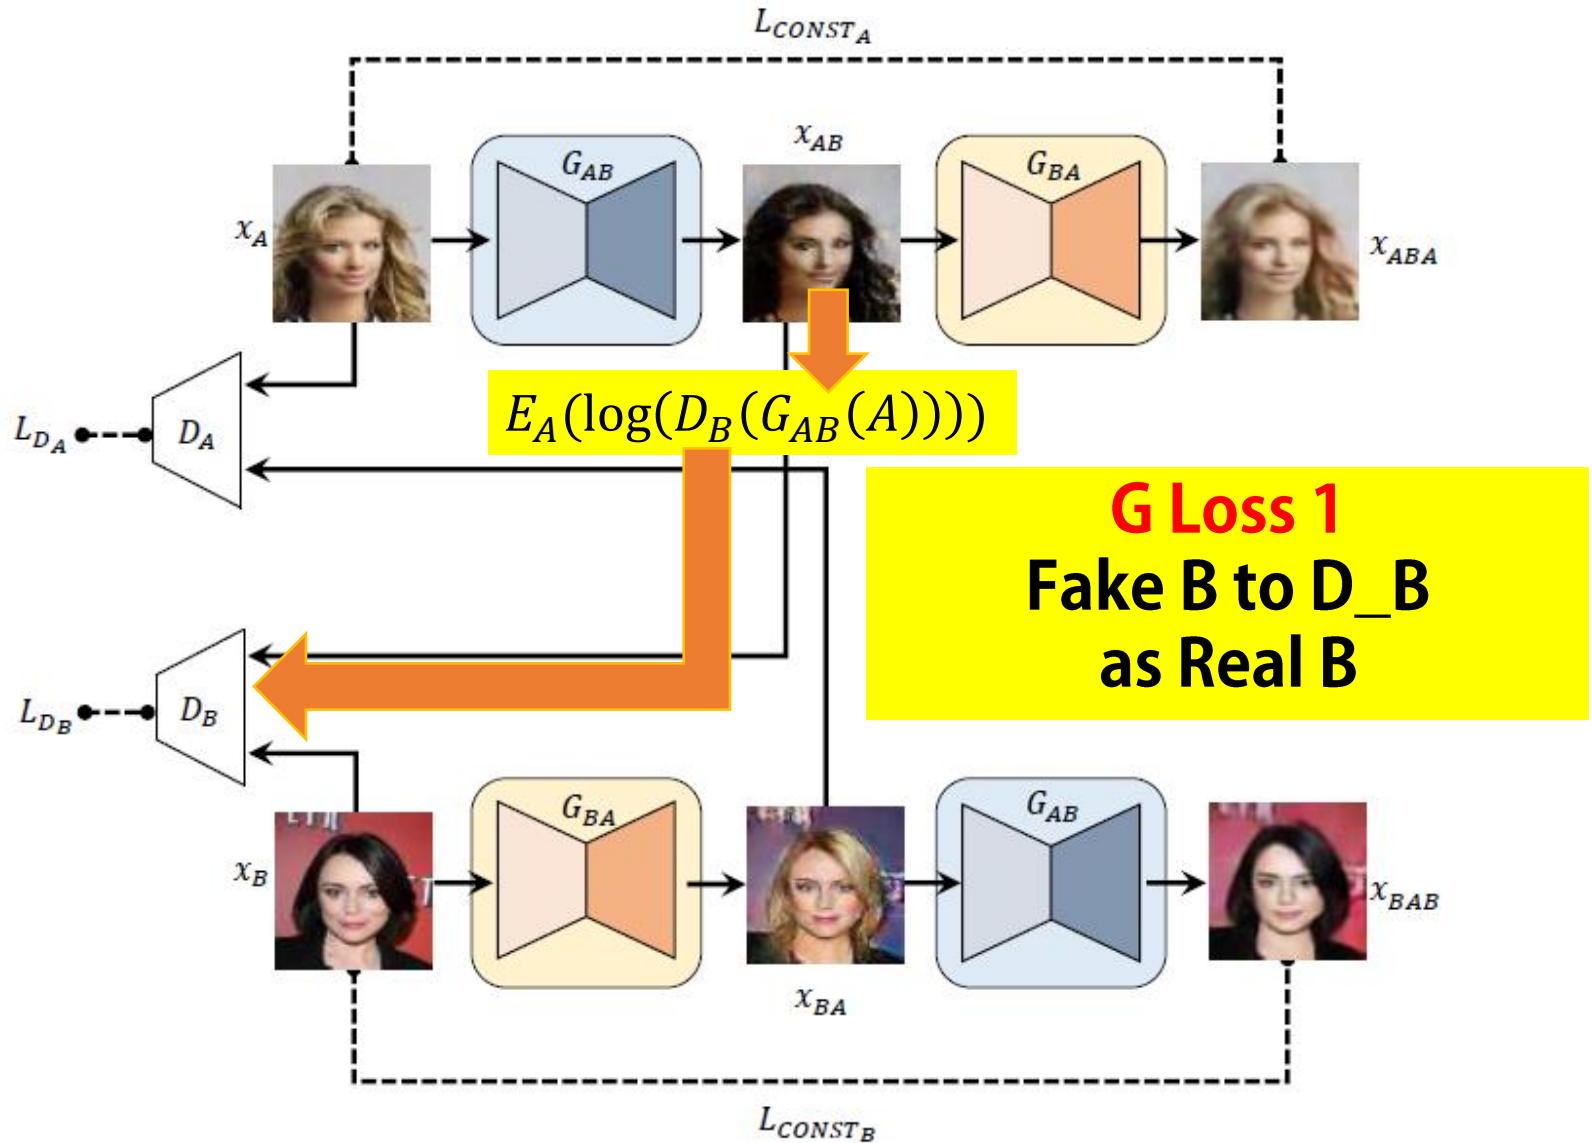


Image to Image Translation 3

Learning to Discover Cross-domain Relations with Generative Adversarial Networks, ICML 2017

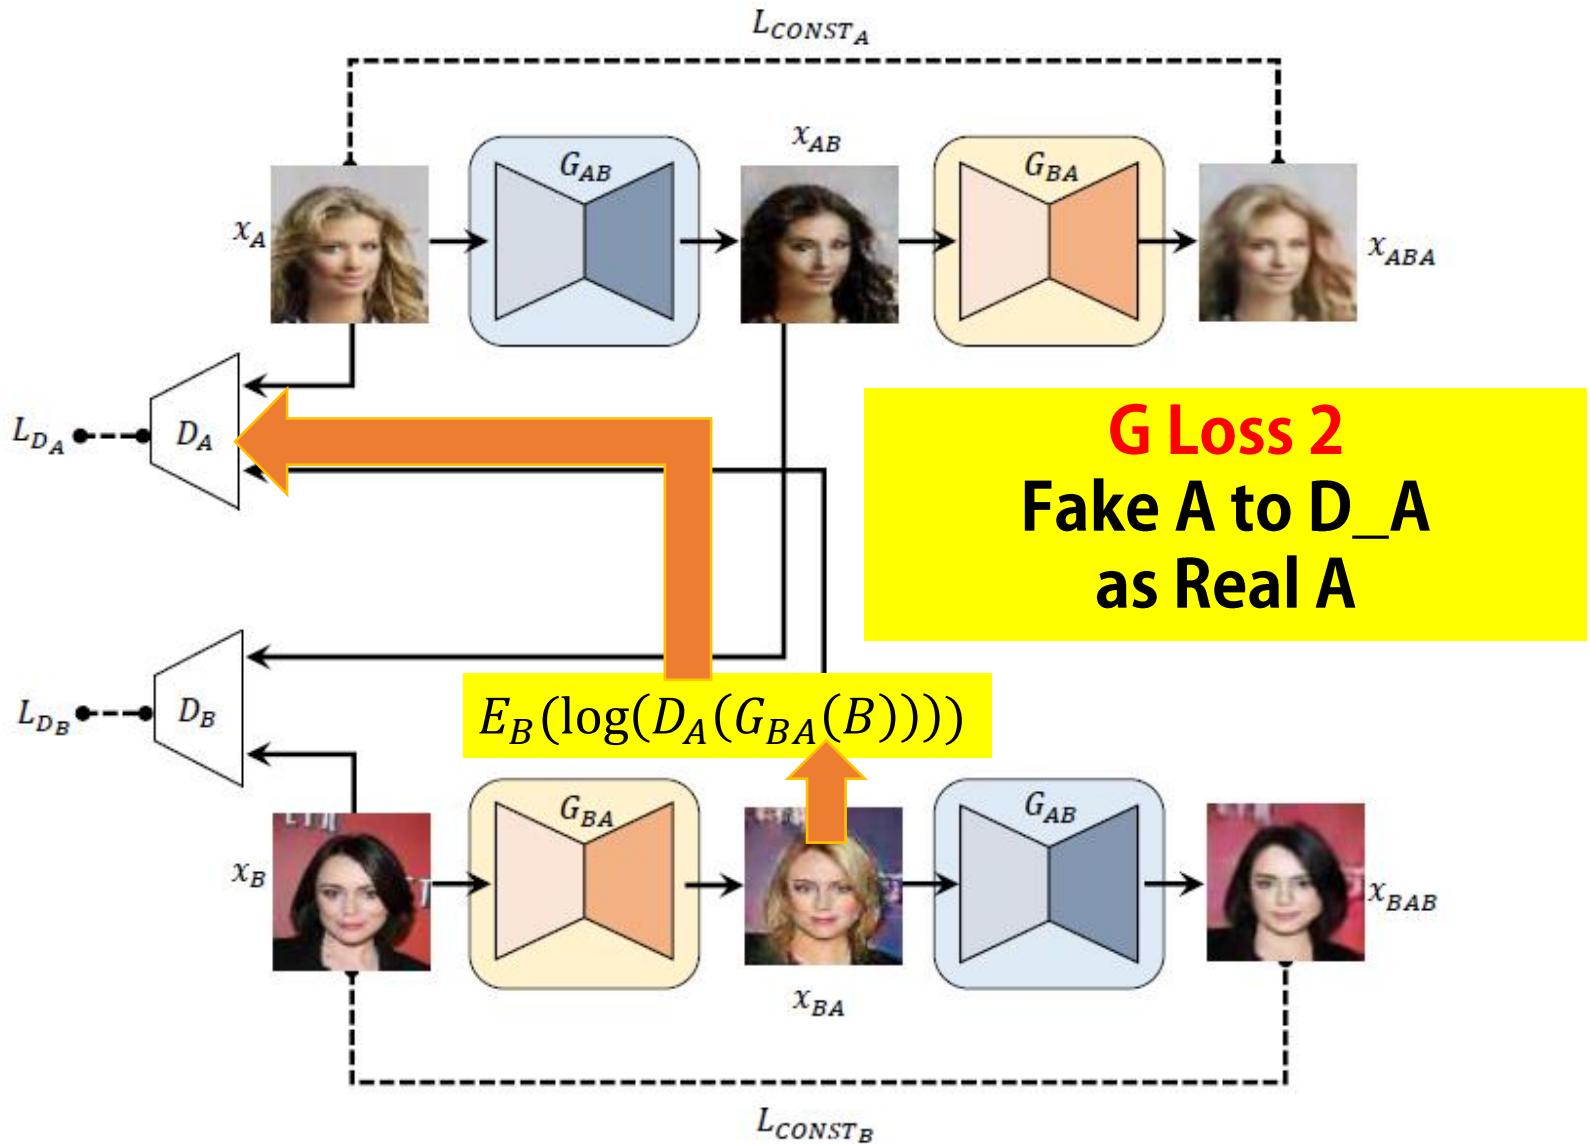


Image to Image Translation 3

Learning to Discover Cross-domain Relations with Generative Adversarial Networks, ICML 2017

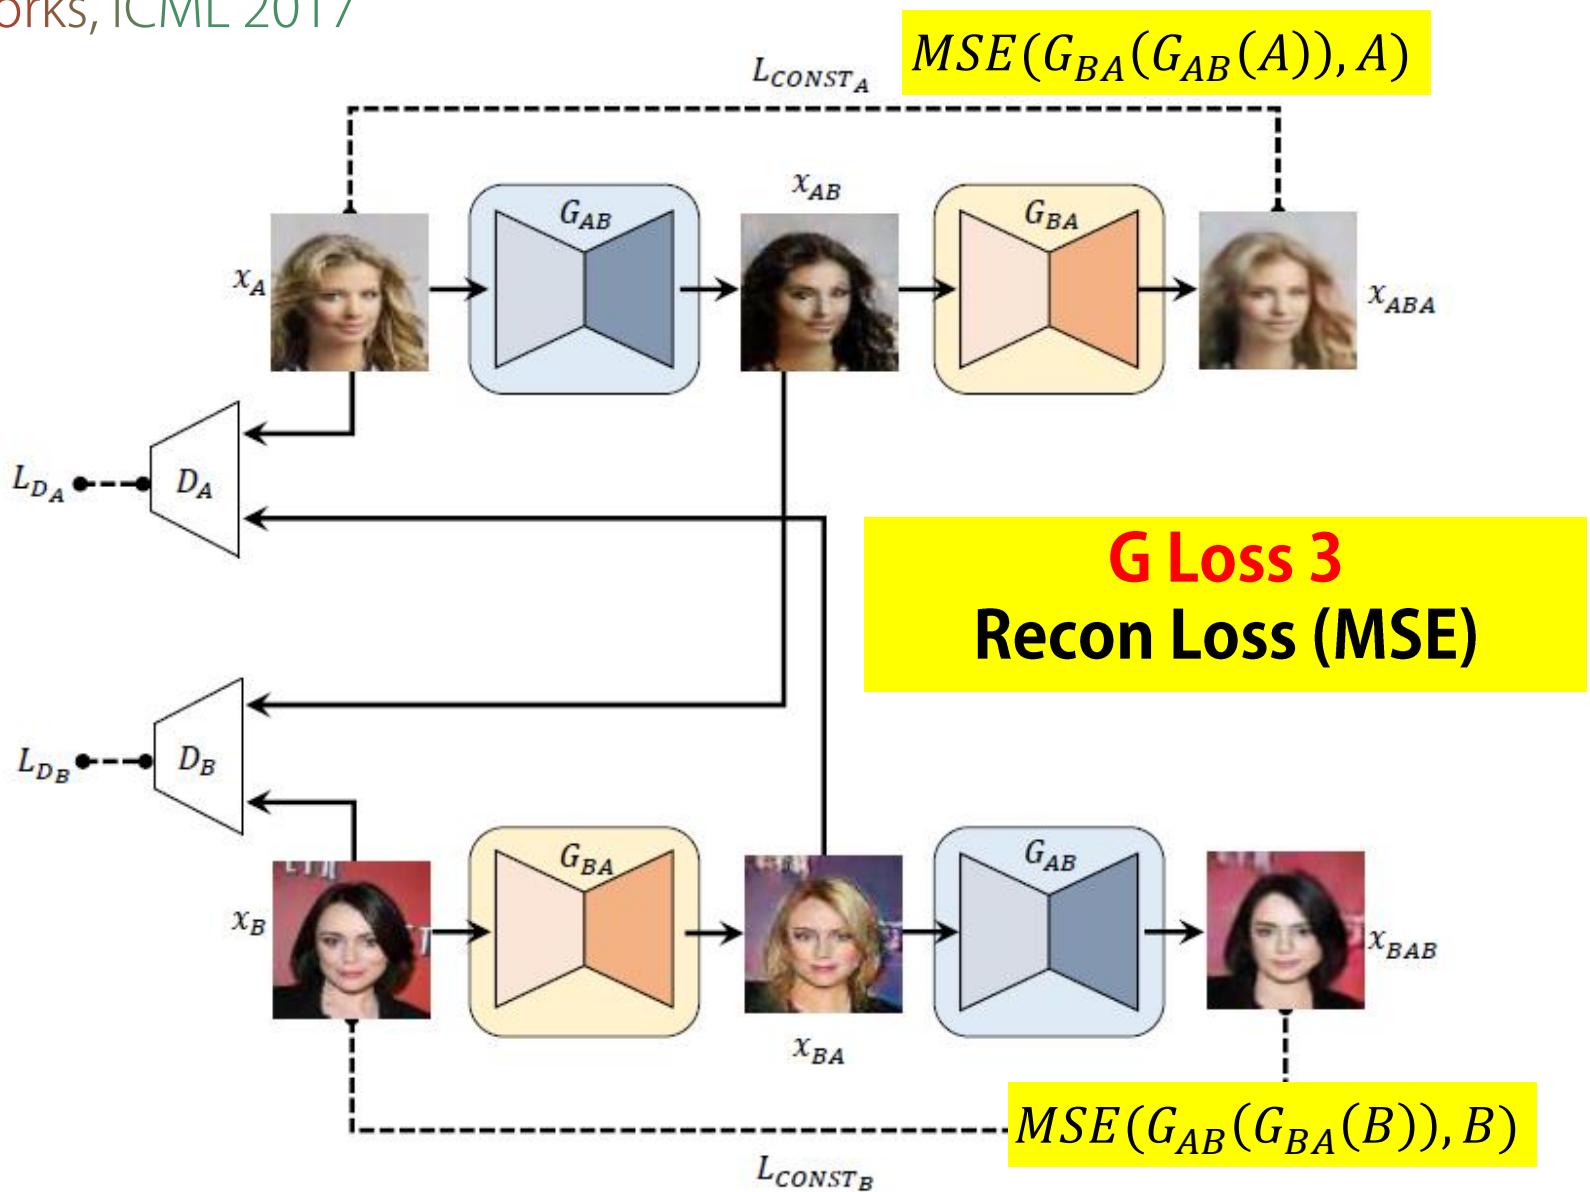


Image to Image Translation 3

Learning to Discover Cross-domain Relations with Generative Adversarial Networks, ICML 2017

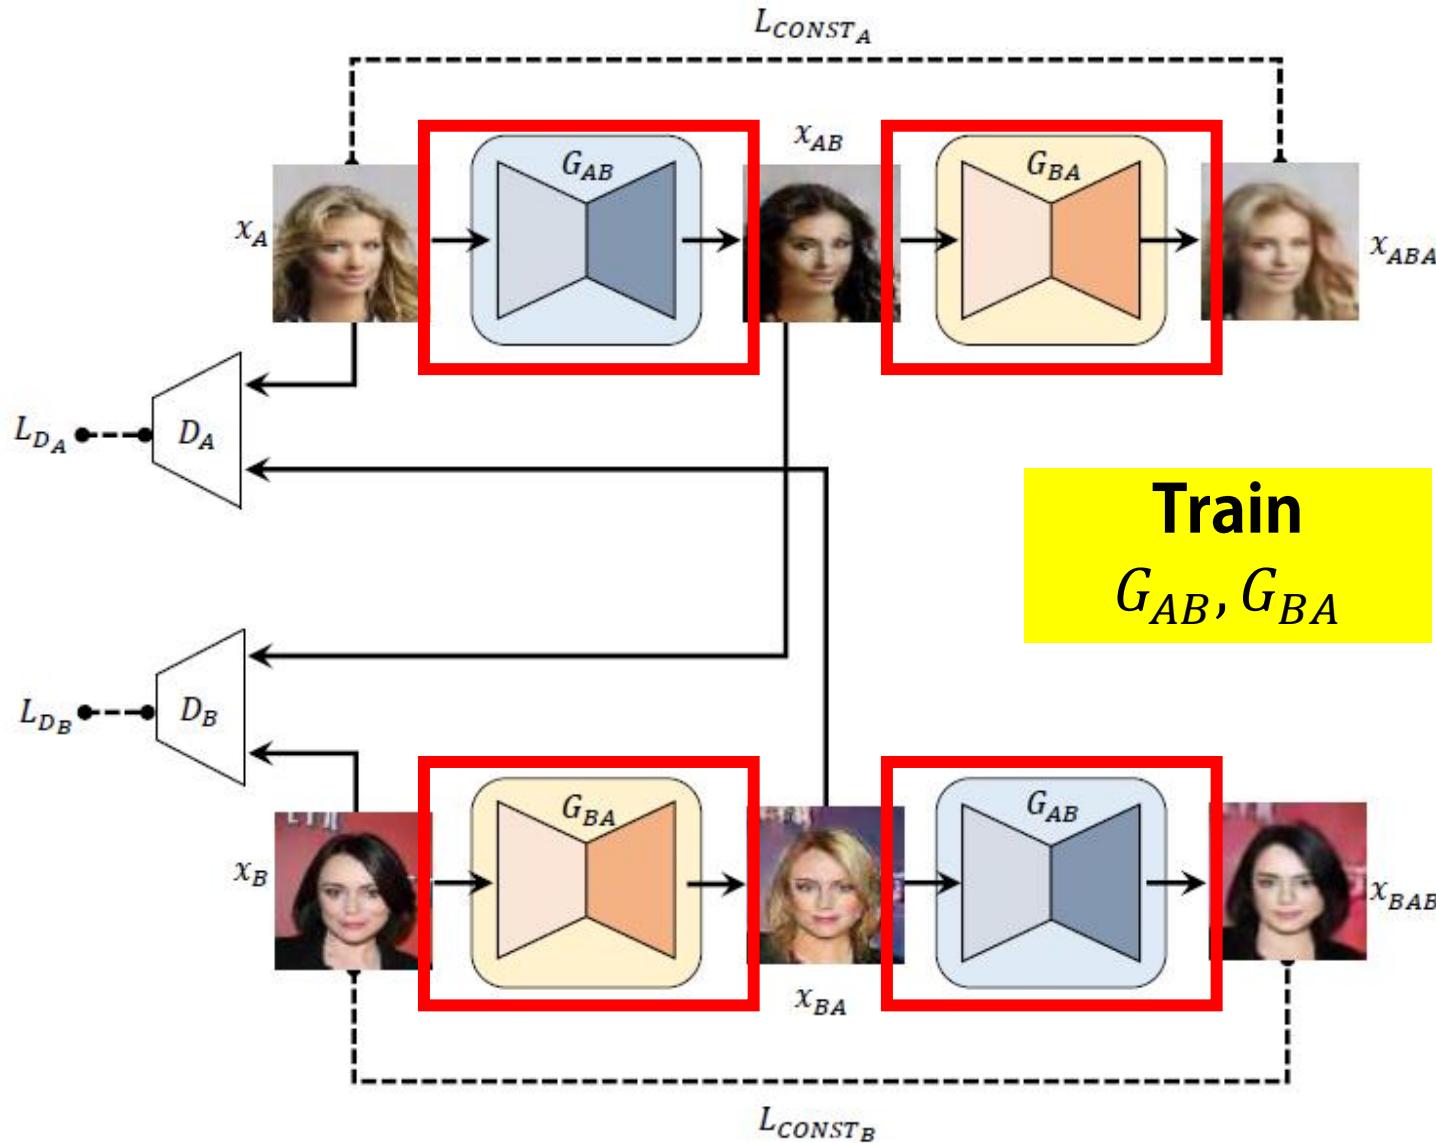


Image to Image Translation 3

Learning to Discover Cross-domain Relations with Generative Adversarial Networks, ICML 2017

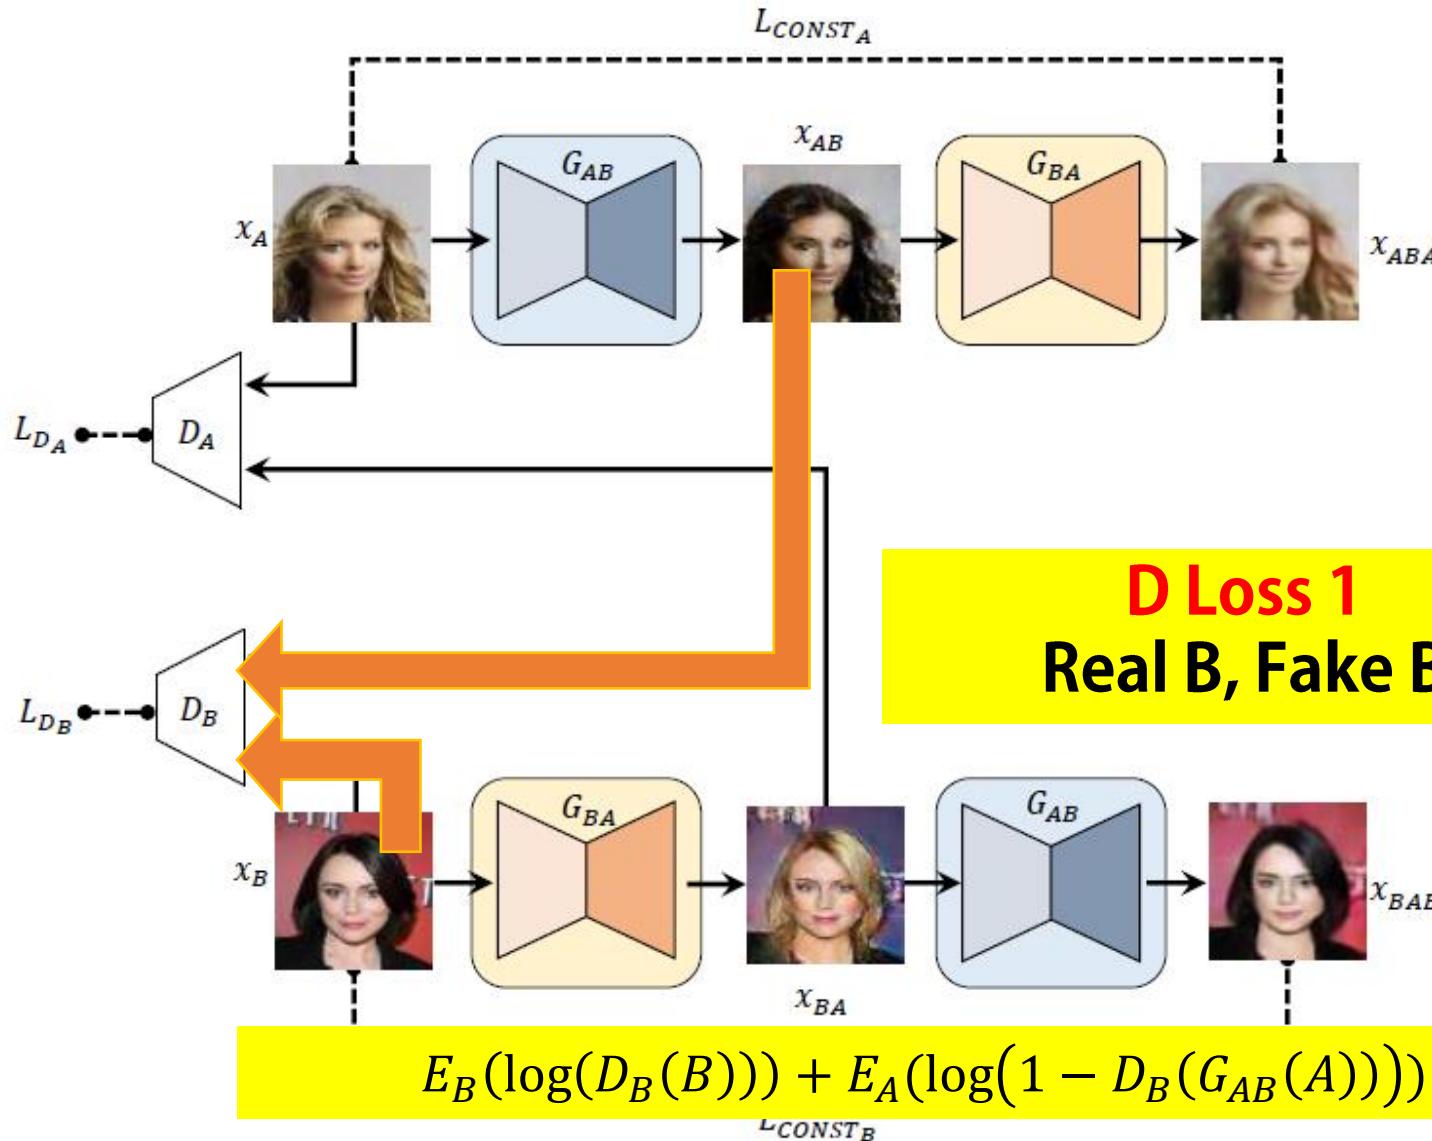
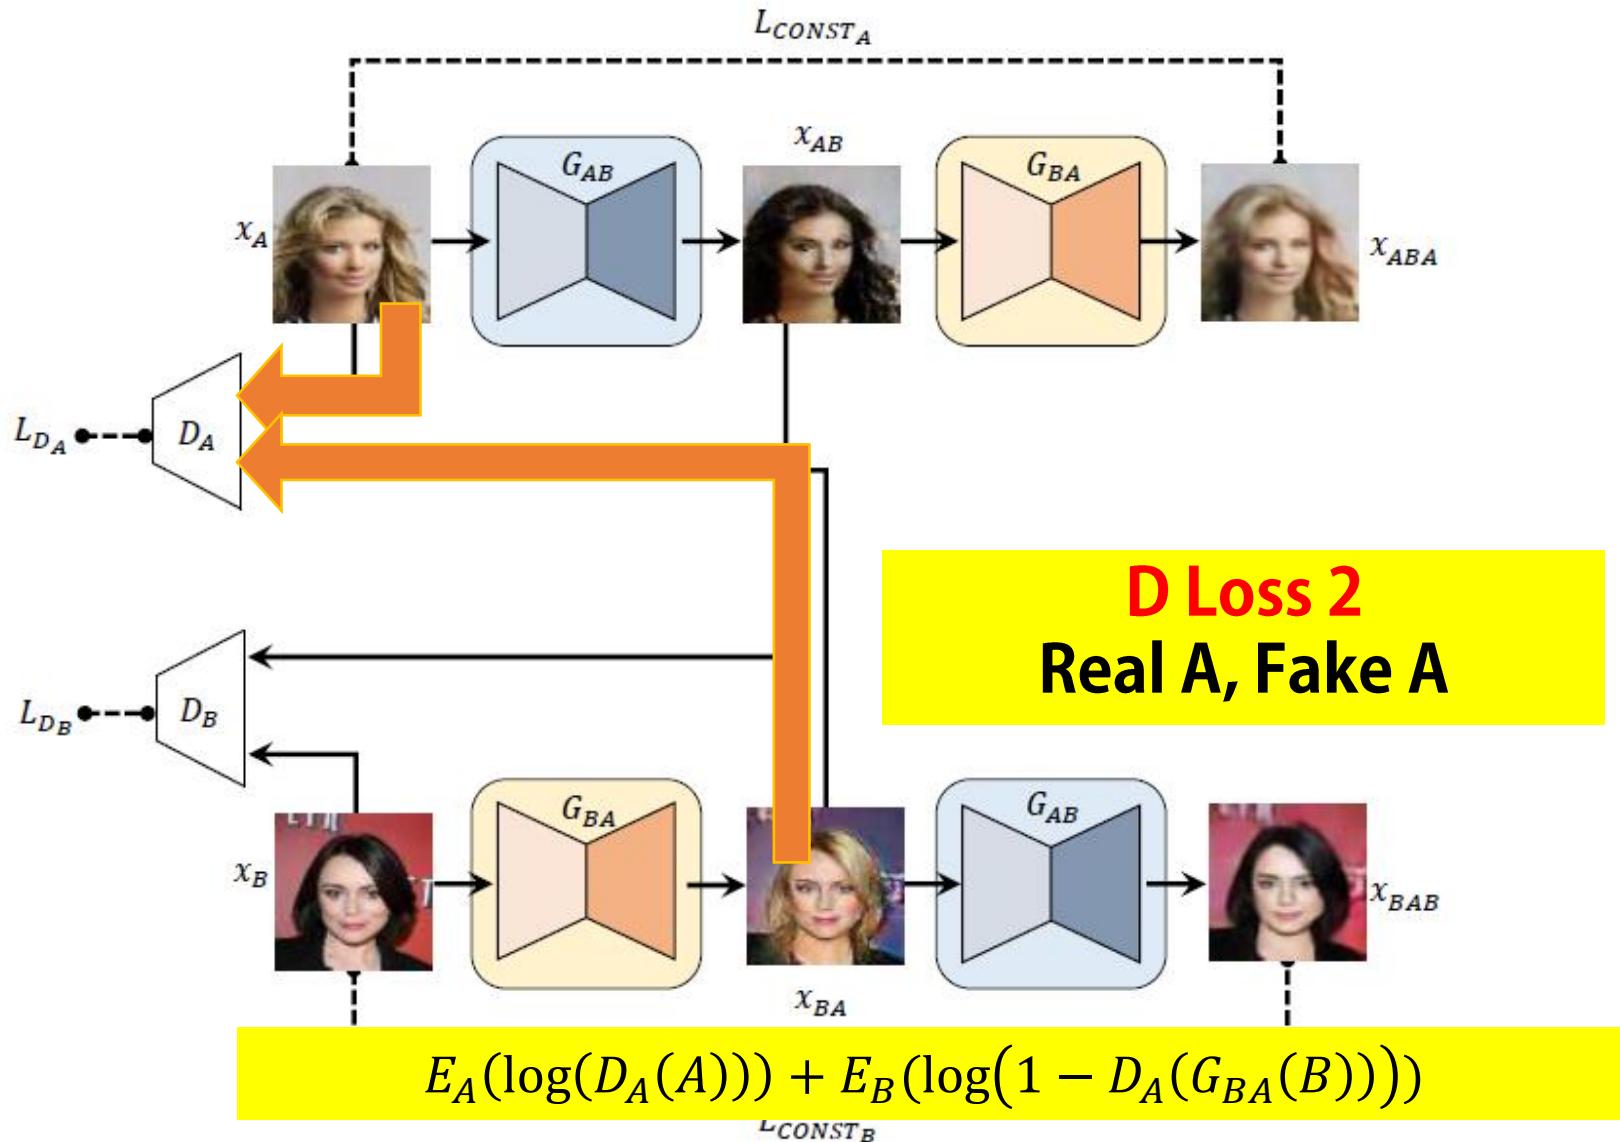
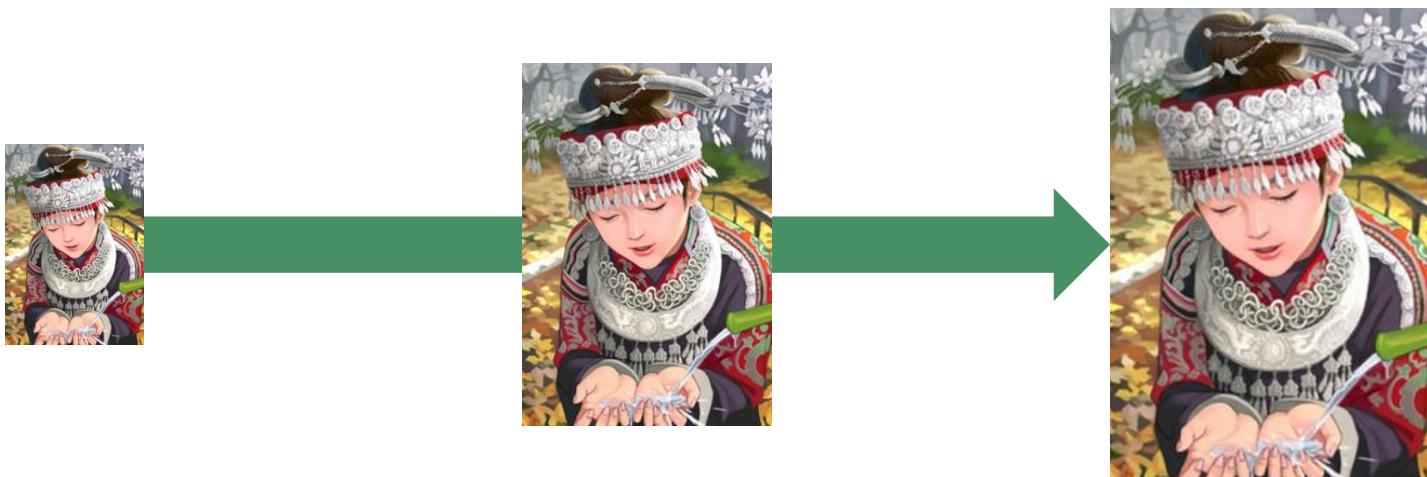


Image to Image Translation 3

Learning to Discover Cross-domain Relations with Generative Adversarial Networks, ICML 2017



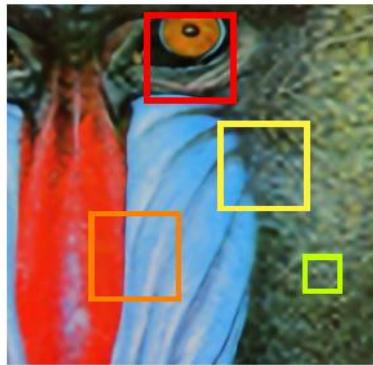
Single Image Super Resolution



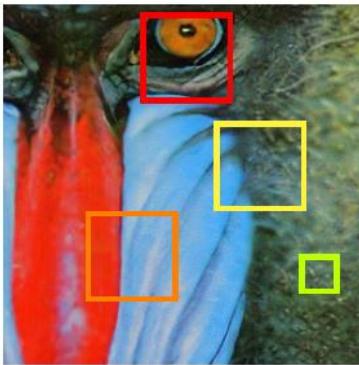
Single Image Super Resolution

Photo-realistic Single Image Super-Resolution using a Generative Adversarial Networks,
a.k.a SRGAN, CVPR 2017

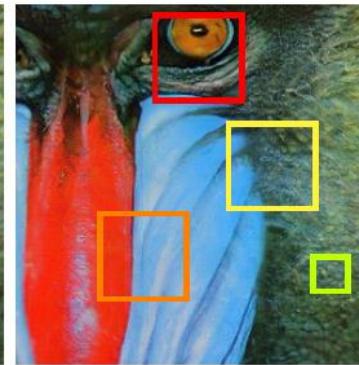
SRResNet



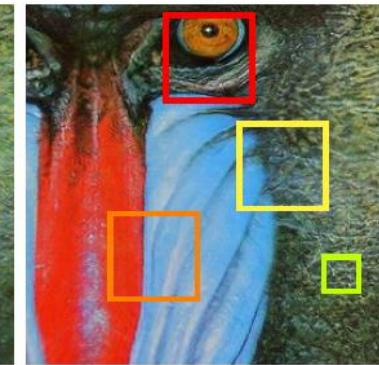
SRGAN-MSE



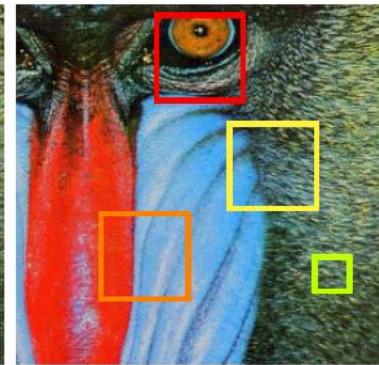
SRGAN-VGG22



SRGAN-VGG54

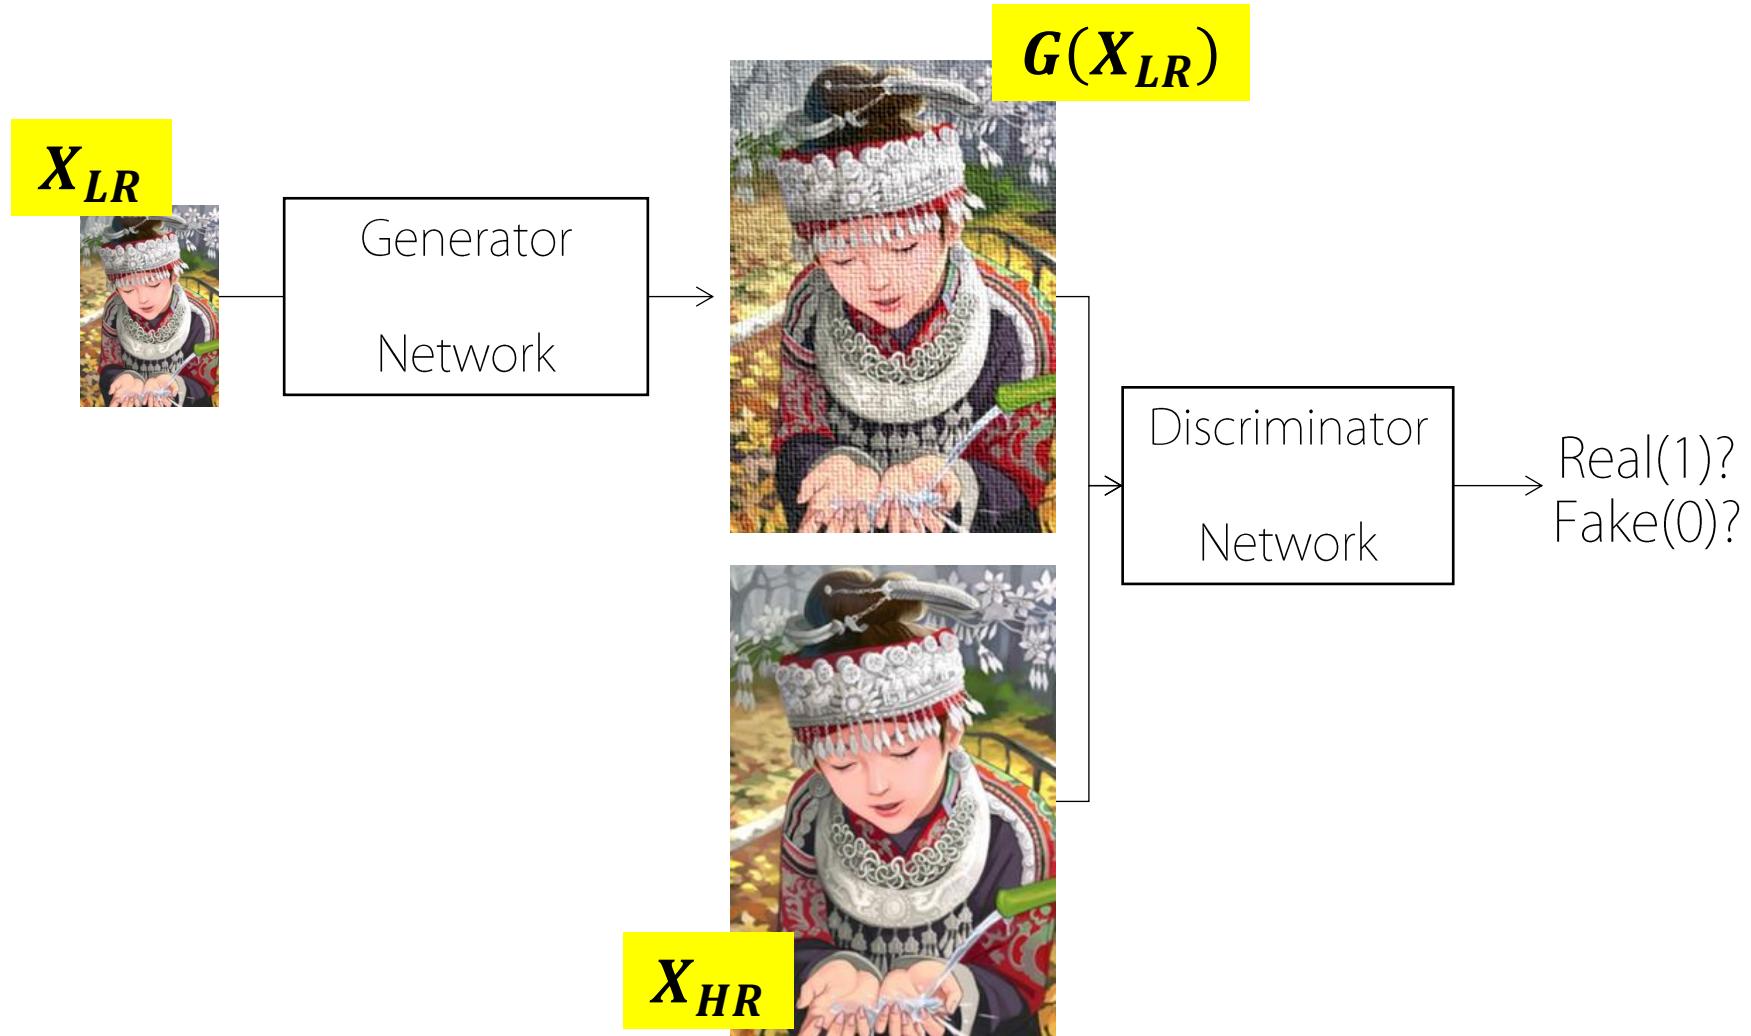


original HR image



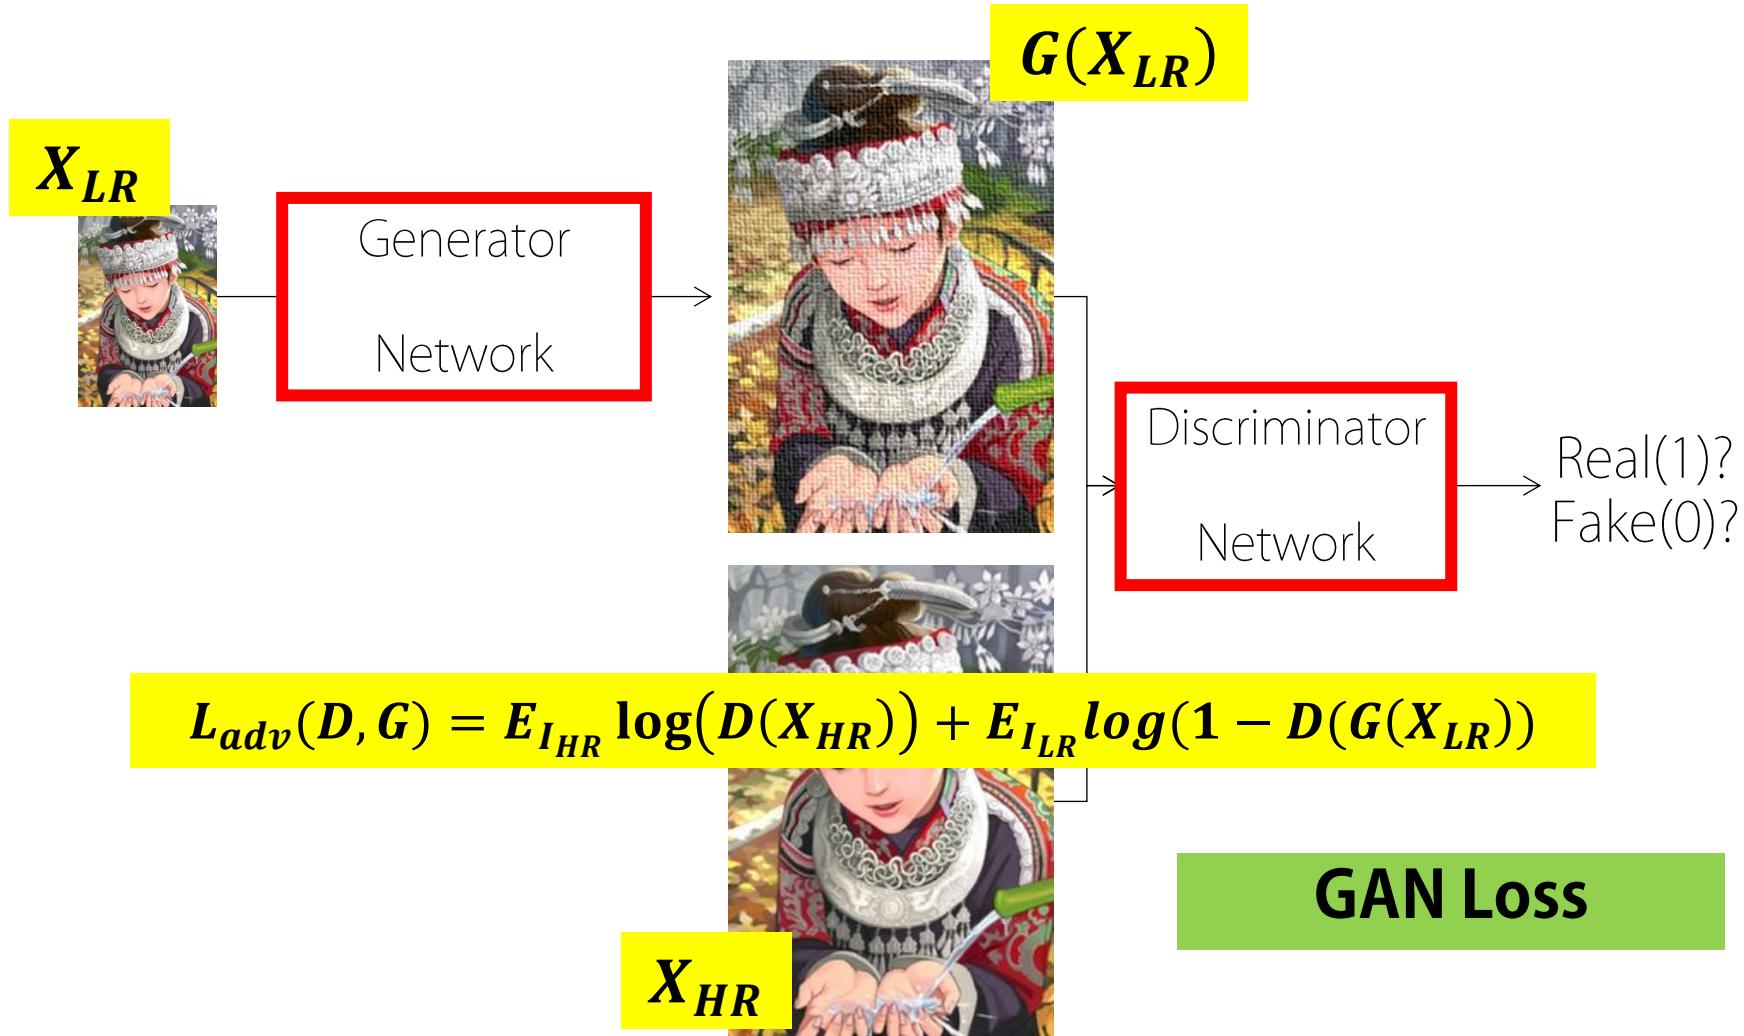
Single Image Super Resolution

Photo-realistic Single Image Super-Resolution using a Generative Adversarial Networks, CVPR 2017



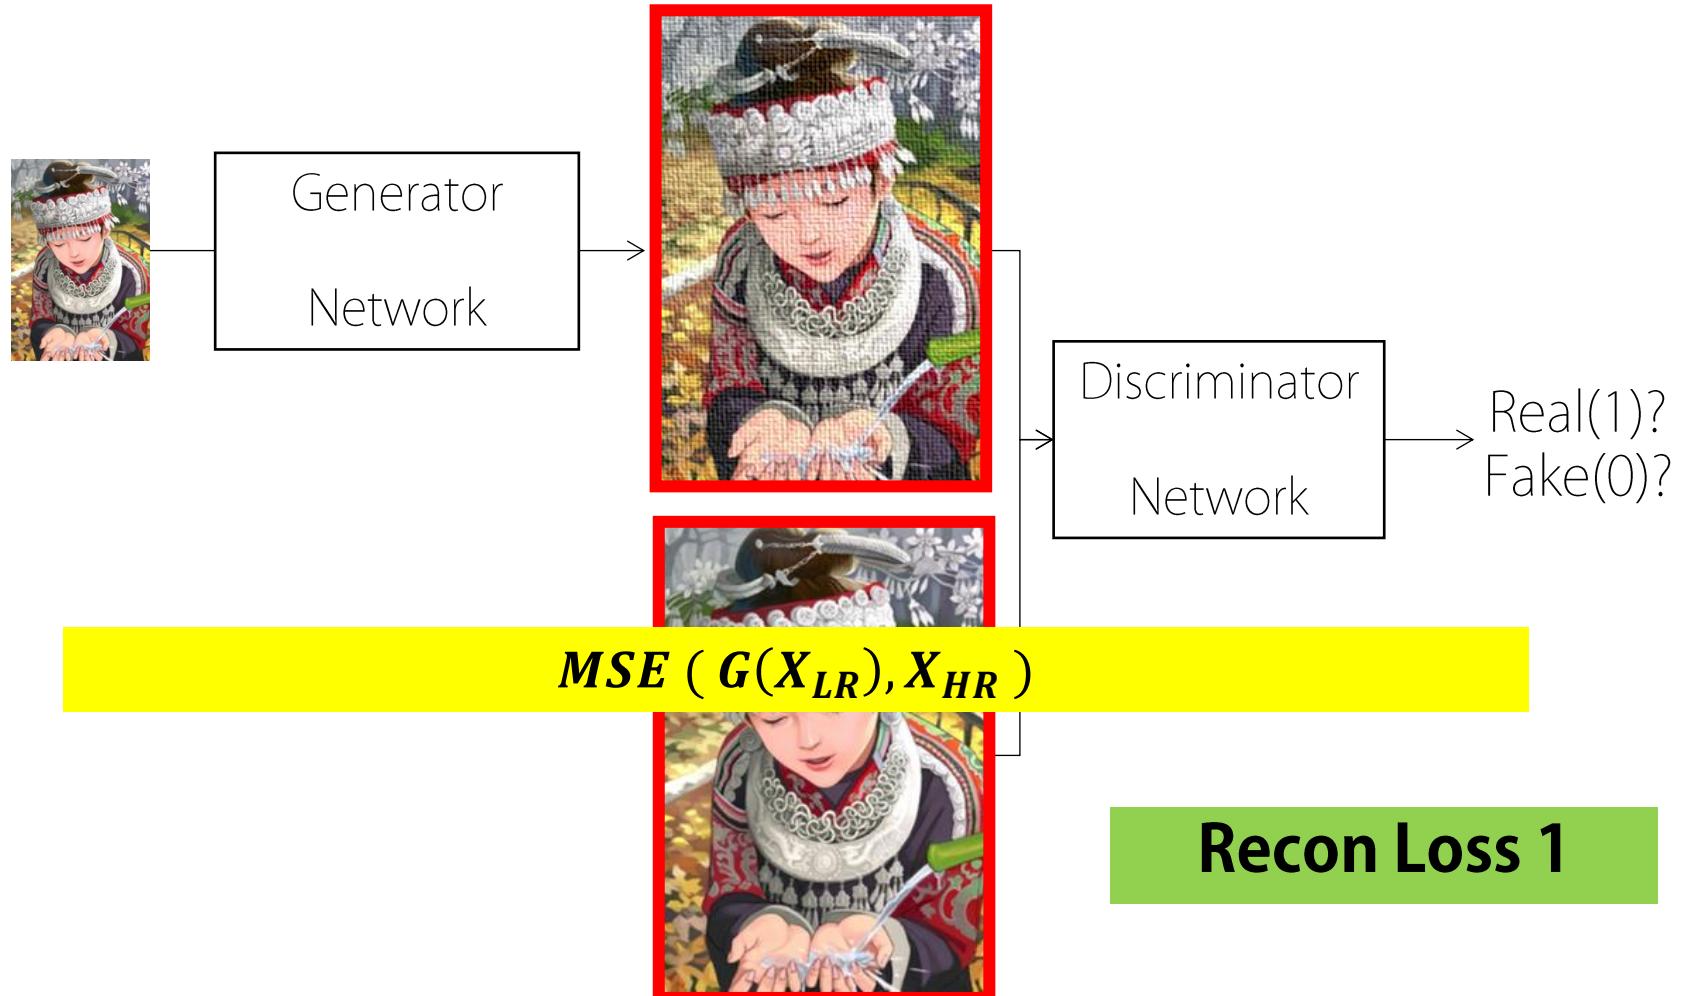
Single Image Super Resolution

Photo-realistic Single Image Super-Resolution using a Generative Adversarial Networks, CVPR 2017



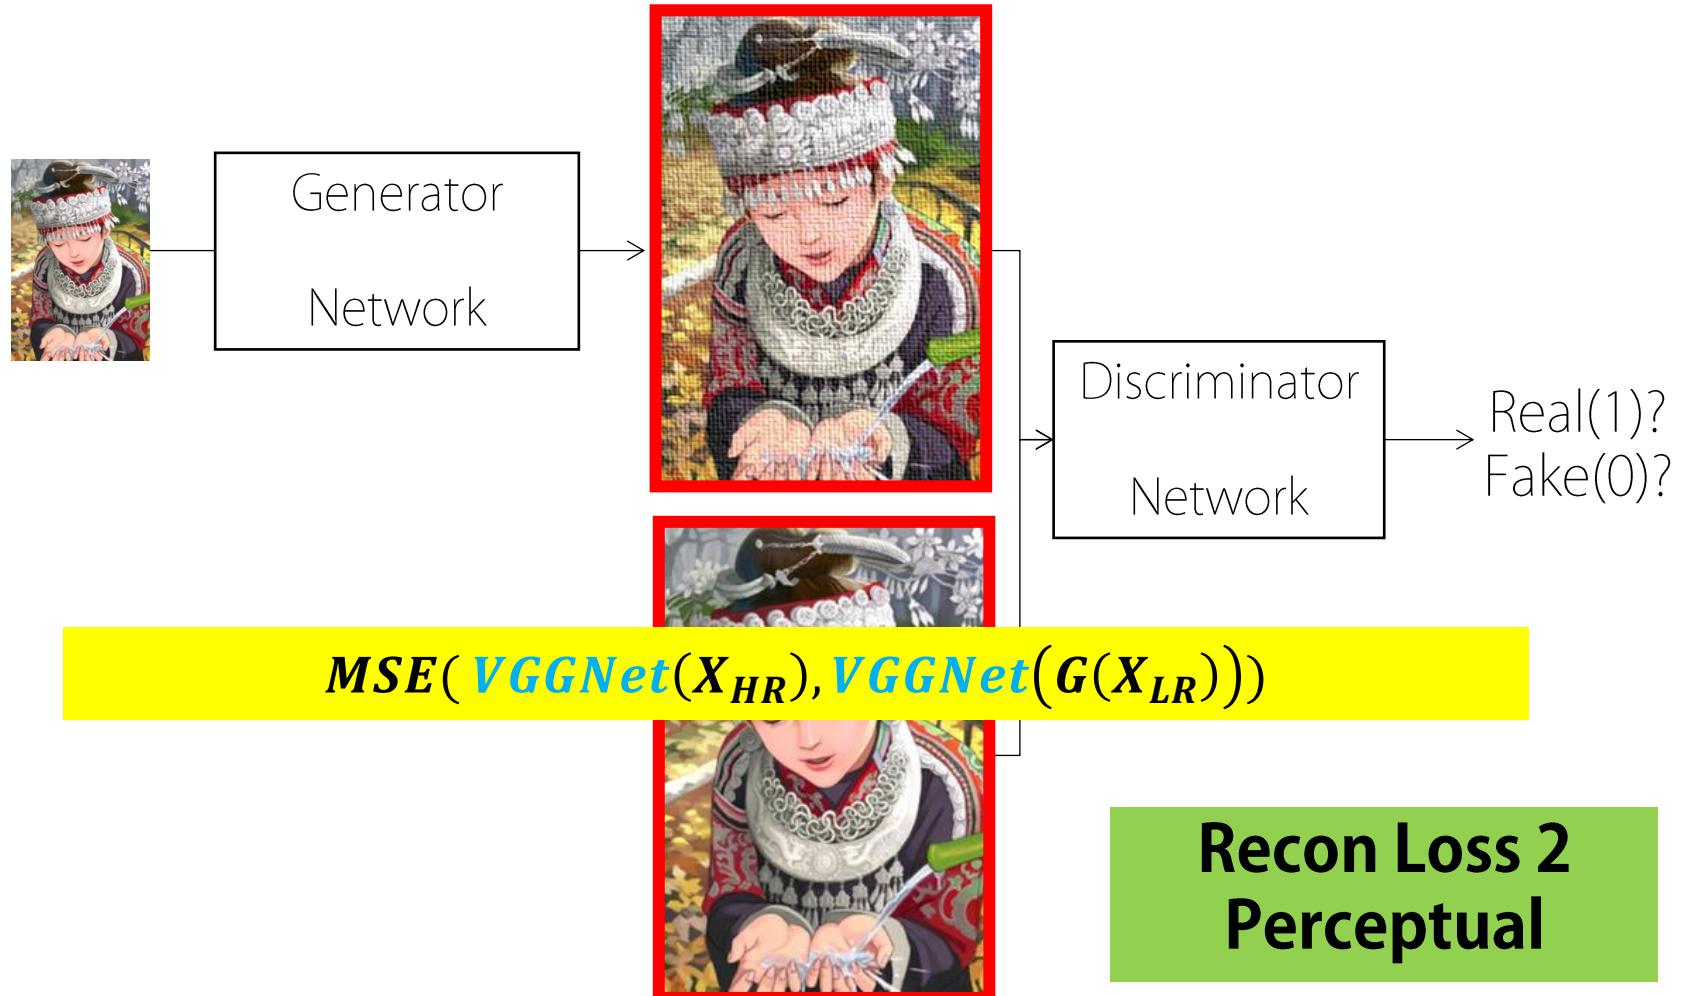
Single Image Super Resolution

Photo-realistic Single Image Super-Resolution using a Generative Adversarial Networks, CVPR 2017



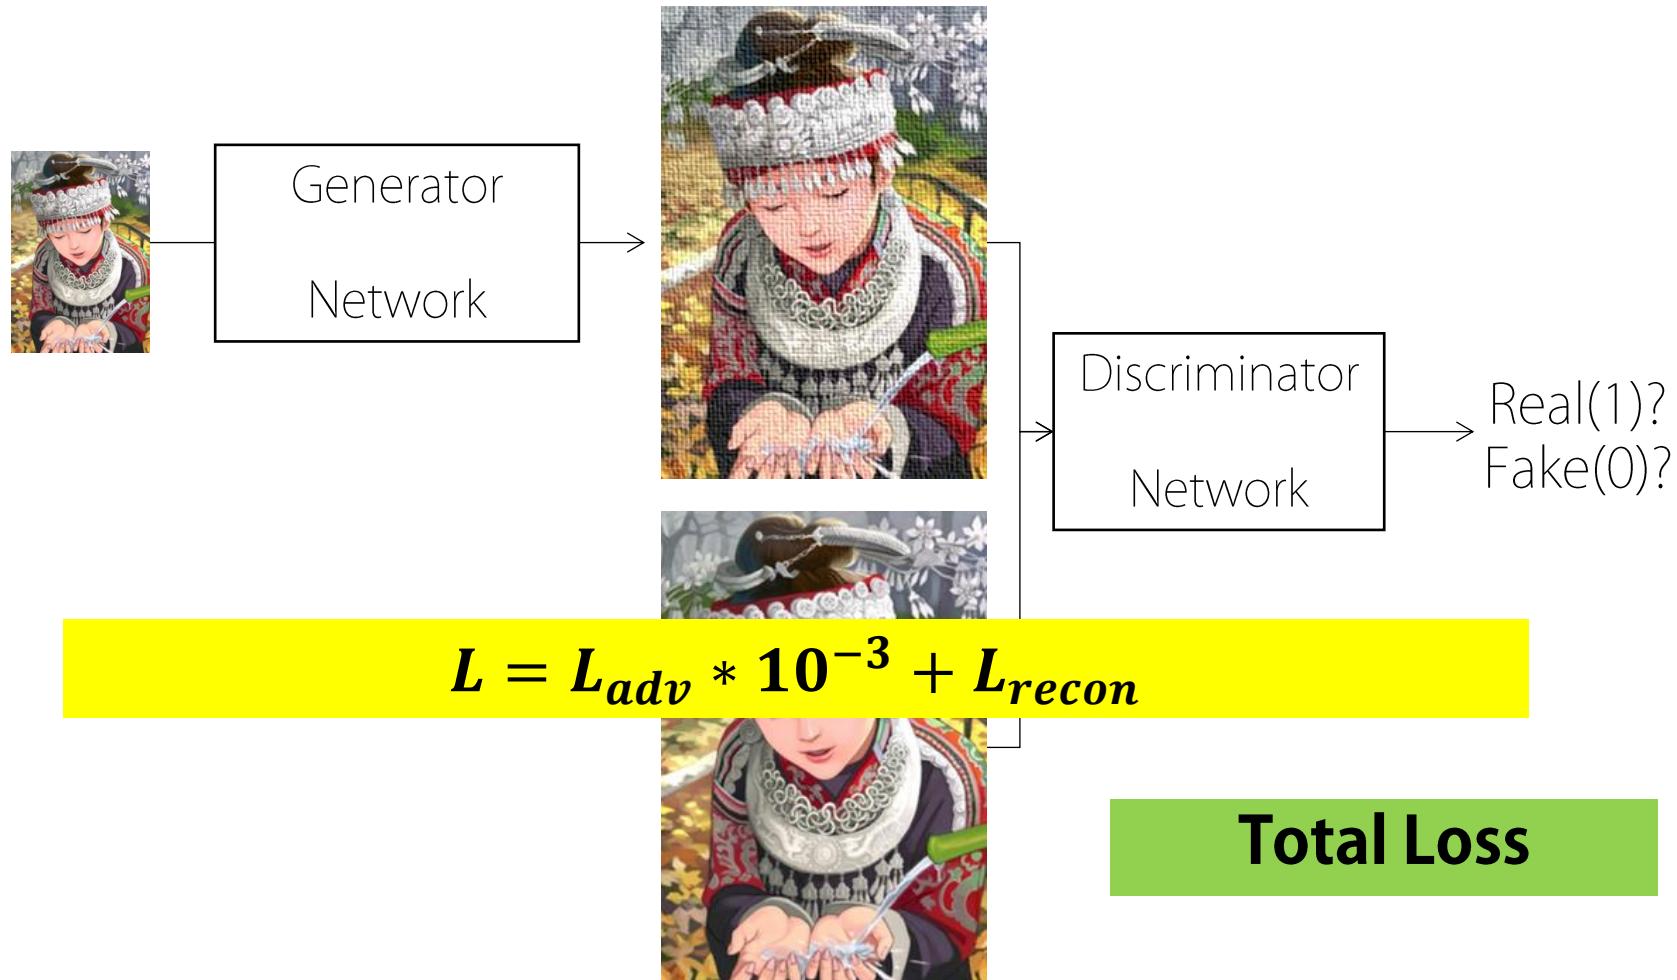
Single Image Super Resolution

Photo-realistic Single Image Super-Resolution using a Generative Adversarial Networks, CVPR 2017



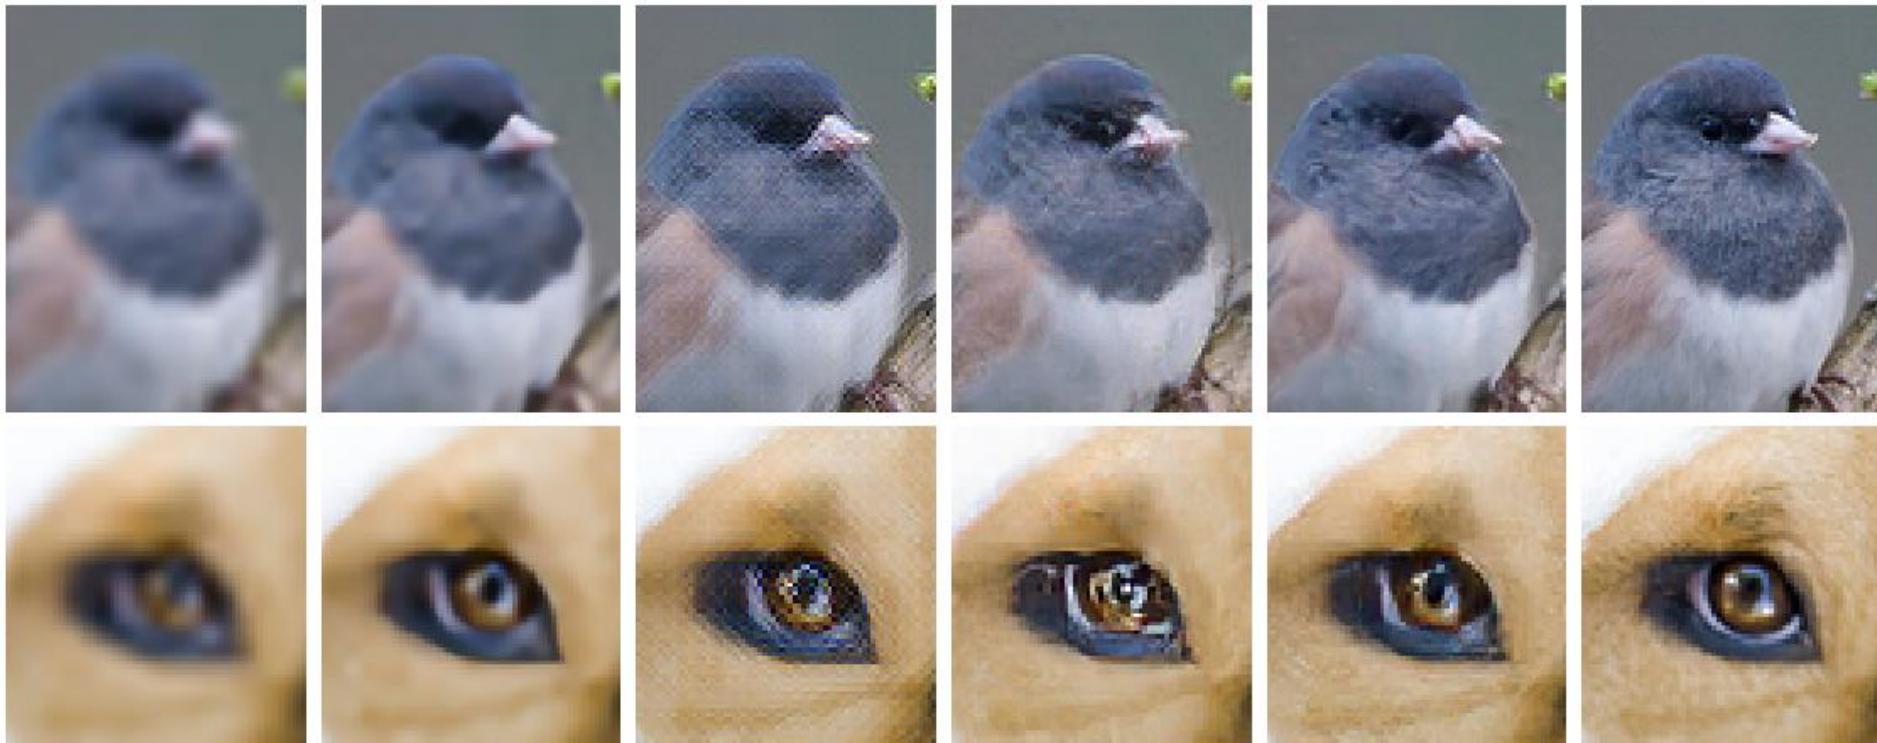
Single Image Super Resolution

Photo-realistic Single Image Super-Resolution using a Generative Adversarial Networks, CVPR 2017



Single Image Super Resolution 2

EnhanceNet: Single Image Super Resolution through Automated Texture
Synthesis,
a.k.a EnhanceNET, ICCV 2017



Bicubic

ENet-E

ENet-P

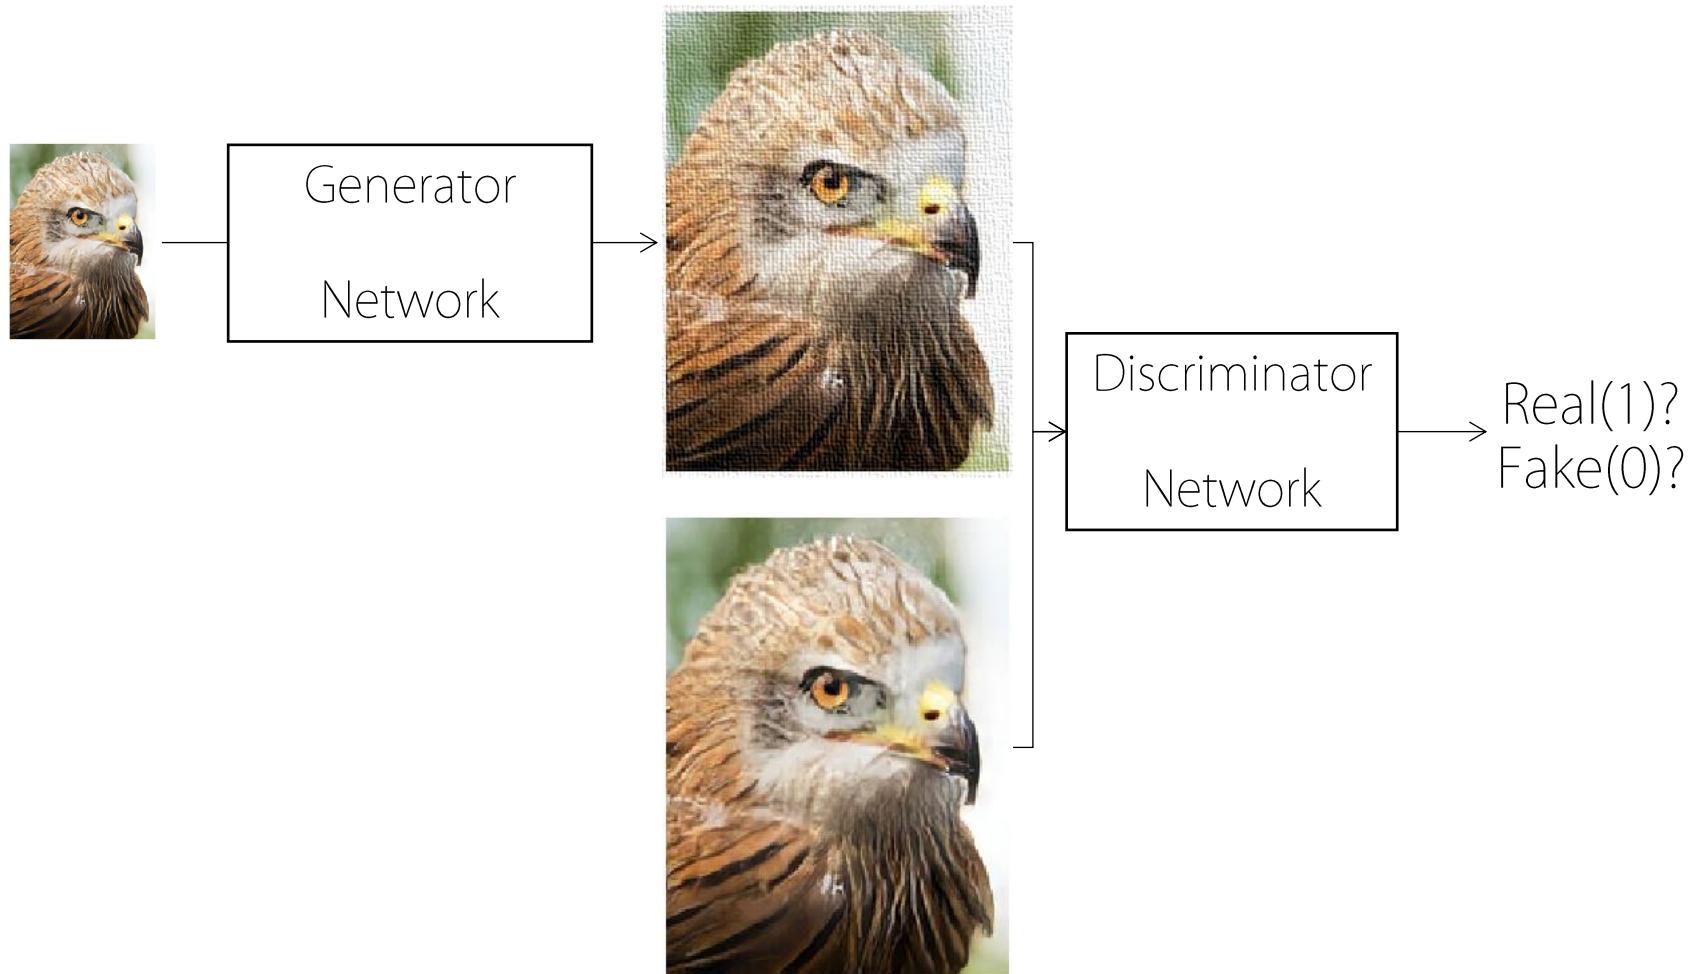
ENet-PA

ENet-PAT

I_{HR}

Single Image Super Resolution 2

EnhanceNet: Single Image Super Resolution through Automated Texture Synthesis, ICCV 2017



Single Image Super Resolution 2

EnhanceNet: Single Image Super Resolution through Automated Texture Synthesis, ICCV 2017

$$L_A = E_{I_{HR}} \log(D(X_{HR})) + E_{I_{LR}} \log(1 - D(G(X_{LR})))$$



GAN Loss

$$L_P = MSE(VGGNet(X_{HR}), VGGNet(G(X_{LR})))$$



Network



Recon Loss

$$L_T = MSE(GRAM(VGG(X_{HR})), GRAM(VGG(G(X_{LR}))))$$

→ Real(1)?
Fake(0)?

Network

Texture Loss



$$L = L_A + L_P + L_T$$



Total Loss

Small Object Detection Feature Upscaling

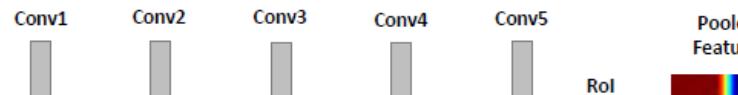


Object Detection

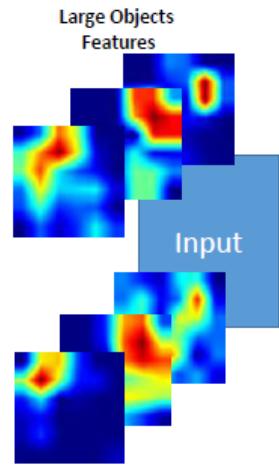
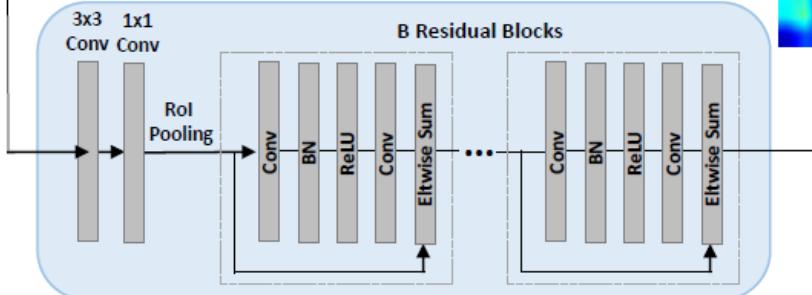
Perceptual Generative Adversarial Networks for Small Object Detection
CVPR 2017

Large Object Features

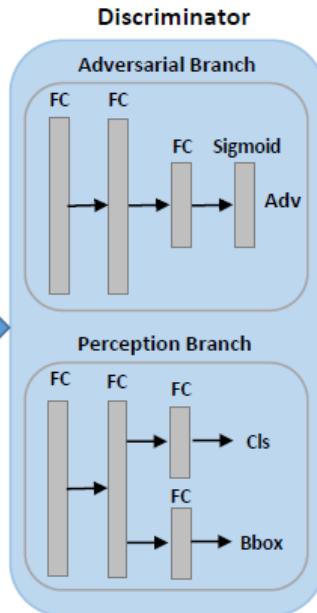
Generator Network



Generator



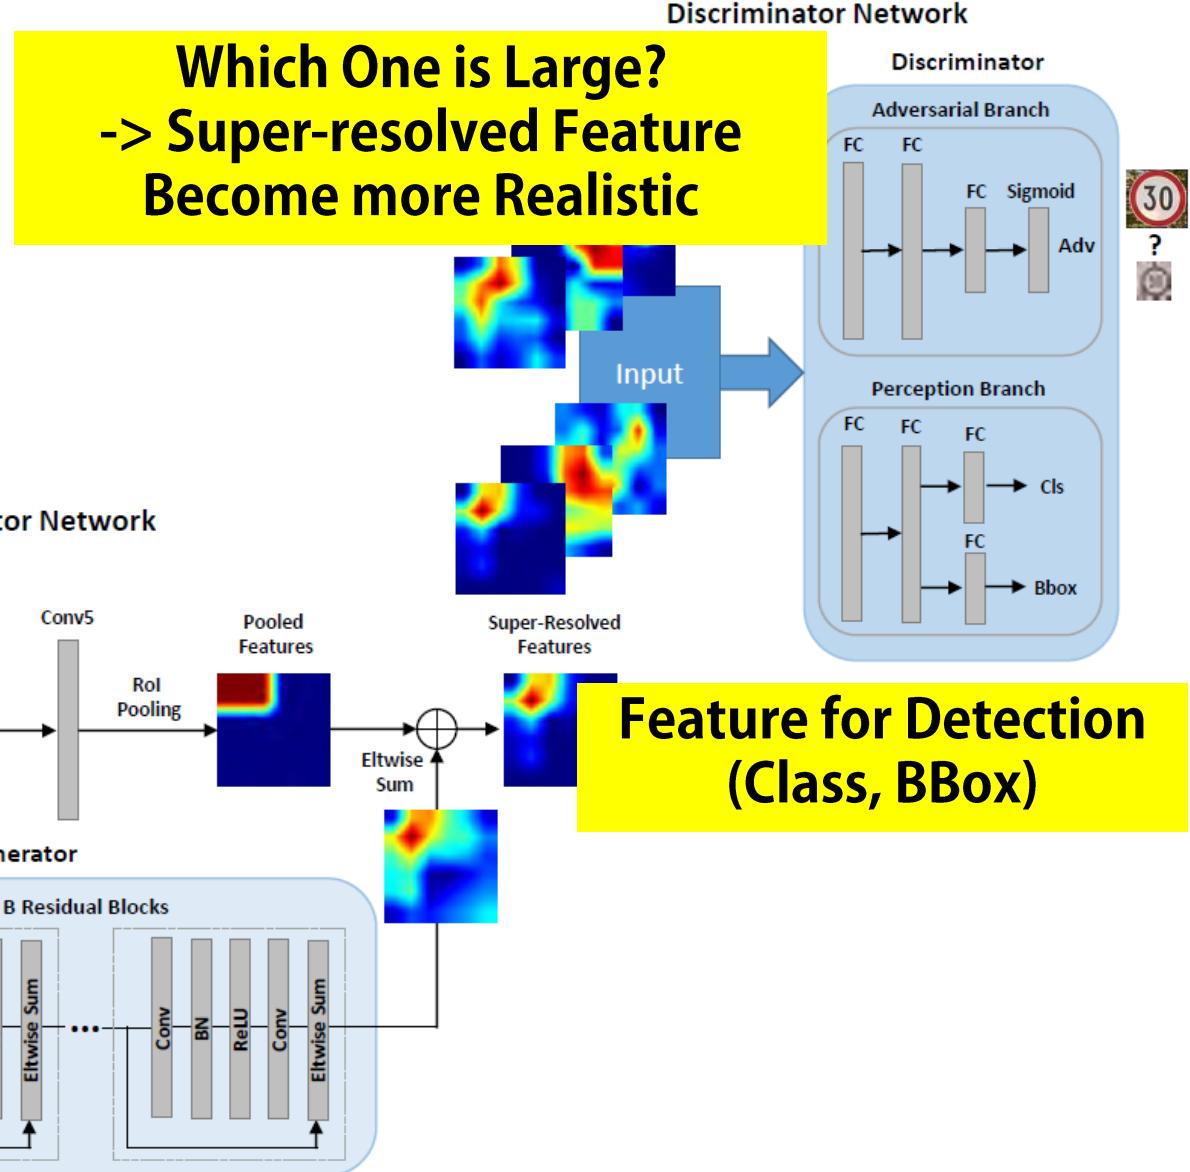
Discriminator Network



Residual Feature Super-Resolution

Object Detection

Perceptual Generative Adversarial Networks for Small Object Detection
CVPR 2017



3D Reconstruction

Reconstruction with DropoutNet

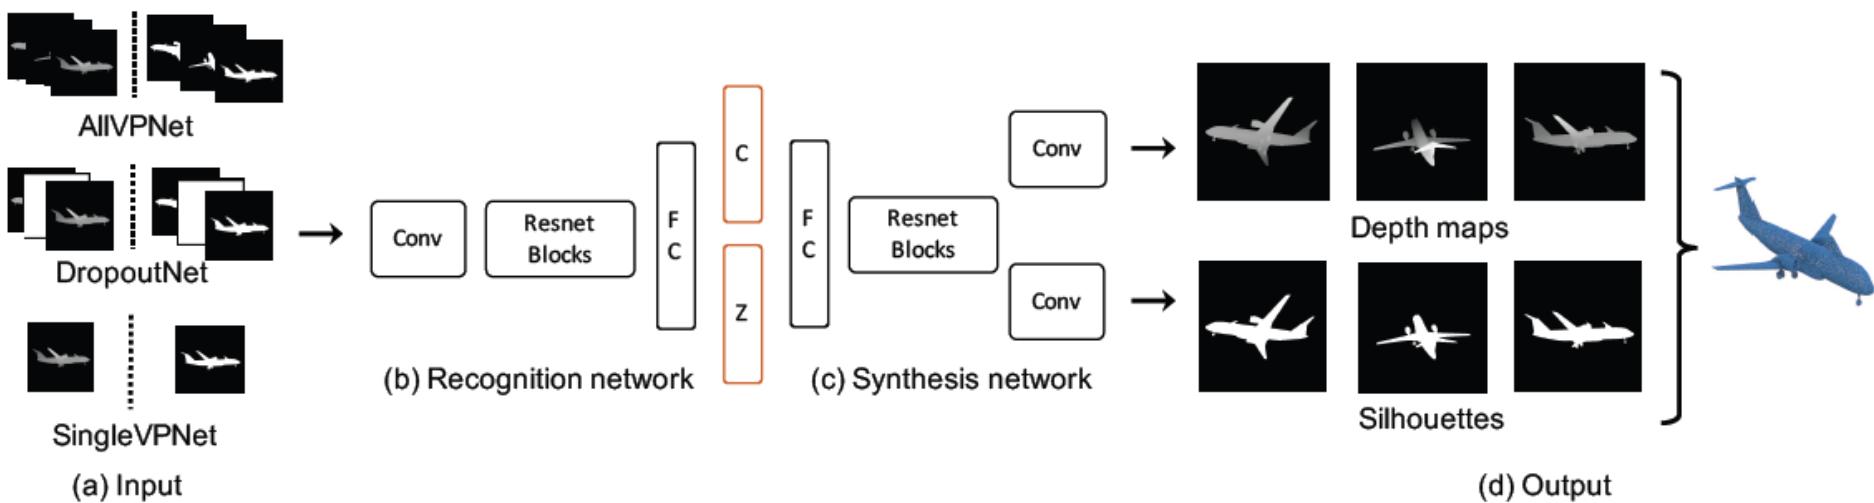


Reconstruction with SingleVPNet



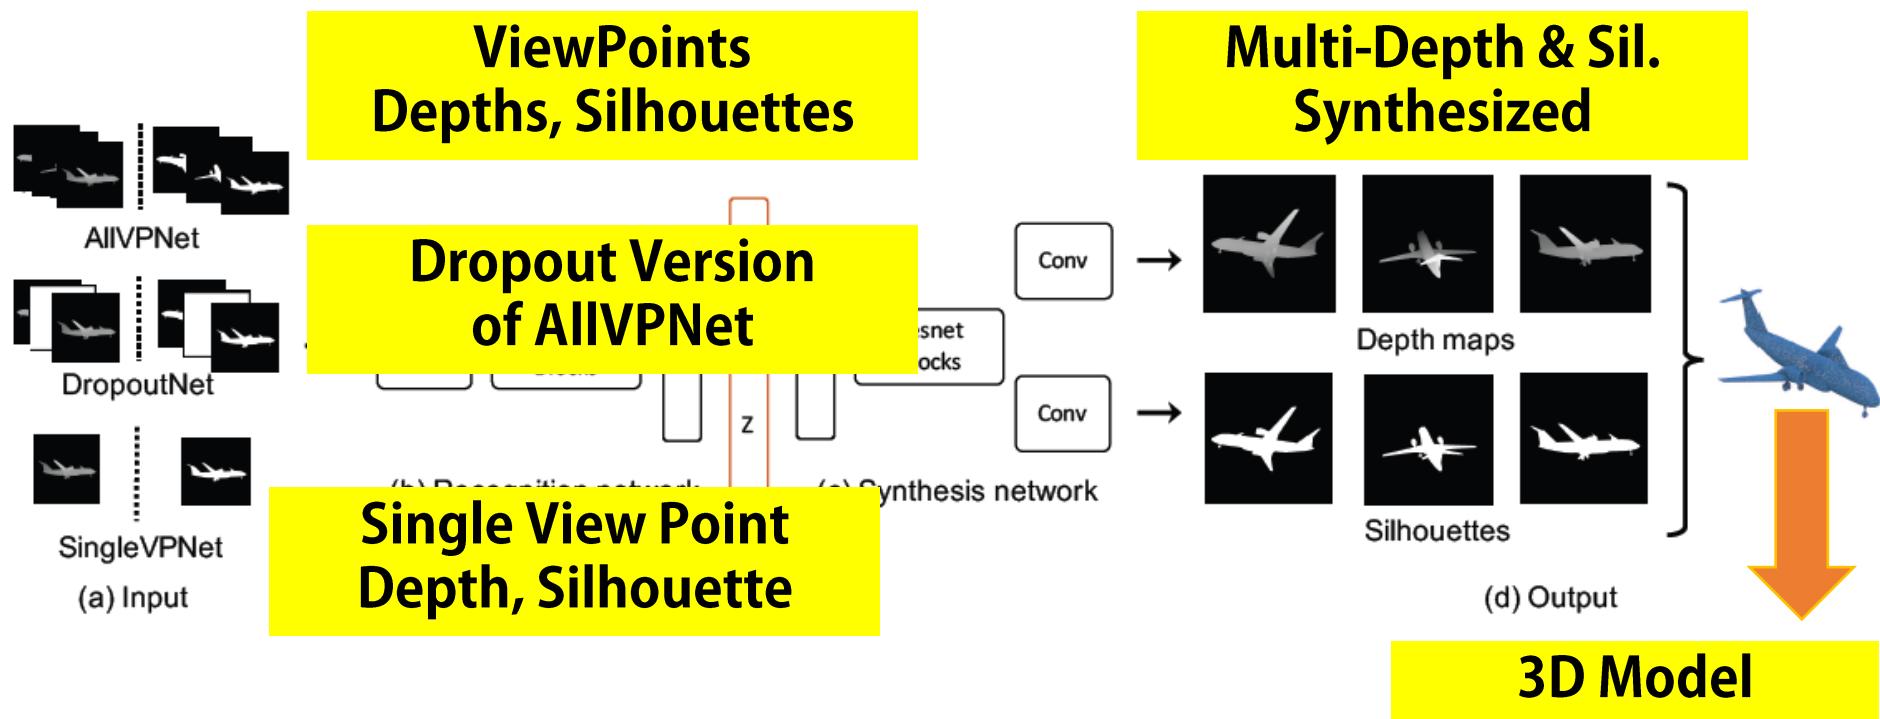
3D Reconstruction

Synthesizing 3D Shapes via Modeling Multi-View Depth Maps and Silhouettes with Deep Generative Networks, CVPR 2017



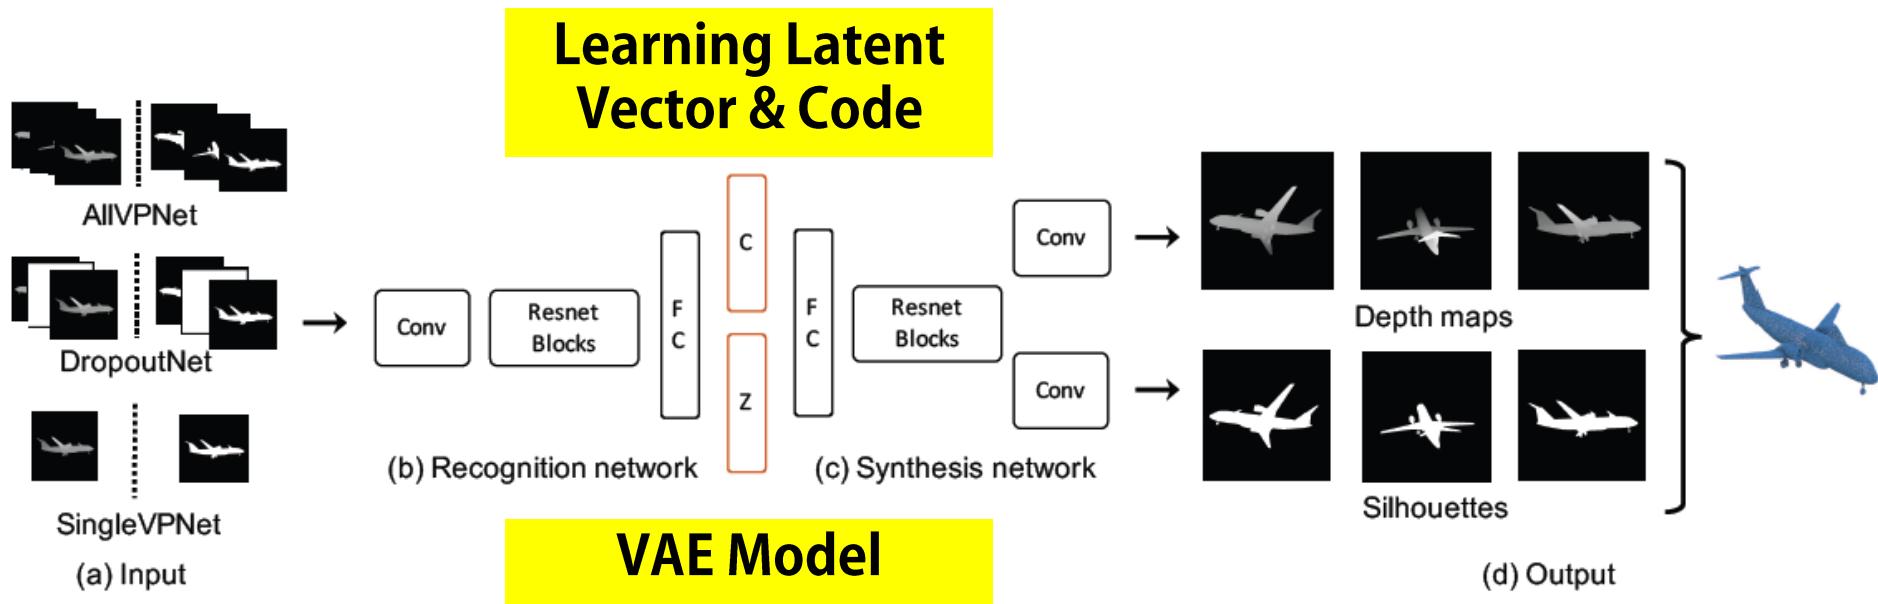
3D Reconstruction

Synthesizing 3D Shapes via Modeling Multi-View Depth Maps and Silhouettes with Deep Generative Networks, CVPR 2017



3D Reconstruction

Synthesizing 3D Shapes via Modeling Multi-View Depth Maps and Silhouettes with Deep Generative Networks, CVPR 2017



Varing c & z → Varing ViewPoints

GAN as Data Augmentation



For Advanced Supervised Learning

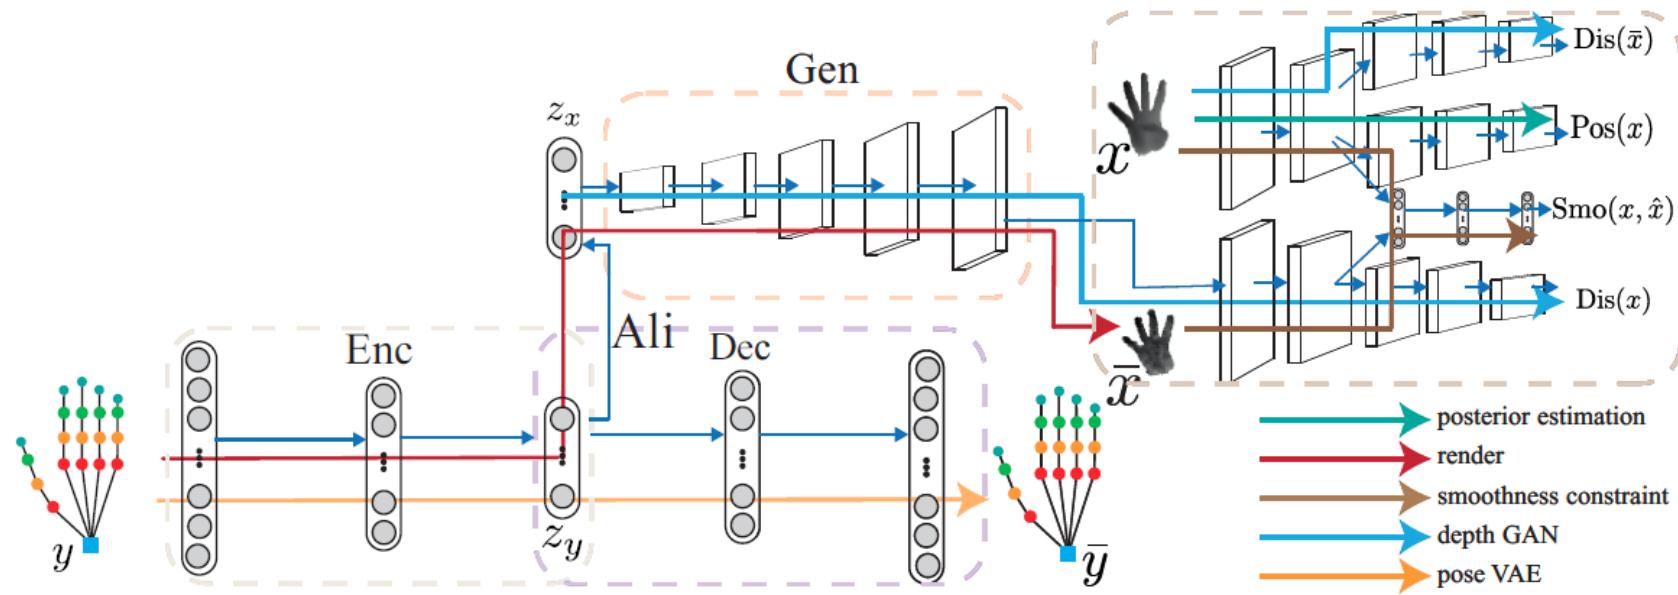
Pose Estimation



Figure 1. Random walk in the learned shared latent space. A set of points is sampled on the connecting line between two points in the shared latent space. The pose and corresponding depth map are then reconstructed through our network. Our method generates meaningful and realistic interpolations in both pose and appearance space.

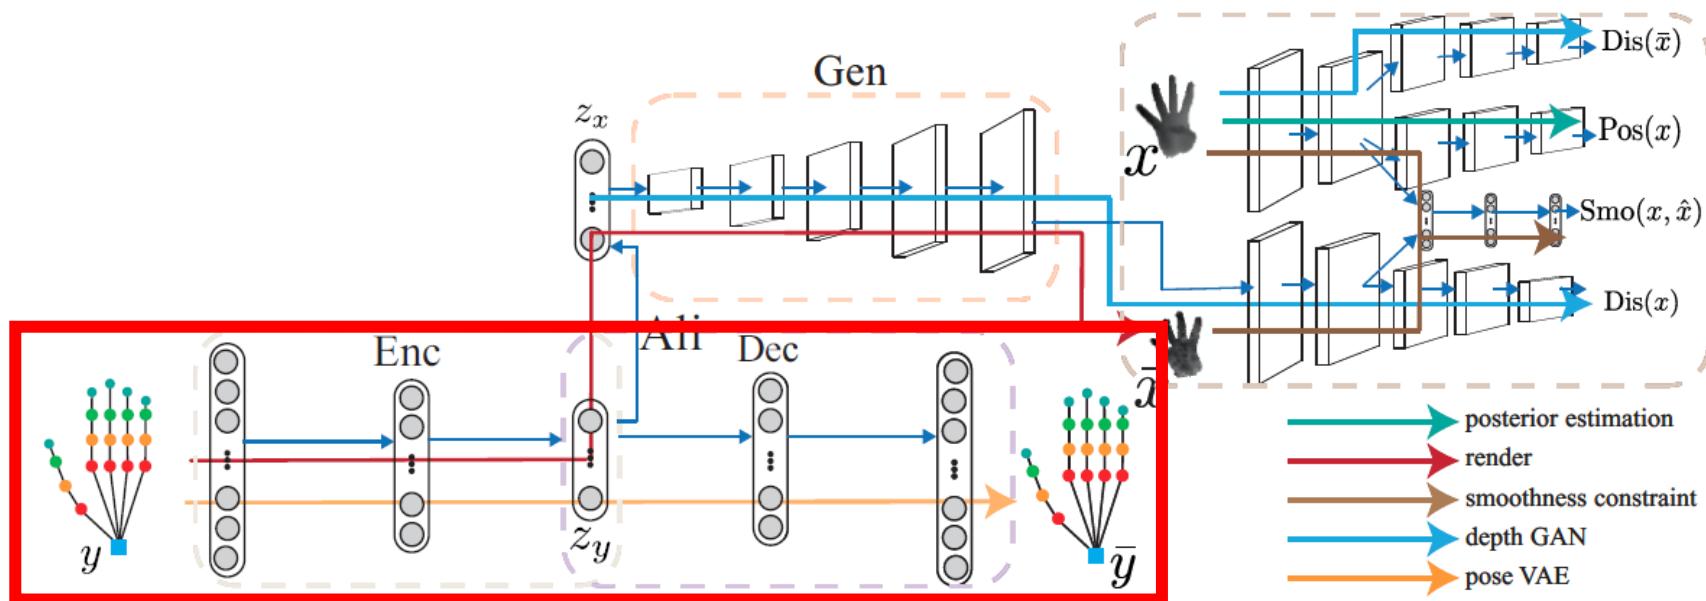
Pose Estimation

Crossing Nets: Combining GANs and VAEs with a Shared Latent Space for Hand Pose Estimation, CVPR 2017



Pose Estimation

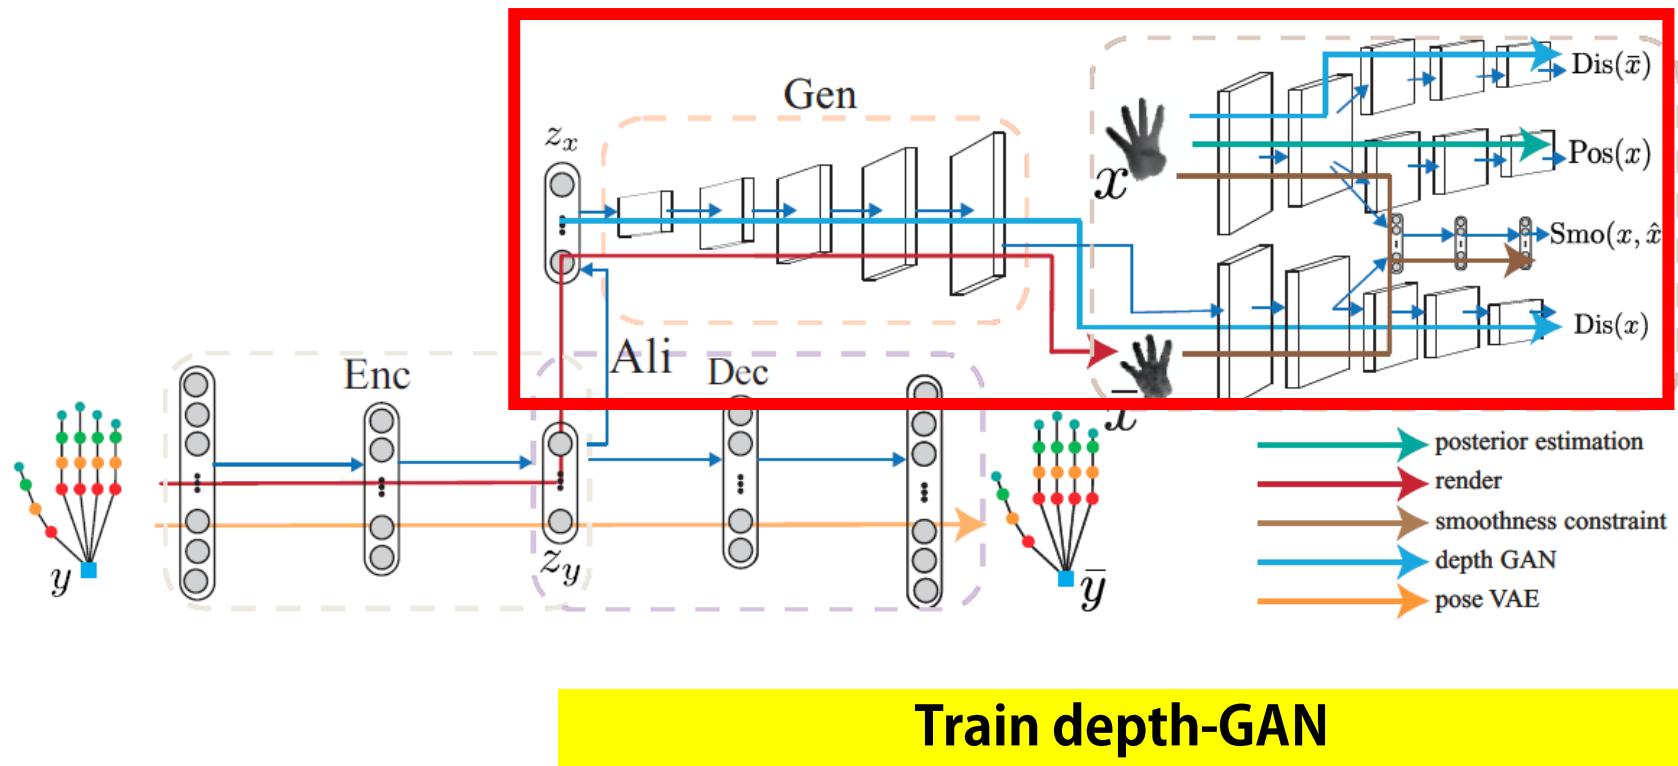
Crossing Nets: Combining GANs and VAEs with a Shared Latent Space for Hand Pose Estimation, CVPR 2017



Train pose-VAE

Pose Estimation

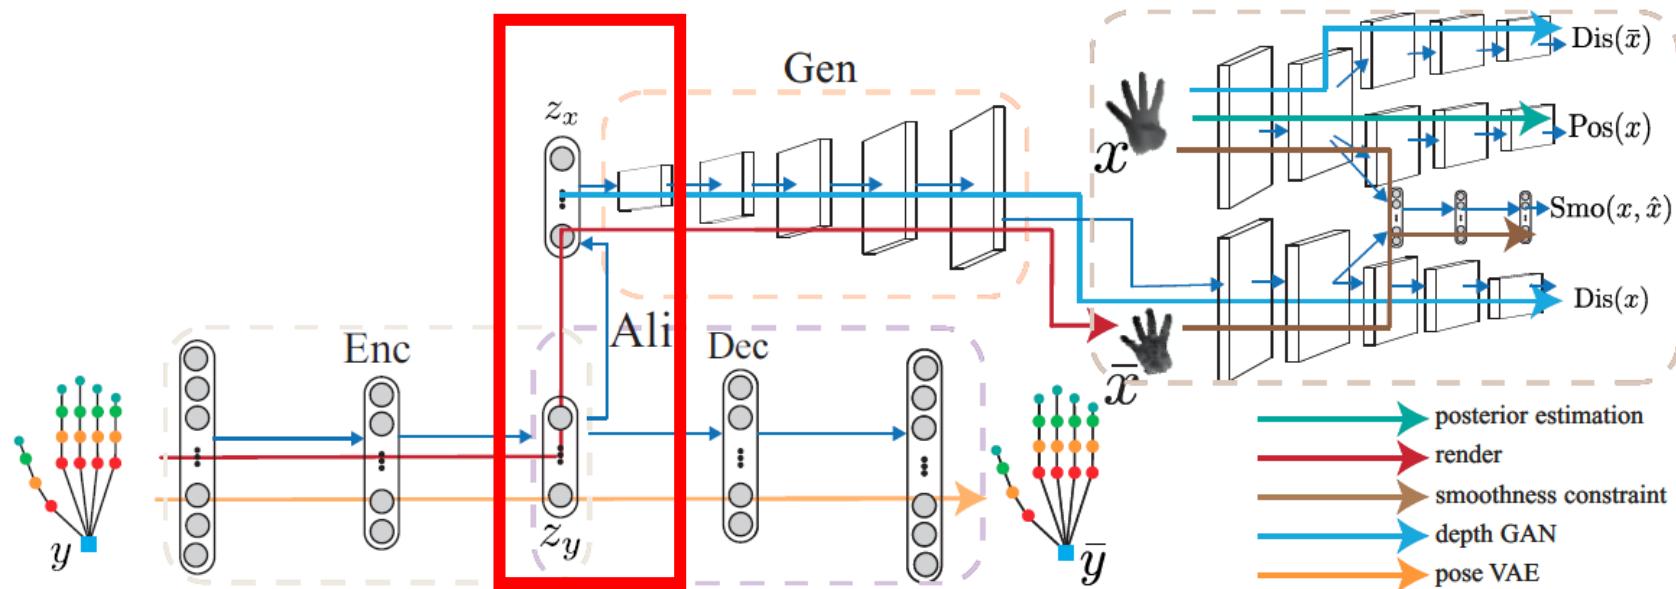
Crossing Nets: Combining GANs and VAEs with a Shared Latent Space for Hand Pose Estimation, CVPR 2017



Pose Estimation

Crossing Nets: Combining GANs and VAEs with a Shared Latent Space for Hand Pose Estimation, CVPR 2017

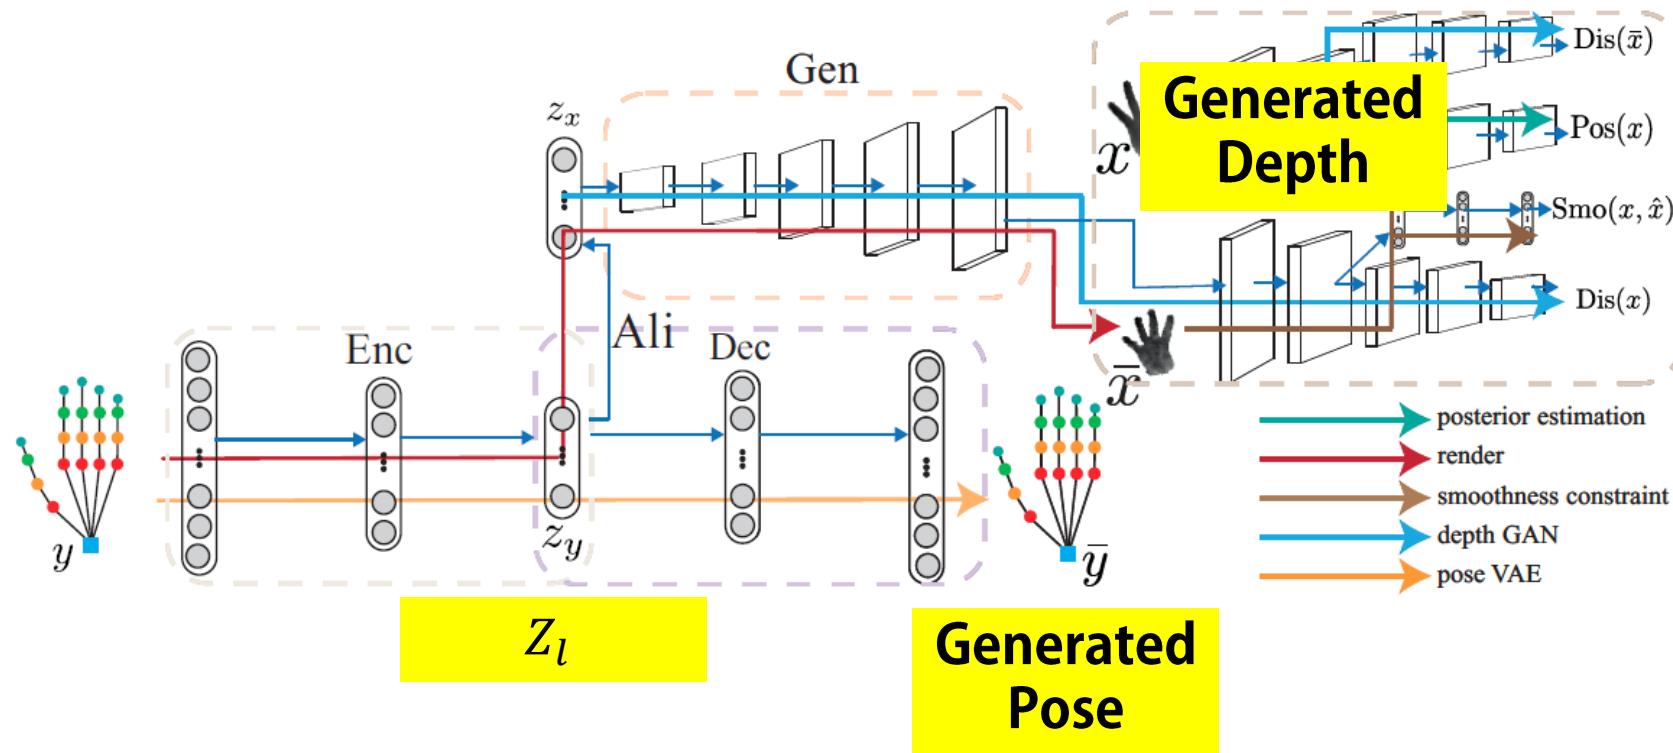
$$L_{recons} = Loss(x, Gen(Ali(Enc(y))))$$



Train Latent Space Mapping

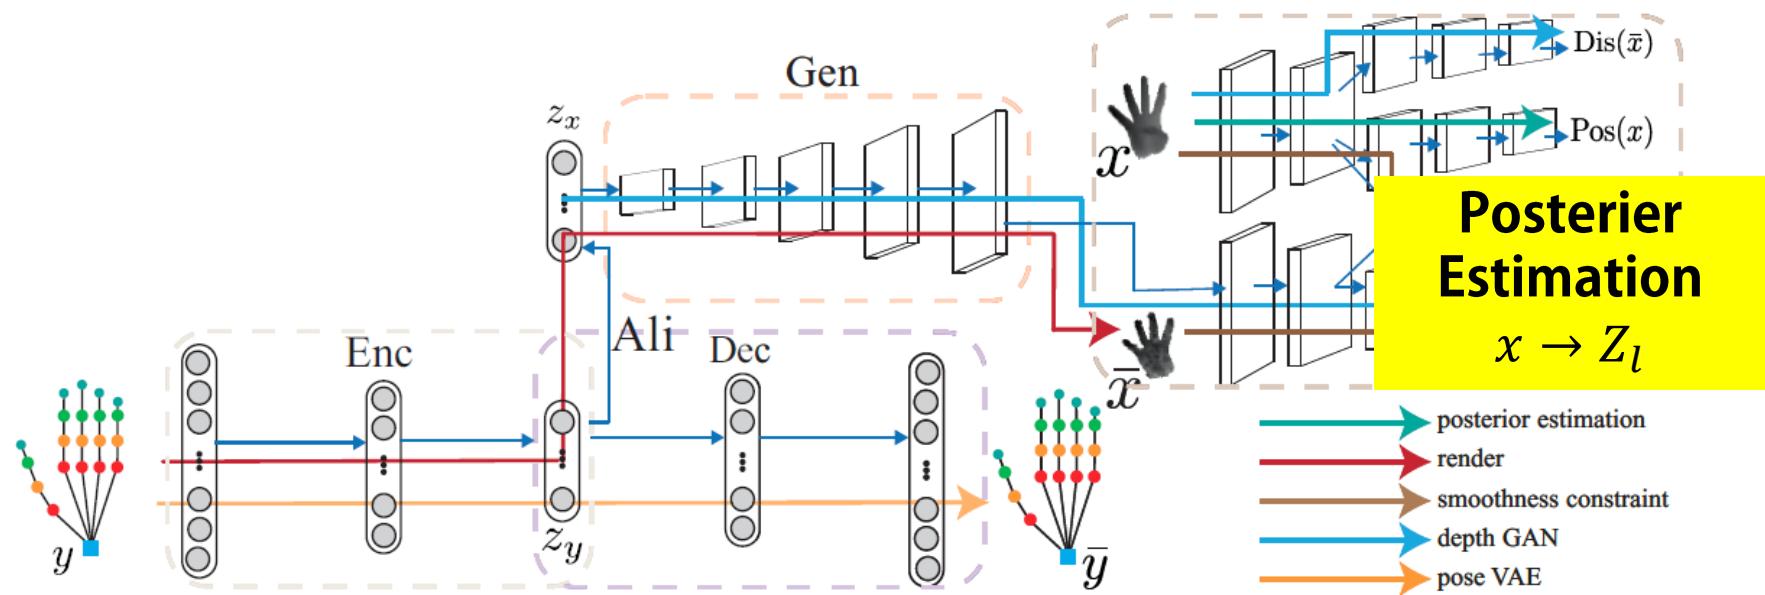
Pose Estimation

Crossing Nets: Combining GANs and VAEs with a Shared Latent Space for Hand Pose Estimation, CVPR 2017



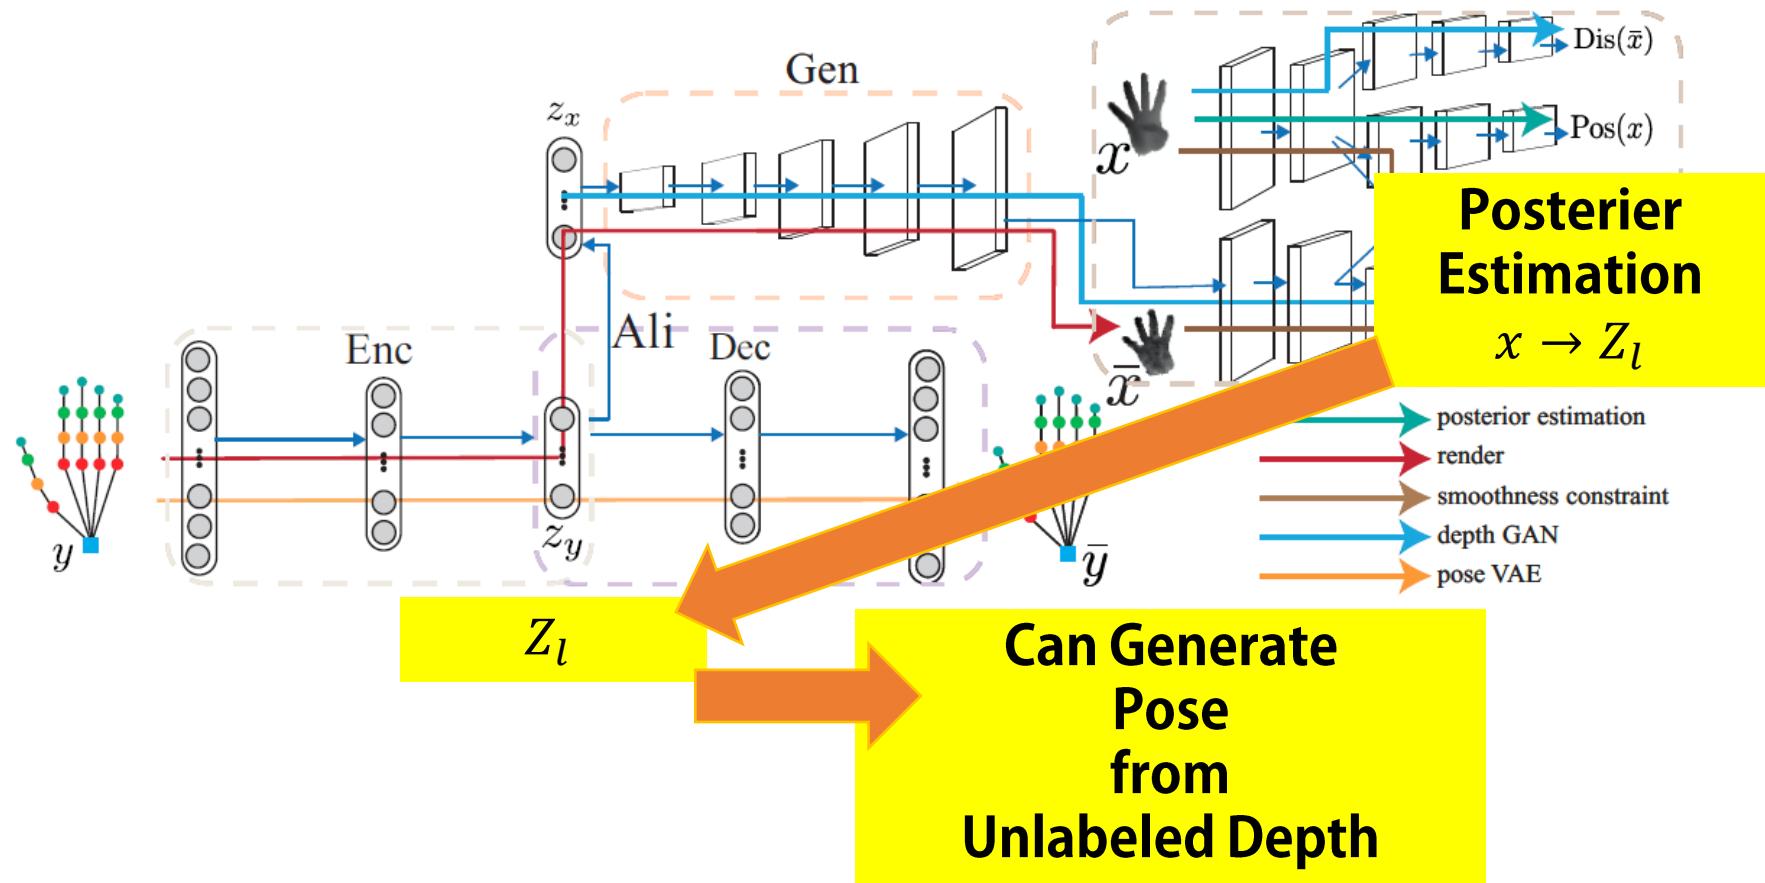
Pose Estimation

Crossing Nets: Combining GANs and VAEs with a Shared Latent Space for Hand Pose Estimation, CVPR 2017



Pose Estimation

Crossing Nets: Combining GANs and VAEs with a Shared Latent Space for Hand Pose Estimation, CVPR 2017

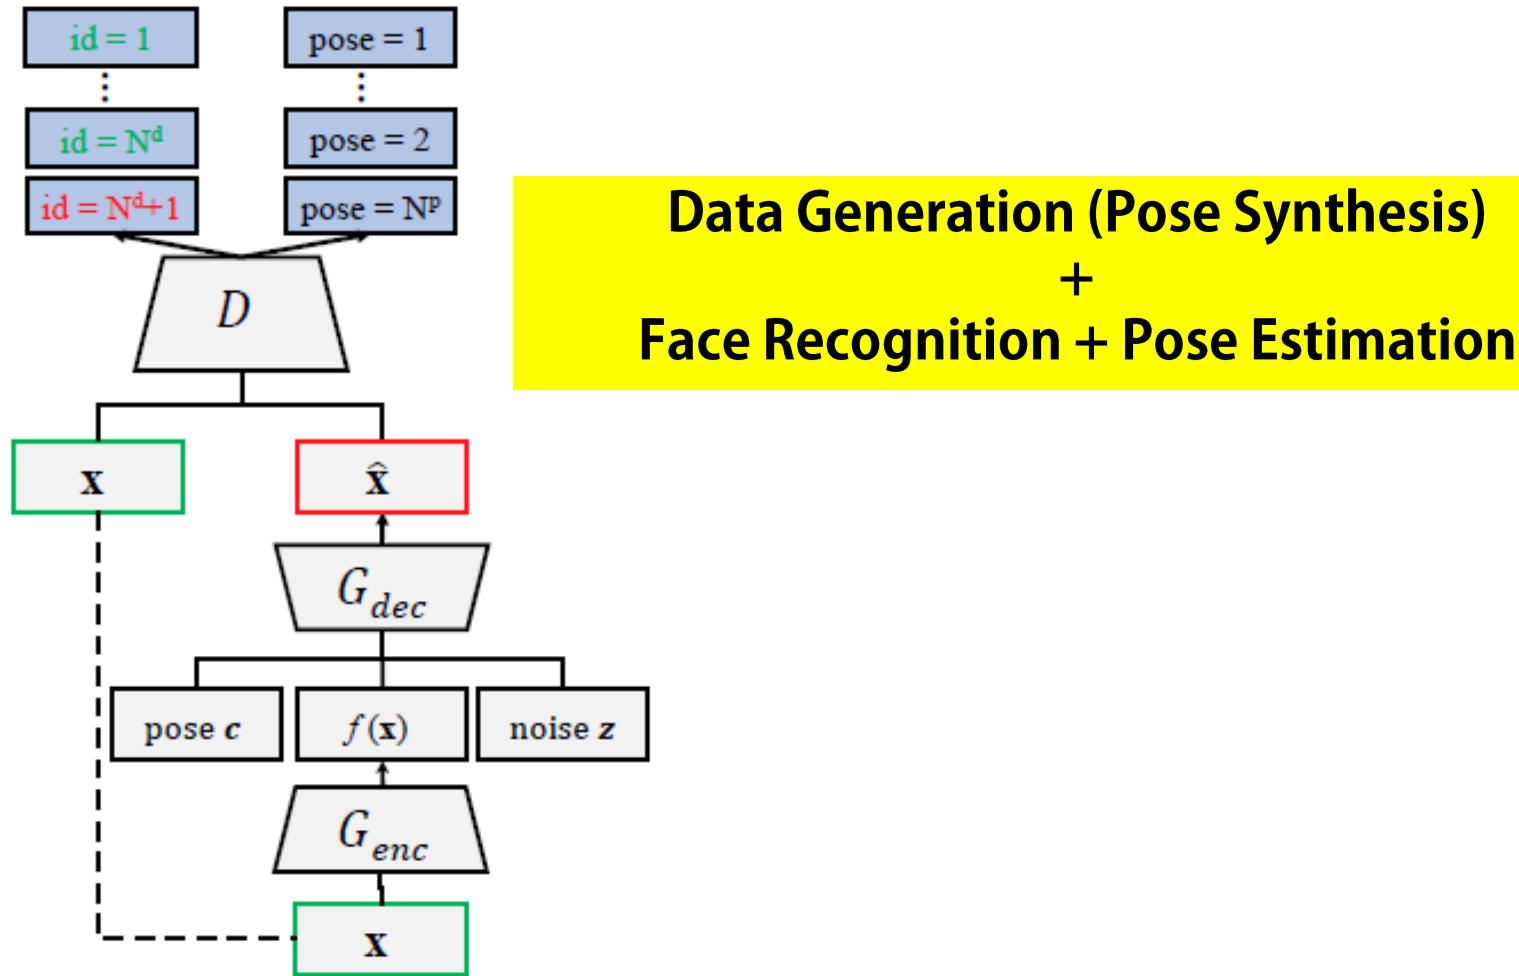


Face Recognition

Input	0°	15°	30°	45°	60°	
						L2
					DR-GAN	
						GT
						L2
					DR-GAN	
						GT

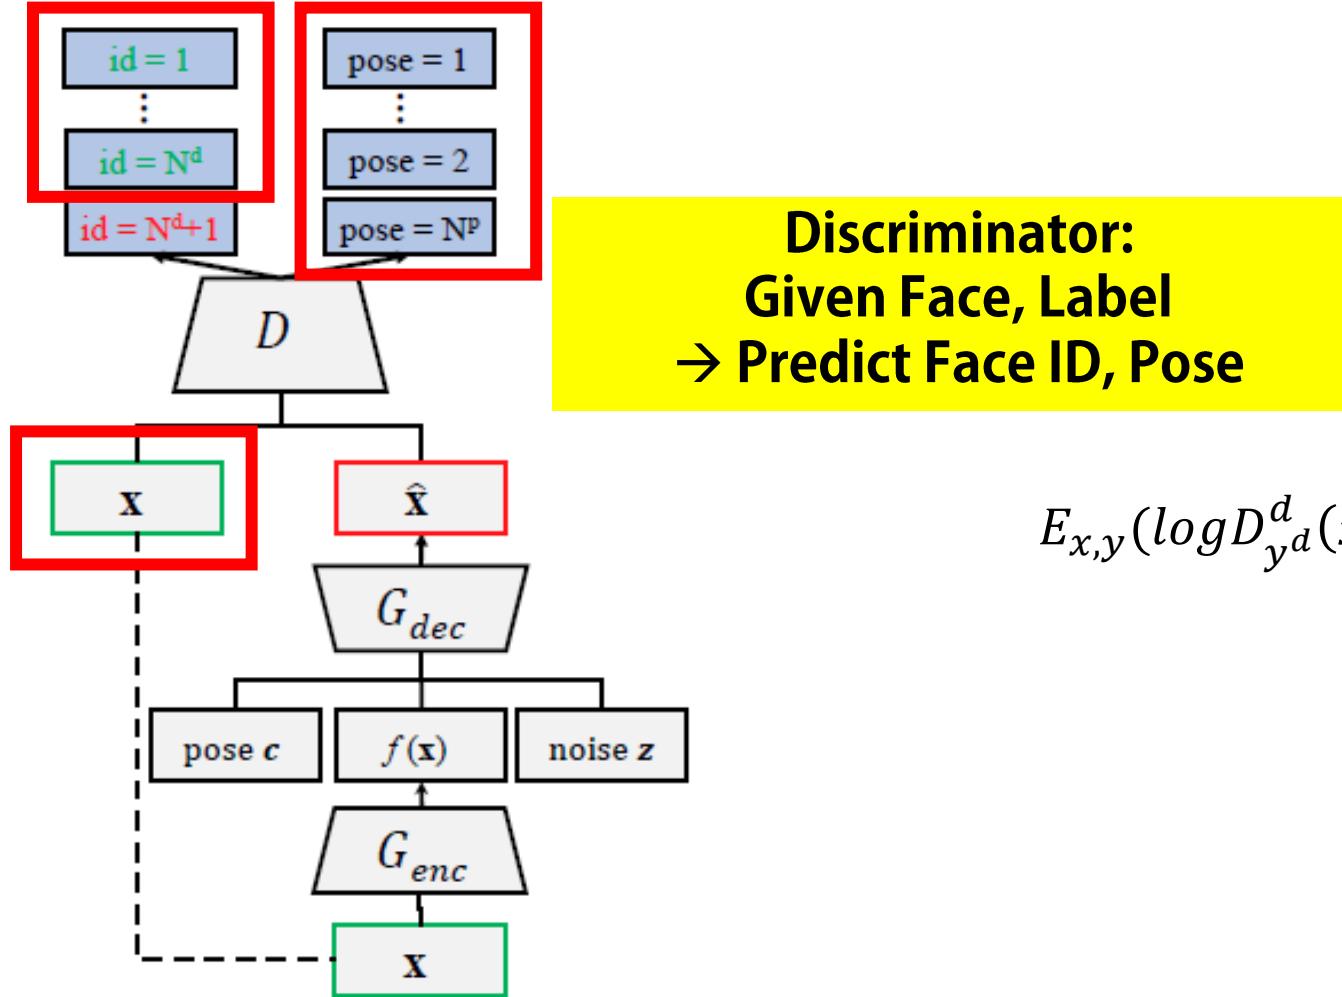
Face Recognition

Distangled Representation Learning GAN for Pose-Invariant Face Recognition
CVPR 2017



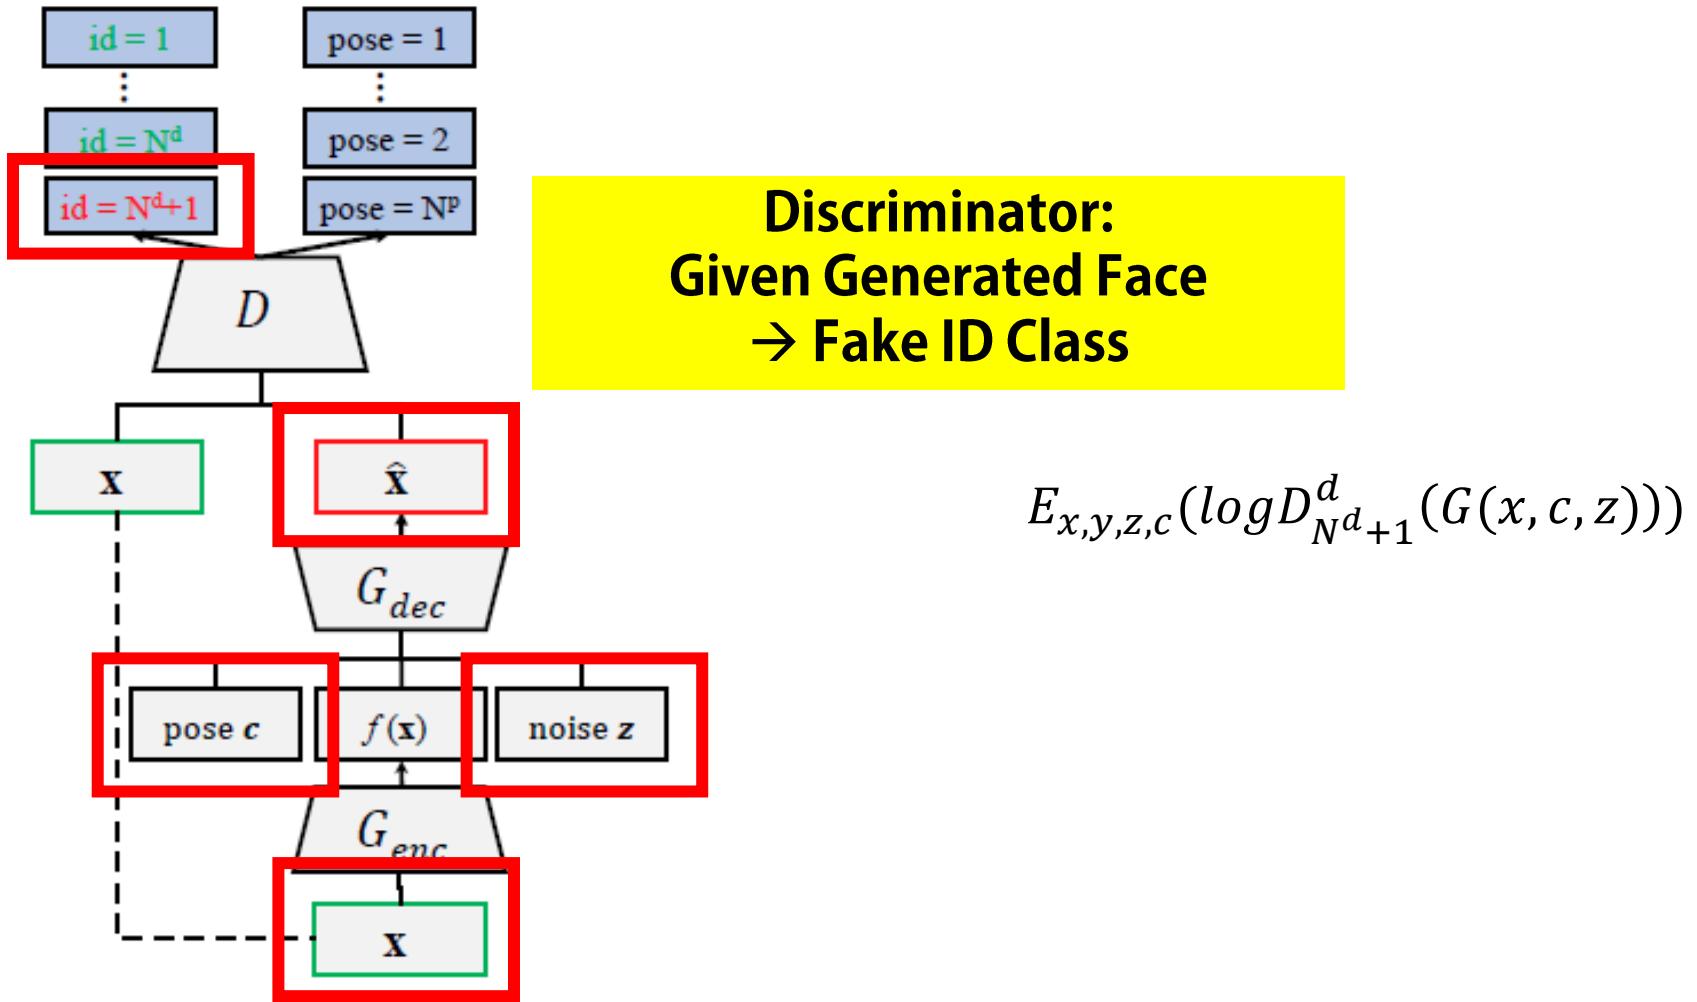
Face Recognition

Distangled Representation Learning GAN for Pose-Invariant Face Recognition
CVPR 2017



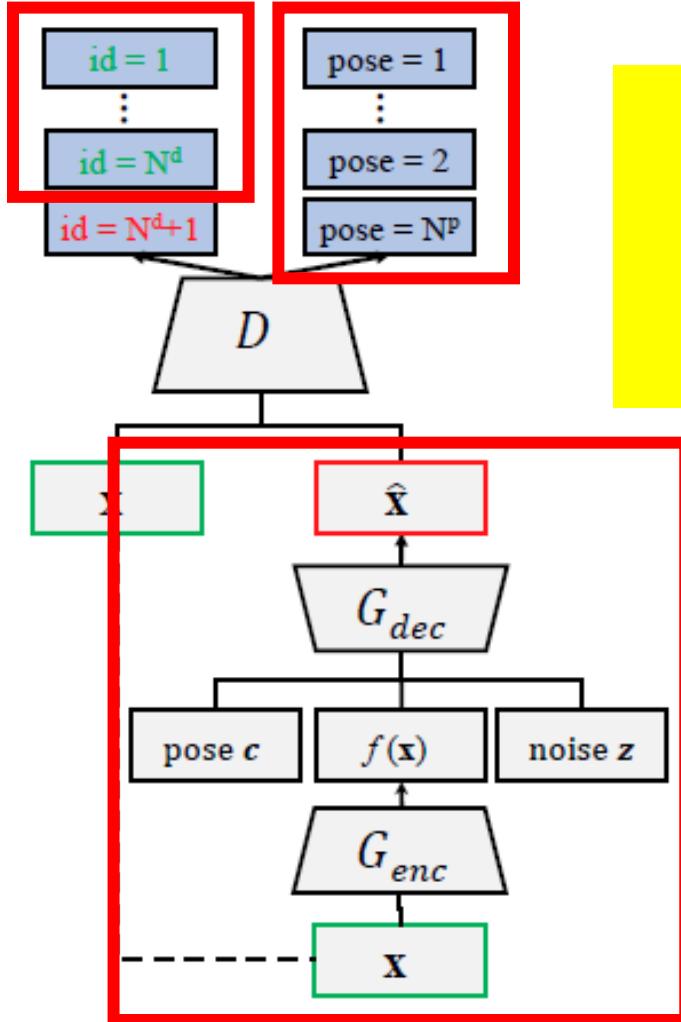
Face Recognition

Distangled Representation Learning GAN for Pose-Invariant Face Recognition
CVPR 2017



Face Recognition

Distangled Representation Learning GAN for Pose-Invariant Face Recognition
CVPR 2017

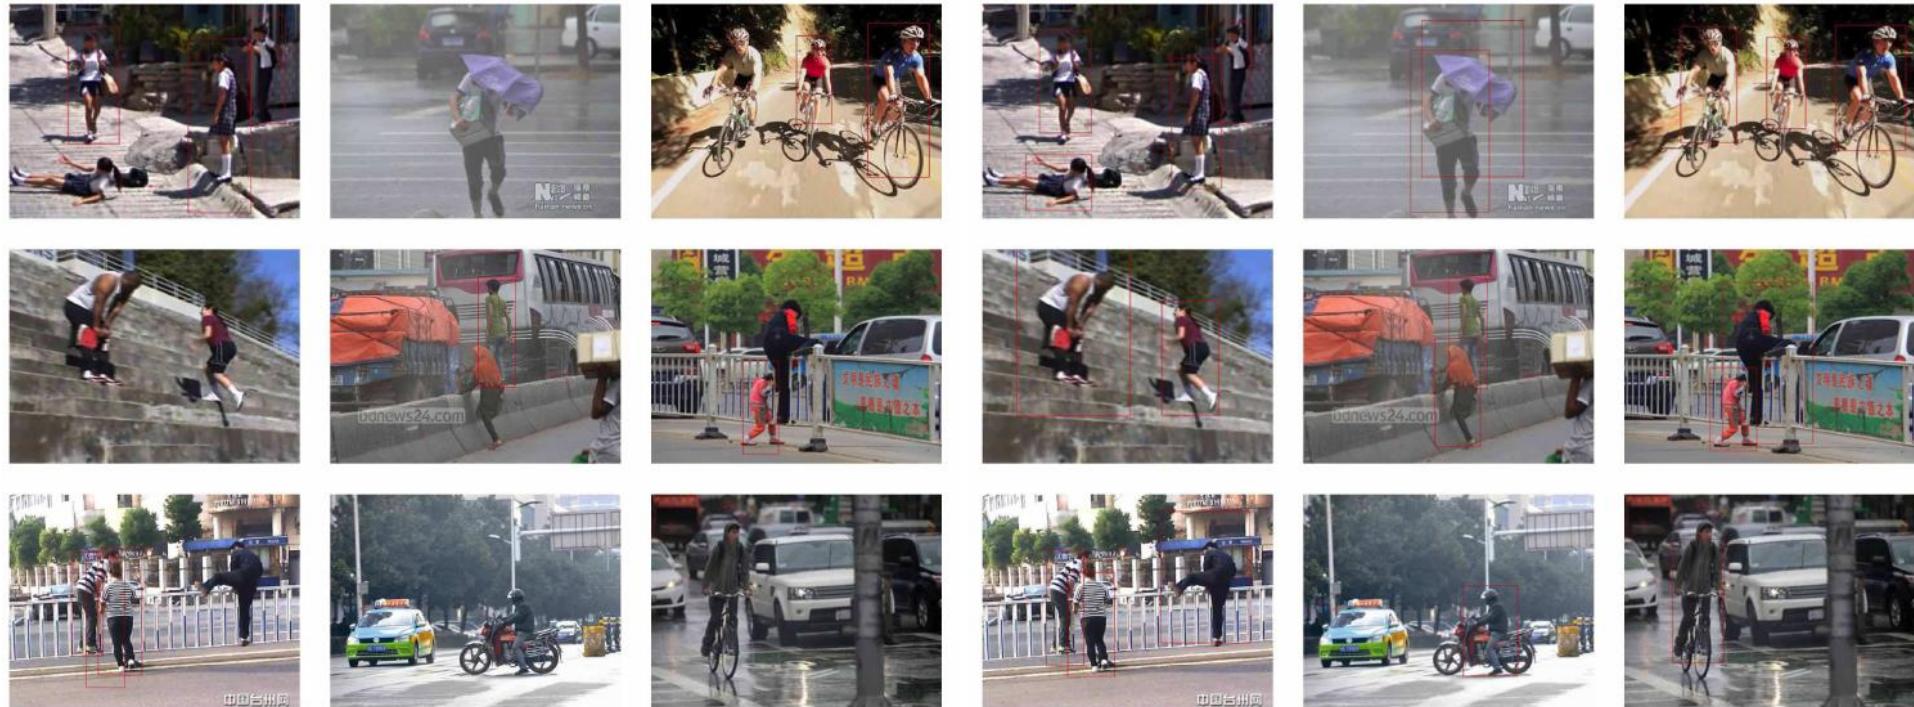


Generator:
Realistic Face Image
and Specific Pose

from Given Face

$$E_{x,y,z,c}(\log D_{y^d}^d(G(x, c, z)) + \log(D_{y^t}^p(G(x, c, z))))$$

Pedestrian Detection

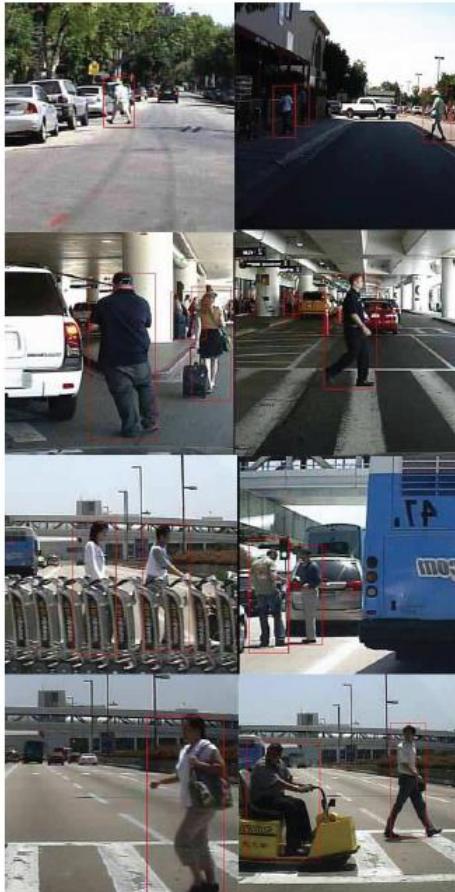


(a) RPN+Caltech

(b) Ours

Pedestrian Detection

Expecting the Unexpected: Training Detectors for Unusual Pedestrians with Adversarial Imposters, CVPR 2017



Caltech



Precarious



Synthesis

Pedestrian Detection

Expecting the Unexpected: Training Detectors for Unusual Pedestrians with Adversarial Imposters, CVPR 2017

	Range
Number of 3D models	[4, 8]
Index of background images	[0, 1726)
Index of 3D models	[0, 20)
Position of 3D models	Within the field of vision
Index of Animations	[0, <i>maxnumber</i>)
Time of animation	[0, 1]
Model's angle on the x axis	[-90°, 90°]
Model's angle on the y axis	[-180°, 180°]
Model's angle on the z axis	[-90°, 90°]
Light intensity	[0.5, 2]
Light's angle on the x axis	[-45°, 45°]
Light's angle on the y axis	[-45°, 45°]



Discriminator

Pedestrian Detection

Expecting the Unexpected: Training Detectors for Unusual Pedestrians with Adversarial Imposters, CVPR 2017

	Range
Number of 3D models	[4, 8]
Index of background images	[0, 1726)
Index of 3D models	[0, 20)
Position of 3D models	Within the field of vision
Index of Animations	[0, maxnumber)
Time of animation	[0, 1]
Model's angle on the x axis	[-90°, 90°]
Model's angle on the y axis	[-180°, 180°]
Model's angle on the z axis	[-90°, 90°]
Light intensity	[0.5, 2]
Light's angle on the x axis	[-45°, 45°]
Light's angle on the y axis	[-45°, 45°]



3D Engine

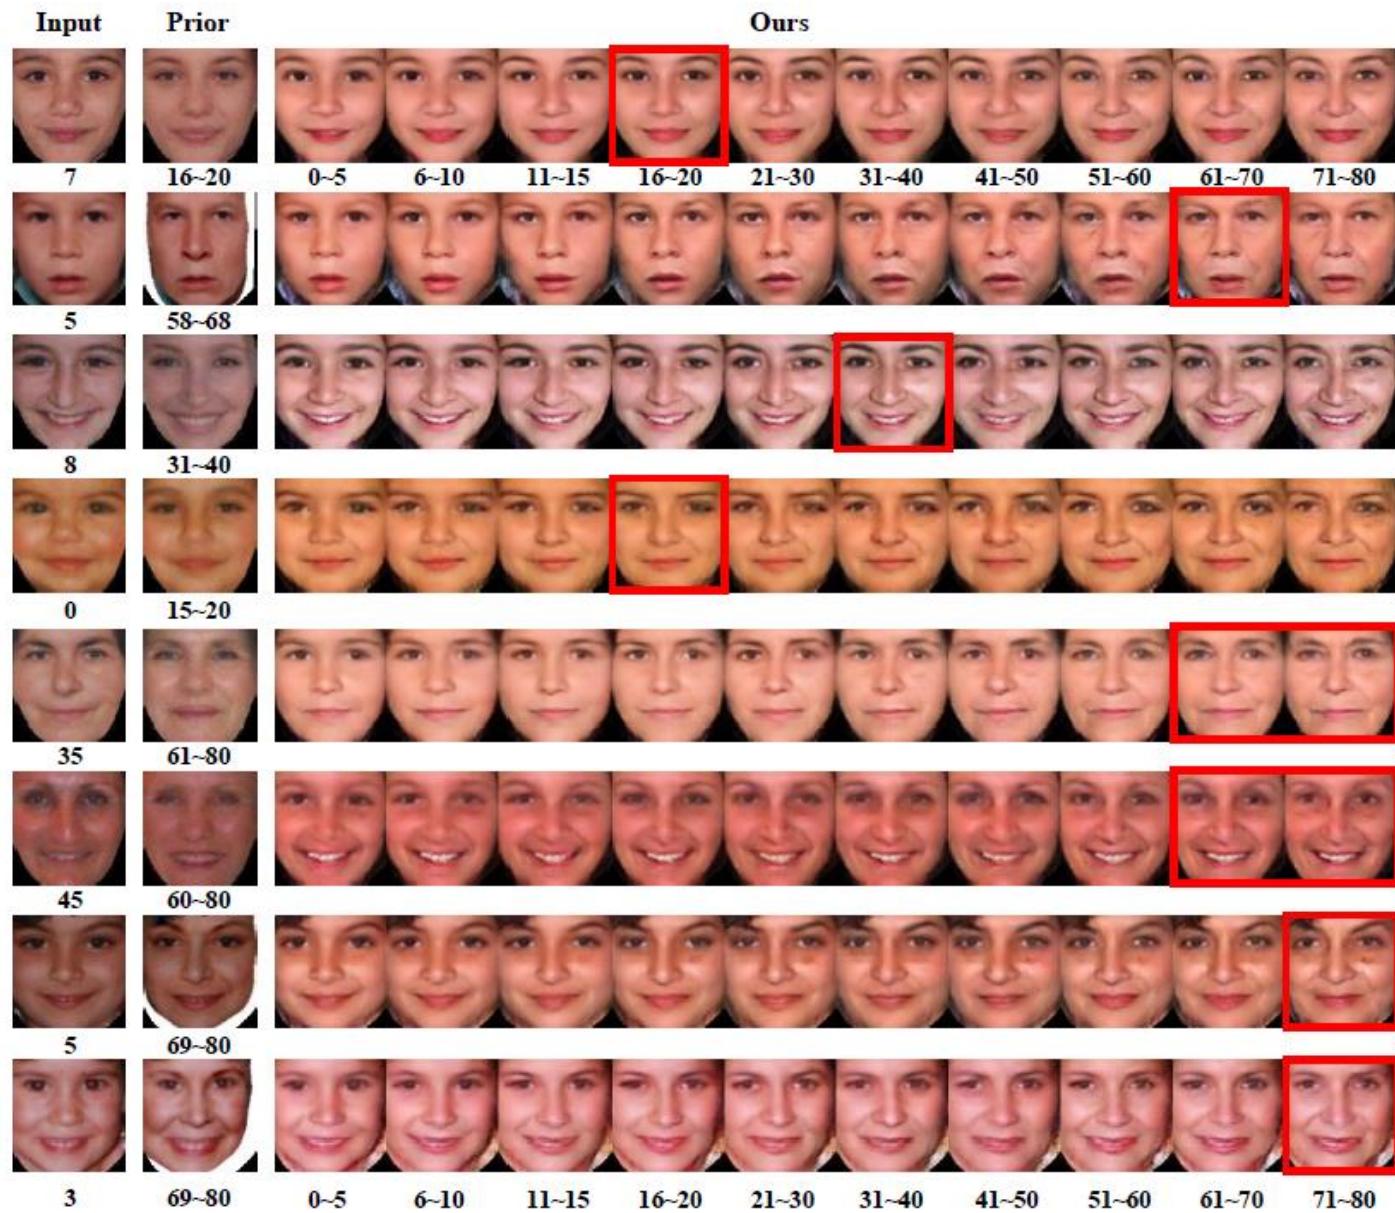
Train
Sampling
Parameter
Distribution



Discriminator

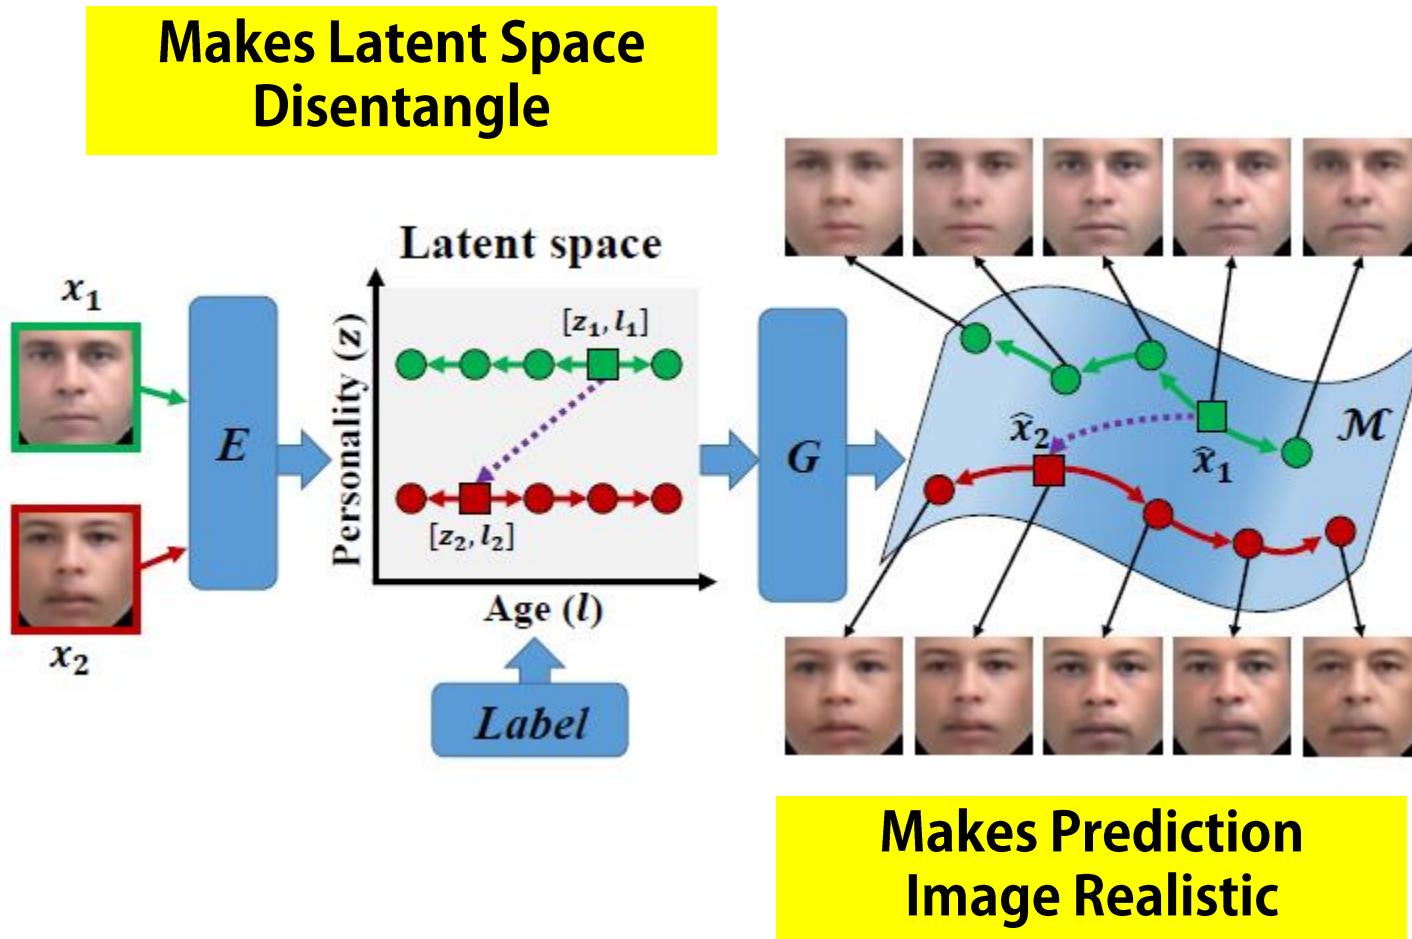
Other Approaches..

Modeling / Regression



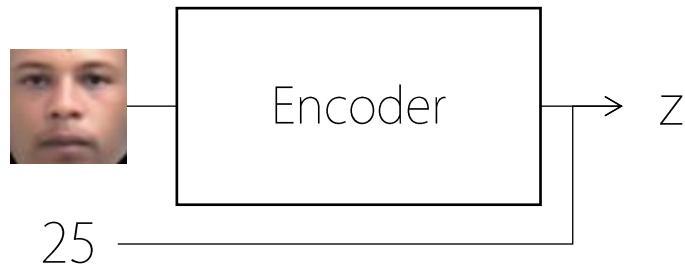
Face Aging Prediction

Age Prediction/Regression by Conditional Adversarial Autoencoder
CVPR 2017



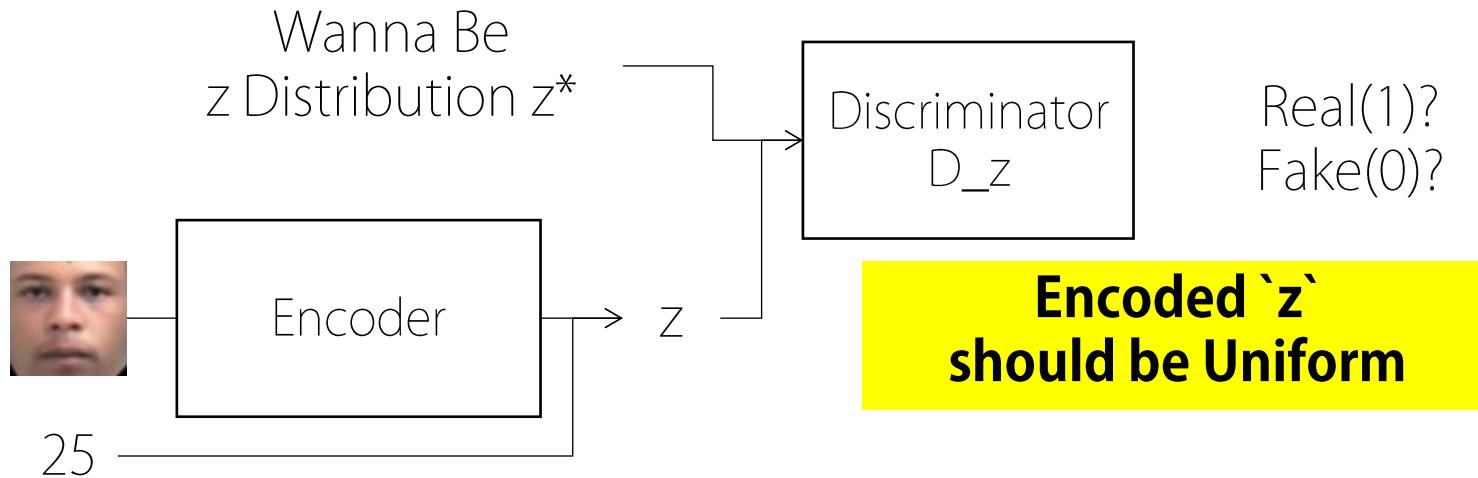
Face Aging Prediction

Age Prediction/Regression by Conditional Adversarial Autoencoder
CVPR 2017



Face Aging Prediction

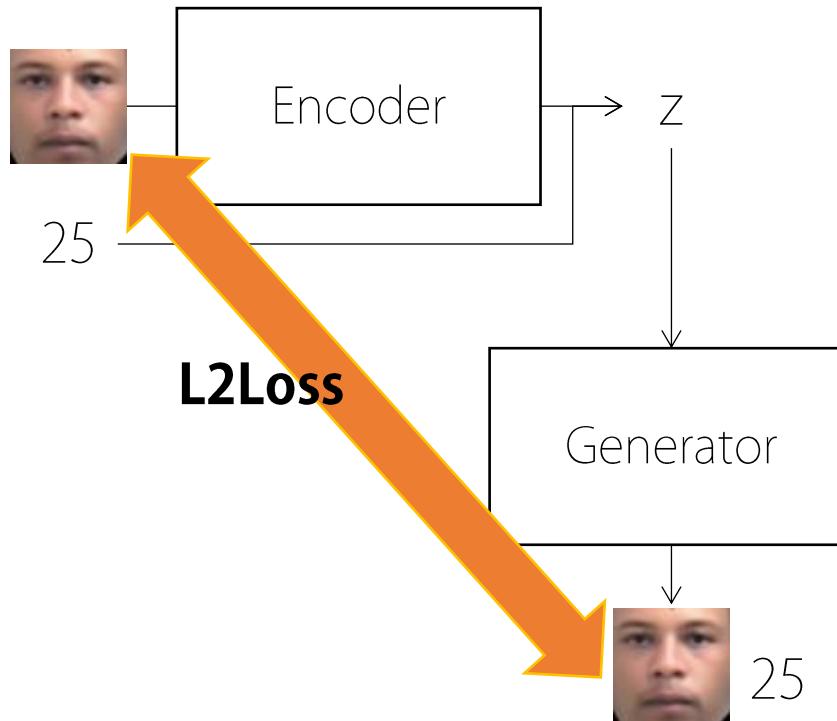
Age Prediction/Regression by Conditional Adversarial Autoencoder
CVPR 2017



$$E_{z^*}(\log D_z(z^*)) + E_x \left(\log \left(1 - D_z(E(x)) \right) \right)$$

Face Aging Prediction

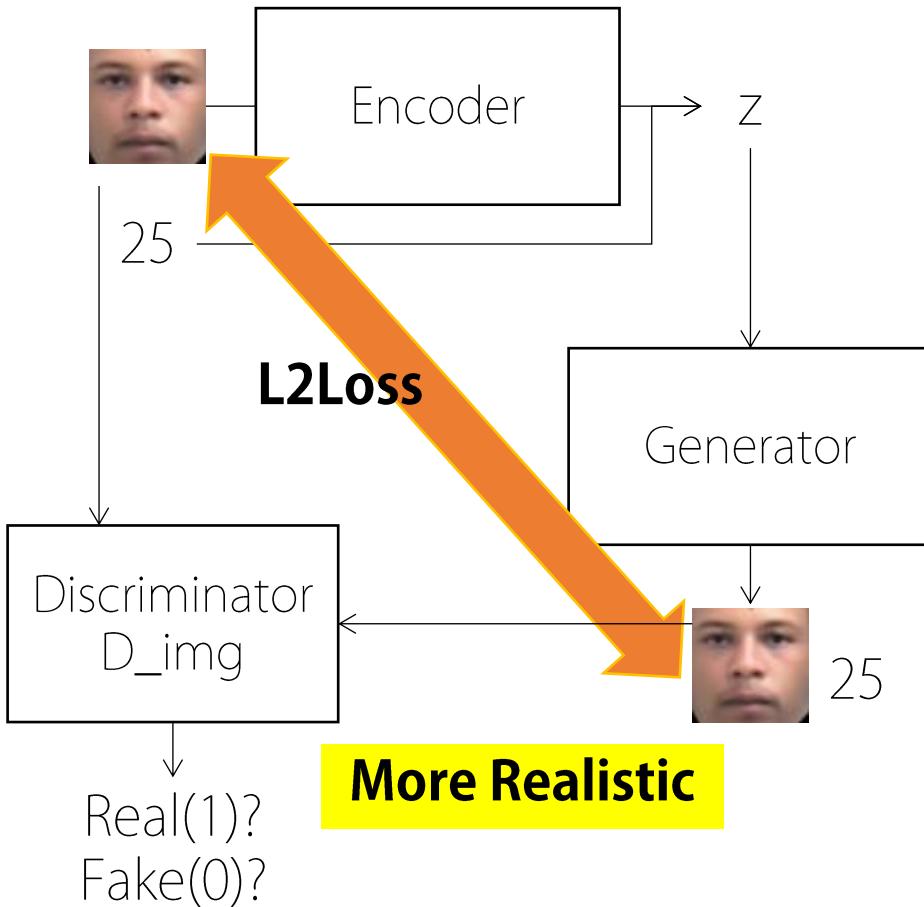
Age Prediction/Regression by Conditional Adversarial Autoencoder
CVPR 2017



$$L_2(x, G(E(x), l)) + TV(G(E(x)), l)$$

Face Aging Prediction

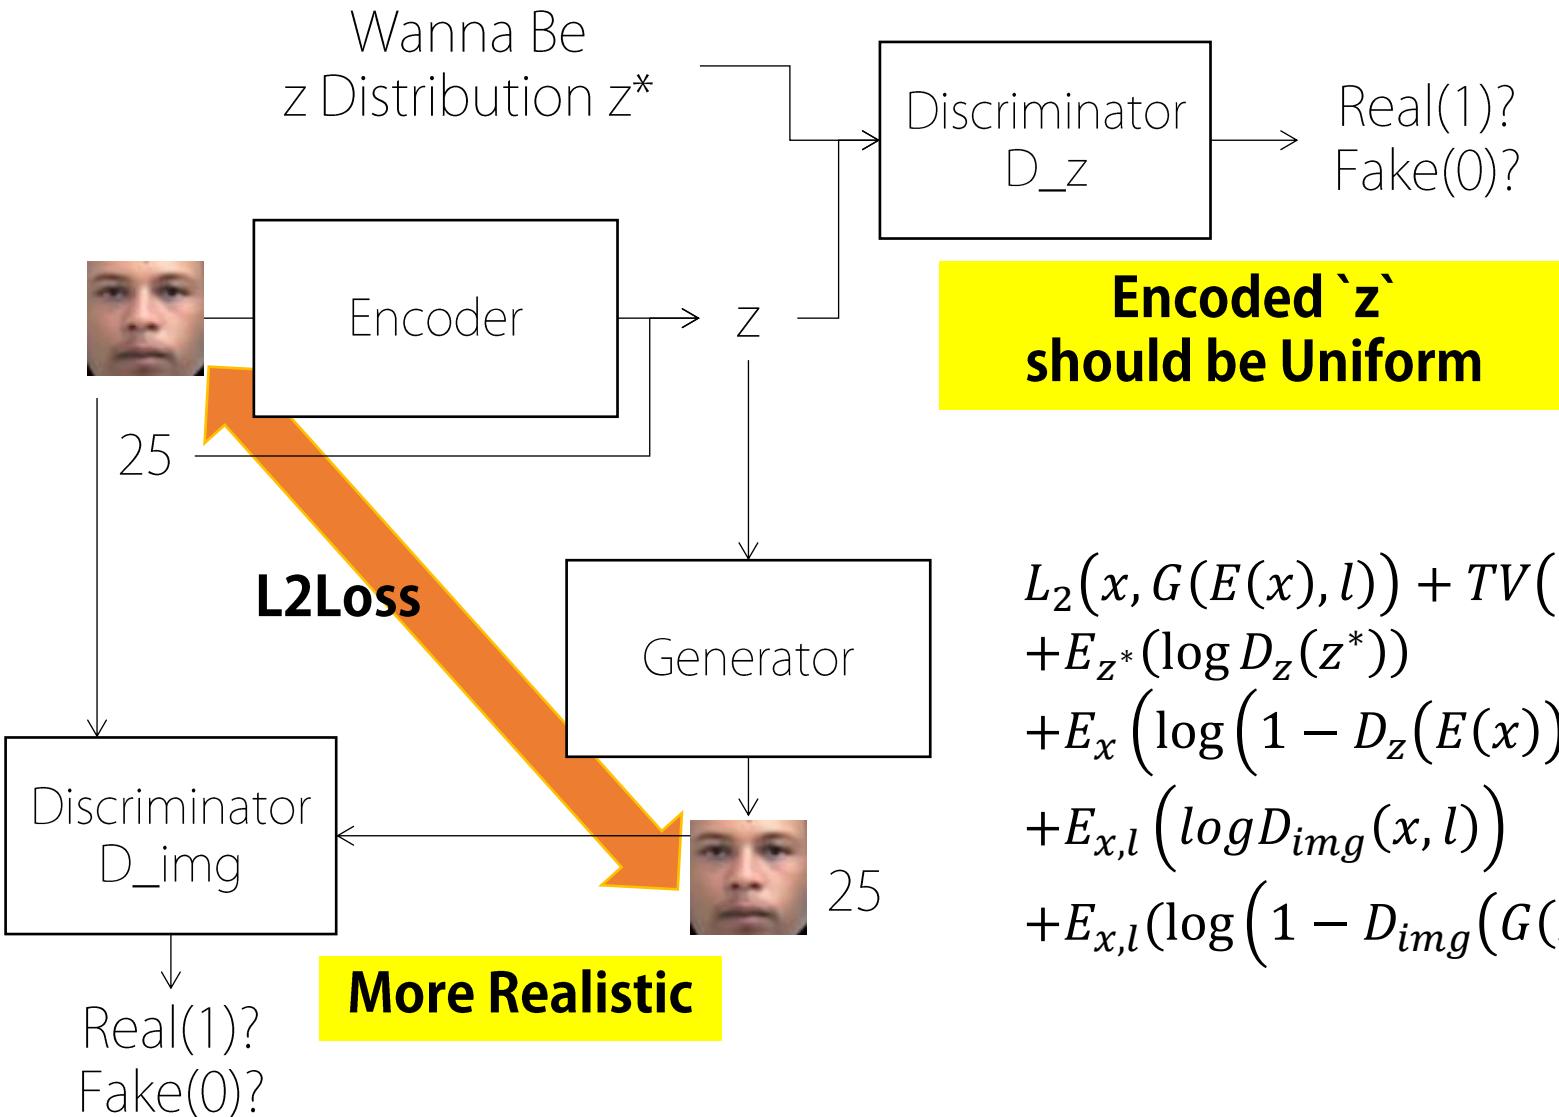
Age Prediction/Regression by Conditional Adversarial Autoencoder
CVPR 2017



$$E_{x,l} \left(\log D_{img}(x, l) \right) + E_{x,l} \left(\log \left(1 - D_{img}(G(E(x), l)) \right) \right)$$

Face Aging Prediction

Age Prediction/Regression by Conditional Adversarial Autoencoder
CVPR 2017

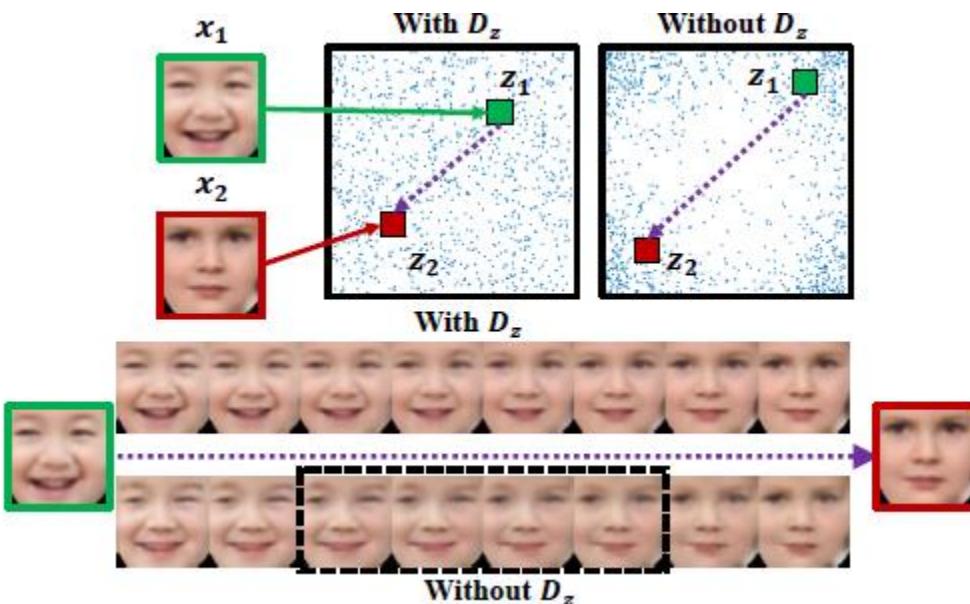


$$\begin{aligned} & L_2(x, G(E(x), l)) + TV(G(E(x)), l) \\ & + E_{z^*}(\log D_z(z^*)) \\ & + E_x(\log(1 - D_z(E(x)))) \\ & + E_{x,l}(\log D_{img}(x, l)) \\ & + E_{x,l}(\log(1 - D_{img}(G(E(x), l)))) \end{aligned}$$

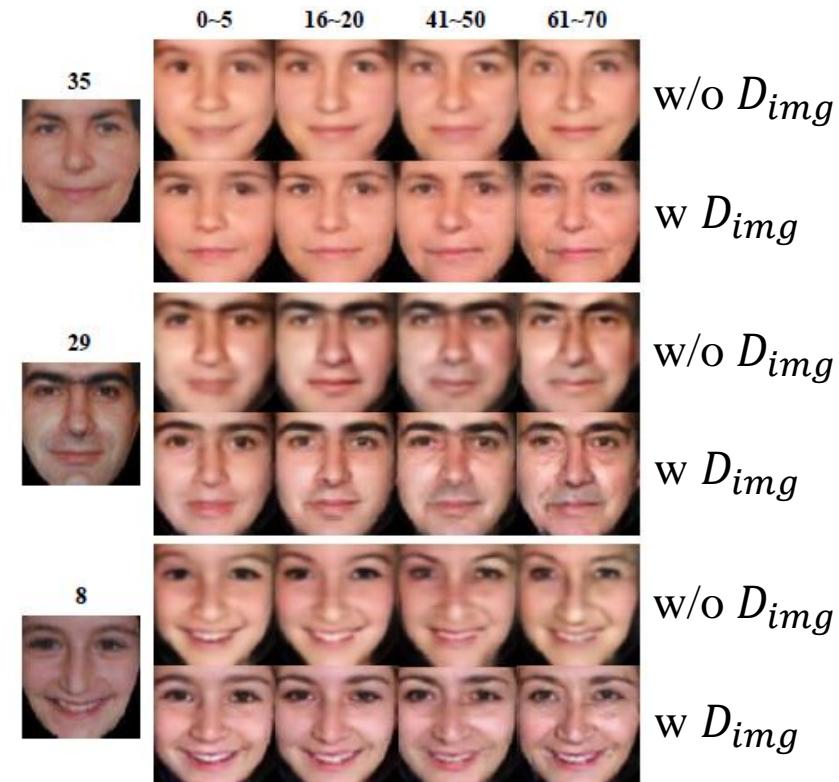
Face Aging Prediction

Age Prediction/Regression by Conditional Adversarial Autoencoder

CVPR 2017



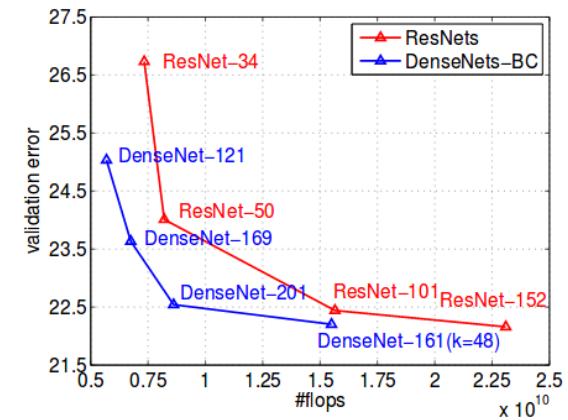
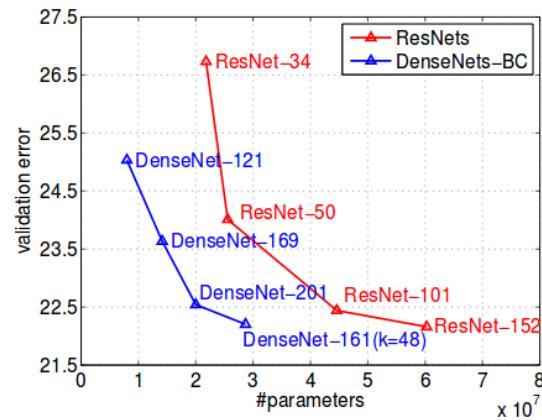
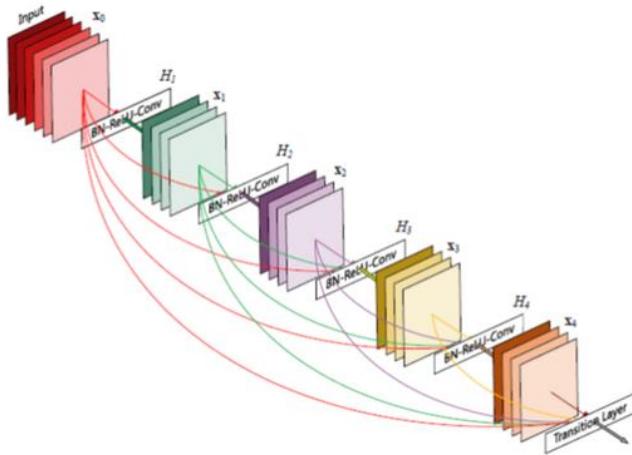
Effect of D_z



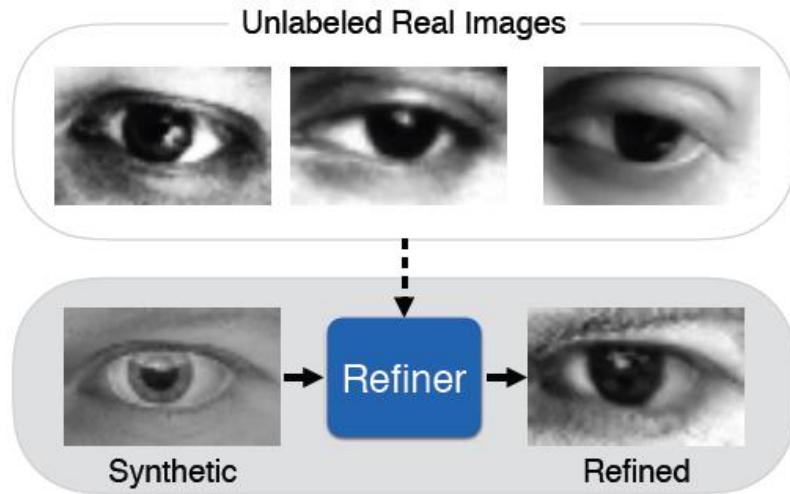
Effect of D_{img}

One More Thing

CVPR 2017 Best Paper



DenseNet, Facebook



S+U Learning,
Apple



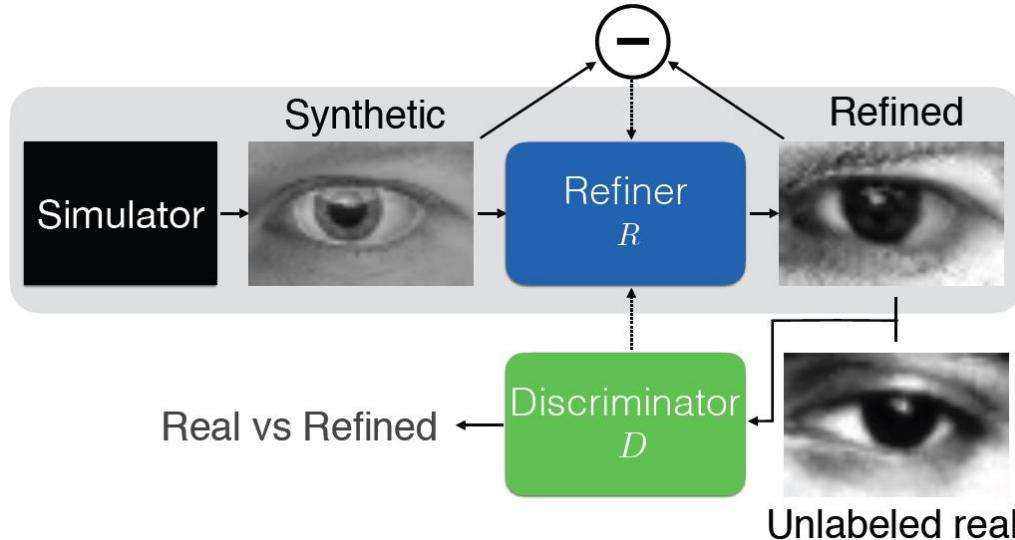
Apple Machine Learning Journal

Improving the Realism of Synthetic Images

Vol. 1, Issue 1 • July 2017

Most successful examples of neural nets today are trained with supervision. However, to achieve high accuracy, the training sets need to be large, diverse, and accurately annotated, which is costly. An alternative to labelling huge amounts of data is to use synthetic images from a simulator. This is cheap as there is no labeling cost, but the synthetic images may not be realistic enough, resulting in poor generalization on real test images. To help close this performance gap, we've developed a method for refining synthetic images to make them look more realistic. We show that training models on these refined images leads to significant improvements in accuracy on various machine learning tasks.

[Read the article >](#)



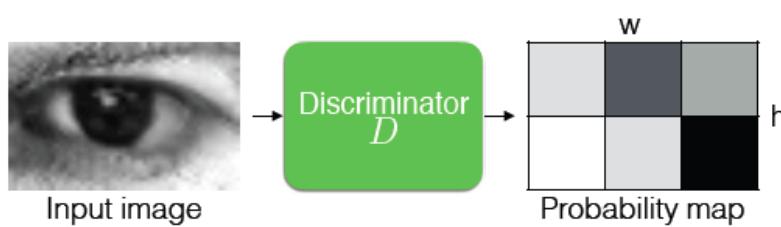
Key Idea:

Synthetic Sample with Label \rightarrow Bad Quality
using Unlabeled Real Image + **GAN**
 \rightarrow Good Quality Synthetic Sample
 \rightarrow Data Augmentation / Semi-Supervised L

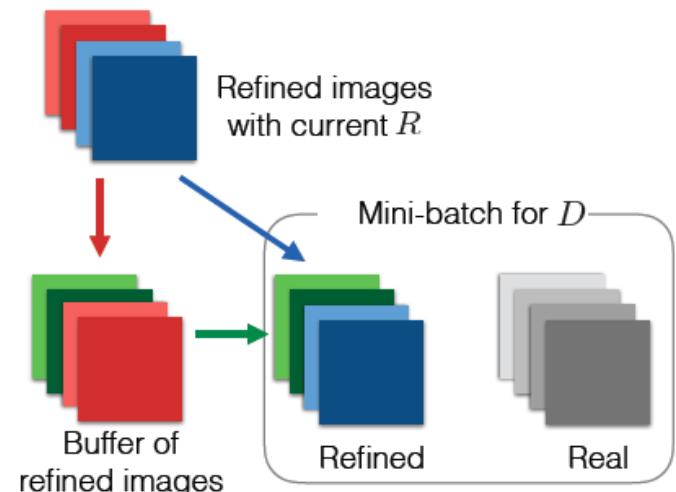
$$L_D = \log(D(X_{refined})) + \log(1 - D(X_{unlabeled}))$$

$$L_R = \log(1 - D(X_{refined})) + L1(CNN(X_{refined}) - CNN(X_{syn}))$$

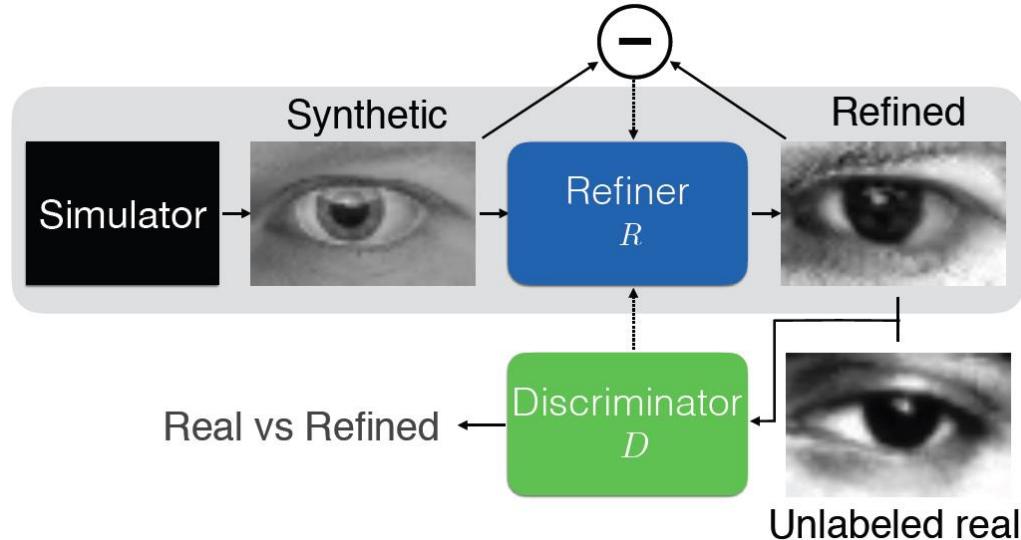
1. Perceptual Loss? (Self-Reg. Term)



2. PatchGAN? (Local Adversarial)



3. Minibatch Discrimination? (using History)



One Contribution
: Applied GAN \rightarrow Synthetic Refinement

Conclusion

GAN → Adversarial Training → Reinforcement

Generative adversarial nets

I Goodfellow, J Pouget-Abadie, M Mirza... - Advances in neural ... , 2014 - papers.nips.cc

Abstract We propose a new framework for estimating generative models via adversarial nets, in which we simultaneously train two models: a generative model G that captures the data distribution, and a discriminative model D that estimates the probability that a sample

1021회 인용 관련 학술자료 전체 18개의 버전 인용 저장

Introduction to GANs

Semi-supervised learning with GANs

Practical tricks and tips for training GANs

Plug and Play Generative Models

Latent variable inference for GANs

Approximate Nash equilibria in neural net GANs

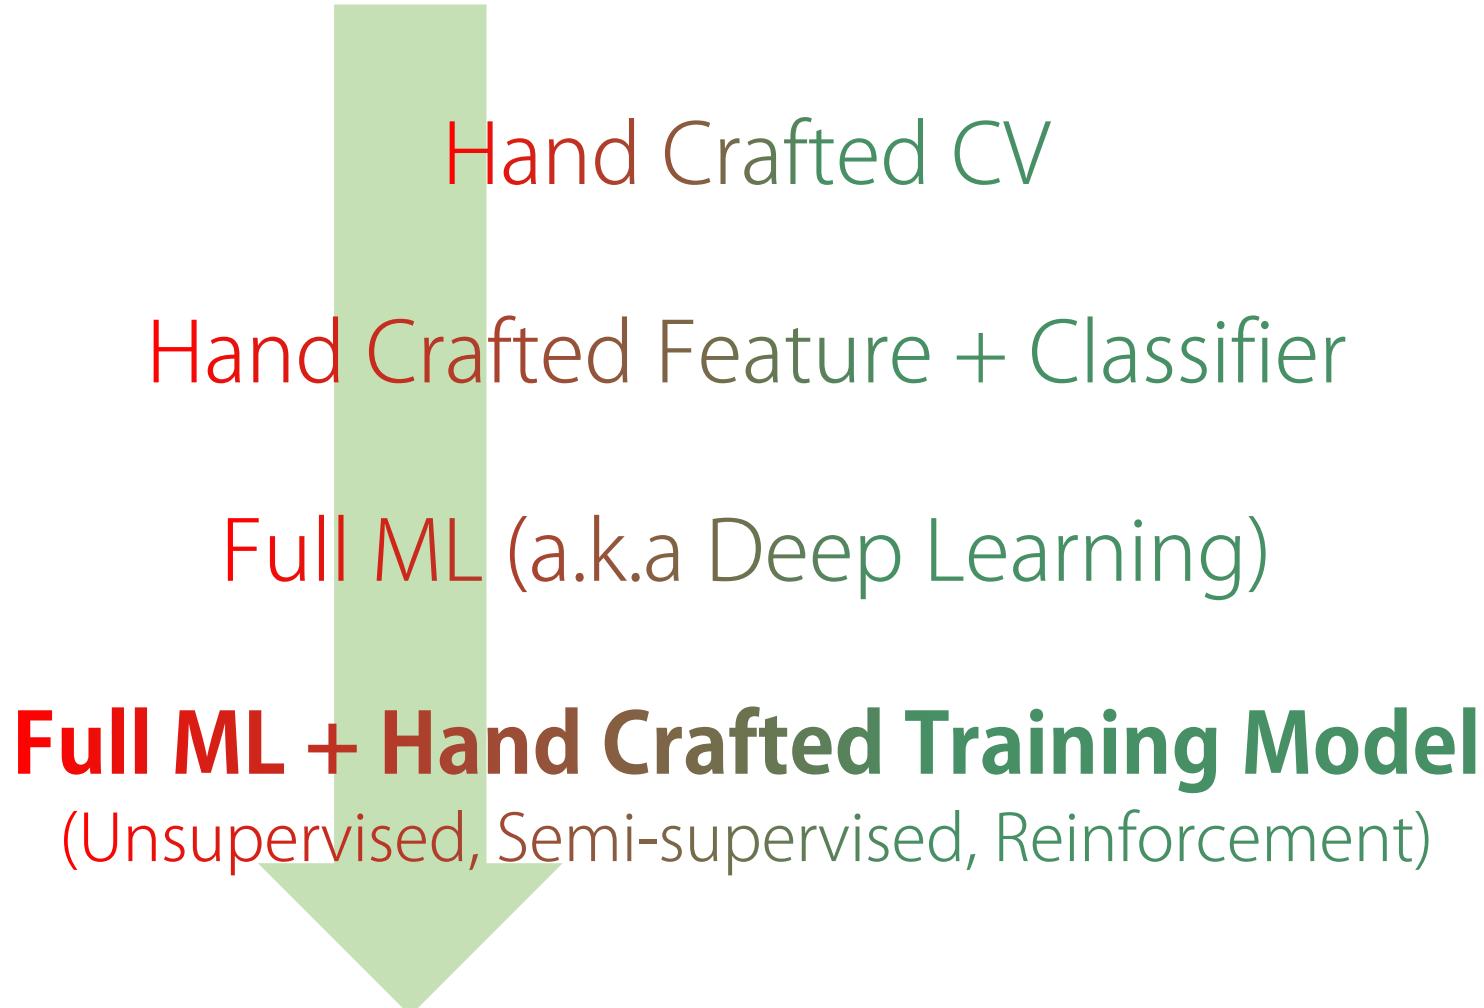
Evaluating generative models

Interactive image generation with GANs

Connections between adversarial training and reinforcement learning

Ian Goodfellow's GAN
Tutorial
Will be held in ICCV2017
22th October

Importance of GANs



Importance of GANs

Generative Models + Adversarial Training

GMM, HMM
Naïve Bayes
RBM, Autoencoders

Core Idea

References

* Image to Image Translation

- Image to Image Translation with Conditional Adversarial Networks, CVPR 2017
- Unpaired Image to Image Translation using Cycle-consistent Adversarial Networks, ICCV 2017
- Learning to Discover Cross-domain Relations with Generative Adversarial Networks, ICML 2017

* Single Image Super-Resolution

- Photo-realistic Single Image Super-Resolution using a Generative Adversarial Networks, CVPR 2017
- EnhanceNet: Single Image Super Resolution through Automated Texture Synthesis, ICCV 2017

* Feature Upscaling & 3D Reconstruction

- Perceptual Generative Adversarial Networks for Small Object Detection, CVPR 2017
- Synthesizing 3D Shapes via Modeling Multi-View Depth Maps and Silhouettes with Deep Generative Networks, CVPR 2017

* Pose Estimation, Face Recognition, Pedestrian Detection using Data Augmentation

- Crossing Nets: Combining GANs and VAEs with a Shared Latent Space for Hand Pose Estimation, CVPR 2017
- Distangled Representation Learning GAN for Pose-Invariant Face Recognition, CVPR 2017
- Expecting the Unexpected: Training Detectors for Unusual Pedestrians with Adversarial Imposters, CVPR 2017

* Modeling and Feature Extraction

- Age Prediction/Regression by Conditional Adversarial Autoencoder, CVPR 2017

-

* One More Thing

- Learning from Simulated and Unsupervised Images through Adversarial Training, CVPR 2017 Best Paper

More Information:

<https://taeoh-kim.github.io/blog>

Thank you!

Any Questions?